

Assignment01

August 7, 2024

Name:- Sk Fardeen Hossain

Enrollment ID:- 2021CSB023

G-Suite ID:- 2021csb023.sk@students.iiests.ac.in

1 Assignment 01

```
[1]: ## install pandas dataframe library
```

```
!pip install pandas
```

```
Requirement already satisfied: pandas in
/home/fardeen/Desktop/MachineLearningLab/venv/lib/python3.12/site-packages
(2.2.2)
Requirement already satisfied: numpy>=1.26.0 in
/home/fardeen/Desktop/MachineLearningLab/venv/lib/python3.12/site-packages (from
pandas) (2.0.1)
Requirement already satisfied: python-dateutil>=2.8.2 in
/home/fardeen/Desktop/MachineLearningLab/venv/lib/python3.12/site-packages (from
pandas) (2.9.0.post0)
Requirement already satisfied: pytz>=2020.1 in
/home/fardeen/Desktop/MachineLearningLab/venv/lib/python3.12/site-packages (from
pandas) (2024.1)
Requirement already satisfied: tzdata>=2022.7 in
/home/fardeen/Desktop/MachineLearningLab/venv/lib/python3.12/site-packages (from
pandas) (2024.1)
Requirement already satisfied: six>=1.5 in
/home/fardeen/Desktop/MachineLearningLab/venv/lib/python3.12/site-packages (from
python-dateutil>=2.8.2->pandas) (1.16.0)
```

1.1 Question 01 / 02

```
[6]: import pandas as pd
TRAIN_DATASET_PATH = "../ML_DRIVE/Assignment01/train.csv"

houses_dataframe = pd.read_csv(TRAIN_DATASET_PATH)
```

```
houses_dataframe.head()
```

```
[6]:
```

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	\
0	1	60	RL	65.0	8450	Pave	NaN	Reg	
1	2	20	RL	80.0	9600	Pave	NaN	Reg	
2	3	60	RL	68.0	11250	Pave	NaN	IR1	
3	4	70	RL	60.0	9550	Pave	NaN	IR1	
4	5	60	RL	84.0	14260	Pave	NaN	IR1	

	LandContour	Utilities	...	PoolArea	PoolQC	Fence	MiscFeature	MiscVal	MoSold	\
0	Lvl	AllPub	...	0	NaN	NaN	NaN	0	2	
1	Lvl	AllPub	...	0	NaN	NaN	NaN	0	5	
2	Lvl	AllPub	...	0	NaN	NaN	NaN	0	9	
3	Lvl	AllPub	...	0	NaN	NaN	NaN	0	2	
4	Lvl	AllPub	...	0	NaN	NaN	NaN	0	12	

	YrSold	SaleType	SaleCondition	SalePrice
0	2008	WD	Normal	208500
1	2007	WD	Normal	181500
2	2008	WD	Normal	223500
3	2006	WD	Abnorml	140000
4	2008	WD	Normal	250000

[5 rows x 81 columns]

1.2 Estimate missing values in the dataframe

First we will drop the columns having at least one 'NaN' value using the `dropna()` function. Here we will be using the `interpolate()` method provided by the pandas library to fill the missing values, method used will be forward average for equally spaced values

```
[7]: houses_dataframe=houses_dataframe.dropna(axis=1,thresh=500)
houses_dataframe=houses_dataframe.interpolate(method="linear",
limit_direction="forward")
houses_dataframe.columns
```

```
[7]: Index(['Id', 'MSSubClass', 'MSZoning', 'LotFrontage', 'LotArea', 'Street',
'LotShape', 'LandContour', 'Utilities', 'LotConfig', 'LandSlope',
'Neighborhood', 'Condition1', 'Condition2', 'BldgType', 'HouseStyle',
'OverallQual', 'OverallCond', 'YearBuilt', 'YearRemodAdd', 'RoofStyle',
'RoofMatl', 'Exterior1st', 'Exterior2nd', 'MasVnrType', 'MasVnrArea',
'ExterQual', 'ExterCond', 'Foundation', 'BsmtQual', 'BsmtCond',
'BsmtExposure', 'BsmtFinType1', 'BsmtFinSF1', 'BsmtFinType2',
'BsmtFinSF2', 'BsmtUnfSF', 'TotalBsmtSF', 'Heating', 'HeatingQC',
'CentralAir', 'Electrical', '1stFlrSF', '2ndFlrSF', 'LowQualFinSF',
'GrLivArea', 'BsmtFullBath', 'BsmtHalfBath', 'FullBath', 'HalfBath',
'BedroomAbvGr', 'KitchenAbvGr', 'KitchenQual', 'TotRmsAbvGrd',
```

```

'Functional', 'Fireplaces', 'FireplaceQu', 'GarageType', 'GarageYrBlt',
'GarageFinish', 'GarageCars', 'GarageArea', 'GarageQual', 'GarageCond',
'PavedDrive', 'WoodDeckSF', 'OpenPorchSF', 'EnclosedPorch', '3SsnPorch',
'ScreenPorch', 'PoolArea', 'MiscVal', 'MoSold', 'YrSold', 'SaleType',
'SaleCondition', 'SalePrice'],
dtype='object')

```

1.3 Question 03

1.4 Estimating Saleprice based on LotArea using Linear Regression

```

[13]: ## Compress the data-frame into ['saleprice','lotarea']

sale_price_lot_area_df = houses_dataframe[['LotArea','SalePrice']]
sale_price_lot_area_df.head()

```

```

[13]:
   LotArea  SalePrice
0     8450    208500
1     9600    181500
2    11250    223500
3     9550    140000
4    14260    250000

```

1.4.1 Before Performing prediction for new values, we need to define our Linear Regression Model

We will import and use the sklearn library for this purpose

```

[14]: !pip install scikit-learn

Requirement already satisfied: scikit-learn in
/home/fardeen/Desktop/MachineLearningLab/venv/lib/python3.12/site-packages
(1.5.1)
Requirement already satisfied: numpy>=1.19.5 in
/home/fardeen/Desktop/MachineLearningLab/venv/lib/python3.12/site-packages (from
scikit-learn) (2.0.1)
Requirement already satisfied: scipy>=1.6.0 in
/home/fardeen/Desktop/MachineLearningLab/venv/lib/python3.12/site-packages (from
scikit-learn) (1.14.0)
Requirement already satisfied: joblib>=1.2.0 in
/home/fardeen/Desktop/MachineLearningLab/venv/lib/python3.12/site-packages (from
scikit-learn) (1.4.2)
Requirement already satisfied: threadpoolctl>=3.1.0 in
/home/fardeen/Desktop/MachineLearningLab/venv/lib/python3.12/site-packages (from
scikit-learn) (3.5.0)

```


```

[15]: ## Plot the saleprice vs lotarea
import matplotlib.pyplot as plt

```

```
plt.
    ↳scatter(x=sale_price_lot_area_df['LotArea'],y=sale_price_lot_area_df['SalePrice'])
plt.xlabel('LotArea(in sq.feet)')
plt.ylabel('SalePrice')
```


```
[15]: Text(0, 0.5, 'SalePrice')
```



We can observe that the ranges of the X and Y axis are not same, so we can normalise the values of the saleprice and lotarea.

```
[11]: from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler() ## (x-xmin)/(max-xmin)
sale_price_lot_area_df=scaler.fit_transform(sale_price_lot_area_df)
sale_price_lot_area_df=pd.
    ↳DataFrame(sale_price_lot_area_df,columns=['LotArea', 'SalePrice'])
plt.
    ↳scatter(x=sale_price_lot_area_df['LotArea'],y=sale_price_lot_area_df['SalePrice'])
plt.xlabel('LotArea(in sq.feet)')
plt.ylabel('SalePrice')
```

```
[11]: Text(0, 0.5, 'SalePrice')
```



1.4.2 Implementing a utility K-fold cross validator for train-test split

```
[16]: # Here we will be considering 5 folds
from sklearn.model_selection import KFold
from sklearn.metrics import mean_squared_error

## R2-Score = 1-( RSS/TSS ) ; RSS --> Residual Sum of Squares , TSS --> Total
↳ Sum of Squares
'''
    @brief:- Utility function to perform KfoldCross Validation on a dataset
    @param:- model - Estimator
              X_data - Independant Variable
              y_data - Dependant Variable

    @return:- Tuple of size 2, 1st being the average R2 score and the 2nd one
↳ being average MSE
'''

def Kfold_util(model,X_data,y_data):
    s_kfold = KFold(n_splits=5)
    model_r2_score = 0
```

```

model_MSE_score = 0
for i,(train_index,test_index) in enumerate(s_kfold.split(X_data,y_data)):
    print(f"Fold: {i+1}")
    model.fit(X_data[train_index],y_data[train_index])
    r2_score = model.score(X_data[test_index],y_data[test_index])
    y_pred = model.predict(X_data[test_index])
    MSE_score = mean_squared_error(y_data[test_index],y_pred)
    print(f"R2 score is: {r2_score} and MSE is: {MSE_score}")
    model_r2_score += r2_score
    model_MSE_score += MSE_score

return (model_r2_score/5,model_MSE_score/5)

```

1.5 Implement the model for Linear Regression

```

[17]: ## Note since the y-values will be scaled we will not be using the scaled
      ↪ values for our model
from sklearn.linear_model import LinearRegression

model = LinearRegression()

X = sale_price_lot_area_df[['LotArea']]
Y = sale_price_lot_area_df[['SalePrice']]

avg_r2_score,avg_MSE = Kfold_util(model,X_data=X.to_numpy(),y_data=Y.to_numpy())

print(f"Average R2-score is :- {avg_r2_score}")
print(f"Average MSE is :- {avg_MSE}")

```

```

Fold: 1
R2 score is: 0.0811255840664582 and MSE is: 5083791990.222536
Fold: 2
R2 score is: -0.09677051342182441 and MSE is: 7206516872.48441
Fold: 3
R2 score is: 0.09034415298115273 and MSE is: 6893188266.548201
Fold: 4
R2 score is: 0.04542629407540688 and MSE is: 4835731884.893979
Fold: 5
R2 score is: 0.06551930527251171 and MSE is: 6266038369.012835
Average R2-score is :- 0.03712896459474102
Average MSE is :- 6057053476.632393

```

Print the slope and intercept

```

[18]: print(f'Saleprice = {model.coef_[0][0]}*LotArea + {model.intercept_[0]}')

Saleprice = 1.9554707455489049*LotArea + 160015.64126957097

```

1.6 Question 04 / 05

```
[19]: ## Meta-data checking for non-null values
houses_dataframe['LotFrontage'].info()
houses_dataframe['OverallQual'].info()
houses_dataframe['OverallCond'].info()
houses_dataframe['1stFlrSF'].info()
houses_dataframe['GrLivArea'].info()
```

```
<class 'pandas.core.series.Series'>
RangeIndex: 1460 entries, 0 to 1459
Series name: LotFrontage
Non-Null Count  Dtype
-----
1460 non-null   float64
dtypes: float64(1)
memory usage: 11.5 KB
<class 'pandas.core.series.Series'>
RangeIndex: 1460 entries, 0 to 1459
Series name: OverallQual
Non-Null Count  Dtype
-----
1460 non-null   int64
dtypes: int64(1)
memory usage: 11.5 KB
<class 'pandas.core.series.Series'>
RangeIndex: 1460 entries, 0 to 1459
Series name: OverallCond
Non-Null Count  Dtype
-----
1460 non-null   int64
dtypes: int64(1)
memory usage: 11.5 KB
<class 'pandas.core.series.Series'>
RangeIndex: 1460 entries, 0 to 1459
Series name: 1stFlrSF
Non-Null Count  Dtype
-----
1460 non-null   int64
dtypes: int64(1)
memory usage: 11.5 KB
<class 'pandas.core.series.Series'>
RangeIndex: 1460 entries, 0 to 1459
Series name: GrLivArea
Non-Null Count  Dtype
-----
1460 non-null   int64
dtypes: int64(1)
```

memory usage: 11.5 KB

1.6.1 Model 1:- SalePrice based on LotFrontage and LotArea

```
[20]: X_model1 = houses_dataframe[['LotArea','LotFrontage']]
      Y_model1 = houses_dataframe[['SalePrice']]

      model1 = LinearRegression()

      model1_avg_r2_score,model1_avg_MSE = Kfold_util(model1,X_data=X_model1.
      ↪to_numpy(),y_data=Y_model1.to_numpy())

      print(f"Average R2-score of Model 1 is :- {model1_avg_r2_score}")
      print(f"Average MSE of Model 1 is :- {model1_avg_MSE}")
```

```
Fold: 1
R2 score is: 0.16662248538019597 and MSE is: 4610769284.887944
Fold: 2
R2 score is: 0.06022044420881212 and MSE is: 6174981130.825078
Fold: 3
R2 score is: 0.12987082205255585 and MSE is: 6593663152.350717
Fold: 4
R2 score is: 0.052979792241372836 and MSE is: 4797466959.203123
Fold: 5
R2 score is: 0.0954056556554943 and MSE is: 6065639346.040892
Average R2-score of Model 1 is :- 0.10101983990768622
Average MSE of Model 1 is :- 5648503974.6615505
```

Print the slope and intercept

```
[21]: print(f'Saleprice = {model1.coef_[0][0]}*LotArea + {model1.
      ↪coef_[0][1]}*LotFrontage + {model1.intercept_[0]}')
```

```
Saleprice = 1.4451006676919291*LotArea + 963.9269485159833*LotFrontage +
97914.36673373333
```

1.6.2 Model 2:- SalePrice based on LotFrontage, LotArea, OverallQual, OverallCond

```
[22]: model2 = LinearRegression()

      X_model2 = ↵
      ↪houses_dataframe[['LotArea','LotFrontage','OverallQual','OverallCond']]
      Y_model2 = houses_dataframe[['SalePrice']]

      model2_avg_r2_score,model2_avg_MSE = Kfold_util(model2,X_data=X_model2.
      ↪to_numpy(),y_data=Y_model2.to_numpy())

      print(f"Average R2-score of Model 2 is :- {model2_avg_r2_score}")
      print(f"Average MSE of Model 2 is :- {model2_avg_MSE}")
```



```

Fold: 1
R2 score is: 0.7113617232553542 and MSE is: 1596928735.7894995
Fold: 2
R2 score is: 0.6575961773354837 and MSE is: 2249822451.4958005
Fold: 3
R2 score is: 0.6498657567034967 and MSE is: 2653246571.786378
Fold: 4
R2 score is: 0.6433965130847672 and MSE is: 1806501521.2943513
Fold: 5
R2 score is: 0.6227651796454086 and MSE is: 2529498866.9175787
Average R2-score of Model 2 is :- 0.6569970700049022
Average MSE of Model 2 is :- 2167199629.4567213

```

Print the slope and intercept

```

[23]: print(f'Saleprice = {model2.coef_[0][0]}*LotArea + {model2.
      ↪coef_[0][1]}*LotFrontage + {model2.coef_[0][2]}*OverallQual +{model2.
      ↪coef_[0][3]}*OverallCond +{model2.intercept_[0]}')

```

```

Saleprice = 1.189847212962575*LotArea + 418.4665742151931*LotFrontage +
42994.96807363269*OverallQual +-867.0796471136888*OverallCond
+-118705.73313785144

```

1.6.3 Model 3:- SalePrice based on LotFrontage, LotArea, OverallQual, OverallCond, 1stFlrSF and GrLivArea

```

[24]: model3 = LinearRegression()

X_model3 =
    ↪houses_dataframe[['LotArea', 'LotFrontage', 'OverallQual', 'OverallCond', '1stFlrSF', 'GrLivArea']
Y_model3 = houses_dataframe[['SalePrice']]

model3_avg_r2_score,model3_avg_MSE = Kfold_util(model3,X_data=X_model3.
    ↪to_numpy(),y_data=Y_model3.to_numpy())

print(f"Average R2-score of Model 3 is :- {model3_avg_r2_score}")
print(f"Average MSE of Model 3 is :- {model3_avg_MSE}")

```

```

Fold: 1
R2 score is: 0.7879932254919686 and MSE is: 1172955001.7146173
Fold: 2
R2 score is: 0.7569992404082071 and MSE is: 1596677748.530303
Fold: 3
R2 score is: 0.7612416434303056 and MSE is: 1809262599.0810008
Fold: 4
R2 score is: 0.7326812554452906 and MSE is: 1354197971.775218
Fold: 5
R2 score is: 0.6423236527060292 and MSE is: 2398352077.7665286
Average R2-score of Model 3 is :- 0.7362478034963603

```

Average MSE of Model 3 is :- 1666289079.7735336

Print the slope and intercept

```
[25]: print(f'Saleprice = {model3.coef_[0][0]}*LotArea + {model3.coef_[0][1]}*LotFrontage + {model3.coef_[0][2]}*OverallQual + {model3.coef_[0][3]}*OverallCond + {model3.coef_[0][4]}*1stFlrSF + {model3.coef_[0][5]}*GrLivArea + {model3.intercept_[0]}')
```

```
Saleprice = 0.6505447246251344*LotArea + 121.80300773461254*LotFrontage + 29613.59960437697*OverallQual + 1309.826685552478*OverallCond + 38.79274996548665*1stFlrSF + 45.046292101878805*GrLivArea +-135266.21756742758
```

1.7 Question 06

1.8 Categorical Value Regression

1.8.1 Model 4:- SalePrice based on LotArea and Street

```
[26]: street_set = set(houses_dataframe['Street'])
street_set
```

```
[26]: {'Grvl', 'Pave'}
```

Since we have two values for the Street feature we can consider a numpy array as a one-hot encoded input feature for estimating SalePrice

```
[27]: ## Utility One Hot Encoder

from sklearn.preprocessing import OneHotEncoder

'''
    @brief:- Function transforms the dataframe by converting the categorical_
    ↪value cols into one-hot encoded form
    @param:- X: Dataframe, col_name:- Categorical Column name
    @return:- Transformed DataFrame
'''
def one_hot_encoder(X,col_name):
    encoder = OneHotEncoder()

    new_df = pd.DataFrame(
        encoder.fit_transform(X[[col_name]]).toarray(),
        columns = encoder.get_feature_names_out()
    )

    X = X.join(new_df)
    X = X.drop(col_name,axis=1) # drop the col_name from the dataframe

    return X
```

```
[28]: houses_dataframe_pd1 = one_hot_encoder(houses_dataframe, 'Street')
      houses_dataframe_pd1.head()
```

```
[28]:
```

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	LotShape	LandContour	\
0	1	60	RL	65.0	8450	Reg	Lvl	
1	2	20	RL	80.0	9600	Reg	Lvl	
2	3	60	RL	68.0	11250	IR1	Lvl	
3	4	70	RL	60.0	9550	IR1	Lvl	
4	5	60	RL	84.0	14260	IR1	Lvl	

	Utilities	LotConfig	LandSlope	...	ScreenPorch	PoolArea	MiscVal	MoSold	\
0	AllPub	Inside	Gtl	...	0	0	0	2	
1	AllPub	FR2	Gtl	...	0	0	0	5	
2	AllPub	Inside	Gtl	...	0	0	0	9	
3	AllPub	Corner	Gtl	...	0	0	0	2	
4	AllPub	FR2	Gtl	...	0	0	0	12	

	YrSold	SaleType	SaleCondition	SalePrice	Street_Grvl	Street_Pave
0	2008	WD	Normal	208500	0.0	1.0
1	2007	WD	Normal	181500	0.0	1.0
2	2008	WD	Normal	223500	0.0	1.0
3	2006	WD	Abnorml	140000	0.0	1.0
4	2008	WD	Normal	250000	0.0	1.0

[5 rows x 78 columns]

```
[29]: ## Street_Grvl and Street_Pave are the two transformed col values
      model4 = LinearRegression()

      X_model4 = houses_dataframe_pd1[['LotArea', 'Street_Grvl', 'Street_Pave']]
      Y_model4 = houses_dataframe_pd1[['SalePrice']]

      model4_avg_r2_score, model4_avg_MSE = Kfold_util(model4, X_data=X_model4.
      ↪to_numpy(), y_data=Y_model4.to_numpy())

      print(f"Average R2-score of Model 4 is :- {model4_avg_r2_score}")
      print(f"Average MSE of Model 4 is :- {model4_avg_MSE}")
```

```
Fold: 1
R2 score is: 0.08083987276169624 and MSE is: 5085372725.1058655
Fold: 2
R2 score is: -0.06216483704535136 and MSE is: 6979134400.3638525
Fold: 3
R2 score is: 0.102774623160658 and MSE is: 6798992674.369186
Fold: 4
R2 score is: 0.056929075521841566 and MSE is: 4777460463.1480665
Fold: 5
R2 score is: 0.06650223035586045 and MSE is: 6259447493.116865
```

Average R2-score of Model 4 is :- 0.04897619295094098

Average MSE of Model 4 is :- 5980081551.220767

Print the slope and intercept

```
[30]: print(f'Saleprice = {model4.coef_[0][0]} * LotArea + {model4.coef_[0][1]} *  
      ↪ Street_Grvl + {model4.coef_[0][2]} * Street_Pave + {model4.intercept_[0]}')
```

```
Saleprice = 2.110142385499554 * LotArea + -64695.324094630705 * Street_Grvl +  
64695.32409463071 * Street_Pave + 94246.82591973609
```

1.8.2 Model 5 :- SalePrice based on LotArea, OverallCond, Street, Neighbourhood

```
[31]: houses_dataframe[['LotArea', 'OverallCond', 'Neighborhood']]
```

```
[31]:      LotArea  OverallCond Neighborhood  
0         8450             5      CollgCr  
1         9600             8      Veenker  
2        11250             5      CollgCr  
3         9550             5      Crawfor  
4        14260             5      NoRidge  
...      ...             ...      ...  
1455        7917             5      Gilbert  
1456       13175             6      NWAmes  
1457        9042             9      Crawfor  
1458        9717             6      NAmes  
1459        9937             6      Edwards
```

[1460 rows x 3 columns]

Neighbourhood is another categorical value that needs to be one-hot encoded

```
[32]: print(len(set(houses_dataframe['Neighborhood'])))
```

25

```
[33]: houses_dataframe_pd2 = one_hot_encoder(houses_dataframe, 'Street')  
      houses_dataframe_pd2 = one_hot_encoder(houses_dataframe_pd2, 'Neighborhood')  
      houses_dataframe_pd2.head()
```

```
[33]:      Id  MSSubClass MSZoning  LotFrontage  LotArea  LotShape  LandContour  \  
0      1           60        RL           65.0      8450        Reg           Lvl  
1      2           20        RL           80.0      9600        Reg           Lvl  
2      3           60        RL           68.0     11250        IR1           Lvl  
3      4           70        RL           60.0      9550        IR1           Lvl  
4      5           60        RL           84.0     14260        IR1           Lvl  
  
      Utilities  LotConfig  LandSlope  ...  Neighborhood_NoRidge  \  
0      AllPub      Inside      Gtl  ...              0.0  
1      AllPub      FR2      Gtl  ...              0.0
```

2	AllPub	Inside	Gtl ...	0.0
3	AllPub	Corner	Gtl ...	0.0
4	AllPub	FR2	Gtl ...	1.0

	Neighborhood_NridgHt	Neighborhood_OldTown	Neighborhood_SWISU	\
0	0.0	0.0	0.0	
1	0.0	0.0	0.0	
2	0.0	0.0	0.0	
3	0.0	0.0	0.0	
4	0.0	0.0	0.0	

	Neighborhood_Sawyer	Neighborhood_SawyerW	Neighborhood_Somerst	\
0	0.0	0.0	0.0	
1	0.0	0.0	0.0	
2	0.0	0.0	0.0	
3	0.0	0.0	0.0	
4	0.0	0.0	0.0	

	Neighborhood_StoneBr	Neighborhood_Timber	Neighborhood_Veenker
0	0.0	0.0	0.0
1	0.0	0.0	1.0
2	0.0	0.0	0.0
3	0.0	0.0	0.0
4	0.0	0.0	0.0

[5 rows x 102 columns]

```
[34]: neighborhood_cols = houses_dataframe_pd2.filter(regex='^Neighborhood_')
neighborhood_cols.head()
```

```
[34]:
```

	Neighborhood_Blmgtn	Neighborhood_Blueste	Neighborhood_BrDale	\
0	0.0	0.0	0.0	
1	0.0	0.0	0.0	
2	0.0	0.0	0.0	
3	0.0	0.0	0.0	
4	0.0	0.0	0.0	

	Neighborhood_BrkSide	Neighborhood_ClearCr	Neighborhood_CollgCr	\
0	0.0	0.0	1.0	
1	0.0	0.0	0.0	
2	0.0	0.0	1.0	
3	0.0	0.0	0.0	
4	0.0	0.0	0.0	

	Neighborhood_Crawfor	Neighborhood_Edwards	Neighborhood_Gilbert	\
0	0.0	0.0	0.0	
1	0.0	0.0	0.0	

2	0.0	0.0	0.0
3	1.0	0.0	0.0
4	0.0	0.0	0.0

	Neighborhood_IDOTRR	...	Neighborhood_NoRidge	Neighborhood_NridgHt	\
0	0.0	...	0.0	0.0	
1	0.0	...	0.0	0.0	
2	0.0	...	0.0	0.0	
3	0.0	...	0.0	0.0	
4	0.0	...	1.0	0.0	

	Neighborhood_OldTown	Neighborhood_SWISU	Neighborhood_Sawyer	\
0	0.0	0.0	0.0	
1	0.0	0.0	0.0	
2	0.0	0.0	0.0	
3	0.0	0.0	0.0	
4	0.0	0.0	0.0	

	Neighborhood_SawyerW	Neighborhood_Somerst	Neighborhood_StoneBr	\
0	0.0	0.0	0.0	
1	0.0	0.0	0.0	
2	0.0	0.0	0.0	
3	0.0	0.0	0.0	
4	0.0	0.0	0.0	

	Neighborhood_Timber	Neighborhood_Veenker
0	0.0	0.0
1	0.0	1.0
2	0.0	0.0
3	0.0	0.0
4	0.0	0.0

[5 rows x 25 columns]

```
[35]: houses_dataframe_pd2 = pd.
      ↪concat([houses_dataframe_pd2[['LotArea', 'OverallCond', 'Street_Grvl', 'Street_Pave']],neighbo
houses_dataframe_pd2.head()
```

	LotArea	OverallCond	Street_Grvl	Street_Pave	Neighborhood_Blmgtn	\
0	8450	5	0.0	1.0	0.0	
1	9600	8	0.0	1.0	0.0	
2	11250	5	0.0	1.0	0.0	
3	9550	5	0.0	1.0	0.0	
4	14260	5	0.0	1.0	0.0	

	Neighborhood_Blueste	Neighborhood_BrDale	Neighborhood_BrkSide	\
0	0.0	0.0	0.0	

1	0.0	0.0	0.0
2	0.0	0.0	0.0
3	0.0	0.0	0.0
4	0.0	0.0	0.0

	Neighborhood_ClearCr	Neighborhood_CollgCr	...	Neighborhood_NoRidge	\
0	0.0	1.0	...	0.0	
1	0.0	0.0	...	0.0	
2	0.0	1.0	...	0.0	
3	0.0	0.0	...	0.0	
4	0.0	0.0	...	1.0	

	Neighborhood_NridgHt	Neighborhood_OldTown	Neighborhood_SWISU	\
0	0.0	0.0	0.0	
1	0.0	0.0	0.0	
2	0.0	0.0	0.0	
3	0.0	0.0	0.0	
4	0.0	0.0	0.0	

	Neighborhood_Sawyer	Neighborhood_SawyerW	Neighborhood_Somerst	\
0	0.0	0.0	0.0	
1	0.0	0.0	0.0	
2	0.0	0.0	0.0	
3	0.0	0.0	0.0	
4	0.0	0.0	0.0	

	Neighborhood_StoneBr	Neighborhood_Timber	Neighborhood_Veenker
0	0.0	0.0	0.0
1	0.0	0.0	1.0
2	0.0	0.0	0.0
3	0.0	0.0	0.0
4	0.0	0.0	0.0

[5 rows x 29 columns]

```
[36]: model5 = LinearRegression()

X_model5 = houses_dataframe_pd2
Y_model5 = houses_dataframe_pd1[['SalePrice']]

model5_avg_r2_score,model5_avg_MSE = Kfold_util(model5,X_data=X_model5.
↪to_numpy(),y_data=Y_model5.to_numpy())

print(f"Average R2-score of Model 5 is :- {model5_avg_r2_score}")
print(f"Average MSE of Model 5 is :- {model5_avg_MSE}")
```

Fold: 1

R2 score is: 0.5841021043157792 and MSE is: 2301009097.833854

```

Fold: 2
R2 score is: 0.5261617473778633 and MSE is: 3113434688.9909825
Fold: 3
R2 score is: 0.5665378038610315 and MSE is: 3284688966.9422917
Fold: 4
R2 score is: 0.5250330812935398 and MSE is: 2406113492.1310244
Fold: 5
R2 score is: 0.545123175604145 and MSE is: 3050117194.4170914
Average R2-score of Model 5 is :- 0.5493915824904716
Average MSE of Model 5 is :- 2831072688.0630484

```

1.8.3 Model 6 :- SalePrice based on LotArea, OverallCond, Street, 1stFlrSF, Neighborhood, Year

```
[37]: houses_dataframe.columns
```

```

## Hmm `Year` is not there but we can use the continuous values of YrSold, ↵
↵YearBuilt and YearRemodAdd

```

```

[37]: Index(['Id', 'MSSubClass', 'MSZoning', 'LotFrontage', 'LotArea', 'Street',
        'LotShape', 'LandContour', 'Utilities', 'LotConfig', 'LandSlope',
        'Neighborhood', 'Condition1', 'Condition2', 'BldgType', 'HouseStyle',
        'OverallQual', 'OverallCond', 'YearBuilt', 'YearRemodAdd', 'RoofStyle',
        'RoofMatl', 'Exterior1st', 'Exterior2nd', 'MasVnrType', 'MasVnrArea',
        'ExterQual', 'ExterCond', 'Foundation', 'BsmtQual', 'BsmtCond',
        'BsmtExposure', 'BsmtFinType1', 'BsmtFinSF1', 'BsmtFinType2',
        'BsmtFinSF2', 'BsmtUnfSF', 'TotalBsmtSF', 'Heating', 'HeatingQC',
        'CentralAir', 'Electrical', '1stFlrSF', '2ndFlrSF', 'LowQualFinSF',
        'GrLivArea', 'BsmtFullBath', 'BsmtHalfBath', 'FullBath', 'HalfBath',
        'BedroomAbvGr', 'KitchenAbvGr', 'KitchenQual', 'TotRmsAbvGrd',
        'Functional', 'Fireplaces', 'FireplaceQu', 'GarageType', 'GarageYrBlt',
        'GarageFinish', 'GarageCars', 'GarageArea', 'GarageQual', 'GarageCond',
        'PavedDrive', 'WoodDeckSF', 'OpenPorchSF', 'EnclosedPorch', '3SsnPorch',
        'ScreenPorch', 'PoolArea', 'MiscVal', 'MoSold', 'YrSold', 'SaleType',
        'SaleCondition', 'SalePrice'],
        dtype='object')

```

```

[38]: houses_dataframe_pd3 = houses_dataframe_pd2
      houses_dataframe_pd3 = pd.
      ↵concat([houses_dataframe_pd3, houses_dataframe[['1stFlrSF', 'YearBuilt', 'YearRemodAdd', 'YrSold',
      houses_dataframe_pd3.head()

```

```

[38]:   LotArea  OverallCond  Street_Grvl  Street_Pave  Neighborhood_Blmngtn  \
0      8450             5           0.0           1.0                0.0
1      9600             8           0.0           1.0                0.0
2     11250             5           0.0           1.0                0.0
3      9550             5           0.0           1.0                0.0

```


4	14260	5	0.0	1.0	0.0
---	-------	---	-----	-----	-----

	Neighborhood_Blueste	Neighborhood_BrDale	Neighborhood_BrkSide	\
0	0.0	0.0	0.0	
1	0.0	0.0	0.0	
2	0.0	0.0	0.0	
3	0.0	0.0	0.0	
4	0.0	0.0	0.0	

	Neighborhood_ClearCr	Neighborhood_CollgCr	...	Neighborhood_Sawyer	\
0	0.0	1.0	...	0.0	
1	0.0	0.0	...	0.0	
2	0.0	1.0	...	0.0	
3	0.0	0.0	...	0.0	
4	0.0	0.0	...	0.0	

	Neighborhood_SawyerW	Neighborhood_Somerst	Neighborhood_StoneBr	\
0	0.0	0.0	0.0	
1	0.0	0.0	0.0	
2	0.0	0.0	0.0	
3	0.0	0.0	0.0	
4	0.0	0.0	0.0	

	Neighborhood_Timber	Neighborhood_Veenker	1stFlrSF	YearBuilt	\
0	0.0	0.0	856	2003	
1	0.0	1.0	1262	1976	
2	0.0	0.0	920	2001	
3	0.0	0.0	961	1915	
4	0.0	0.0	1145	2000	

	YearRemodAdd	YrSold
0	2003	2008
1	1976	2007
2	2002	2008
3	1970	2006
4	2000	2008

[5 rows x 33 columns]

```
[39]: model6 = LinearRegression()

X_model6 = houses_dataframe_pd3
Y_model6 = houses_dataframe_pd1[['SalePrice']]

model6_avg_r2_score,model6_avg_MSE = Kfold_util(model6,X_data=X_model6.
↳to_numpy(),y_data=Y_model6.to_numpy())
```

```
print(f"Average R2-score of Model 6 is :- {model6_avg_r2_score}")
print(f"Average MSE of Model 6 is :- {model6_avg_MSE}")
```

Fold: 1

R2 score is: 0.7095584627129405 and MSE is: 1606905508.8314323

Fold: 2

R2 score is: 0.7138242036094097 and MSE is: 1880366658.2457325

Fold: 3

R2 score is: 0.6985593789138491 and MSE is: 2284256138.342621

Fold: 4

R2 score is: 0.6829225245459274 and MSE is: 1606268482.484338

Fold: 5

R2 score is: 0.6451907294515968 and MSE is: 2379127268.740516

Average R2-score of Model 6 is :- 0.6900110598467447

Average MSE of Model 6 is :- 1951384811.328928

1.9 Question 07

1.9.1 Comparison of the feature “LotArea” weights/coefficients for all the six trained models

```
[40]: x_axis = range(1,7)
lot_area_weights = [model1.coef_[0][0],model2.coef_[0][0],model3.
    ↪coef_[0][0],model4.coef_[0][0],model5.coef_[0][0],model6.coef_[0][0]]

plt.plot(x_axis,lot_area_weights,".-")
plt.xlabel("Model Number")
plt.ylabel("LotArea Weighths")
plt.title("Variation of LotArea Weights with the models")
plt.show()
```



1.10 Question 08

1.10.1 Bonus - Using Polynomial Regression for predicting SalePrice based on LotArea and applying it on the training and test dataset

```
[41]: sale_price = houses_dataframe[['SalePrice']]
lot_area = houses_dataframe[['LotArea']]
plt.plot(lot_area,sale_price,'.',label="original data")

## define a domain
x_axis = range(0,200000,200)

## Linear Model (degree 1)

linear_model_degree_1 = LinearRegression()
X_data_1 = lot_area[['LotArea']]
Y_data_1 = sale_price[['SalePrice']]
Kfold_util(linear_model_degree_1,X_data_1.to_numpy(),Y_data_1.to_numpy())
linear_model_equation = [linear_model_degree_1.coef_[0][0]*x +
↪ linear_model_degree_1.intercept_[0] for x in x_axis]
```

```

plt.plot(x_axis,linear_model_equation,'red',label='linear')

curve_model_degree_2 = LinearRegression()
X_data_2 = X_data_1
X_data_2[['LotArea_squared']]=X_data_1[['LotArea']]**2
Y_data_2 = Y_data_1
Kfold_util(curve_model_degree_2,X_data_2.to_numpy(),Y_data_2.to_numpy())
curve_model_2_equation = [curve_model_degree_2.coef_[0][0]*x +
    ↳curve_model_degree_2.coef_[0][1]*(x**2) + curve_model_degree_2.intercept_[0]
    ↳for x in x_axis]
plt.plot(x_axis,curve_model_2_equation,'green',label='degree-2-polynomial')

curve_model_degree_3 = LinearRegression()
X_data_3 = X_data_2
X_data_3[['LotArea_cubed']] = X_data_1[['LotArea']]**3
Y_data_3 = Y_data_2
Kfold_util(curve_model_degree_3,X_data_3.to_numpy(),Y_data_3.to_numpy())
curve_model_3_equation = [curve_model_degree_3.coef_[0][0]*x +
    ↳curve_model_degree_3.coef_[0][1]*(x**2) + curve_model_degree_3.
    ↳coef_[0][2]*(x**3) + curve_model_degree_3.intercept_[0] for x in x_axis]
plt.plot(x_axis,curve_model_3_equation,'purple',label='degree-3-polynomial')

plt.title("Saleprice based on various curvilinear regression models")
plt.legend()
plt.xlabel('LotArea')
plt.ylabel('SalePrice')
plt.show()

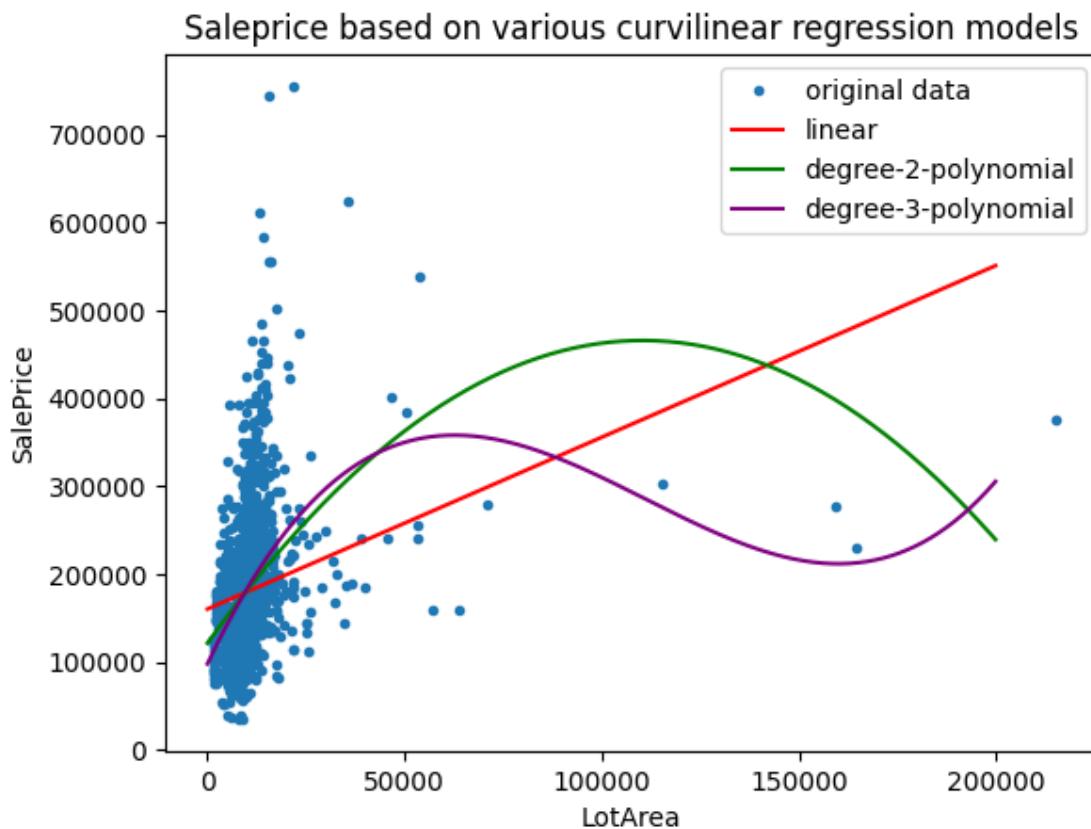
```

```

Fold: 1
R2 score is: 0.0811255840664582 and MSE is: 5083791990.222536
Fold: 2
R2 score is: -0.09677051342182441 and MSE is: 7206516872.48441
Fold: 3
R2 score is: 0.09034415298115273 and MSE is: 6893188266.548201
Fold: 4
R2 score is: 0.04542629407540688 and MSE is: 4835731884.893979
Fold: 5
R2 score is: 0.06551930527251171 and MSE is: 6266038369.012835
Fold: 1
R2 score is: 0.1880996406892882 and MSE is: 4491944135.074409
Fold: 2
R2 score is: -0.13326083373331588 and MSE is: 7446282717.562326
Fold: 3
R2 score is: 0.17210416157128516 and MSE is: 6273627436.226067
Fold: 4
R2 score is: 0.08756692413546696 and MSE is: 4622253567.644948
Fold: 5

```

R2 score is: 0.10154612807937147 and MSE is: 6024465209.401183
Fold: 1
R2 score is: 0.21551510546124142 and MSE is: 4340264517.273444
Fold: 2
R2 score is: -0.15932525137145825 and MSE is: 7617543398.973908
Fold: 3
R2 score is: 0.18848183298735877 and MSE is: 6149520750.374883
Fold: 4
R2 score is: 0.08050638806129573 and MSE is: 4658021218.90792
Fold: 5
R2 score is: 0.13134466316099958 and MSE is: 5824655020.47448



We can clearly see that the higher order polynomials(here cubic) capture the training data more precisely, but it should not be a good indicator for model selection as it may hint at chances of overfitting.

```
[42]: TEST_DATA_PATH = "../ML_DRIVE/Assignment01/test.csv"

test_df = pd.read_csv(TEST_DATA_PATH)
X_data = test_df[['LotArea']]
```

```

## Linear Model Prediction
test_df[['Linear_SalePrice']] = linear_model_degree_1.predict(X_data.to_numpy())

X_data_2 = X_data
X_data_2[['LotArea_squared']] = X_data[['LotArea']]**2

## Quadratic Model Prediction
test_df[['Quadratic_SalePrice']] = curve_model_degree_2.predict(X_data_2.
    ↪to_numpy())

X_data_3 = X_data_2
X_data_3[['LotArea_cubed']] = X_data[['LotArea']]**3

## Cubic Model Prediction
test_df[['Cubic_SalePrice']] = curve_model_degree_3.predict(X_data_3.to_numpy())
test_df[['Linear_SalePrice', 'Quadratic_SalePrice', 'Cubic_SalePrice']]

```

```

[42]:
      Linear_SalePrice  Quadratic_SalePrice  Cubic_SalePrice
0      182742.122274      189993.344744      194862.710030
1      187914.342396      204567.098660      213283.237485
2      187059.801681      202186.557959      210330.004739
3      179527.328369      180735.550583      182745.955495
4      169802.772351      151800.366145      142881.730423
...
1454      163801.432633      133245.300499      115795.352440
1455      163719.302862      132987.673554      115411.178384
1456      199125.056181      234796.704797      248815.620205
1457      180432.711324      183358.293784      186210.825113
1458      178840.958137      178739.167815      180091.665582

```

```

[1459 rows x 3 columns]

```