

Assignment 03

Name:- Sk Fardeen Hossain

Roll:- 2021CSB023

Department:- Computer Science and Technology

G-Suite ID:- 2021csb023.sk@students.iiests.ac.in

Question 1:

Download the [Forest Cover Type Dataset](#) and preprocess the dummy variables to create training, test and development set. Reduce the train data size if the system is unable to process the whole dataset.

```
In [6]: import pandas as pd
forest_cover_df = pd.read_csv('../ML_DRIVE/Assignment03/covtype.csv')
forest_cover_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 581012 entries, 0 to 581011  
Data columns (total 55 columns):
```

#	Column	Non-Null Count	Dtype
0	Elevation	581012 non-null	int64
1	Aspect	581012 non-null	int64
2	Slope	581012 non-null	int64
3	Horizontal_Distance_To_Hydrology	581012 non-null	int64
4	Vertical_Distance_To_Hydrology	581012 non-null	int64
5	Horizontal_Distance_To_Roadways	581012 non-null	int64
6	Hillshade_9am	581012 non-null	int64
7	Hillshade_Noon	581012 non-null	int64
8	Hillshade_3pm	581012 non-null	int64
9	Horizontal_Distance_To_Fire_Points	581012 non-null	int64
10	Wilderness_Area1	581012 non-null	int64
11	Wilderness_Area2	581012 non-null	int64
12	Wilderness_Area3	581012 non-null	int64
13	Wilderness_Area4	581012 non-null	int64
14	Soil_Type1	581012 non-null	int64
15	Soil_Type2	581012 non-null	int64
16	Soil_Type3	581012 non-null	int64
17	Soil_Type4	581012 non-null	int64
18	Soil_Type5	581012 non-null	int64
19	Soil_Type6	581012 non-null	int64
20	Soil_Type7	581012 non-null	int64
21	Soil_Type8	581012 non-null	int64
22	Soil_Type9	581012 non-null	int64
23	Soil_Type10	581012 non-null	int64
24	Soil_Type11	581012 non-null	int64
25	Soil_Type12	581012 non-null	int64
26	Soil_Type13	581012 non-null	int64
27	Soil_Type14	581012 non-null	int64
28	Soil_Type15	581012 non-null	int64
29	Soil_Type16	581012 non-null	int64
30	Soil_Type17	581012 non-null	int64
31	Soil_Type18	581012 non-null	int64
32	Soil_Type19	581012 non-null	int64
33	Soil_Type20	581012 non-null	int64
34	Soil_Type21	581012 non-null	int64
35	Soil_Type22	581012 non-null	int64
36	Soil_Type23	581012 non-null	int64
37	Soil_Type24	581012 non-null	int64
38	Soil_Type25	581012 non-null	int64
39	Soil_Type26	581012 non-null	int64
40	Soil_Type27	581012 non-null	int64

```

41 Soil_Type28          581012 non-null int64
42 Soil_Type29          581012 non-null int64
43 Soil_Type30          581012 non-null int64
44 Soil_Type31          581012 non-null int64
45 Soil_Type32          581012 non-null int64
46 Soil_Type33          581012 non-null int64
47 Soil_Type34          581012 non-null int64
48 Soil_Type35          581012 non-null int64
49 Soil_Type36          581012 non-null int64
50 Soil_Type37          581012 non-null int64
51 Soil_Type38          581012 non-null int64
52 Soil_Type39          581012 non-null int64
53 Soil_Type40          581012 non-null int64
54 Cover_Type          581012 non-null int64
dtypes: int64(55)
memory usage: 243.8 MB

```

As we can see above, all the columns are important since all have non-null values

```
In [7]: forest_cover_df.head()
```

```

Out[7]:   Elevation  Aspect  Slope  Horizontal_Distance_To_Hydrology  Vertical_Distance_To_Hydrology  Horizontal_Distance_To_Roadways  Hillshade_9am  Hillst
0      2596      51      3                258                      0                510                221
1      2590      56      2                212                     -6                390                220
2      2804     139      9                268                     65               3180                234
3      2785     155     18                242                    118               3090                238
4      2595      45      2                153                     -1                391                220

```

5 rows × 55 columns

Standardisation of non-binary data columns

```

In [8]: from sklearn.preprocessing import StandardScaler

scaler=StandardScaler()

...
    @param:- [df:pd.DataFrame, col_name: string]
    @brief:- Function returns a standardised dataframe on the column `col_name`
    @return:- Standardised Dataframe
...

```

```
def standardise(df:"pd.DataFrame",col_name:"str")->"pd.DataFrame":
    df[[col_name]]=pd.DataFrame(
        data = scaler.fit_transform(df[[col_name]]),
        columns=[col_name],
        index=df.index
    )

    return df
```

```
In [9]: '''
Elevation                581012 non-null  int64
1  Aspect                581012 non-null  int64
2  Slope                 581012 non-null  int64
3  Horizontal_Distance_To_Hydrology  581012 non-null  int64
4  Vertical_Distance_To_Hydrology    581012 non-null  int64
5  Horizontal_Distance_To_Roadways   581012 non-null  int64
6  Hillshade_9am             581012 non-null  int64
7  Hillshade_Noon            581012 non-null  int64
8  Hillshade_3pm             581012 non-null  int64
9  Horizontal_Distance_To_Fire_Points 581012 non-null  int64
'''

columns_to_standardise = ['Elevation','Aspect','Slope','Horizontal_Distance_To_Hydrology','Vertical_Distance_To_Hydrology',
                          'Horizontal_Distance_To_Roadways','Hillshade_9am','Hillshade_Noon','Hillshade_3pm','Horizontal_Distan

for col_name in columns_to_standardise:
    forest_cover_df=standardise(forest_cover_df,col_name)

forest_cover_df.head()
```

```
Out[9]:
```

	Elevation	Aspect	Slope	Horizontal_Distance_To_Hydrology	Vertical_Distance_To_Hydrology	Horizontal_Distance_To_Roadways	Hillshade_9am
0	-1.297805	-0.935157	-1.482820	-0.053767	-0.796273	-1.180146	0.330743
1	-1.319235	-0.890480	-1.616363	-0.270188	-0.899197	-1.257106	0.293388
2	-0.554907	-0.148836	-0.681563	-0.006719	0.318742	0.532212	0.816364
3	-0.622768	-0.005869	0.520322	-0.129044	1.227908	0.474492	0.965786
4	-1.301377	-0.988770	-1.616363	-0.547771	-0.813427	-1.256464	0.293388

5 rows × 55 columns

```
In [10]: ## Number of target classes
```

```
forest_cover_df[['Cover_Type']].value_counts()
```

```
Out[10]: Cover_Type
2         283301
1         211840
3         35754
7         20510
6         17367
5          9493
4          2747
Name: count, dtype: int64
```

Since there high number of samples especially from classes 1 , 2 , 3 we need to perform random sampling across population points

```
In [11]: forest_cover_df_rough = forest_cover_df.sample(frac=0.1, random_state=5)
forest_cover_df_rough[['Cover_Type']].value_counts()
```

```
Out[11]: Cover_Type
2         28366
1         21062
3          3611
7         2068
6         1724
5          985
4          285
Name: count, dtype: int64
```

```
In [12]: y=forest_cover_df_rough[['Cover_Type']]
X=forest_cover_df_rough.drop(columns=['Cover_Type'])
X
```

Out[12]:	Elevation	Aspect	Slope	Horizontal_Distance_To_Hydrology	Vertical_Distance_To_Hydrology	Horizontal_Distance_To_Roadways	Hillshade
456056	1.223763	1.566773	0.653865	1.842264	3.029086	0.569409	-1.49
456659	0.280854	1.816965	-0.280934	-0.067882	-0.213034	-0.778672	-0.37
210876	1.095185	-0.827931	-0.414477	-0.703029	-0.264496	2.261886	0.66
490911	-0.097739	-0.827931	0.653865	0.962468	1.159292	-1.083305	0.77
267218	-2.994328	1.575708	1.321579	-1.126460	-0.659040	-0.806890	-2.06
...
509165	0.777310	-1.194285	0.386780	0.938944	0.816211	-0.506747	-0.19
261011	0.266567	1.057451	0.119694	-0.373693	0.130048	-1.256464	-1.20
278278	-2.165710	1.745482	1.989293	-0.561885	0.181510	-1.316750	-2.06
494618	0.923747	1.736546	-0.013849	-0.105520	-0.110110	-1.171167	-0.67
24166	0.373716	-0.184578	-1.082191	-1.126460	-0.727656	2.867304	0.66

58101 rows × 54 columns

```
In [13]: # 80% -> train | 10% -> validation | 10% -> testing
from sklearn.model_selection import train_test_split

X_train,X_rest,y_train,y_rest = train_test_split(X,y,random_state=5,test_size=0.2)
X_val,X_test,y_val,y_test = train_test_split(X_rest,y_rest,random_state=5,test_size=0.5)
```

Question 02:

Consider only two features and three classes and train Logistic Regression 3- class classifier (any three class) to show the training and test area in a 2D plane, using matplotlib.

Here we are choosing the features `Elevation` and `Slope` which are to be mapped to classes `1`, `2`, `3`

```
In [14]: X_train_subset = X_train[['Elevation','Slope']]
X_train_subset
```

Out [14]:

	Elevation	Slope
257482	0.366573	0.520322
264358	-0.597766	1.588665
1767	-0.358467	0.787408
58906	-0.447758	-0.147392
390314	0.902317	-1.215734
...
53980	0.173705	0.119694
162900	0.988036	0.520322
41957	-1.065649	0.787408
15266	-0.719202	-0.280934
474993	-0.229889	1.989293

46480 rows × 2 columns

```
In [15]: train_subset = X_train_subset.join(y_train)
train_subset = train_subset[train_subset['Cover_Type'].isin([1,2,3])]
train_subset
```

```
Out[15]:
```

	Elevation	Slope	Cover_Type
257482	0.366573	0.520322	1
264358	-0.597766	1.588665	2
1767	-0.358467	0.787408	2
58906	-0.447758	-0.147392	2
390314	0.902317	-1.215734	1
...
53980	0.173705	0.119694	1
162900	0.988036	0.520322	2
41957	-1.065649	0.787408	2
15266	-0.719202	-0.280934	2
474993	-0.229889	1.989293	2

42451 rows × 3 columns

```
In [16]: from sklearn.linear_model import LogisticRegression

LR_model = LogisticRegression(
    solver='saga',
    max_iter=10000,
)
LR_model.fit(train_subset.iloc[:,0:2],train_subset.iloc[:,2])
```


```
Out[16]:
```

▼
LogisticRegression
i ?

```
LogisticRegression(max_iter=10000, solver='saga')
```

```
In [17]: # Training Data
from sklearn.inspection import DecisionBoundaryDisplay

display_fig = DecisionBoundaryDisplay.from_estimator(
    LR_model,
    train_subset.iloc[:,0:2],
    xlabel = 'Elevation',
    ylabel = 'Slope',
    grid_resolution = 6000,
```

```
In [18]: X_test_subset = X_test[['Elevation', 'Slope']]

test_subset = X_test_subset.join(y_test)
test_subset = test_subset[test_subset['Cover_Type'].isin([1,2,3])]
test_subset
```


Out[18]:

	Elevation	Slope	Cover_Type
402158	0.005839	-1.349277	2
271152	-2.097849	-1.215734	3
318946	0.305855	0.119694	2
554660	-1.008503	0.787408	2
450417	1.366628	0.119694	1
...
521316	0.748737	0.920951	1
540121	0.327285	-0.013849	2
153065	-0.176314	0.520322	2
192414	0.927318	-1.349277	1
445729	0.852314	0.386780	2

5311 rows × 3 columns

```
In [19]: # Test Data

display_fig = DecisionBoundaryDisplay.from_estimator(
    LR_model,
    test_subset.iloc[:,0:2],
    xlabel = 'Elevation',
    ylabel = 'Slope',
    grid_resolution = 6000
)
```



Question 03:

Analyze and control the overfitting by varying the inverse of regularization strength parameter (0.1, 0.25, 0.5, 0.75, 0.9) and plot the accuracy graph for the test set.

```
In [20]: from sklearn.metrics import accuracy_score

...
@params: - X_train,y_train,X_val,y_val,X_test,y_test,solver,inv_reg_str,penalty
@return:- list([accuracy_before_finetune,accuracy_after_finetune])
@brief:- Trains and tests a LR model on training, testing and validation dataset
...

def LR_model_util(
    X_train: "pd.DataFrame",
    y_train: "pd.DataFrame",
    X_val: "pd.DataFrame",
    y_val: "pd.DataFrame",
```

```

X_test: "pd.DataFrame",
y_test: "pd.DataFrame",
solver: "str",
inv_reg_str: "double",
penalty: "str"
)->"list":
    model = LogisticRegression(
        solver=solver,
        C=inv_reg_str,
        max_iter=10000,
        penalty=penalty
    )
    model.fit(X_train,y_train)
    y_pred_before_finetune = model.predict(X_val)
    accuracy_before_finetuning = accuracy_score(y_val,y_pred_before_finetune)
    model.fit(X_val,y_val)
    y_pred_after_finetune = model.predict(X_test)
    accuracy_after_finetuning= accuracy_score(y_test,y_pred_after_finetune)

    return [accuracy_before_finetuning,accuracy_after_finetuning]

```

In [21]: `X_val_subset = X_val[['Elevation','Slope']]`

```

val_subset = X_val_subset.join(y_val)
val_subset = val_subset[val_subset['Cover_Type'].isin([1,2,3])]
val_subset

```

Out[21]:

	Elevation	Slope	Cover_Type
--	-----------	-------	------------

286194	-0.072737	-0.013849	2
519090	-0.069166	-0.147392	2
429189	-0.069166	0.386780	1
175544	0.898745	-0.414477	2
320041	-1.554962	-0.414477	2
...
120920	-0.154885	-0.681563	2
451493	1.013037	-0.548020	1
391006	0.716592	0.386780	1
508515	0.948748	-0.147392	1
108165	-1.165655	-0.147392	2

5277 rows × 3 columns

In [22]: `import matplotlib.pyplot as plt`

```
inverse_regularization_strength_list = [0.1,0.25,0.5,0.75,0.9]
```

```
train_subset_X = train_subset.iloc[:,0:2]
```

```
train_subset_y = train_subset.iloc[:,2]
```

```
val_subset_X = val_subset.iloc[:,0:2]
```

```
val_subset_y = val_subset.iloc[:,2]
```

```
test_subset_X = test_subset.iloc[:,0:2]
```

```
test_subset_y = test_subset.iloc[:,2]
```

```
accuracy_before_finetune_list=[]
```

```
accuracy_after_finetune_list=[]
```

```
for C in inverse_regularization_strength_list:
```

```
    [acc_b_f,acc_a_f] = LR_model_util(train_subset_X,train_subset_y,val_subset_X,val_subset_y,  
                                     test_subset_X,test_subset_y,"saga",C,"l1")
```

```
    accuracy_before_finetune_list.append(acc_b_f)
```

```
    accuracy_after_finetune_list.append(acc_a_f)
```

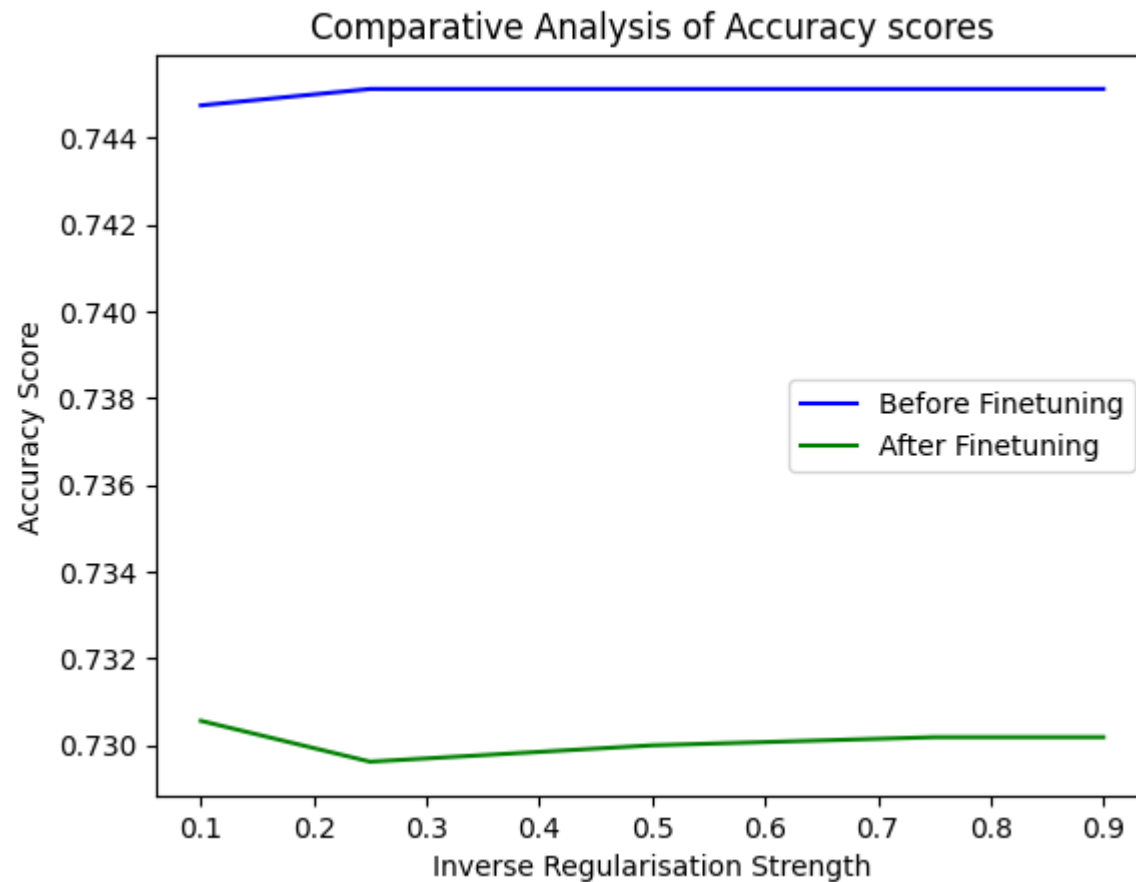
```
plt.plot(inverse_regularization_strength_list,accuracy_before_finetune_list,"blue",label="Before Finetuning")
```

```
plt.plot(inverse_regularization_strength_list,accuracy_after_finetune_list,"green",label="After Finetuning")
```

```
plt.title("Comparative Analysis of Accuracy scores")
plt.xlabel('Inverse Regularisation Strength')
plt.ylabel('Accuracy Score')

plt.legend()

plt.show()
```



Question 04:

Apply multiclass classification in Support Vector Machine (SVM) using Forest Cover Type dataset.

```
In [39]: # Using a Support vector machine to perform multi-class classification
from sklearn.svm import LinearSVC
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
```

```

...
@params:- X_train,y_train,X_val,y_val,X_test,y_test,penalty,inverse regularisation strength,loss
@return:- list[penalty,inv_reg_str,loss,accuracy_before_finetune,accuracy_after_finetune,precision_before_finetune,precision_a
          recall_before_finetune,recall_after_finetune,f1score_before_finetune,f1score_after_finetune]
@brief:- Function creates Linear SVC(Support Vector Classifier) based on tuned parameters and returns its accuracy metrics
...
def SVC_util(
    X_train: "pd.DataFrame",
    y_train: "pd.DataFrame",
    X_val: "pd.DataFrame",
    y_val: "pd.DataFrame",
    X_test: "pd.DataFrame",
    y_test: "pd.DataFrame",
    penalty: 'str',
    inv_reg_str: "double",
    loss: 'str'
)->"list":
    model = LinearSVC(
        max_iter=10000,
        C = inv_reg_str,
        loss=loss
    ).fit(X_train,y_train)
    y_pred = model.predict(X_val)
    accuracy_before_finetune = accuracy_score(y_val,y_pred)
    precision_before_finetune = precision_score(y_val,y_pred,average='macro').item()
    recall_before_finetune = recall_score(y_val,y_pred,average='macro').item()
    f1score_before_finetune = f1_score(y_val,y_pred,average='macro').item()
    model.fit(X_val,y_val)
    y_pred = model.predict(X_test)
    accuracy_after_finetune = accuracy_score(y_test,y_pred)
    precision_after_finetune = precision_score(y_test,y_pred,average='macro').item()
    recall_after_finetune = recall_score(y_test,y_pred,average='macro').item()
    f1score_after_finetune = f1_score(y_test,y_pred,average='macro').item()
    return [penalty,inv_reg_str,loss,accuracy_before_finetune,accuracy_after_finetune,precision_before_finetune,precision_a
            recall_before_finetune,recall_after_finetune,f1score_before_finetune,f1score_after_finetune]

```

In [41]: kernels = ['linear','rbf','poly','sigmoid','precomputed']

```

linear_1 = SVC_util(X_train,y_train,X_val,y_val,X_test,y_test,'l1',0.1,'hinge')
linear_2 = SVC_util(X_train,y_train,X_val,y_val,X_test,y_test,'l1',0.1,'squared_hinge')
linear_3 = SVC_util(X_train,y_train,X_val,y_val,X_test,y_test,'l2',0.1,'hinge')
linear_4 = SVC_util(X_train,y_train,X_val,y_val,X_test,y_test,'l2',0.1,'squared_hinge')

linear_5 = SVC_util(X_train,y_train,X_val,y_val,X_test,y_test,'l1',0.25,'hinge')
linear_6 = SVC_util(X_train,y_train,X_val,y_val,X_test,y_test,'l1',0.25,'squared_hinge')

```

```

linear_7 = SVC_util(X_train,y_train,X_val,y_val,X_test,y_test,'l2',0.25,'hinge')
linear_8 = SVC_util(X_train,y_train,X_val,y_val,X_test,y_test,'l2',0.25,'squared_hinge')

linear_9 = SVC_util(X_train,y_train,X_val,y_val,X_test,y_test,'l1',0.5,'hinge')
linear_10 = SVC_util(X_train,y_train,X_val,y_val,X_test,y_test,'l1',0.5,'squared_hinge')
linear_11 = SVC_util(X_train,y_train,X_val,y_val,X_test,y_test,'l2',0.5,'hinge')
linear_12 = SVC_util(X_train,y_train,X_val,y_val,X_test,y_test,'l2',0.5,'squared_hinge')

linear_13 = SVC_util(X_train,y_train,X_val,y_val,X_test,y_test,'l1',0.75,'hinge')
linear_14 = SVC_util(X_train,y_train,X_val,y_val,X_test,y_test,'l1',0.75,'squared_hinge')
linear_15 = SVC_util(X_train,y_train,X_val,y_val,X_test,y_test,'l2',0.75,'hinge')
linear_16 = SVC_util(X_train,y_train,X_val,y_val,X_test,y_test,'l2',0.75,'squared_hinge')

linear_17 = SVC_util(X_train,y_train,X_val,y_val,X_test,y_test,'l1',0.9,'hinge')
linear_18 = SVC_util(X_train,y_train,X_val,y_val,X_test,y_test,'l1',0.9,'squared_hinge')
linear_19 = SVC_util(X_train,y_train,X_val,y_val,X_test,y_test,'l2',0.9,'hinge')
linear_20 = SVC_util(X_train,y_train,X_val,y_val,X_test,y_test,'l2',0.9,'squared_hinge')

SVC_report = pd.DataFrame(
    columns=['penalty','inv_reg_strength','loss','accuracy_before_finetune','accuracy_after_finetune','precision_before_finetune',
            'recall_before_finetune','recall_after_finetune','f1score_before_finetune','f1score_after_finetune'],
    data = [linear_1,linear_2,linear_3,linear_4,linear_5,linear_6,linear_7,linear_8,linear_9,linear_10,
            linear_11,linear_12,linear_13,linear_14,linear_15,linear_16,linear_17,linear_18,linear_19,
            linear_20]
)

SVC_report

```


Out[41]:	penalty	inv_reg_strength	loss	accuracy_before_finetune	accuracy_after_finetune	precision_before_finetune	precision_after_finetune	recall
	0	l1	0.10	hinge	0.713597	0.700396	0.521611	0.365776
	1	l1	0.10	squared_hinge	0.713769	0.705903	0.523590	0.425119
	2	l2	0.10	hinge	0.713597	0.700568	0.551474	0.413395
	3	l2	0.10	squared_hinge	0.713769	0.705903	0.523590	0.425119
	4	l1	0.25	hinge	0.712565	0.706763	0.461092	0.456528
	5	l1	0.25	squared_hinge	0.713941	0.708828	0.679575	0.429829
	6	l2	0.25	hinge	0.712565	0.706419	0.488530	0.406084
	7	l2	0.25	squared_hinge	0.713941	0.708828	0.679575	0.429829
	8	l1	0.50	hinge	0.713253	0.708828	0.530967	0.442193
	9	l1	0.50	squared_hinge	0.714114	0.709000	0.679629	0.576261
	10	l2	0.50	hinge	0.712392	0.707796	0.482432	0.407406
	11	l2	0.50	squared_hinge	0.714114	0.709000	0.679629	0.576261
	12	l1	0.75	hinge	0.713597	0.710377	0.548501	0.452501
	13	l1	0.75	squared_hinge	0.714286	0.707968	0.682814	0.574935
	14	l2	0.75	hinge	0.714114	0.709861	0.548739	0.429317
	15	l2	0.75	squared_hinge	0.714286	0.707968	0.682814	0.574935
	16	l1	0.90	hinge	0.713425	0.709861	0.528923	0.531003
	17	l1	0.90	squared_hinge	0.714114	0.708312	0.682751	0.574776
	18	l2	0.90	hinge	0.713425	0.708828	0.492822	0.480583
	19	l2	0.90	squared_hinge	0.714114	0.708312	0.682751	0.574776


From the above table based on f1-score, we have seen that the LinearSVC model with l2 penalty, 0.9 inverse_reg_strength and squared_hinge_loss performs the best

Question 05:

Plot and analyze the confusion matrix for the above applied SVM method.

```
In [43]: from sklearn.metrics import ConfusionMatrixDisplay, confusion_matrix
```

```
model = LinearSVC(  
    penalty='l2',  
    C=0.9,  
    loss='squared_hinge',  
    max_iter=10000  
) .fit(X_train,y_train)  
  
y_pred = model.predict(X_test)  
matrix = confusion_matrix(y_test,y_pred,labels=model.classes_)  
disp = ConfusionMatrixDisplay(confusion_matrix=matrix, display_labels= model.classes_)  
  
disp.plot()  
  
plt.show()
```



From the above confusion matrix we can conclude that the major misclassification occur among groups:-

1. (1,2)

2. (5,2)
3. (6,3)
4. (7,1)

Label 2 is most accurately classified from the given model among all other labels

Question 06:

Download [Titanic Dataset](#) and do the initial preprocessing and train a Decision Tree classifier and vary the maximum depth of the tree to train at least 5 classifiers to analyze the effectiveness.

```
In [23]: TRAIN_DATA_PATH = "../ML_DRIVE/Assignment03/train.csv"
TEST_DATA_PATH = "../ML_DRIVE/Assignment03/test.csv"

train_df = pd.read_csv(TRAIN_DATA_PATH)

test_df = pd.read_csv(TEST_DATA_PATH) # only for prediction since final label 'Survived' is not given

train_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   PassengerId  891 non-null    int64
 1   Survived     891 non-null    int64
 2   Pclass       891 non-null    int64
 3   Name         891 non-null    object
 4   Sex          891 non-null    object
 5   Age          714 non-null    float64
 6   SibSp        891 non-null    int64
 7   Parch        891 non-null    int64
 8   Ticket       891 non-null    object
 9   Fare         891 non-null    float64
10   Cabin        204 non-null    object
11   Embarked     889 non-null    object
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB
```

```
In [24]: train_df.columns
```

```
Out[24]: Index(['PassengerId', 'Survived', 'Pclass', 'Name', 'Sex', 'Age', 'SibSp',  
              'Parch', 'Ticket', 'Fare', 'Cabin', 'Embarked'],  
              dtype='object')
```

```
In [25]: # Dropping passengerId, name, cabin, Ticket  
train_df = train_df.drop("PassengerId",axis=1)  
train_df = train_df.drop("Name",axis=1)  
train_df = train_df.drop("Cabin",axis=1)  
train_df = train_df.drop("Ticket",axis=1)
```

```
In [26]: # Dropping null age values  
train_df = train_df[train_df['Age'].notna()]  
train_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
Index: 714 entries, 0 to 890  
Data columns (total 8 columns):  
#   Column      Non-Null Count  Dtype  
---  ---  
0   Survived    714 non-null    int64  
1   Pclass      714 non-null    int64  
2   Sex         714 non-null    object  
3   Age         714 non-null    float64  
4   SibSp       714 non-null    int64  
5   Parch       714 non-null    int64  
6   Fare        714 non-null    float64  
7   Embarked    712 non-null    object  
dtypes: float64(2), int64(4), object(2)  
memory usage: 50.2+ KB
```

```
In [27]: train_df.head()
```

```
Out[27]:
```

	Survived	Pclass	Sex	Age	SibSp	Parch	Fare	Embarked
0	0	3	male	22.0	1	0	7.2500	S
1	1	1	female	38.0	1	0	71.2833	C
2	1	3	female	26.0	0	0	7.9250	S
3	1	1	female	35.0	1	0	53.1000	S
4	0	3	male	35.0	0	0	8.0500	S

Preprocessing the dataset is done !!!!

```
In [28]: ## One-hot encode `Embarked`, `Pclass` and `Sex`
from sklearn.preprocessing import OneHotEncoder

encoder = OneHotEncoder()

encoded_df = pd.DataFrame(
    encoder.fit_transform(train_df[['Embarked']]).toarray(),
    index = train_df.index,
    columns = encoder.get_feature_names_out()
)

train_df = train_df.join(encoded_df)
train_df = train_df.drop('Embarked',axis=1)

encoded_df = pd.DataFrame(
    encoder.fit_transform(train_df[['Pclass']]).toarray(),
    index = train_df.index,
    columns = encoder.get_feature_names_out()
)

train_df = train_df.join(encoded_df)
train_df = train_df.drop('Pclass',axis=1)

encoded_df = pd.DataFrame(
    encoder.fit_transform(train_df[['Sex']]).toarray(),
    index = train_df.index,
    columns = encoder.get_feature_names_out()
)

train_df = train_df.join(encoded_df)
train_df = train_df.drop('Sex',axis=1)
```

```
In [29]: X = train_df.drop(columns=['Survived'])

y = train_df[['Survived']]
```

```
In [30]: X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.4)

from sklearn.tree import DecisionTreeClassifier

'''
@params:- X_train,y_train,X_test,y_test,max_depth
@return:- Accuracy score of the model
@brief:- Trains a decision tree classifier and then returns its accuracy score
'''
```

```

def DT_util(
    X_train: "pd.DataFrame",
    y_train: "pd.DataFrame",
    X_test: "pd.DataFrame",
    y_test: "pd.DataFrame",
    max_depth: "integer"
) -> "float":
    model = DecisionTreeClassifier(max_depth=max_depth).fit(X_train,y_train)
    y_pred = model.predict(X_test)
    return accuracy_score(y_test,y_pred)

```

In [31]: max_depths = range(1,50)

```

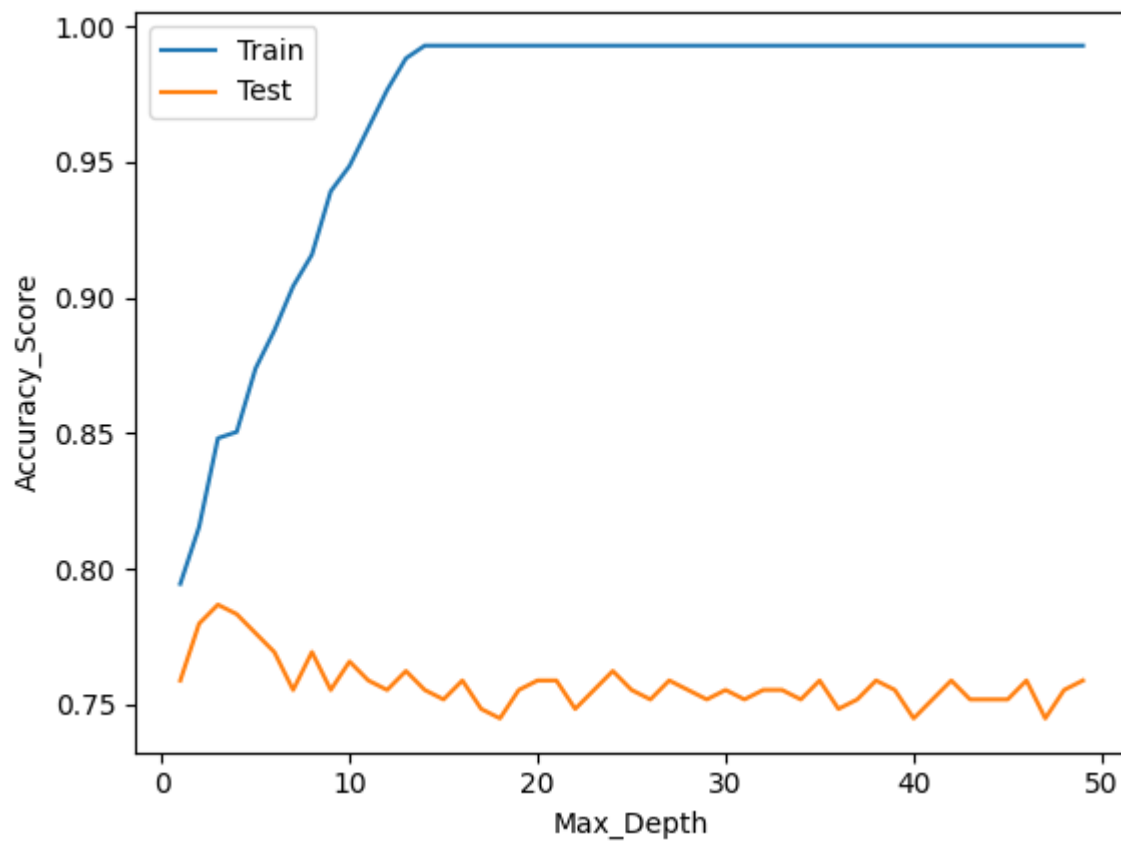
training_accuracies = [DT_util(X_train,y_train,X_train,y_train,depth) for depth in max_depths]
testing_accuracies = [DT_util(X_train,y_train,X_test,y_test,depth) for depth in max_depths]

plt.plot(max_depths,training_accuracies,'-',label='Train')
plt.plot(max_depths,testing_accuracies,'-',label='Test')

plt.xlabel('Max_Depth')
plt.ylabel('Accuracy_Score')

plt.legend()
plt.show()

```



Question 07:

Estimate the average accuracy of the Naïve Bayes Classifier using 5-fold cross- validation using a scikit-learn package in python. Plot the bar graph using matplotlib.

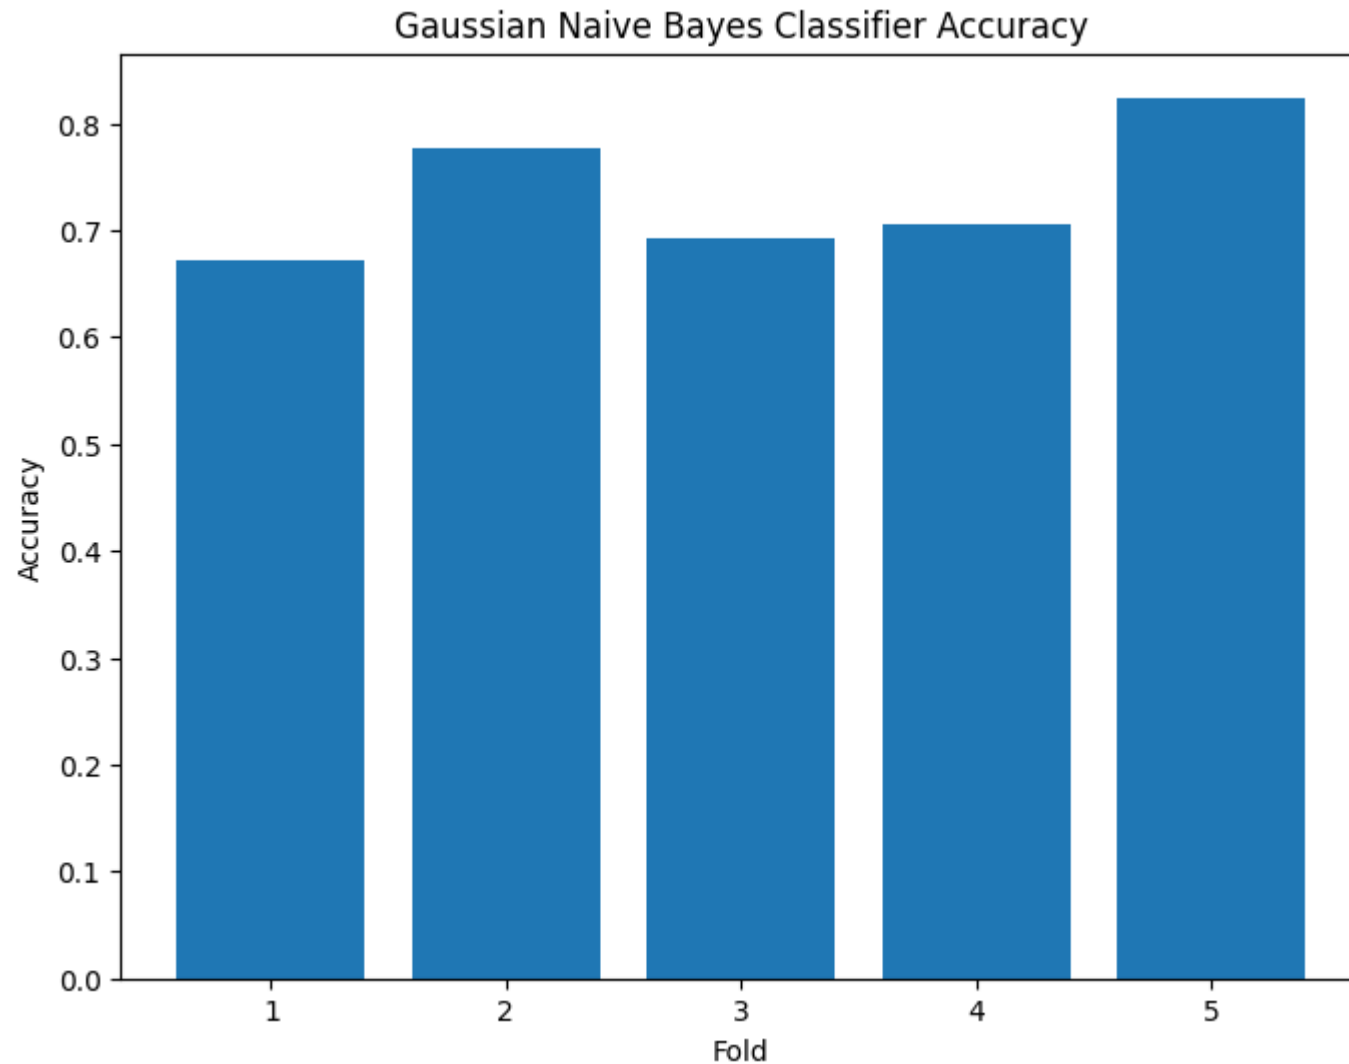
```
In [48]: from sklearn.model_selection import KFold, cross_val_score
from sklearn.naive_bayes import GaussianNB, MultinomialNB, BernoulliNB, ComplementNB

gnb_clf = GaussianNB()
scores = cross_val_score(gnb_clf, X, y, cv=5)
gnb_avg_accuracy = scores.mean()

# Plot the bar graph for Naive Bayes Classifier accuracy
plt.figure(figsize=(8, 6))
plt.bar(range(1, 6), scores)
plt.xlabel('Fold')
plt.ylabel('Accuracy')
plt.title('Gaussian Naive Bayes Classifier Accuracy')
```

```
plt.show()
```

```
print(f'Average Accuracy of Gaussian Naive Bayes Classifier: {gnb_avg_accuracy:.2f}')
```



Average Accuracy of Gaussian Naive Bayes Classifier: 0.73

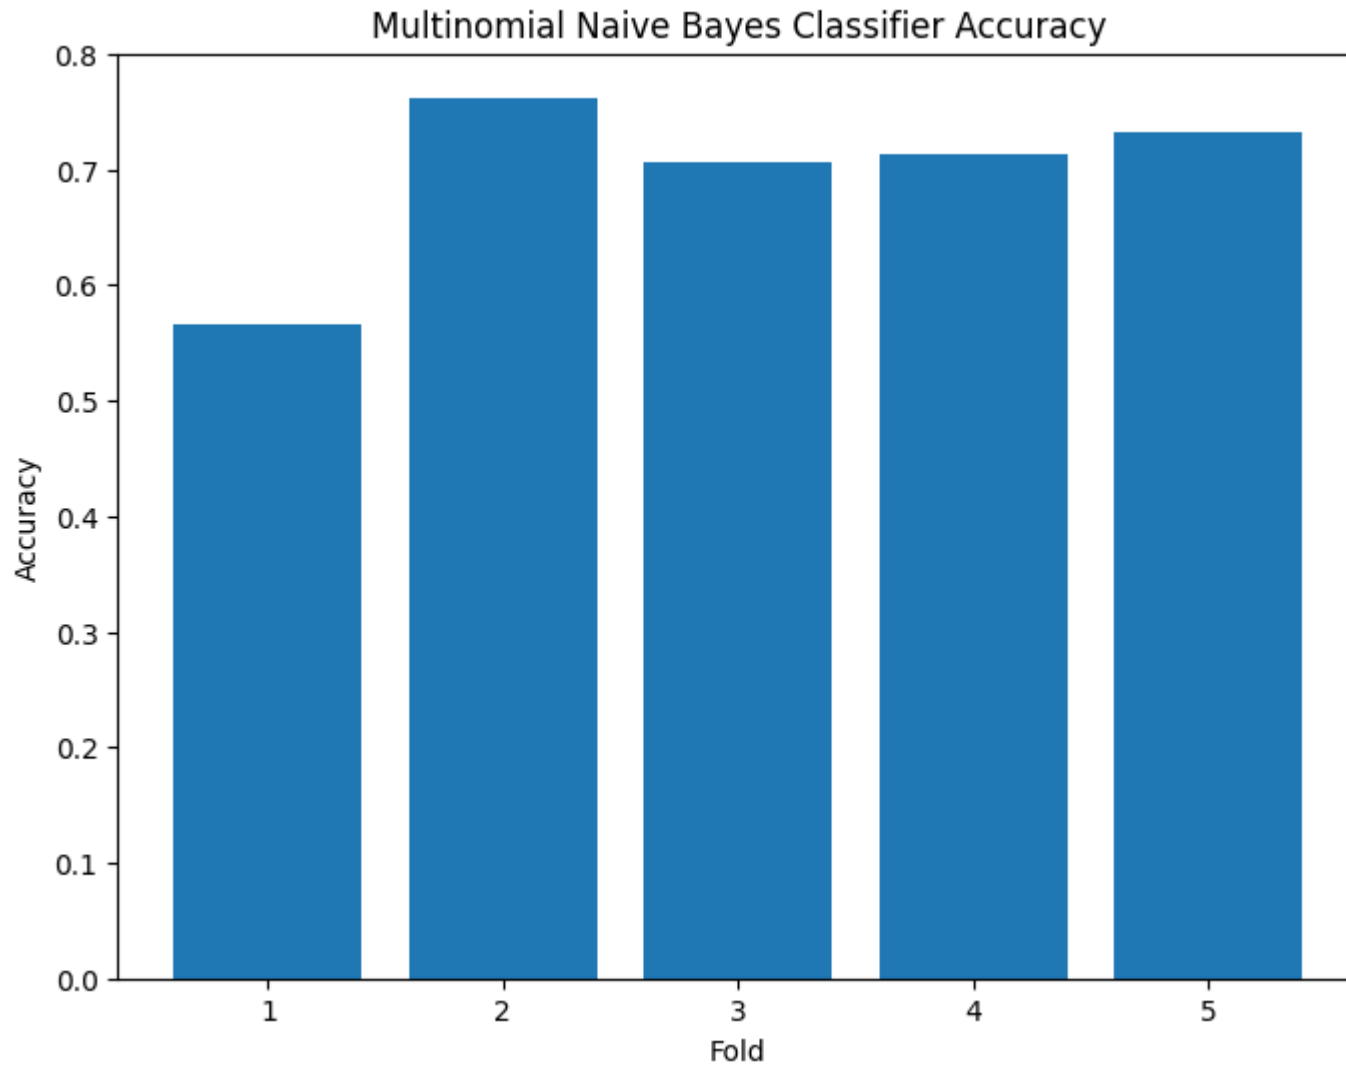
```
In [49]: mnb_clf = MultinomialNB()
scores = cross_val_score(mnb_clf, X, y, cv=5)
mnb_avg_accuracy = scores.mean()

# Plot the bar graph for Naive Bayes Classifier accuracy
plt.figure(figsize=(8, 6))
plt.bar(range(1, 6), scores)
plt.xlabel('Fold')
```



```
plt.ylabel('Accuracy')
plt.title('Multinomial Naive Bayes Classifier Accuracy')
plt.show()

print(f'Average Accuracy of Multinomial Naive Bayes Classifier: {mnb_avg_accuracy:.2f}')
```



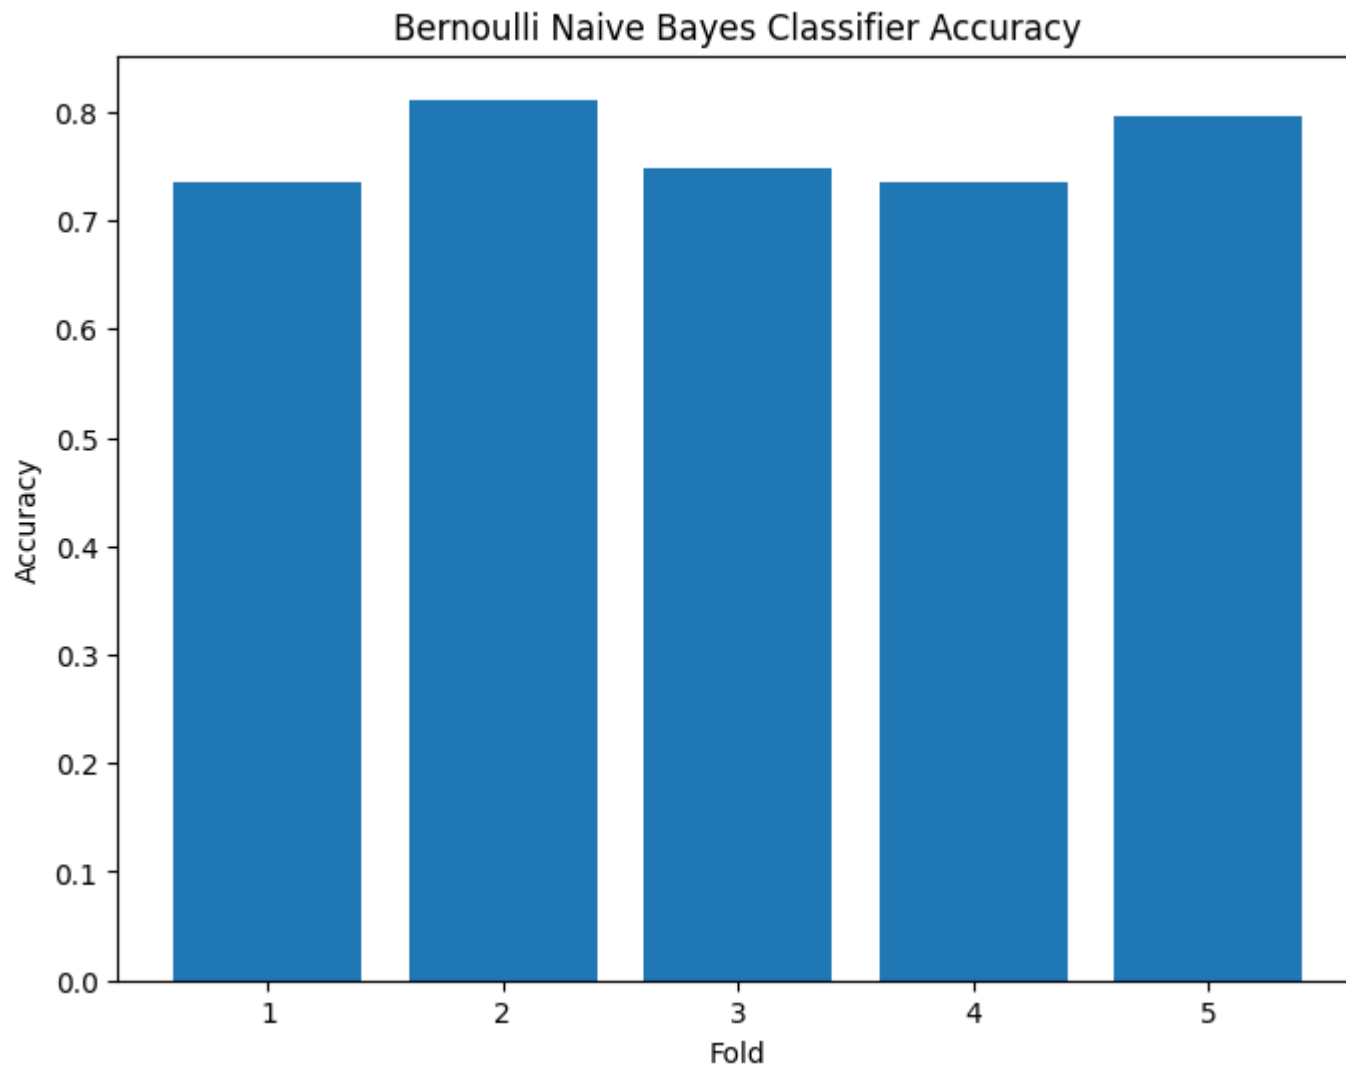
Average Accuracy of Multinomial Naive Bayes Classifier: 0.70

```
In [50]: bnb_clf = BernoulliNB()
scores = cross_val_score(bnb_clf, X, y, cv=5)
bnb_avg_accuracy = scores.mean()

# Plot the bar graph for Naive Bayes Classifier accuracy
plt.figure(figsize=(8, 6))
```

```
plt.bar(range(1, 6), scores)
plt.xlabel('Fold')
plt.ylabel('Accuracy')
plt.title('Bernoulli Naive Bayes Classifier Accuracy')
plt.show()

print(f'Average Accuracy of Bernoulli Naive Bayes Classifier: {bnb_avg_accuracy:.2f}')
```

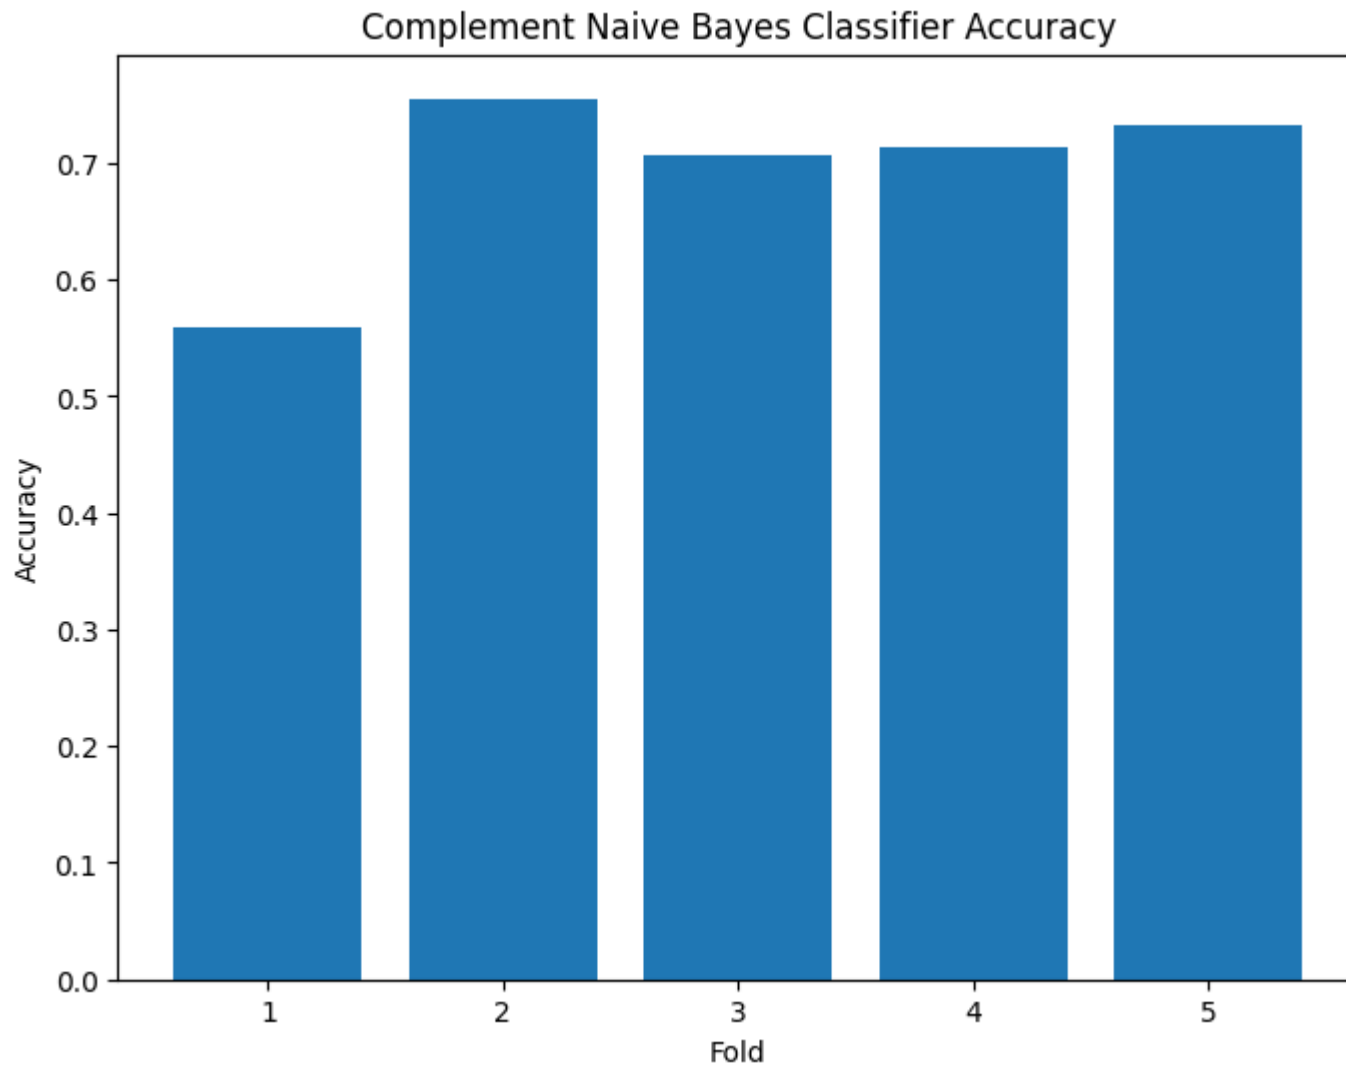


Average Accuracy of Bernoulli Naive Bayes Classifier: 0.76

```
In [51]: cnb_clf = ComplementNB()
scores = cross_val_score(cnb_clf, X, y, cv=5)
cnb_avg_accuracy = scores.mean()
```

```
# Plot the bar graph for Naive Bayes Classifier accuracy
plt.figure(figsize=(8, 6))
plt.bar(range(1, 6), scores)
plt.xlabel('Fold')
plt.ylabel('Accuracy')
plt.title('Complement Naive Bayes Classifier Accuracy')
plt.show()

print(f'Average Accuracy of Complement Naive Bayes Classifier: {cnb_avg_accuracy:.2f}')
```



Average Accuracy of Complement Naive Bayes Classifier: 0.69