

Assignment 04

Name:- Sk Fardeen Hossain

Roll No. :- 2021CSB023

G-Suite Id:- 2021csb023.sk@students.iiests.ac.in

Department:- Computer Science and Technology

Question 01

Download and install TensorFlow from https://www.tensorflow.org/install/install_sources or using command `sudo pip install tensorflow` (Alternatively the Keras library can be used).

```
In [1]: import warnings
warnings.filterwarnings('ignore')
```

```
In [2]: # Install the tensorflow library
!pip install tensorflow
```

Requirement already satisfied: tensorflow in /usr/local/lib/python3.10/dist-packages (2.17.0)
Requirement already satisfied: absl-py>=1.0.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (1.4.0)
Requirement already satisfied: astunparse>=1.6.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (1.6.3)
Requirement already satisfied: flatbuffers>=24.3.25 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (24.3.25)
Requirement already satisfied: gast!=0.5.0,!=0.5.1,!=0.5.2,>=0.2.1 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (0.6.0)
Requirement already satisfied: google-pasta>=0.1.1 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (0.2.0)
Requirement already satisfied: h5py>=3.10.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (3.11.0)
Requirement already satisfied: libclang>=13.0.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (18.1.1)
Requirement already satisfied: ml-dtypes<0.5.0,>=0.3.1 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (0.4.1)
Requirement already satisfied: opt-einsum>=2.3.2 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (3.4.0)
Requirement already satisfied: packaging in /usr/local/lib/python3.10/dist-packages (from tensorflow) (24.1)
Requirement already satisfied: protobuf!=4.21.0,!=4.21.1,!=4.21.2,!=4.21.3,!=4.21.4,!=4.21.5,<5.0.0dev,>=3.20.3 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (3.20.3)
Requirement already satisfied: requests<3,>=2.21.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (2.32.3)
Requirement already satisfied: setuptools in /usr/local/lib/python3.10/dist-packages (from tensorflow) (71.0.4)
Requirement already satisfied: six>=1.12.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (1.16.0)
Requirement already satisfied: termcolor>=1.1.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (2.4.0)
Requirement already satisfied: typing-extensions>=3.6.6 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (4.12.2)
Requirement already satisfied: wrapt>=1.11.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (1.16.0)
Requirement already satisfied: grpcio<2.0,>=1.24.3 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (1.64.1)
Requirement already satisfied: tensorboard<2.18,>=2.17 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (2.17.0)
Requirement already satisfied: keras>=3.2.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (3.4.1)
Requirement already satisfied: tensorflow-io-gcs-filesystem>=0.23.1 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (0.37.1)
Requirement already satisfied: numpy<2.0.0,>=1.23.5 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (1.26.4)
Requirement already satisfied: wheel<1.0,>=0.23.0 in /usr/local/lib/python3.10/dist-packages (from astunparse>=1.6.0->tensorflow) (0.44.0)
Requirement already satisfied: rich in /usr/local/lib/python3.10/dist-packages (from keras>=3.2.0->tensorflow) (13.8.1)
Requirement already satisfied: namex in /usr/local/lib/python3.10/dist-packages (from keras>=3.2.0->tensorflow) (0.0.8)
Requirement already satisfied: optree in /usr/local/lib/python3.10/dist-packages (from keras>=3.2.0->tensorflow) (0.12.1)
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.10/dist-packages (from requests<3,>=2.21.0->tensorflow) (3.3.2)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests<3,>=2.21.0->tensorflow) (3.10)
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.10/dist-packages (from requests<3,>=2.21.0->tensorflow) (2.2.3)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/dist-packages (from requests<3,>=2.21.0->tensorflow) (2024.8.30)
Requirement already satisfied: markdown>=2.6.8 in /usr/local/lib/python3.10/dist-packages (from tensorboard<2.18,>=2.17->tensorflow) (3.7)
Requirement already satisfied: tensorboard-data-server<0.8.0,>=0.7.0 in /usr/local/lib/python3.10/dist-packages (from tensorboard<2.18,>=2.17->tensorflow) (0.7.2)
Requirement already satisfied: werkzeug>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from tensorboard<2.18,>=2.17->tensorflow) (3.0.4)
Requirement already satisfied: MarkupSafe>=2.1.1 in /usr/local/lib/python3.10/dist-packages (from werkzeug>=1.0.1->tensorboard<

2.18,>=2.17->tensorflow)(2.1.5)

Requirement already satisfied: markdown-it-py>=2.2.0 in /usr/local/lib/python3.10/dist-packages (from rich->keras>=3.2.0->tensorflow) (3.0.0)

Requirement already satisfied: pygments<3.0.0,>=2.13.0 in /usr/local/lib/python3.10/dist-packages (from rich->keras>=3.2.0->tensorflow) (2.18.0)

Requirement already satisfied: mdurl~=0.1 in /usr/local/lib/python3.10/dist-packages (from markdown-it-py>=2.2.0->rich->keras>=3.2.0->tensorflow) (0.1.2)

Question 02

Download the MNIST dataset (contains class labels for digits 0 – 9)

```
In [3]: from keras.datasets import mnist  
(X_train, y_train), (X_test, y_test) = mnist.load_data()
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz>
11490434/11490434 ————— 1s 0us/step

```
In [4]: X_train.shape, y_train.shape, X_test.shape, y_test.shape
```

```
Out[4]: ((60000, 28, 28), (60000,), (10000, 28, 28), (10000,))
```

```
In [5]: print(X_train[0])
```

[0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0]							
[0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0]							
[0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0]							
[0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0]							
[0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0]							
[0	0	0	0	0	0	0	0	0	0	0	0	3	18	18	18	126
	175	26	166	255	247	127	0	0	0	0]							136
[0	0	0	0	0	0	0	0	30	36	94	154	170	253	253	253	253
	225	172	253	242	195	64	0	0	0	0]							
[0	0	0	0	0	0	0	49	238	253	253	253	253	253	253	253	251
	93	82	82	56	39	0	0	0	0	0]							
[0	0	0	0	0	0	0	18	219	253	253	253	253	253	198	182	247
	0	0	0	0	0	0	0	0	0	0]							241
[0	0	0	0	0	0	0	0	80	156	107	253	253	205	11	0	43
	0	0	0	0	0	0	0	0	0	0]							154
[0	0	0	0	0	0	0	0	0	14	1	154	253	90	0	0	0
	0	0	0	0	0	0	0	0	0	0]							0
[0	0	0	0	0	0	0	0	0	0	0	139	253	190	2	0	0
	0	0	0	0	0	0	0	0	0	0]							0
[0	0	0	0	0	0	0	0	0	0	0	11	190	253	70	0	0
	0	0	0	0	0	0	0	0	0	0]							0
[0	0	0	0	0	0	0	0	0	0	0	0	35	241	225	160	108
	0	0	0	0	0	0	0	0	0	0]							1
[0	0	0	0	0	0	0	0	0	0	0	0	0	0	81	240	253
	25	0	0	0	0	0	0	0	0	0]							119
[0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	45	186
	150	27	0	0	0	0	0	0	0	0]							253
[0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	16	93
	253	187	0	0	0	0	0	0	0	0]							252
[0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	249
	253	249	64	0	0	0	0	0	0	0]							
[0	0	0	0	0	0	0	0	0	0	0	0	0	0	46	130	183
	253	207	2	0	0	0	0	0	0	0]							

```
[ 0  0  0  0  55 172 226 253 253 253 253 244 133 11  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0]
[ 0  0  0  0 136 253 253 253 212 135 132 16  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0]
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0]
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0]
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0]]
```

As we can see, every image is given as a 28x28 pixel map where every value ranges from 0 to 255 on an intensity scale (0 ---> white and 255 ---> black). So it is basically a gray-scale image.

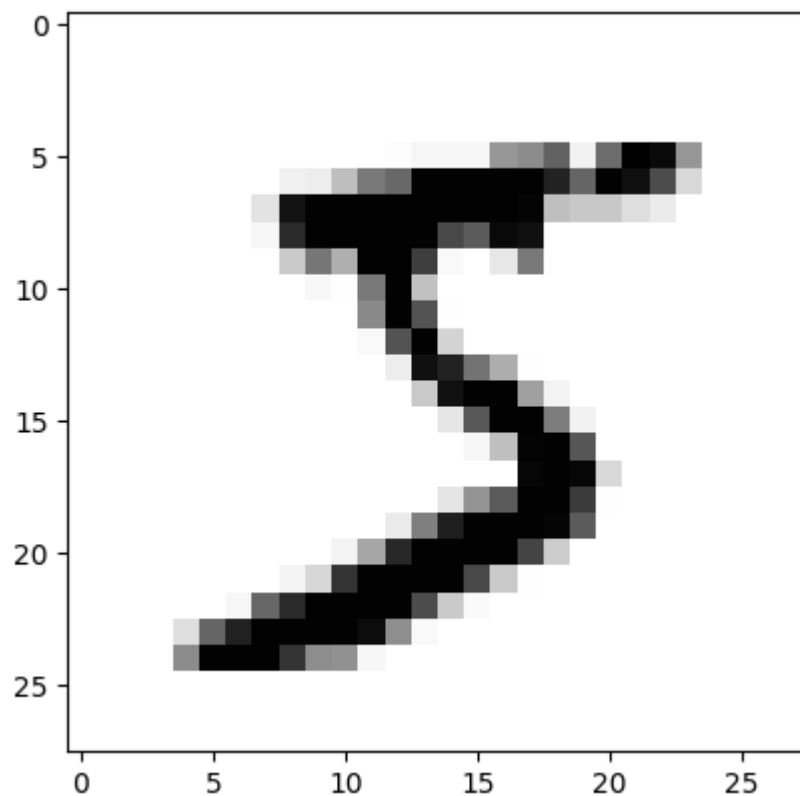
```
In [6]: # Let's see a sample image

import matplotlib.pyplot as plt

plt.imshow(X_train[0], cmap="Greys")

print(f"The image is a {y_train[0]}")
```

The image is a 5



Question 03

Reduce the training size by 1/10 if computation resources are limited. Define Radial Basis Function (RBF) as `def RBF(x, c, s): return np.exp(-np.sum((x - c) ** 2, axis = 1)/(2 * s ** 2))`, where `x` is the actual value, `c` is center (assumed as mean) and `s` is the standard deviation. Converted 28×28 image into 32×32 using rbf and store the new dataset with the labels. Split the dataset as 80% training and 10% validation and 10% test.

In [7]: `import numpy as np`

```
def RBF(x, c, s):
    return np.exp(-np.sum((x-c)**2)/(2*s**2))
```

In [8]: `## Convert a 28x28 image to 32x32 using RBF`

```
# Assuming x_train[0] is a 28x28 image
image = X_train[0]

# Calculate mean and standard deviation of the image
mean = np.mean(image)
```

```
std = np.std(image)

# Vectorize the RBF function to apply it element-wise on the image
RBF_vectorized = np.vectorize(RBF)

x_new = np.linspace(0, 27, 32).astype(int) # Mapping 32 points to 28
y_new = np.linspace(0, 27, 32).astype(int)

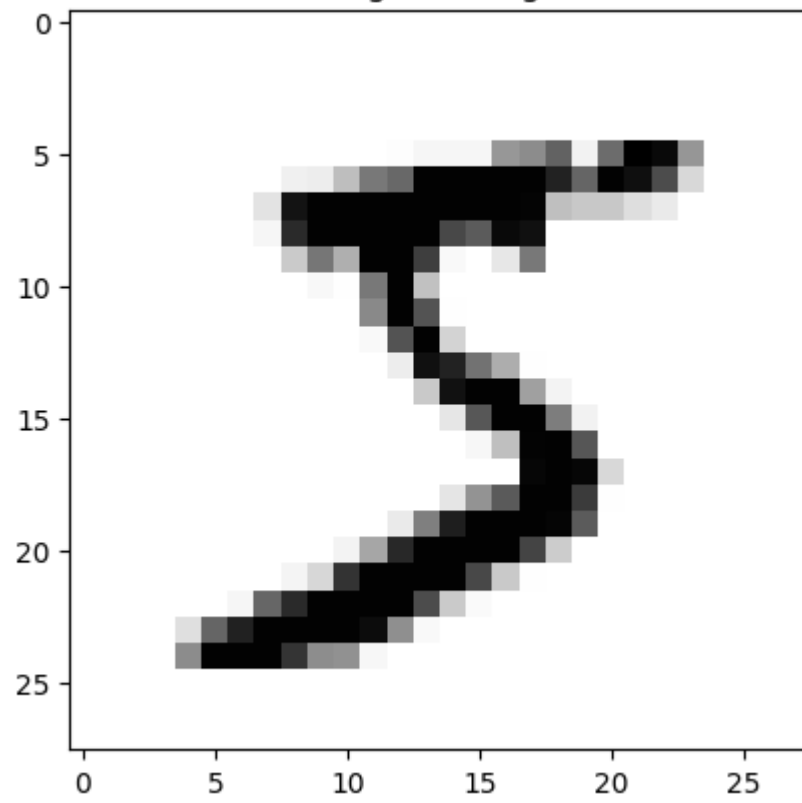
# Resize the image by selecting pixels from the original 28x28 image
new_image = image[np.ix_(x_new, y_new)]

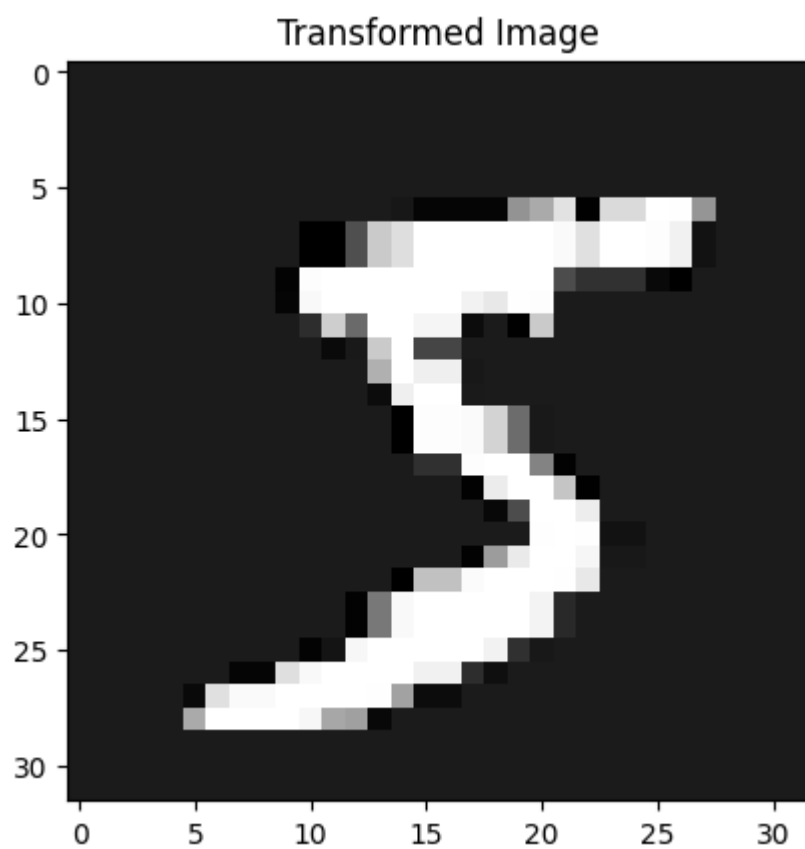
# Apply the RBF kernel on the resized image element-wise
new_image = RBF_vectorized(new_image, mean, std)

plt.imshow(image, cmap="Greys")
plt.title("Original Image")
plt.show()

plt.imshow(new_image, cmap="Greys")
plt.title("Transformed Image")
plt.show()
```

Original Image





```
In [9]: ## Do this iteratively for all the images
        '''
        @params:- A 28x28 image
        @return:- A 32x32 image
        @brief:- Convert a 28x28 image to 32x32 using RBF
        '''
        RBF_vectorized = np.vectorize(RBF)

        def get_transformed_image(image):
            mean = np.mean(image)
            std = np.std(image)
            x_new = np.linspace(0, 27, 32).astype(int) # Mapping 32 points to 28
            y_new = np.linspace(0, 27, 32).astype(int)

            # Resize the image by selecting pixels from the original 28x28 image
            new_image = image[np.ix_(x_new, y_new)]

            # Apply the RBF kernel on the resized image element-wise
            new_image = RBF_vectorized(new_image, mean, std)
```

```
return new_image
```

```
In [10]: # Combine both training and testing data into training for more samples
X_train = np.concatenate((X_train,X_test),axis=0)
y_train = np.concatenate((y_train,y_test),axis=0)
X_train.shape
```

```
Out[10]: (70000, 28, 28)
```

```
In [11]: # Training Data -> 80% , Validation Data -> 10% , Testing Data -> 10%

from sklearn.model_selection import train_test_split

X_train, X_rem, y_train, y_rem = train_test_split(X_train, y_train, test_size=0.2, random_state=4)
X_val, X_test, y_val, y_test = train_test_split(X_rem, y_rem, test_size=0.5, random_state=4)
```

```
In [12]: from tqdm import tqdm
X_train_transformed = np.empty((X_train.shape[0],32,32))

for i,image in tqdm(enumerate(X_train)):
    X_train_transformed[i] = get_transformed_image(image)

X_train_transformed.shape
```

```
56000it [08:56, 104.33it/s]
```

```
Out[12]: (56000, 32, 32)
```

```
In [13]: X_val_transformed = np.empty((X_val.shape[0],32,32))

for i,image in tqdm(enumerate(X_val)):
    X_val_transformed[i] = get_transformed_image(image)

X_val_transformed.shape
```

```
7000it [01:06, 105.69it/s]
```

```
Out[13]: (7000, 32, 32)
```

```
In [14]: X_test_transformed = np.empty((X_test.shape[0],32,32))

for i,image in tqdm(enumerate(X_test)):
    X_test_transformed[i] = get_transformed_image(image)

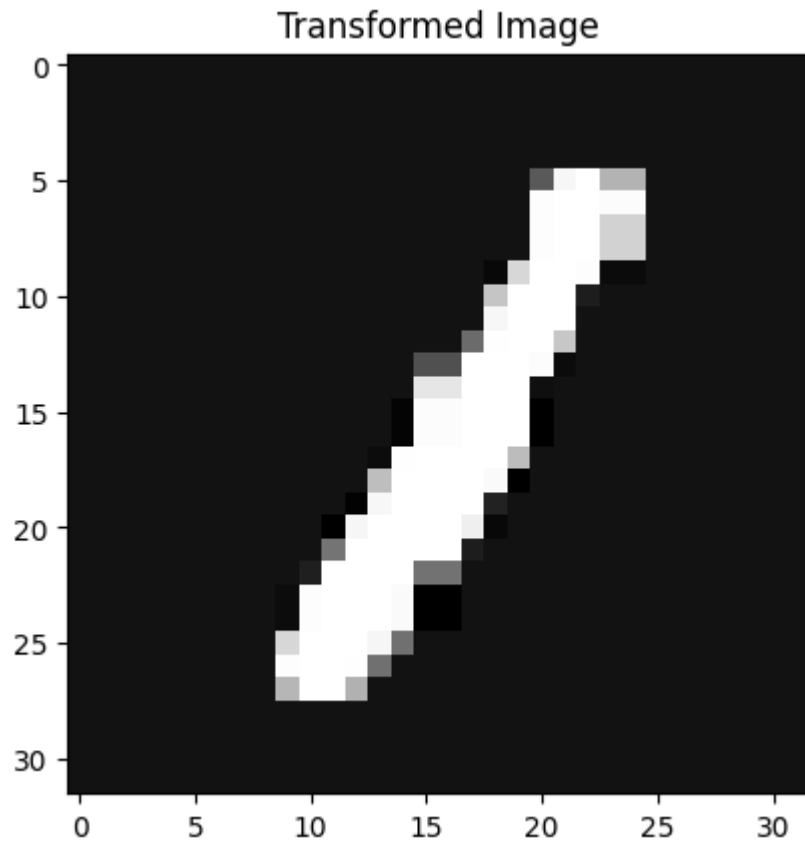
X_test_transformed.shape
```

```
7000it [01:06, 105.69it/s]
```

Out[14]: (7000, 32, 32)

```
In [15]: # Renaming
X_train = X_train_transformed
X_val = X_val_transformed
X_test = X_test_transformed
```

```
In [16]: plt.imshow(X_train[28], cmap="Greys")
plt.title("Transformed Image")
plt.show()
```



```
In [17]: # One hot encode the target labels [0-9]

import pandas as pd

y_train = pd.get_dummies(y_train, dtype='int').to_numpy()
y_val = pd.get_dummies(y_val, dtype='int').to_numpy()
y_test = pd.get_dummies(y_test, dtype='int').to_numpy()
y_train[28]
```

```
Out[17]: array([0, 1, 0, 0, 0, 0, 0, 0, 0, 0])
```

Question 04

Now run the fully connected network after flattening the data by changing the number of hyper-parameters:

- Use Gradient Descent Optimizer (learning rate = 0.001) and Squared Error Loss
- Use Adam Optimizer (learning rate = 0.001) and Categorical Cross Entropy Loss

Hidden Layers	Activation Function	Hidden Neurons
1	Sigmoid	[16]
2	Sigmoid	[16, 32]
3	Sigmoid	[16, 32, 64]

Try all the possible combinations.

```
In [21]: # Import necessary libraries
from tensorflow.keras import Sequential, Input
from tensorflow.keras.layers import Dense, Flatten, Dropout
from tensorflow.keras.losses import MeanSquaredError, CategoricalCrossentropy
from tensorflow.keras.optimizers import Adam, SGD
from tensorflow.keras.callbacks import EarlyStopping # Stop training when a monitored metric has stopped improving
import time # Calculate training time

# Defining Utility functions
def train_model(
    activation_func: 'str',
    hidden_neurons: 'list(int)',
    optimizer: 'tensorflow.keras.optimizers',
    loss_func: 'tensorflow.keras.losses',
    epochs: 'int',
    dropout_rate: 'float' = None,
```

```

learning_rate: 'float' = 0.001,

):
    model = Sequential() # Sequential model (feed forward neural net)
    # Adding the input layer
    model.add(Input(shape=(32,32)))

    # Flatten the input
    model.add(Flatten())

    # Hidden Layer Addition
    for i in range(len(hidden_neurons)):
        model.add(Dense(hidden_neurons[i], activation=activation_func))
        if dropout_rate:
            model.add(Dropout(rate=dropout_rate))

    # Adding the Output Layer
    model.add(Dense(10, activation='softmax')) # Multiclass classification so softmax activation function

    # Model summary
    model.summary()

    # Compile the model
    model.compile(
        optimizer=optimizer(learning_rate=learning_rate),
        loss=loss_func(),
        metrics=['accuracy']
    )

    start=time.time()

    # Fit the model
    history = model.fit(
        x=X_train,
        y=y_train,
        epochs=epochs,
        validation_data=(X_val, y_val),
        callbacks=[
            EarlyStopping(monitor='val_loss', patience=5, restore_best_weights=True)
        ],
        verbose='auto'
    )

    end=time.time()

    duration = end-start

```

```

return model, history, duration

def plot_metrics(
    history: 'tensorflow.keras.callbacks.History',
    activation_func: 'str',
    hidden_neurons: 'list(int)',
    dropout_rate: 'float' = None
):

    plt.plot(history.history['loss'], label='Training Loss')
    plt.plot(history.history['val_loss'], label='Validation Loss')
    plt.ylabel('Loss')
    plt.xlabel('Epoch')
    plt.legend()

    if dropout_rate is None:
        plt.title(f"Loss vs epochs for {activation_func} {hidden_neurons}")
    else:
        plt.title(f"Loss vs epochs for {activation_func} {hidden_neurons} dropout {dropout_rate}")

    plt.legend()
    plt.show()

    plt.plot(history.history['accuracy'], label='Training Accuracy')
    plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
    plt.ylabel('Accuracy')
    plt.xlabel('Epoch')
    plt.legend()

    if dropout_rate is None:
        plt.title(f"Accuracy vs epochs for {activation_func} {hidden_neurons}")
    else:
        plt.title(f"Accuracy vs epochs for {activation_func} {hidden_neurons} dropout {dropout_rate}")

    plt.legend()
    plt.show()

```

In [43]: `import pandas as pd`

```

result_df = pd.DataFrame(
    columns=[
        'Hidden Layers',
        'Activation Function',

```

```

        'Hidden Neurons',
        'Optimizer',
        'Training Time(in seconds)',
        'Test Loss',
        'Test Accuracy'
    ]
)

```

In [44]: *## Model 1 --> SGD, lr =0.001, Mean Squared Error loss, hidden_neurons=[16]*

```

hidden_neurons = [16]
activation_function = 'sigmoid'

model, history, duration = train_model(
    activation_func=activation_function,
    hidden_neurons=hidden_neurons,
    optimizer=SGD,
    loss_func=MeanSquaredError,
    epochs=100,
    dropout_rate=None,
    learning_rate=0.001
)

test_loss, test_acc = model.evaluate(X_test,y_test)

print(f"Test loss={test_loss} Test accuracy = {test_acc}")
print(f"Time required for training the model is {duration} seconds")

result_df.loc[len(result_df.index)] = [
    len(hidden_neurons),
    activation_function,
    str(hidden_neurons),
    'SGD',
    duration,
    test_loss,
    test_acc
]

plot_metrics(history, activation_function, hidden_neurons)

```

Model: "sequential_16"

Layer (type)	Output Shape	Param #
flatten_16 (Flatten)	(None, 1024)	0
dense_56 (Dense)	(None, 16)	16,400
dense_57 (Dense)	(None, 10)	170

Total params: 16,570 (64.73 KB)























Trainable params: 16,570 (64.73 KB)

Non-trainable params: 0 (0.00 B)

Epoch 1/100	1750/1750	3s	2ms/step	- accuracy: 0.0980	- loss: 0.0954	- val_accuracy: 0.0943	- val_loss: 0.0926
Epoch 2/100	1750/1750	6s	2ms/step	- accuracy: 0.0970	- loss: 0.0922	- val_accuracy: 0.0976	- val_loss: 0.0913
Epoch 3/100	1750/1750	3s	2ms/step	- accuracy: 0.1023	- loss: 0.0912	- val_accuracy: 0.1031	- val_loss: 0.0908
Epoch 4/100	1750/1750	5s	2ms/step	- accuracy: 0.1053	- loss: 0.0907	- val_accuracy: 0.1040	- val_loss: 0.0904
Epoch 5/100	1750/1750	3s	2ms/step	- accuracy: 0.1082	- loss: 0.0904	- val_accuracy: 0.1161	- val_loss: 0.0902
Epoch 6/100	1750/1750	5s	2ms/step	- accuracy: 0.1191	- loss: 0.0903	- val_accuracy: 0.1334	- val_loss: 0.0901
Epoch 7/100	1750/1750	5s	2ms/step	- accuracy: 0.1364	- loss: 0.0901	- val_accuracy: 0.1534	- val_loss: 0.0899
Epoch 8/100	1750/1750	3s	2ms/step	- accuracy: 0.1504	- loss: 0.0900	- val_accuracy: 0.1666	- val_loss: 0.0898
Epoch 9/100	1750/1750	5s	2ms/step	- accuracy: 0.1558	- loss: 0.0899	- val_accuracy: 0.1746	- val_loss: 0.0898
Epoch 10/100	1750/1750	5s	2ms/step	- accuracy: 0.1675	- loss: 0.0898	- val_accuracy: 0.1794	- val_loss: 0.0897
Epoch 11/100	1750/1750	6s	2ms/step	- accuracy: 0.1738	- loss: 0.0897	- val_accuracy: 0.1806	- val_loss: 0.0896
Epoch 12/100	1750/1750	4s	2ms/step	- accuracy: 0.1759	- loss: 0.0896	- val_accuracy: 0.1833	- val_loss: 0.0895
Epoch 13/100	1750/1750	6s	2ms/step	- accuracy: 0.1796	- loss: 0.0895	- val_accuracy: 0.1870	- val_loss: 0.0895
Epoch 14/100	1750/1750	5s	3ms/step	- accuracy: 0.1812	- loss: 0.0895	- val_accuracy: 0.1901	- val_loss: 0.0894
Epoch 15/100	1750/1750	4s	2ms/step	- accuracy: 0.1839	- loss: 0.0895	- val_accuracy: 0.1917	- val_loss: 0.0894
Epoch 16/100	1750/1750	3s	2ms/step	- accuracy: 0.1835	- loss: 0.0894	- val_accuracy: 0.1927	- val_loss: 0.0893
Epoch 17/100	1750/1750	8s	4ms/step	- accuracy: 0.1879	- loss: 0.0894	- val_accuracy: 0.1947	- val_loss: 0.0893
Epoch 18/100	1750/1750	7s	2ms/step	- accuracy: 0.1901	- loss: 0.0893	- val_accuracy: 0.1973	- val_loss: 0.0892
Epoch 19/100	1750/1750	5s	2ms/step	- accuracy: 0.1907	- loss: 0.0893	- val_accuracy: 0.1996	- val_loss: 0.0892
Epoch 20/100	1750/1750	3s	2ms/step	- accuracy: 0.1896	- loss: 0.0892	- val_accuracy: 0.2024	- val_loss: 0.0891
Epoch 21/100	1750/1750	3s	2ms/step	- accuracy: 0.1928	- loss: 0.0892	- val_accuracy: 0.2054	- val_loss: 0.0891
Epoch 22/100	1750/1750	6s	2ms/step	- accuracy: 0.1953	- loss: 0.0891	- val_accuracy: 0.2066	- val_loss: 0.0891
Epoch 23/100	1750/1750	3s	2ms/step	- accuracy: 0.1952	- loss: 0.0891	- val_accuracy: 0.2073	- val_loss: 0.0890

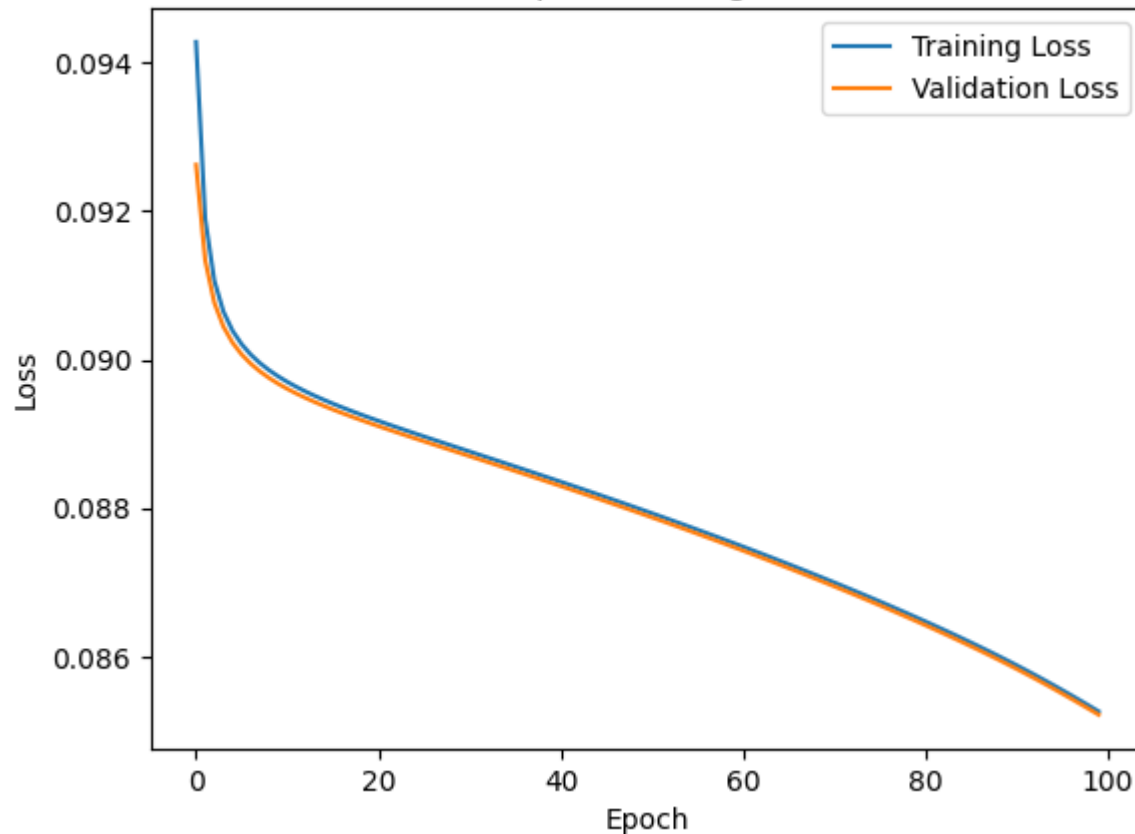
Epoch 24/100	1750/1750	5s	2ms/step	- accuracy: 0.1972	- loss: 0.0891	- val_accuracy: 0.2104	- val_loss: 0.0890
Epoch 25/100	1750/1750	7s	3ms/step	- accuracy: 0.2036	- loss: 0.0890	- val_accuracy: 0.2116	- val_loss: 0.0889
Epoch 26/100	1750/1750	8s	2ms/step	- accuracy: 0.2039	- loss: 0.0890	- val_accuracy: 0.2159	- val_loss: 0.0889
Epoch 27/100	1750/1750	5s	2ms/step	- accuracy: 0.2041	- loss: 0.0889	- val_accuracy: 0.2180	- val_loss: 0.0889
Epoch 28/100	1750/1750	3s	2ms/step	- accuracy: 0.2072	- loss: 0.0889	- val_accuracy: 0.2211	- val_loss: 0.0888
Epoch 29/100	1750/1750	6s	2ms/step	- accuracy: 0.2098	- loss: 0.0889	- val_accuracy: 0.2246	- val_loss: 0.0888
Epoch 30/100	1750/1750	4s	2ms/step	- accuracy: 0.2132	- loss: 0.0888	- val_accuracy: 0.2269	- val_loss: 0.0887
Epoch 31/100	1750/1750	3s	2ms/step	- accuracy: 0.2145	- loss: 0.0888	- val_accuracy: 0.2296	- val_loss: 0.0887
Epoch 32/100	1750/1750	3s	2ms/step	- accuracy: 0.2219	- loss: 0.0887	- val_accuracy: 0.2334	- val_loss: 0.0887
Epoch 33/100	1750/1750	3s	2ms/step	- accuracy: 0.2224	- loss: 0.0887	- val_accuracy: 0.2366	- val_loss: 0.0886
Epoch 34/100	1750/1750	4s	2ms/step	- accuracy: 0.2253	- loss: 0.0886	- val_accuracy: 0.2393	- val_loss: 0.0886
Epoch 35/100	1750/1750	4s	2ms/step	- accuracy: 0.2303	- loss: 0.0886	- val_accuracy: 0.2431	- val_loss: 0.0885
Epoch 36/100	1750/1750	5s	2ms/step	- accuracy: 0.2332	- loss: 0.0886	- val_accuracy: 0.2466	- val_loss: 0.0885
Epoch 37/100	1750/1750	6s	2ms/step	- accuracy: 0.2353	- loss: 0.0885	- val_accuracy: 0.2504	- val_loss: 0.0885
Epoch 38/100	1750/1750	3s	2ms/step	- accuracy: 0.2408	- loss: 0.0885	- val_accuracy: 0.2541	- val_loss: 0.0884
Epoch 39/100	1750/1750	5s	2ms/step	- accuracy: 0.2419	- loss: 0.0885	- val_accuracy: 0.2580	- val_loss: 0.0884
Epoch 40/100	1750/1750	4s	2ms/step	- accuracy: 0.2458	- loss: 0.0884	- val_accuracy: 0.2627	- val_loss: 0.0883
Epoch 41/100	1750/1750	4s	2ms/step	- accuracy: 0.2524	- loss: 0.0884	- val_accuracy: 0.2673	- val_loss: 0.0883
Epoch 42/100	1750/1750	3s	2ms/step	- accuracy: 0.2582	- loss: 0.0883	- val_accuracy: 0.2710	- val_loss: 0.0883
Epoch 43/100	1750/1750	3s	2ms/step	- accuracy: 0.2627	- loss: 0.0883	- val_accuracy: 0.2767	- val_loss: 0.0882
Epoch 44/100	1750/1750	4s	2ms/step	- accuracy: 0.2662	- loss: 0.0882	- val_accuracy: 0.2816	- val_loss: 0.0882
Epoch 45/100	1750/1750	4s	2ms/step	- accuracy: 0.2695	- loss: 0.0882	- val_accuracy: 0.2853	- val_loss: 0.0881
Epoch 46/100	1750/1750	3s	2ms/step	- accuracy: 0.2738	- loss: 0.0881	- val_accuracy: 0.2894	- val_loss: 0.0881

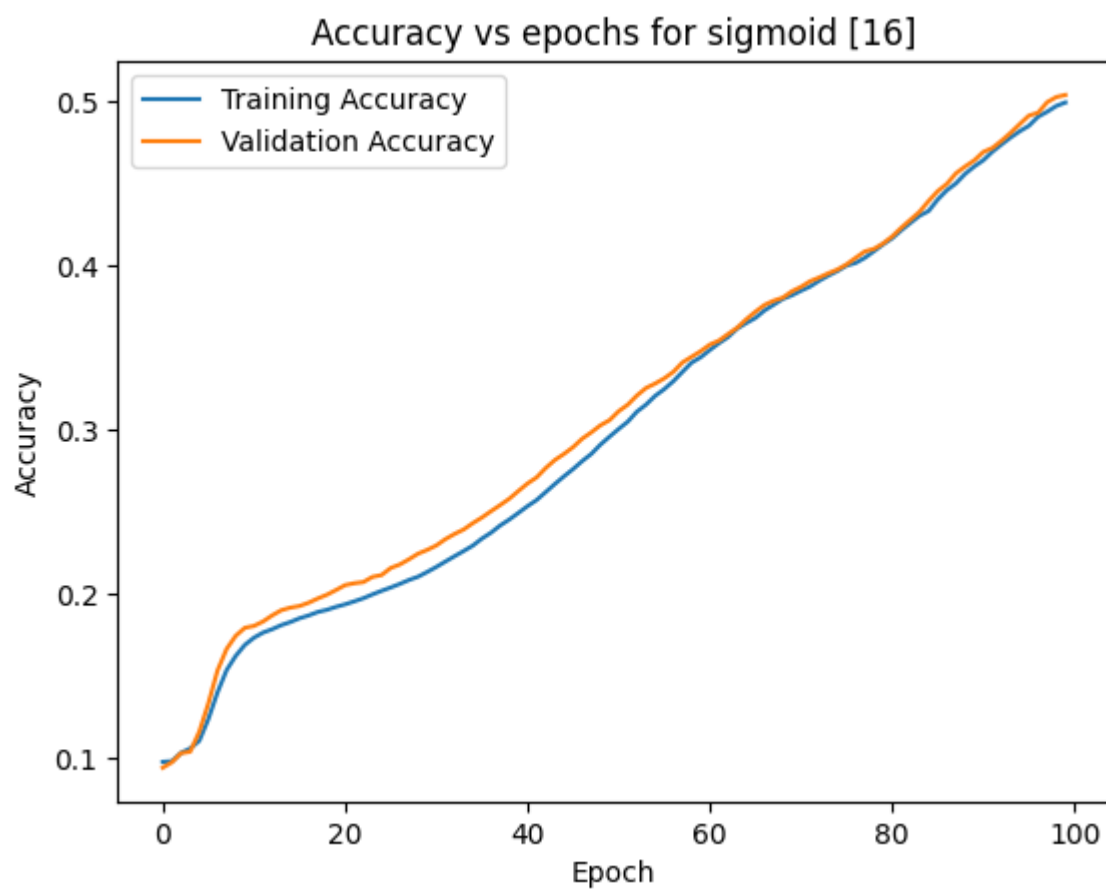
Epoch 47/100	1750/1750	3s	2ms/step	- accuracy: 0.2819	- loss: 0.0881	- val_accuracy: 0.2946	- val_loss: 0.0880
Epoch 48/100	1750/1750	5s	2ms/step	- accuracy: 0.2882	- loss: 0.0880	- val_accuracy: 0.2984	- val_loss: 0.0880
Epoch 49/100	1750/1750	5s	2ms/step	- accuracy: 0.2903	- loss: 0.0880	- val_accuracy: 0.3027	- val_loss: 0.0880
Epoch 50/100	1750/1750	3s	2ms/step	- accuracy: 0.2922	- loss: 0.0880	- val_accuracy: 0.3059	- val_loss: 0.0879
Epoch 51/100	1750/1750	5s	2ms/step	- accuracy: 0.2975	- loss: 0.0880	- val_accuracy: 0.3113	- val_loss: 0.0879
Epoch 52/100	1750/1750	3s	2ms/step	- accuracy: 0.3047	- loss: 0.0879	- val_accuracy: 0.3153	- val_loss: 0.0878
Epoch 53/100	1750/1750	6s	2ms/step	- accuracy: 0.3120	- loss: 0.0878	- val_accuracy: 0.3209	- val_loss: 0.0878
Epoch 54/100	1750/1750	4s	2ms/step	- accuracy: 0.3160	- loss: 0.0878	- val_accuracy: 0.3254	- val_loss: 0.0877
Epoch 55/100	1750/1750	4s	2ms/step	- accuracy: 0.3202	- loss: 0.0878	- val_accuracy: 0.3281	- val_loss: 0.0877
Epoch 56/100	1750/1750	3s	2ms/step	- accuracy: 0.3285	- loss: 0.0877	- val_accuracy: 0.3313	- val_loss: 0.0877
Epoch 57/100	1750/1750	4s	2ms/step	- accuracy: 0.3291	- loss: 0.0877	- val_accuracy: 0.3353	- val_loss: 0.0876
Epoch 58/100	1750/1750	4s	2ms/step	- accuracy: 0.3359	- loss: 0.0876	- val_accuracy: 0.3410	- val_loss: 0.0876
Epoch 59/100	1750/1750	5s	2ms/step	- accuracy: 0.3412	- loss: 0.0876	- val_accuracy: 0.3444	- val_loss: 0.0875
Epoch 60/100	1750/1750	6s	2ms/step	- accuracy: 0.3436	- loss: 0.0875	- val_accuracy: 0.3479	- val_loss: 0.0875
Epoch 61/100	1750/1750	4s	2ms/step	- accuracy: 0.3467	- loss: 0.0875	- val_accuracy: 0.3519	- val_loss: 0.0874
Epoch 62/100	1750/1750	5s	2ms/step	- accuracy: 0.3534	- loss: 0.0874	- val_accuracy: 0.3543	- val_loss: 0.0874
Epoch 63/100	1750/1750	4s	2ms/step	- accuracy: 0.3562	- loss: 0.0874	- val_accuracy: 0.3583	- val_loss: 0.0873
Epoch 64/100	1750/1750	3s	2ms/step	- accuracy: 0.3580	- loss: 0.0874	- val_accuracy: 0.3621	- val_loss: 0.0873
Epoch 65/100	1750/1750	5s	2ms/step	- accuracy: 0.3614	- loss: 0.0873	- val_accuracy: 0.3673	- val_loss: 0.0872
Epoch 66/100	1750/1750	4s	2ms/step	- accuracy: 0.3693	- loss: 0.0873	- val_accuracy: 0.3719	- val_loss: 0.0872
Epoch 67/100	1750/1750	4s	2ms/step	- accuracy: 0.3678	- loss: 0.0872	- val_accuracy: 0.3760	- val_loss: 0.0871
Epoch 68/100	1750/1750	3s	2ms/step	- accuracy: 0.3731	- loss: 0.0872	- val_accuracy: 0.3786	- val_loss: 0.0871
Epoch 69/100	1750/1750	6s	2ms/step	- accuracy: 0.3771	- loss: 0.0871	- val_accuracy: 0.3804	- val_loss: 0.0870

Epoch 70/100		4s	2ms/step	- accuracy: 0.3797	- loss: 0.0871	- val_accuracy: 0.3843	- val_loss: 0.0870
Epoch 71/100		3s	2ms/step	- accuracy: 0.3850	- loss: 0.0870	- val_accuracy: 0.3870	- val_loss: 0.0869
Epoch 72/100		5s	2ms/step	- accuracy: 0.3868	- loss: 0.0870	- val_accuracy: 0.3904	- val_loss: 0.0869
Epoch 73/100		4s	2ms/step	- accuracy: 0.3875	- loss: 0.0869	- val_accuracy: 0.3927	- val_loss: 0.0868
Epoch 74/100		3s	2ms/step	- accuracy: 0.3942	- loss: 0.0868	- val_accuracy: 0.3953	- val_loss: 0.0868
Epoch 75/100		3s	2ms/step	- accuracy: 0.3955	- loss: 0.0868	- val_accuracy: 0.3976	- val_loss: 0.0867
Epoch 76/100		3s	2ms/step	- accuracy: 0.4028	- loss: 0.0867	- val_accuracy: 0.4007	- val_loss: 0.0867
Epoch 77/100		6s	2ms/step	- accuracy: 0.4032	- loss: 0.0867	- val_accuracy: 0.4047	- val_loss: 0.0866
Epoch 78/100		3s	2ms/step	- accuracy: 0.4009	- loss: 0.0867	- val_accuracy: 0.4086	- val_loss: 0.0866
Epoch 79/100		5s	2ms/step	- accuracy: 0.4073	- loss: 0.0866	- val_accuracy: 0.4101	- val_loss: 0.0865
Epoch 80/100		6s	2ms/step	- accuracy: 0.4119	- loss: 0.0865	- val_accuracy: 0.4134	- val_loss: 0.0865
Epoch 81/100		3s	2ms/step	- accuracy: 0.4159	- loss: 0.0865	- val_accuracy: 0.4177	- val_loss: 0.0864
Epoch 82/100		5s	2ms/step	- accuracy: 0.4220	- loss: 0.0864	- val_accuracy: 0.4230	- val_loss: 0.0864
Epoch 83/100		5s	2ms/step	- accuracy: 0.4239	- loss: 0.0864	- val_accuracy: 0.4279	- val_loss: 0.0863
Epoch 84/100		5s	2ms/step	- accuracy: 0.4275	- loss: 0.0863	- val_accuracy: 0.4327	- val_loss: 0.0863
Epoch 85/100		7s	3ms/step	- accuracy: 0.4327	- loss: 0.0863	- val_accuracy: 0.4393	- val_loss: 0.0862
Epoch 86/100		3s	2ms/step	- accuracy: 0.4360	- loss: 0.0862	- val_accuracy: 0.4453	- val_loss: 0.0861
Epoch 87/100		5s	2ms/step	- accuracy: 0.4438	- loss: 0.0861	- val_accuracy: 0.4497	- val_loss: 0.0861
Epoch 88/100		4s	2ms/step	- accuracy: 0.4479	- loss: 0.0861	- val_accuracy: 0.4561	- val_loss: 0.0860
Epoch 89/100		4s	2ms/step	- accuracy: 0.4532	- loss: 0.0860	- val_accuracy: 0.4603	- val_loss: 0.0860
Epoch 90/100		5s	2ms/step	- accuracy: 0.4593	- loss: 0.0859	- val_accuracy: 0.4640	- val_loss: 0.0859
Epoch 91/100		4s	2ms/step	- accuracy: 0.4649	- loss: 0.0859	- val_accuracy: 0.4691	- val_loss: 0.0858
Epoch 92/100		3s	2ms/step	- accuracy: 0.4671	- loss: 0.0858	- val_accuracy: 0.4716	- val_loss: 0.0858

Epoch 93/100 **1750/1750** 5s 2ms/step - accuracy: 0.4722 - loss: 0.0858 - val_accuracy: 0.4760 - val_loss: 0.0857
Epoch 94/100 **1750/1750** 3s 2ms/step - accuracy: 0.4746 - loss: 0.0857 - val_accuracy: 0.4809 - val_loss: 0.0856
Epoch 95/100 **1750/1750** 5s 2ms/step - accuracy: 0.4799 - loss: 0.0856 - val_accuracy: 0.4861 - val_loss: 0.0856
Epoch 96/100 **1750/1750** 3s 2ms/step - accuracy: 0.4830 - loss: 0.0856 - val_accuracy: 0.4913 - val_loss: 0.0855
Epoch 97/100 **1750/1750** 6s 2ms/step - accuracy: 0.4902 - loss: 0.0855 - val_accuracy: 0.4930 - val_loss: 0.0854
Epoch 98/100 **1750/1750** 5s 2ms/step - accuracy: 0.4946 - loss: 0.0854 - val_accuracy: 0.4994 - val_loss: 0.0854
Epoch 99/100 **1750/1750** 3s 2ms/step - accuracy: 0.5003 - loss: 0.0853 - val_accuracy: 0.5026 - val_loss: 0.0853
Epoch 100/100 **1750/1750** 3s 2ms/step - accuracy: 0.4970 - loss: 0.0853 - val_accuracy: 0.5039 - val_loss: 0.0852
219/219 1s 2ms/step - accuracy: 0.5047 - loss: 0.0853
Test loss=0.08531524986028671 Test accuracy = 0.4977142810821533
Time required for training the model is 428.50232338905334 seconds

Loss vs epochs for sigmoid [16]





In [45]: `## Model 2 --> SGD, lr =0.001, Mean Squared Error loss, hidden_neurons=[16,32]`

```
hidden_neurons = [16,32]
activation_function = 'sigmoid'

model, history,duration = train_model(
    activation_func=activation_function,
    hidden_neurons=hidden_neurons,
    optimizer=SGD,
    loss_func=MeanSquaredError,
    epochs=100,
    dropout_rate=None,
    learning_rate=0.001
)

test_loss, test_acc = model.evaluate(X_test,y_test)

print(f"Test loss={test_loss} Test accuracy = {test_acc}")
```

```

print(f"Time required to train the model is {duration} seconds")

result_df.loc[len(result_df.index)] = [
    len(hidden_neurons),
    activation_function,
    str(hidden_neurons),
    'SGD',
    duration,
    test_loss,
    test_acc
]

plot_metrics(history, activation_function, hidden_neurons)

```

Model: "sequential_17"

Layer (type)	Output Shape	Param #
flatten_17 (Flatten)	(None, 1024)	0
dense_58 (Dense)	(None, 16)	16,400
dense_59 (Dense)	(None, 32)	544
dense_60 (Dense)	(None, 10)	330

Total params: 17,274 (67.48 KB)

Trainable params: 17,274 (67.48 KB)

Non-trainable params: 0 (0.00 B)

Epoch 1/100	1750/1750	3s	2ms/step	- accuracy: 0.1012	- loss: 0.0943	- val_accuracy: 0.1037	- val_loss: 0.0939
Epoch 2/100	1750/1750	5s	2ms/step	- accuracy: 0.1008	- loss: 0.0938	- val_accuracy: 0.1056	- val_loss: 0.0935
Epoch 3/100	1750/1750	5s	2ms/step	- accuracy: 0.1024	- loss: 0.0934	- val_accuracy: 0.1083	- val_loss: 0.0931
Epoch 4/100	1750/1750	3s	2ms/step	- accuracy: 0.1018	- loss: 0.0930	- val_accuracy: 0.1090	- val_loss: 0.0928
Epoch 5/100	1750/1750	3s	2ms/step	- accuracy: 0.1051	- loss: 0.0928	- val_accuracy: 0.1101	- val_loss: 0.0926
Epoch 6/100	1750/1750	3s	2ms/step	- accuracy: 0.1031	- loss: 0.0925	- val_accuracy: 0.1100	- val_loss: 0.0923
Epoch 7/100	1750/1750	4s	2ms/step	- accuracy: 0.1040	- loss: 0.0923	- val_accuracy: 0.1113	- val_loss: 0.0922
Epoch 8/100	1750/1750	5s	2ms/step	- accuracy: 0.1077	- loss: 0.0921	- val_accuracy: 0.1114	- val_loss: 0.0920
Epoch 9/100	1750/1750	6s	2ms/step	- accuracy: 0.1100	- loss: 0.0920	- val_accuracy: 0.1127	- val_loss: 0.0919
Epoch 10/100	1750/1750	4s	2ms/step	- accuracy: 0.1125	- loss: 0.0918	- val_accuracy: 0.1136	- val_loss: 0.0917
Epoch 11/100	1750/1750	6s	2ms/step	- accuracy: 0.1102	- loss: 0.0917	- val_accuracy: 0.1127	- val_loss: 0.0916
Epoch 12/100	1750/1750	4s	2ms/step	- accuracy: 0.1138	- loss: 0.0915	- val_accuracy: 0.1157	- val_loss: 0.0915
Epoch 13/100	1750/1750	3s	2ms/step	- accuracy: 0.1153	- loss: 0.0915	- val_accuracy: 0.1153	- val_loss: 0.0914
Epoch 14/100	1750/1750	5s	2ms/step	- accuracy: 0.1155	- loss: 0.0914	- val_accuracy: 0.1153	- val_loss: 0.0914
Epoch 15/100	1750/1750	4s	2ms/step	- accuracy: 0.1160	- loss: 0.0913	- val_accuracy: 0.1166	- val_loss: 0.0913
Epoch 16/100	1750/1750	3s	2ms/step	- accuracy: 0.1132	- loss: 0.0913	- val_accuracy: 0.1173	- val_loss: 0.0912
Epoch 17/100	1750/1750	5s	2ms/step	- accuracy: 0.1177	- loss: 0.0912	- val_accuracy: 0.1174	- val_loss: 0.0912
Epoch 18/100	1750/1750	6s	2ms/step	- accuracy: 0.1173	- loss: 0.0912	- val_accuracy: 0.1173	- val_loss: 0.0911
Epoch 19/100	1750/1750	4s	2ms/step	- accuracy: 0.1166	- loss: 0.0911	- val_accuracy: 0.1173	- val_loss: 0.0911
Epoch 20/100	1750/1750	5s	2ms/step	- accuracy: 0.1187	- loss: 0.0911	- val_accuracy: 0.1177	- val_loss: 0.0910
Epoch 21/100	1750/1750	6s	2ms/step	- accuracy: 0.1224	- loss: 0.0910	- val_accuracy: 0.1187	- val_loss: 0.0910
Epoch 22/100	1750/1750	3s	2ms/step	- accuracy: 0.1217	- loss: 0.0909	- val_accuracy: 0.1190	- val_loss: 0.0909
Epoch 23/100	1750/1750	5s	2ms/step	- accuracy: 0.1234	- loss: 0.0909	- val_accuracy: 0.1189	- val_loss: 0.0909

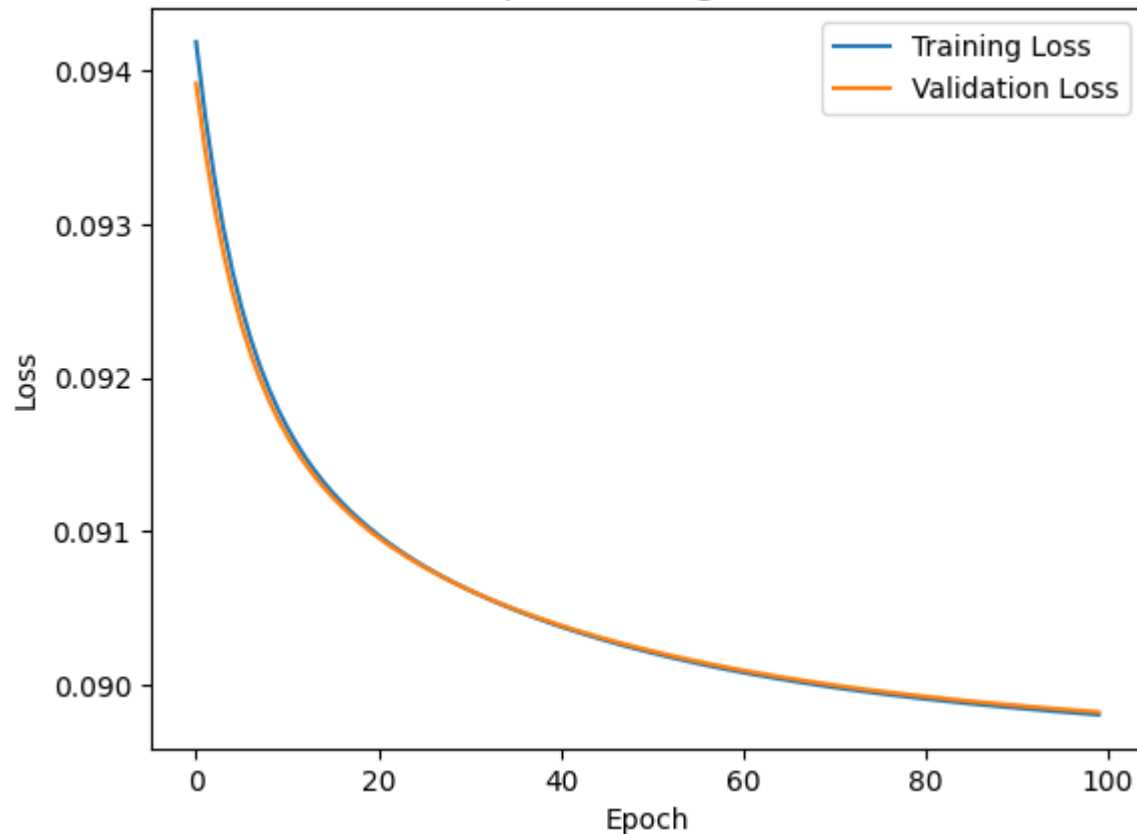
Epoch 24/100	1750/1750	5s	2ms/step	- accuracy: 0.1208	- loss: 0.0909	- val_accuracy: 0.1184	- val_loss: 0.0908
Epoch 25/100	1750/1750	3s	2ms/step	- accuracy: 0.1208	- loss: 0.0908	- val_accuracy: 0.1176	- val_loss: 0.0908
Epoch 26/100	1750/1750	5s	2ms/step	- accuracy: 0.1193	- loss: 0.0908	- val_accuracy: 0.1179	- val_loss: 0.0908
Epoch 27/100	1750/1750	4s	2ms/step	- accuracy: 0.1221	- loss: 0.0907	- val_accuracy: 0.1183	- val_loss: 0.0907
Epoch 28/100	1750/1750	4s	2ms/step	- accuracy: 0.1205	- loss: 0.0907	- val_accuracy: 0.1173	- val_loss: 0.0907
Epoch 29/100	1750/1750	3s	2ms/step	- accuracy: 0.1208	- loss: 0.0907	- val_accuracy: 0.1170	- val_loss: 0.0907
Epoch 30/100	1750/1750	3s	2ms/step	- accuracy: 0.1210	- loss: 0.0906	- val_accuracy: 0.1169	- val_loss: 0.0906
Epoch 31/100	1750/1750	5s	2ms/step	- accuracy: 0.1190	- loss: 0.0906	- val_accuracy: 0.1206	- val_loss: 0.0906
Epoch 32/100	1750/1750	3s	2ms/step	- accuracy: 0.1248	- loss: 0.0906	- val_accuracy: 0.1234	- val_loss: 0.0906
Epoch 33/100	1750/1750	6s	2ms/step	- accuracy: 0.1276	- loss: 0.0906	- val_accuracy: 0.1319	- val_loss: 0.0906
Epoch 34/100	1750/1750	3s	2ms/step	- accuracy: 0.1381	- loss: 0.0905	- val_accuracy: 0.1434	- val_loss: 0.0905
Epoch 35/100	1750/1750	4s	2ms/step	- accuracy: 0.1455	- loss: 0.0905	- val_accuracy: 0.1519	- val_loss: 0.0905
Epoch 36/100	1750/1750	3s	2ms/step	- accuracy: 0.1531	- loss: 0.0905	- val_accuracy: 0.1599	- val_loss: 0.0905
Epoch 37/100	1750/1750	4s	2ms/step	- accuracy: 0.1586	- loss: 0.0905	- val_accuracy: 0.1631	- val_loss: 0.0905
Epoch 38/100	1750/1750	4s	2ms/step	- accuracy: 0.1656	- loss: 0.0905	- val_accuracy: 0.1669	- val_loss: 0.0905
Epoch 39/100	1750/1750	5s	2ms/step	- accuracy: 0.1707	- loss: 0.0904	- val_accuracy: 0.1703	- val_loss: 0.0904
Epoch 40/100	1750/1750	4s	2ms/step	- accuracy: 0.1722	- loss: 0.0904	- val_accuracy: 0.1711	- val_loss: 0.0904
Epoch 41/100	1750/1750	4s	2ms/step	- accuracy: 0.1732	- loss: 0.0904	- val_accuracy: 0.1731	- val_loss: 0.0904
Epoch 42/100	1750/1750	3s	2ms/step	- accuracy: 0.1757	- loss: 0.0904	- val_accuracy: 0.1730	- val_loss: 0.0904
Epoch 43/100	1750/1750	6s	2ms/step	- accuracy: 0.1745	- loss: 0.0903	- val_accuracy: 0.1710	- val_loss: 0.0904
Epoch 44/100	1750/1750	4s	2ms/step	- accuracy: 0.1729	- loss: 0.0903	- val_accuracy: 0.1700	- val_loss: 0.0903
Epoch 45/100	1750/1750	5s	2ms/step	- accuracy: 0.1720	- loss: 0.0903	- val_accuracy: 0.1699	- val_loss: 0.0903
Epoch 46/100	1750/1750	6s	2ms/step	- accuracy: 0.1741	- loss: 0.0903	- val_accuracy: 0.1683	- val_loss: 0.0903

Epoch 47/100	<div><div></div></div>	4s	2ms/step	-	accuracy: 0.1725	-	loss: 0.0903	-	val_accuracy: 0.1666	-	val_loss: 0.0903
Epoch 48/100	<div><div></div></div>	3s	2ms/step	-	accuracy: 0.1736	-	loss: 0.0903	-	val_accuracy: 0.1631	-	val_loss: 0.0903
Epoch 49/100	<div><div></div></div>	6s	2ms/step	-	accuracy: 0.1658	-	loss: 0.0903	-	val_accuracy: 0.1614	-	val_loss: 0.0903
Epoch 50/100	<div><div></div></div>	3s	2ms/step	-	accuracy: 0.1678	-	loss: 0.0902	-	val_accuracy: 0.1593	-	val_loss: 0.0902
Epoch 51/100	<div><div></div></div>	5s	2ms/step	-	accuracy: 0.1649	-	loss: 0.0902	-	val_accuracy: 0.1574	-	val_loss: 0.0902
Epoch 52/100	<div><div></div></div>	6s	2ms/step	-	accuracy: 0.1649	-	loss: 0.0902	-	val_accuracy: 0.1560	-	val_loss: 0.0902
Epoch 53/100	<div><div></div></div>	5s	2ms/step	-	accuracy: 0.1642	-	loss: 0.0902	-	val_accuracy: 0.1549	-	val_loss: 0.0902
Epoch 54/100	<div><div></div></div>	6s	2ms/step	-	accuracy: 0.1584	-	loss: 0.0902	-	val_accuracy: 0.1537	-	val_loss: 0.0902
Epoch 55/100	<div><div></div></div>	4s	2ms/step	-	accuracy: 0.1581	-	loss: 0.0901	-	val_accuracy: 0.1530	-	val_loss: 0.0902
Epoch 56/100	<div><div></div></div>	5s	2ms/step	-	accuracy: 0.1586	-	loss: 0.0901	-	val_accuracy: 0.1524	-	val_loss: 0.0902
Epoch 57/100	<div><div></div></div>	6s	2ms/step	-	accuracy: 0.1563	-	loss: 0.0901	-	val_accuracy: 0.1511	-	val_loss: 0.0901
Epoch 58/100	<div><div></div></div>	4s	2ms/step	-	accuracy: 0.1547	-	loss: 0.0901	-	val_accuracy: 0.1490	-	val_loss: 0.0901
Epoch 59/100	<div><div></div></div>	3s	2ms/step	-	accuracy: 0.1541	-	loss: 0.0901	-	val_accuracy: 0.1480	-	val_loss: 0.0901
Epoch 60/100	<div><div></div></div>	6s	2ms/step	-	accuracy: 0.1532	-	loss: 0.0901	-	val_accuracy: 0.1471	-	val_loss: 0.0901
Epoch 61/100	<div><div></div></div>	4s	2ms/step	-	accuracy: 0.1519	-	loss: 0.0901	-	val_accuracy: 0.1463	-	val_loss: 0.0901
Epoch 62/100	<div><div></div></div>	6s	2ms/step	-	accuracy: 0.1496	-	loss: 0.0901	-	val_accuracy: 0.1450	-	val_loss: 0.0901
Epoch 63/100	<div><div></div></div>	5s	2ms/step	-	accuracy: 0.1506	-	loss: 0.0901	-	val_accuracy: 0.1443	-	val_loss: 0.0901
Epoch 64/100	<div><div></div></div>	3s	2ms/step	-	accuracy: 0.1502	-	loss: 0.0900	-	val_accuracy: 0.1437	-	val_loss: 0.0901
Epoch 65/100	<div><div></div></div>	3s	2ms/step	-	accuracy: 0.1504	-	loss: 0.0901	-	val_accuracy: 0.1433	-	val_loss: 0.0901
Epoch 66/100	<div><div></div></div>	6s	2ms/step	-	accuracy: 0.1466	-	loss: 0.0901	-	val_accuracy: 0.1427	-	val_loss: 0.0900
Epoch 67/100	<div><div></div></div>	5s	2ms/step	-	accuracy: 0.1496	-	loss: 0.0900	-	val_accuracy: 0.1426	-	val_loss: 0.0900
Epoch 68/100	<div><div></div></div>	6s	2ms/step	-	accuracy: 0.1482	-	loss: 0.0900	-	val_accuracy: 0.1423	-	val_loss: 0.0900
Epoch 69/100	<div><div></div></div>	4s	2ms/step	-	accuracy: 0.1465	-	loss: 0.0900	-	val_accuracy: 0.1419	-	val_loss: 0.0900

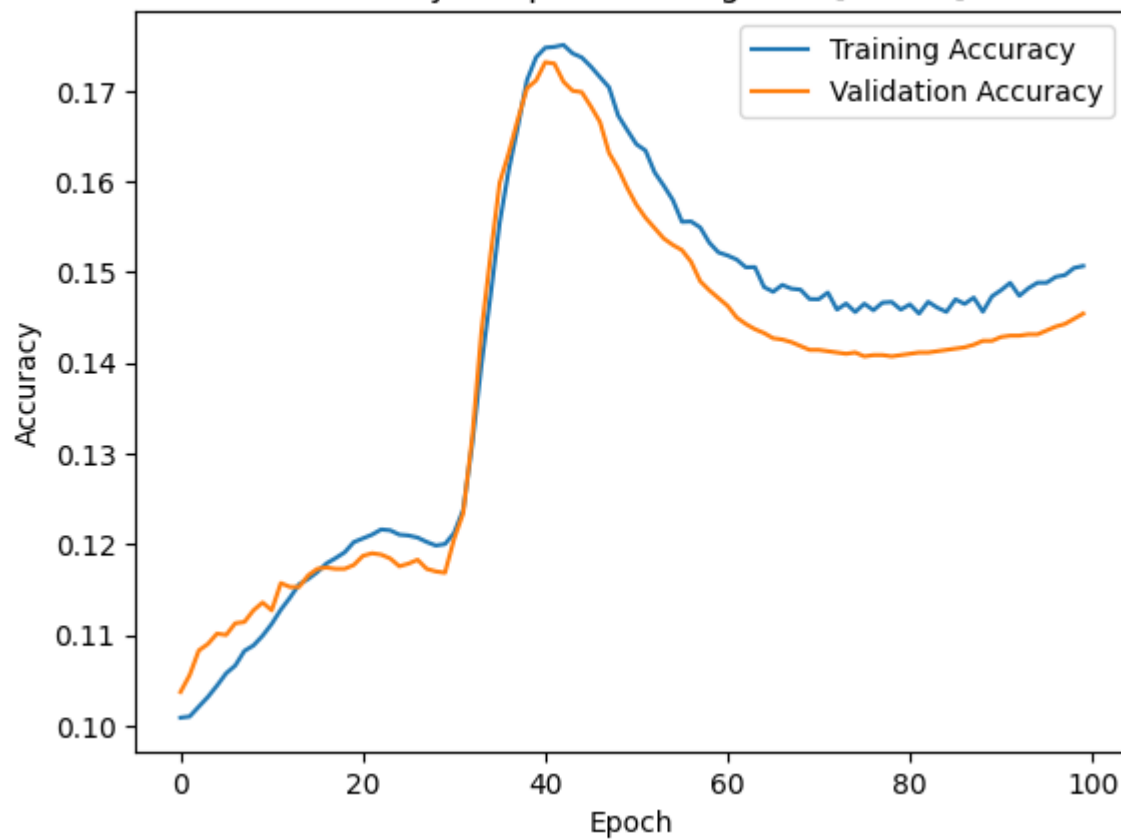
Epoch 70/100	1750/1750	5s	2ms/step	- accuracy: 0.1482	- loss: 0.0900	- val_accuracy: 0.1414	- val_loss: 0.0900
Epoch 71/100	1750/1750	6s	2ms/step	- accuracy: 0.1502	- loss: 0.0900	- val_accuracy: 0.1414	- val_loss: 0.0900
Epoch 72/100	1750/1750	5s	2ms/step	- accuracy: 0.1443	- loss: 0.0900	- val_accuracy: 0.1413	- val_loss: 0.0900
Epoch 73/100	1750/1750	6s	2ms/step	- accuracy: 0.1460	- loss: 0.0900	- val_accuracy: 0.1411	- val_loss: 0.0900
Epoch 74/100	1750/1750	3s	2ms/step	- accuracy: 0.1463	- loss: 0.0900	- val_accuracy: 0.1410	- val_loss: 0.0900
Epoch 75/100	1750/1750	5s	2ms/step	- accuracy: 0.1469	- loss: 0.0899	- val_accuracy: 0.1411	- val_loss: 0.0900
Epoch 76/100	1750/1750	6s	2ms/step	- accuracy: 0.1477	- loss: 0.0899	- val_accuracy: 0.1407	- val_loss: 0.0900
Epoch 77/100	1750/1750	4s	2ms/step	- accuracy: 0.1475	- loss: 0.0899	- val_accuracy: 0.1409	- val_loss: 0.0900
Epoch 78/100	1750/1750	5s	2ms/step	- accuracy: 0.1462	- loss: 0.0899	- val_accuracy: 0.1409	- val_loss: 0.0899
Epoch 79/100	1750/1750	4s	2ms/step	- accuracy: 0.1428	- loss: 0.0899	- val_accuracy: 0.1407	- val_loss: 0.0899
Epoch 80/100	1750/1750	4s	2ms/step	- accuracy: 0.1474	- loss: 0.0899	- val_accuracy: 0.1409	- val_loss: 0.0899
Epoch 81/100	1750/1750	5s	2ms/step	- accuracy: 0.1448	- loss: 0.0899	- val_accuracy: 0.1410	- val_loss: 0.0899
Epoch 82/100	1750/1750	4s	2ms/step	- accuracy: 0.1472	- loss: 0.0899	- val_accuracy: 0.1411	- val_loss: 0.0899
Epoch 83/100	1750/1750	3s	2ms/step	- accuracy: 0.1471	- loss: 0.0899	- val_accuracy: 0.1411	- val_loss: 0.0899
Epoch 84/100	1750/1750	3s	2ms/step	- accuracy: 0.1472	- loss: 0.0899	- val_accuracy: 0.1413	- val_loss: 0.0899
Epoch 85/100	1750/1750	3s	2ms/step	- accuracy: 0.1447	- loss: 0.0899	- val_accuracy: 0.1414	- val_loss: 0.0899
Epoch 86/100	1750/1750	3s	2ms/step	- accuracy: 0.1454	- loss: 0.0899	- val_accuracy: 0.1416	- val_loss: 0.0899
Epoch 87/100	1750/1750	5s	2ms/step	- accuracy: 0.1473	- loss: 0.0899	- val_accuracy: 0.1417	- val_loss: 0.0899
Epoch 88/100	1750/1750	3s	2ms/step	- accuracy: 0.1474	- loss: 0.0899	- val_accuracy: 0.1420	- val_loss: 0.0899
Epoch 89/100	1750/1750	3s	2ms/step	- accuracy: 0.1456	- loss: 0.0899	- val_accuracy: 0.1424	- val_loss: 0.0899
Epoch 90/100	1750/1750	5s	2ms/step	- accuracy: 0.1450	- loss: 0.0899	- val_accuracy: 0.1424	- val_loss: 0.0899
Epoch 91/100	1750/1750	3s	2ms/step	- accuracy: 0.1484	- loss: 0.0899	- val_accuracy: 0.1429	- val_loss: 0.0899
Epoch 92/100	1750/1750	3s	2ms/step	- accuracy: 0.1486	- loss: 0.0898	- val_accuracy: 0.1430	- val_loss: 0.0899

Epoch 93/100 **1750/1750** ————— 3s 2ms/step - accuracy: 0.1480 - loss: 0.0898 - val_accuracy: 0.1430 - val_loss: 0.0899
Epoch 94/100 **1750/1750** ————— 4s 2ms/step - accuracy: 0.1499 - loss: 0.0898 - val_accuracy: 0.1431 - val_loss: 0.0899
Epoch 95/100 **1750/1750** ————— 5s 2ms/step - accuracy: 0.1469 - loss: 0.0898 - val_accuracy: 0.1431 - val_loss: 0.0898
Epoch 96/100 **1750/1750** ————— 4s 2ms/step - accuracy: 0.1476 - loss: 0.0898 - val_accuracy: 0.1436 - val_loss: 0.0898
Epoch 97/100 **1750/1750** ————— 4s 2ms/step - accuracy: 0.1464 - loss: 0.0898 - val_accuracy: 0.1440 - val_loss: 0.0898
Epoch 98/100 **1750/1750** ————— 3s 2ms/step - accuracy: 0.1494 - loss: 0.0898 - val_accuracy: 0.1443 - val_loss: 0.0898
Epoch 99/100 **1750/1750** ————— 6s 2ms/step - accuracy: 0.1501 - loss: 0.0898 - val_accuracy: 0.1449 - val_loss: 0.0898
Epoch 100/100 **1750/1750** ————— 4s 2ms/step - accuracy: 0.1498 - loss: 0.0898 - val_accuracy: 0.1454 - val_loss: 0.0898
219/219 ————— 0s 1ms/step - accuracy: 0.1576 - loss: 0.0899
Test loss=0.08983103930950165 Test accuracy = 0.15728572010993958
Time required to train the model is 436.4306344985962 seconds

Loss vs epochs for sigmoid [16, 32]



Accuracy vs epochs for sigmoid [16, 32]



In [46]: `## Model 3 --> SGD, lr =0.001, Mean Squared Error loss, hidden_neurons=[16,32,64]`

```
hidden_neurons = [16,32,64]
activation_function = 'sigmoid'

model, history, duration = train_model(
    activation_func=activation_function,
    hidden_neurons=hidden_neurons,
    optimizer=SGD,
    loss_func=MeanSquaredError,
    epochs=100,
    dropout_rate=None,
    learning_rate=0.001
)

test_loss, test_acc = model.evaluate(X_test,y_test)

print(f"Test loss={test_loss} Test accuracy = {test_acc}")
```

```

print(f"Time required to train the model is {duration} seconds")

result_df.loc[len(result_df.index)] = [
    len(hidden_neurons),
    activation_function,
    str(hidden_neurons),
    'SGD',
    duration,
    test_loss,
    test_acc
]

plot_metrics(history, activation_function, hidden_neurons)

```

Model: "sequential_18"

Layer (type)	Output Shape	Param #
flatten_18 (Flatten)	(None, 1024)	0
dense_61 (Dense)	(None, 16)	16,400
dense_62 (Dense)	(None, 32)	544
dense_63 (Dense)	(None, 64)	2,112
dense_64 (Dense)	(None, 10)	650

Total params: 19,706 (76.98 KB)

Trainable params: 19,706 (76.98 KB)

Non-trainable params: 0 (0.00 B)

Epoch 1/100	1750/1750	4s	2ms/step	- accuracy: 0.0966	- loss: 0.0969	- val_accuracy: 0.0943	- val_loss: 0.0956
Epoch 2/100	1750/1750	4s	2ms/step	- accuracy: 0.0975	- loss: 0.0950	- val_accuracy: 0.0943	- val_loss: 0.0945
Epoch 3/100	1750/1750	3s	2ms/step	- accuracy: 0.0970	- loss: 0.0942	- val_accuracy: 0.0943	- val_loss: 0.0938
Epoch 4/100	1750/1750	6s	2ms/step	- accuracy: 0.0984	- loss: 0.0935	- val_accuracy: 0.0943	- val_loss: 0.0934
Epoch 5/100	1750/1750	3s	2ms/step	- accuracy: 0.0977	- loss: 0.0931	- val_accuracy: 0.0943	- val_loss: 0.0930
Epoch 6/100	1750/1750	5s	2ms/step	- accuracy: 0.0994	- loss: 0.0927	- val_accuracy: 0.0943	- val_loss: 0.0928
Epoch 7/100	1750/1750	3s	2ms/step	- accuracy: 0.0968	- loss: 0.0926	- val_accuracy: 0.0943	- val_loss: 0.0925
Epoch 8/100	1750/1750	4s	2ms/step	- accuracy: 0.0976	- loss: 0.0923	- val_accuracy: 0.0943	- val_loss: 0.0923
Epoch 9/100	1750/1750	3s	2ms/step	- accuracy: 0.0958	- loss: 0.0921	- val_accuracy: 0.0943	- val_loss: 0.0921
Epoch 10/100	1750/1750	3s	2ms/step	- accuracy: 0.0996	- loss: 0.0919	- val_accuracy: 0.0943	- val_loss: 0.0920
Epoch 11/100	1750/1750	6s	2ms/step	- accuracy: 0.0963	- loss: 0.0918	- val_accuracy: 0.0943	- val_loss: 0.0918
Epoch 12/100	1750/1750	4s	2ms/step	- accuracy: 0.0965	- loss: 0.0917	- val_accuracy: 0.0943	- val_loss: 0.0917
Epoch 13/100	1750/1750	3s	2ms/step	- accuracy: 0.0983	- loss: 0.0915	- val_accuracy: 0.0943	- val_loss: 0.0916
Epoch 14/100	1750/1750	3s	2ms/step	- accuracy: 0.0971	- loss: 0.0914	- val_accuracy: 0.0943	- val_loss: 0.0915
Epoch 15/100	1750/1750	6s	2ms/step	- accuracy: 0.0983	- loss: 0.0914	- val_accuracy: 0.0943	- val_loss: 0.0914
Epoch 16/100	1750/1750	4s	2ms/step	- accuracy: 0.0978	- loss: 0.0913	- val_accuracy: 0.0937	- val_loss: 0.0914
Epoch 17/100	1750/1750	5s	2ms/step	- accuracy: 0.0972	- loss: 0.0912	- val_accuracy: 0.0859	- val_loss: 0.0913
Epoch 18/100	1750/1750	6s	2ms/step	- accuracy: 0.0926	- loss: 0.0911	- val_accuracy: 0.0716	- val_loss: 0.0912
Epoch 19/100	1750/1750	3s	2ms/step	- accuracy: 0.0768	- loss: 0.0911	- val_accuracy: 0.0699	- val_loss: 0.0912
Epoch 20/100	1750/1750	5s	2ms/step	- accuracy: 0.0768	- loss: 0.0910	- val_accuracy: 0.0866	- val_loss: 0.0911
Epoch 21/100	1750/1750	4s	3ms/step	- accuracy: 0.0925	- loss: 0.0910	- val_accuracy: 0.1006	- val_loss: 0.0911
Epoch 22/100	1750/1750	3s	2ms/step	- accuracy: 0.1068	- loss: 0.0909	- val_accuracy: 0.1074	- val_loss: 0.0910
Epoch 23/100	1750/1750	3s	2ms/step	- accuracy: 0.1093	- loss: 0.0909	- val_accuracy: 0.1083	- val_loss: 0.0910

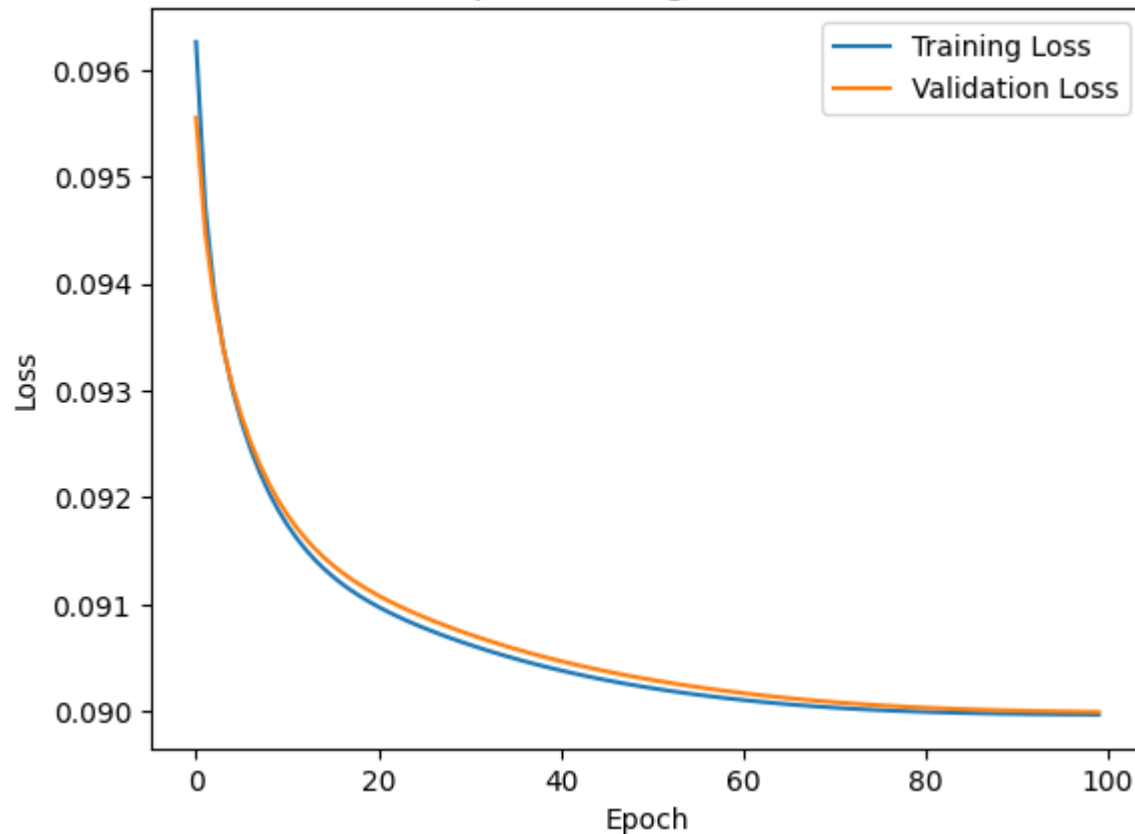
Epoch 24/100	1750/1750	6s	2ms/step	- accuracy: 0.1109	- loss: 0.0908	- val_accuracy: 0.1110	- val_loss: 0.0910
Epoch 25/100	1750/1750	4s	2ms/step	- accuracy: 0.1110	- loss: 0.0908	- val_accuracy: 0.1093	- val_loss: 0.0909
Epoch 26/100	1750/1750	3s	2ms/step	- accuracy: 0.1067	- loss: 0.0908	- val_accuracy: 0.1093	- val_loss: 0.0909
Epoch 27/100	1750/1750	5s	2ms/step	- accuracy: 0.1055	- loss: 0.0908	- val_accuracy: 0.0976	- val_loss: 0.0908
Epoch 28/100	1750/1750	5s	2ms/step	- accuracy: 0.1018	- loss: 0.0908	- val_accuracy: 0.1236	- val_loss: 0.0908
Epoch 29/100	1750/1750	3s	2ms/step	- accuracy: 0.1233	- loss: 0.0907	- val_accuracy: 0.1081	- val_loss: 0.0908
Epoch 30/100	1750/1750	6s	2ms/step	- accuracy: 0.1122	- loss: 0.0907	- val_accuracy: 0.1080	- val_loss: 0.0907
Epoch 31/100	1750/1750	4s	2ms/step	- accuracy: 0.1134	- loss: 0.0906	- val_accuracy: 0.1080	- val_loss: 0.0907
Epoch 32/100	1750/1750	3s	2ms/step	- accuracy: 0.1108	- loss: 0.0906	- val_accuracy: 0.1080	- val_loss: 0.0907
Epoch 33/100	1750/1750	3s	2ms/step	- accuracy: 0.1124	- loss: 0.0906	- val_accuracy: 0.1080	- val_loss: 0.0907
Epoch 34/100	1750/1750	3s	2ms/step	- accuracy: 0.1127	- loss: 0.0905	- val_accuracy: 0.1080	- val_loss: 0.0906
Epoch 35/100	1750/1750	4s	3ms/step	- accuracy: 0.1131	- loss: 0.0905	- val_accuracy: 0.1080	- val_loss: 0.0906
Epoch 36/100	1750/1750	3s	2ms/step	- accuracy: 0.1138	- loss: 0.0905	- val_accuracy: 0.1080	- val_loss: 0.0906
Epoch 37/100	1750/1750	5s	2ms/step	- accuracy: 0.1118	- loss: 0.0905	- val_accuracy: 0.1080	- val_loss: 0.0906
Epoch 38/100	1750/1750	4s	2ms/step	- accuracy: 0.1120	- loss: 0.0905	- val_accuracy: 0.1080	- val_loss: 0.0905
Epoch 39/100	1750/1750	4s	2ms/step	- accuracy: 0.1139	- loss: 0.0904	- val_accuracy: 0.1080	- val_loss: 0.0905
Epoch 40/100	1750/1750	5s	2ms/step	- accuracy: 0.1110	- loss: 0.0904	- val_accuracy: 0.1080	- val_loss: 0.0905
Epoch 41/100	1750/1750	4s	2ms/step	- accuracy: 0.1128	- loss: 0.0904	- val_accuracy: 0.1080	- val_loss: 0.0905
Epoch 42/100	1750/1750	4s	2ms/step	- accuracy: 0.1124	- loss: 0.0904	- val_accuracy: 0.1080	- val_loss: 0.0904
Epoch 43/100	1750/1750	5s	2ms/step	- accuracy: 0.1116	- loss: 0.0903	- val_accuracy: 0.1080	- val_loss: 0.0904
Epoch 44/100	1750/1750	6s	2ms/step	- accuracy: 0.1121	- loss: 0.0903	- val_accuracy: 0.1080	- val_loss: 0.0904
Epoch 45/100	1750/1750	4s	2ms/step	- accuracy: 0.1102	- loss: 0.0903	- val_accuracy: 0.1080	- val_loss: 0.0904
Epoch 46/100	1750/1750	6s	2ms/step	- accuracy: 0.1124	- loss: 0.0903	- val_accuracy: 0.1080	- val_loss: 0.0904

Epoch 47/100	<div><div></div></div>	5s	2ms/step	- accuracy: 0.1131 - loss: 0.0903 - val_accuracy: 0.1080 - val_loss: 0.0904
Epoch 48/100	<div><div></div></div>	3s	2ms/step	- accuracy: 0.1124 - loss: 0.0903 - val_accuracy: 0.1080 - val_loss: 0.0903
Epoch 49/100	<div><div></div></div>	6s	2ms/step	- accuracy: 0.1151 - loss: 0.0902 - val_accuracy: 0.1080 - val_loss: 0.0903
Epoch 50/100	<div><div></div></div>	4s	2ms/step	- accuracy: 0.1123 - loss: 0.0902 - val_accuracy: 0.1080 - val_loss: 0.0903
Epoch 51/100	<div><div></div></div>	5s	2ms/step	- accuracy: 0.1122 - loss: 0.0902 - val_accuracy: 0.1080 - val_loss: 0.0903
Epoch 52/100	<div><div></div></div>	4s	2ms/step	- accuracy: 0.1121 - loss: 0.0902 - val_accuracy: 0.1080 - val_loss: 0.0903
Epoch 53/100	<div><div></div></div>	4s	2ms/step	- accuracy: 0.1141 - loss: 0.0902 - val_accuracy: 0.1080 - val_loss: 0.0903
Epoch 54/100	<div><div></div></div>	5s	2ms/step	- accuracy: 0.1122 - loss: 0.0902 - val_accuracy: 0.1080 - val_loss: 0.0903
Epoch 55/100	<div><div></div></div>	6s	2ms/step	- accuracy: 0.1114 - loss: 0.0902 - val_accuracy: 0.1080 - val_loss: 0.0902
Epoch 56/100	<div><div></div></div>	3s	2ms/step	- accuracy: 0.1150 - loss: 0.0901 - val_accuracy: 0.1080 - val_loss: 0.0902
Epoch 57/100	<div><div></div></div>	5s	2ms/step	- accuracy: 0.1119 - loss: 0.0901 - val_accuracy: 0.1080 - val_loss: 0.0902
Epoch 58/100	<div><div></div></div>	5s	2ms/step	- accuracy: 0.1106 - loss: 0.0901 - val_accuracy: 0.1080 - val_loss: 0.0902
Epoch 59/100	<div><div></div></div>	3s	2ms/step	- accuracy: 0.1115 - loss: 0.0901 - val_accuracy: 0.1080 - val_loss: 0.0902
Epoch 60/100	<div><div></div></div>	3s	2ms/step	- accuracy: 0.1117 - loss: 0.0901 - val_accuracy: 0.1080 - val_loss: 0.0902
Epoch 61/100	<div><div></div></div>	4s	2ms/step	- accuracy: 0.1135 - loss: 0.0901 - val_accuracy: 0.1080 - val_loss: 0.0902
Epoch 62/100	<div><div></div></div>	4s	2ms/step	- accuracy: 0.1120 - loss: 0.0901 - val_accuracy: 0.1080 - val_loss: 0.0902
Epoch 63/100	<div><div></div></div>	3s	2ms/step	- accuracy: 0.1120 - loss: 0.0901 - val_accuracy: 0.1080 - val_loss: 0.0901
Epoch 64/100	<div><div></div></div>	6s	2ms/step	- accuracy: 0.1110 - loss: 0.0901 - val_accuracy: 0.1080 - val_loss: 0.0901
Epoch 65/100	<div><div></div></div>	4s	2ms/step	- accuracy: 0.1106 - loss: 0.0901 - val_accuracy: 0.1080 - val_loss: 0.0901
Epoch 66/100	<div><div></div></div>	5s	2ms/step	- accuracy: 0.1128 - loss: 0.0901 - val_accuracy: 0.1080 - val_loss: 0.0901
Epoch 67/100	<div><div></div></div>	4s	2ms/step	- accuracy: 0.1154 - loss: 0.0900 - val_accuracy: 0.1080 - val_loss: 0.0901
Epoch 68/100	<div><div></div></div>	4s	2ms/step	- accuracy: 0.1122 - loss: 0.0900 - val_accuracy: 0.1080 - val_loss: 0.0901
Epoch 69/100	<div><div></div></div>	3s	2ms/step	- accuracy: 0.1148 - loss: 0.0900 - val_accuracy: 0.1080 - val_loss: 0.0901

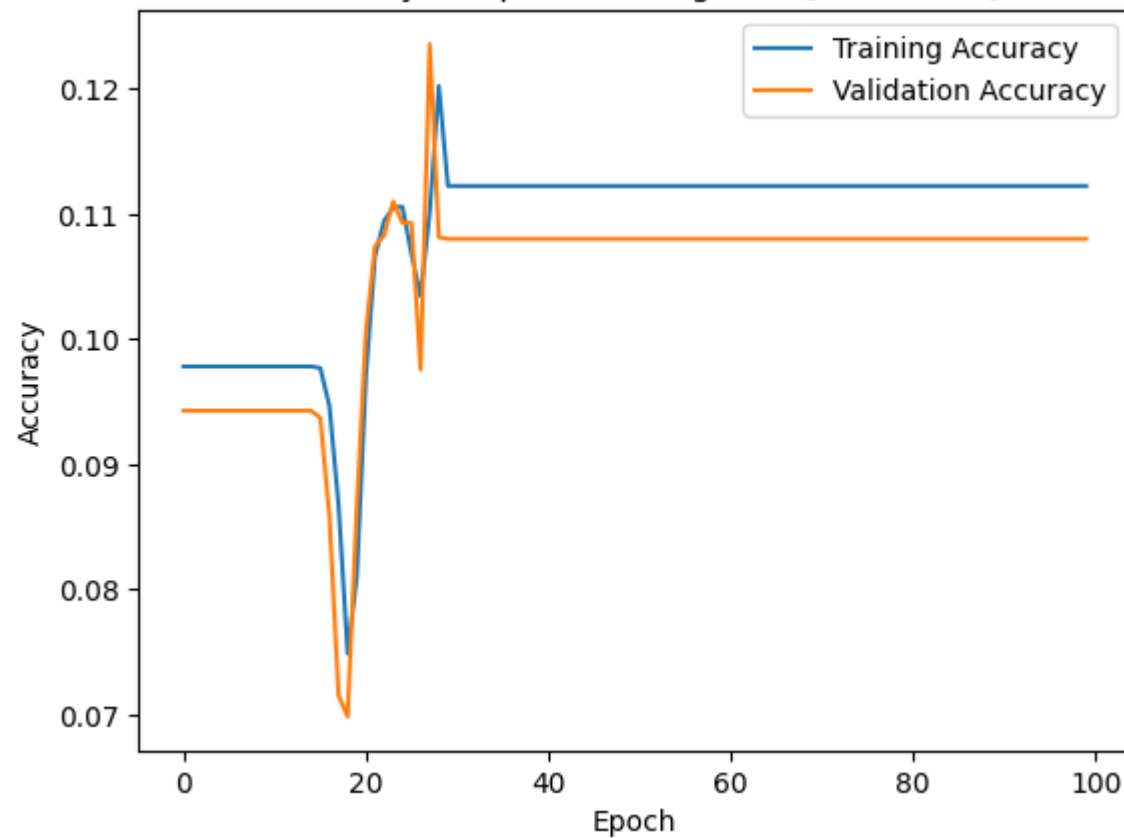
Epoch 70/100	<div><div></div></div>	3s	2ms/step	- accuracy: 0.1139 - loss: 0.0900 - val_accuracy: 0.1080 - val_loss: 0.0901
Epoch 71/100	<div><div></div></div>	5s	2ms/step	- accuracy: 0.1120 - loss: 0.0900 - val_accuracy: 0.1080 - val_loss: 0.0901
Epoch 72/100	<div><div></div></div>	5s	2ms/step	- accuracy: 0.1125 - loss: 0.0900 - val_accuracy: 0.1080 - val_loss: 0.0901
Epoch 73/100	<div><div></div></div>	4s	2ms/step	- accuracy: 0.1117 - loss: 0.0900 - val_accuracy: 0.1080 - val_loss: 0.0901
Epoch 74/100	<div><div></div></div>	4s	2ms/step	- accuracy: 0.1126 - loss: 0.0900 - val_accuracy: 0.1080 - val_loss: 0.0901
Epoch 75/100	<div><div></div></div>	5s	2ms/step	- accuracy: 0.1108 - loss: 0.0900 - val_accuracy: 0.1080 - val_loss: 0.0901
Epoch 76/100	<div><div></div></div>	4s	2ms/step	- accuracy: 0.1087 - loss: 0.0900 - val_accuracy: 0.1080 - val_loss: 0.0901
Epoch 77/100	<div><div></div></div>	4s	2ms/step	- accuracy: 0.1118 - loss: 0.0900 - val_accuracy: 0.1080 - val_loss: 0.0900
Epoch 78/100	<div><div></div></div>	5s	2ms/step	- accuracy: 0.1142 - loss: 0.0900 - val_accuracy: 0.1080 - val_loss: 0.0900
Epoch 79/100	<div><div></div></div>	6s	2ms/step	- accuracy: 0.1159 - loss: 0.0900 - val_accuracy: 0.1080 - val_loss: 0.0900
Epoch 80/100	<div><div></div></div>	3s	2ms/step	- accuracy: 0.1110 - loss: 0.0900 - val_accuracy: 0.1080 - val_loss: 0.0900
Epoch 81/100	<div><div></div></div>	5s	2ms/step	- accuracy: 0.1115 - loss: 0.0900 - val_accuracy: 0.1080 - val_loss: 0.0900
Epoch 82/100	<div><div></div></div>	5s	2ms/step	- accuracy: 0.1125 - loss: 0.0900 - val_accuracy: 0.1080 - val_loss: 0.0900
Epoch 83/100	<div><div></div></div>	5s	2ms/step	- accuracy: 0.1129 - loss: 0.0900 - val_accuracy: 0.1080 - val_loss: 0.0900
Epoch 84/100	<div><div></div></div>	3s	2ms/step	- accuracy: 0.1104 - loss: 0.0900 - val_accuracy: 0.1080 - val_loss: 0.0900
Epoch 85/100	<div><div></div></div>	5s	2ms/step	- accuracy: 0.1142 - loss: 0.0900 - val_accuracy: 0.1080 - val_loss: 0.0900
Epoch 86/100	<div><div></div></div>	5s	2ms/step	- accuracy: 0.1105 - loss: 0.0900 - val_accuracy: 0.1080 - val_loss: 0.0900
Epoch 87/100	<div><div></div></div>	6s	2ms/step	- accuracy: 0.1118 - loss: 0.0900 - val_accuracy: 0.1080 - val_loss: 0.0900
Epoch 88/100	<div><div></div></div>	3s	2ms/step	- accuracy: 0.1121 - loss: 0.0900 - val_accuracy: 0.1080 - val_loss: 0.0900
Epoch 89/100	<div><div></div></div>	5s	2ms/step	- accuracy: 0.1131 - loss: 0.0900 - val_accuracy: 0.1080 - val_loss: 0.0900
Epoch 90/100	<div><div></div></div>	4s	2ms/step	- accuracy: 0.1123 - loss: 0.0900 - val_accuracy: 0.1080 - val_loss: 0.0900
Epoch 91/100	<div><div></div></div>	4s	2ms/step	- accuracy: 0.1127 - loss: 0.0900 - val_accuracy: 0.1080 - val_loss: 0.0900
Epoch 92/100	<div><div></div></div>	5s	2ms/step	- accuracy: 0.1136 - loss: 0.0900 - val_accuracy: 0.1080 - val_loss: 0.0900

Epoch 93/100 **1750/1750** 4s 2ms/step - accuracy: 0.1110 - loss: 0.0900 - val_accuracy: 0.1080 - val_loss: 0.0900
Epoch 94/100 **1750/1750** 3s 2ms/step - accuracy: 0.1131 - loss: 0.0900 - val_accuracy: 0.1080 - val_loss: 0.0900
Epoch 95/100 **1750/1750** 3s 2ms/step - accuracy: 0.1130 - loss: 0.0900 - val_accuracy: 0.1080 - val_loss: 0.0900
Epoch 96/100 **1750/1750** 3s 2ms/step - accuracy: 0.1139 - loss: 0.0900 - val_accuracy: 0.1080 - val_loss: 0.0900
Epoch 97/100 **1750/1750** 6s 2ms/step - accuracy: 0.1094 - loss: 0.0900 - val_accuracy: 0.1080 - val_loss: 0.0900
Epoch 98/100 **1750/1750** 3s 2ms/step - accuracy: 0.1114 - loss: 0.0900 - val_accuracy: 0.1080 - val_loss: 0.0900
Epoch 99/100 **1750/1750** 5s 2ms/step - accuracy: 0.1119 - loss: 0.0900 - val_accuracy: 0.1080 - val_loss: 0.0900
Epoch 100/100 **1750/1750** 5s 3ms/step - accuracy: 0.1128 - loss: 0.0900 - val_accuracy: 0.1080 - val_loss: 0.0900
219/219 0s 1ms/step - accuracy: 0.1219 - loss: 0.0900
Test loss=0.08996734023094177 Test accuracy = 0.11942857503890991
Time required to train the model is 428.4792711734772 seconds

Loss vs epochs for sigmoid [16, 32, 64]



Accuracy vs epochs for sigmoid [16, 32, 64]



In [47]: `## Model 4 --> Adam, lr =0.001, CCE loss, hidden_neurons=[16]`

```
hidden_neurons = [16]
activation_function = 'sigmoid'

model, history, duration = train_model(
    activation_func=activation_function,
    hidden_neurons=hidden_neurons,
    optimizer=Adam,
    loss_func=CategoricalCrossentropy,
    epochs=100,
    dropout_rate=None,
    learning_rate=0.001
)

test_loss, test_acc = model.evaluate(X_test, y_test)

print(f"Test loss={test_loss} Test accuracy = {test_acc}")
```

```
print(f"Time required to train the model is {duration} seconds")

result_df.loc[len(result_df.index)] = [
    len(hidden_neurons),
    activation_function,
    str(hidden_neurons),
    'Adam',
    duration,
    test_loss,
    test_acc
]

plot_metrics(history, activation_function, hidden_neurons)
```

















Model: "sequential_19"

Layer (type)	Output Shape	Param #
flatten_19 (Flatten)	(None, 1024)	0
dense_65 (Dense)	(None, 16)	16,400
dense_66 (Dense)	(None, 10)	170

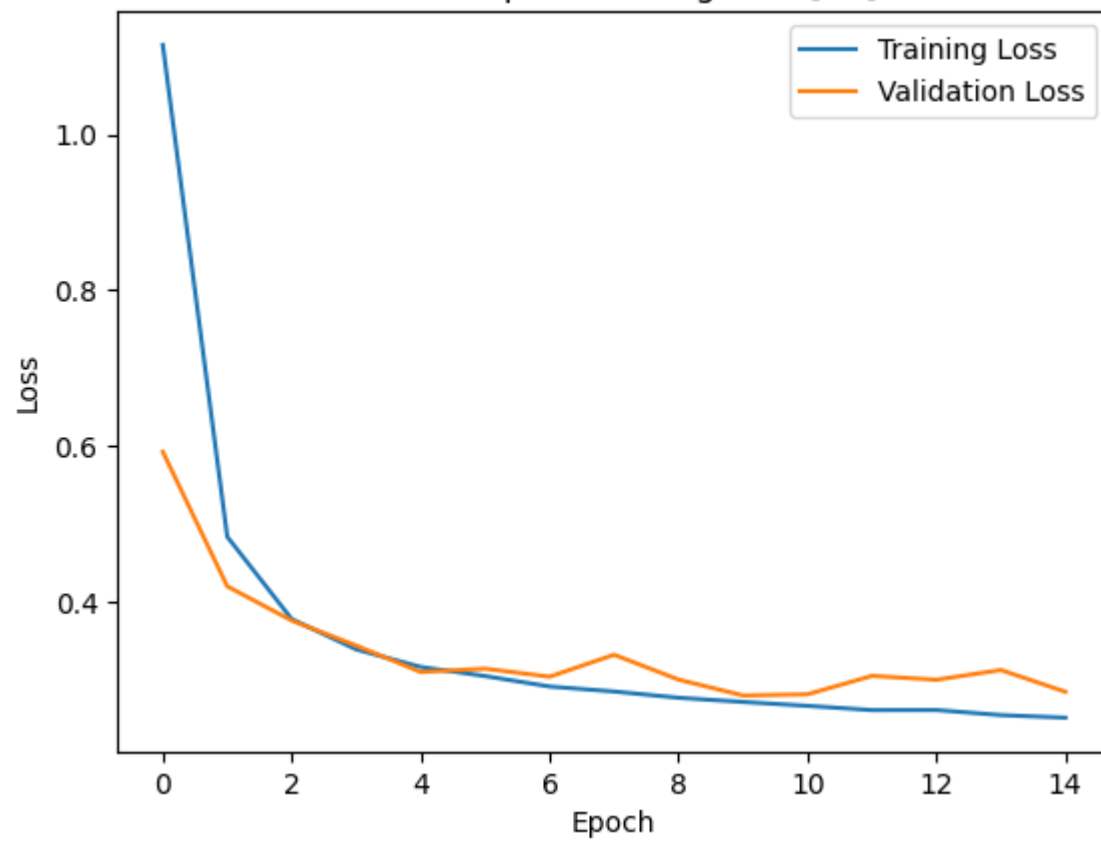
Total params: 16,570 (64.73 KB)

Trainable params: 16,570 (64.73 KB)

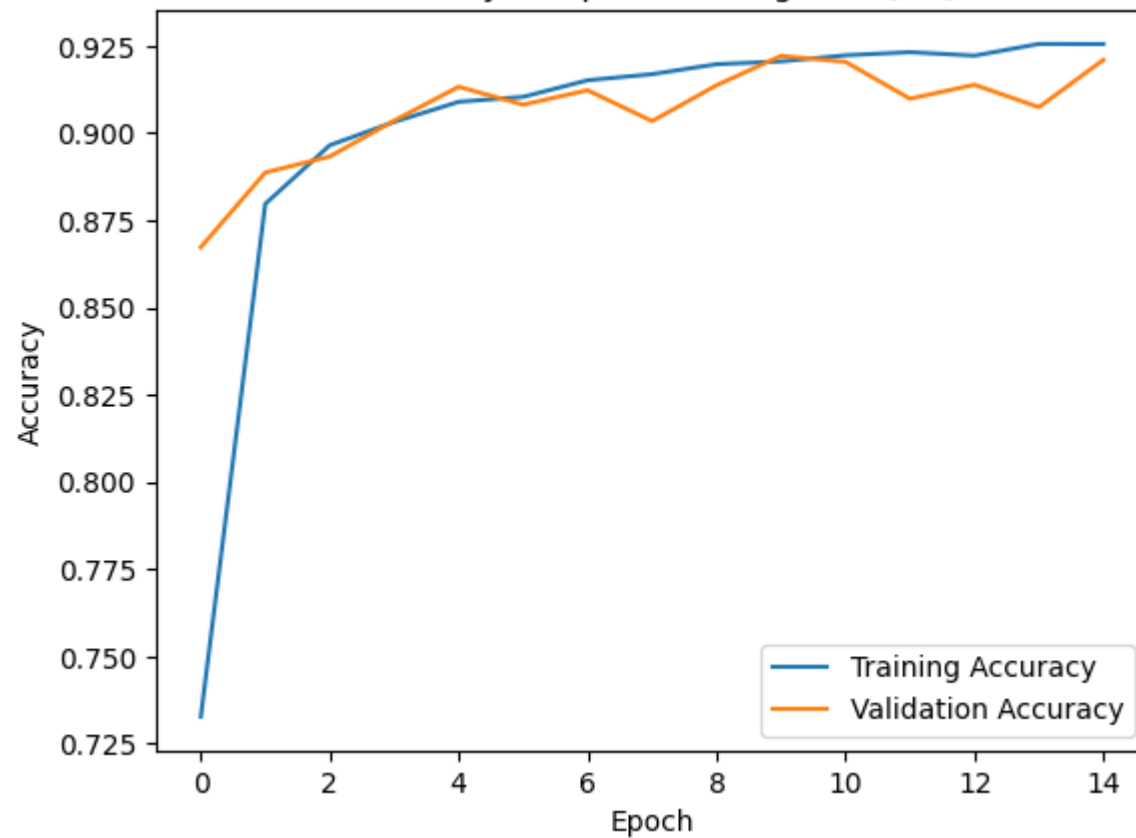
Non-trainable params: 0 (0.00 B)

Epoch 1/100
1750/1750  4s 2ms/step - accuracy: 0.5717 - loss: 1.5646 - val_accuracy: 0.8673 - val_loss: 0.5924
Epoch 2/100
1750/1750  5s 2ms/step - accuracy: 0.8753 - loss: 0.5274 - val_accuracy: 0.8887 - val_loss: 0.4196
Epoch 3/100
1750/1750  4s 2ms/step - accuracy: 0.8940 - loss: 0.3933 - val_accuracy: 0.8933 - val_loss: 0.3756
Epoch 4/100
1750/1750  4s 2ms/step - accuracy: 0.9047 - loss: 0.3397 - val_accuracy: 0.9036 - val_loss: 0.3434
Epoch 5/100
1750/1750  6s 2ms/step - accuracy: 0.9084 - loss: 0.3175 - val_accuracy: 0.9133 - val_loss: 0.3095
Epoch 6/100
1750/1750  4s 2ms/step - accuracy: 0.9134 - loss: 0.2966 - val_accuracy: 0.9081 - val_loss: 0.3137
Epoch 7/100
1750/1750  4s 2ms/step - accuracy: 0.9127 - loss: 0.2955 - val_accuracy: 0.9123 - val_loss: 0.3034
Epoch 8/100
1750/1750  3s 2ms/step - accuracy: 0.9158 - loss: 0.2891 - val_accuracy: 0.9034 - val_loss: 0.3316
Epoch 9/100
1750/1750  5s 2ms/step - accuracy: 0.9200 - loss: 0.2760 - val_accuracy: 0.9137 - val_loss: 0.2997
Epoch 10/100
1750/1750  5s 2ms/step - accuracy: 0.9220 - loss: 0.2679 - val_accuracy: 0.9221 - val_loss: 0.2793
Epoch 11/100
1750/1750  3s 2ms/step - accuracy: 0.9226 - loss: 0.2666 - val_accuracy: 0.9204 - val_loss: 0.2808
Epoch 12/100
1750/1750  5s 2ms/step - accuracy: 0.9230 - loss: 0.2583 - val_accuracy: 0.9099 - val_loss: 0.3045
Epoch 13/100
1750/1750  3s 2ms/step - accuracy: 0.9219 - loss: 0.2611 - val_accuracy: 0.9139 - val_loss: 0.2995
Epoch 14/100
1750/1750  3s 2ms/step - accuracy: 0.9259 - loss: 0.2533 - val_accuracy: 0.9074 - val_loss: 0.3121
Epoch 15/100
1750/1750  4s 2ms/step - accuracy: 0.9238 - loss: 0.2526 - val_accuracy: 0.9210 - val_loss: 0.2840
219/219  0s 1ms/step - accuracy: 0.9234 - loss: 0.2686
Test loss=0.28263649344444275 Test accuracy = 0.9197142720222473
Time required to train the model is 63.880293130874634 seconds

Loss vs epochs for sigmoid [16]



Accuracy vs epochs for sigmoid [16]



In [48]: `## Model 5 --> Adam, lr =0.001, CCE loss, hidden_neurons=[16,32]`

```
hidden_neurons = [16,32]
activation_function = 'sigmoid'

model, history,duration = train_model(
    activation_func=activation_function,
    hidden_neurons=hidden_neurons,
    optimizer=Adam,
    loss_func=CategoricalCrossentropy,
    epochs=100,
    dropout_rate=None,
    learning_rate=0.001
)

test_loss, test_acc = model.evaluate(X_test,y_test)

print(f"Test loss={test_loss} Test accuracy = {test_acc}")
```



```

print(f"Time taken to train the model is {duration} seconds")

result_df.loc[len(result_df.index)] = [
    len(hidden_neurons),
    activation_function,
    str(hidden_neurons),
    'Adam',
    duration,
    test_loss,
    test_acc
]

plot_metrics(history, activation_function, hidden_neurons)

```



















Model: "sequential_20"

Layer (type)	Output Shape	Param #
flatten_20 (Flatten)	(None, 1024)	0
dense_67 (Dense)	(None, 16)	16,400
dense_68 (Dense)	(None, 32)	544
dense_69 (Dense)	(None, 10)	330

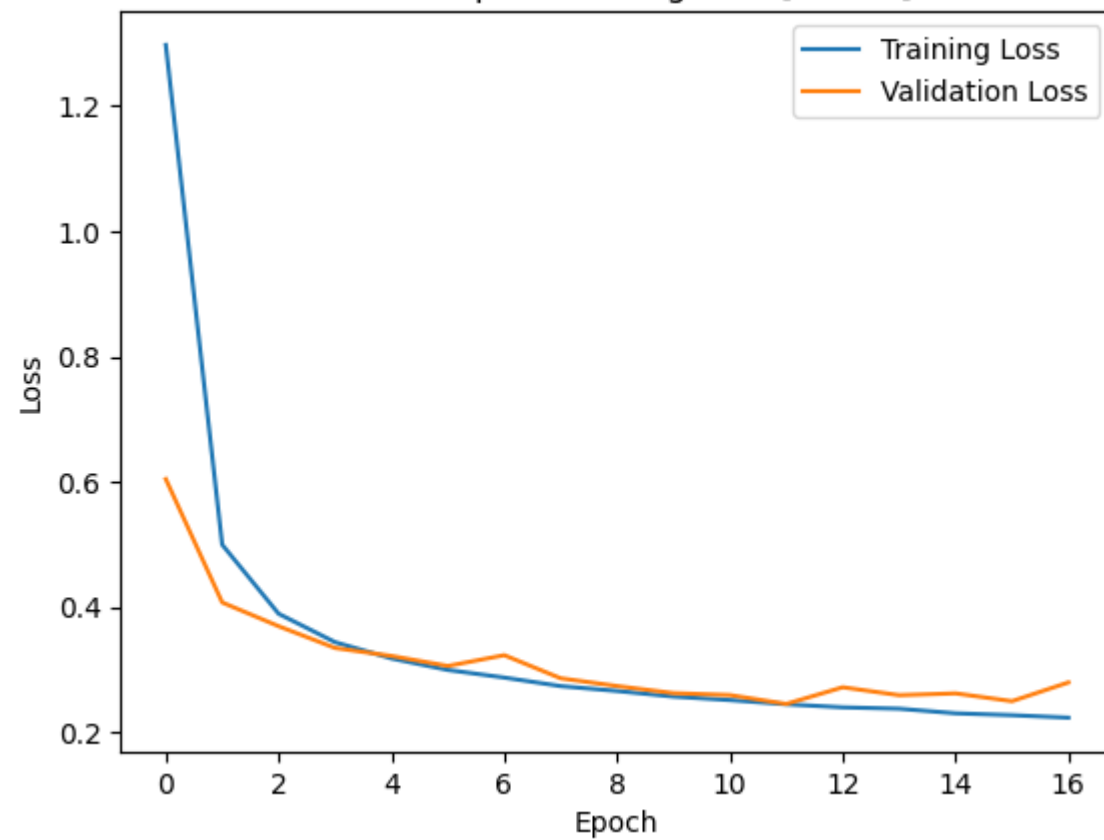
Total params: 17,274 (67.48 KB)

Trainable params: 17,274 (67.48 KB)

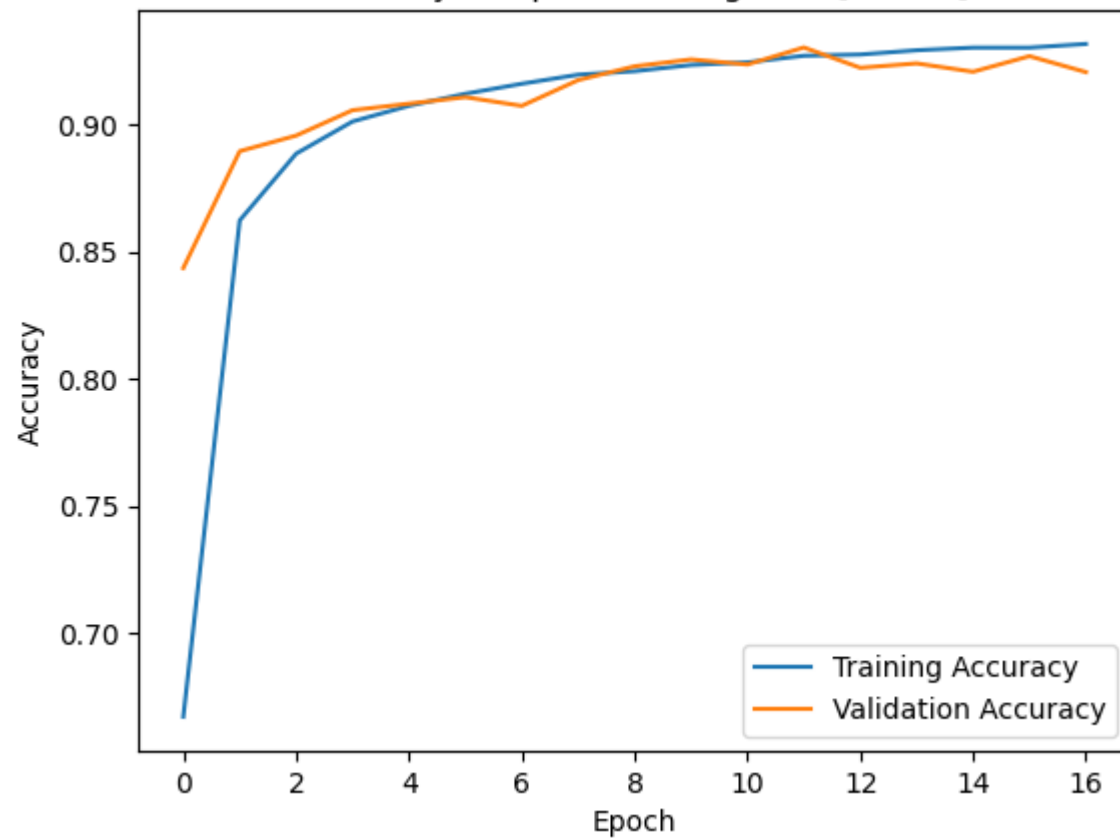
Non-trainable params: 0 (0.00 B)

Epoch 1/100
1750/1750  4s 2ms/step - accuracy: 0.4717 - loss: 1.8169 - val_accuracy: 0.8436 - val_loss: 0.6044
Epoch 2/100
1750/1750  5s 2ms/step - accuracy: 0.8507 - loss: 0.5469 - val_accuracy: 0.8896 - val_loss: 0.4074
Epoch 3/100
1750/1750  5s 2ms/step - accuracy: 0.8866 - loss: 0.3982 - val_accuracy: 0.8957 - val_loss: 0.3695
Epoch 4/100
1750/1750  3s 2ms/step - accuracy: 0.8988 - loss: 0.3551 - val_accuracy: 0.9057 - val_loss: 0.3348
Epoch 5/100
1750/1750  3s 2ms/step - accuracy: 0.9074 - loss: 0.3195 - val_accuracy: 0.9083 - val_loss: 0.3220
Epoch 6/100
1750/1750  5s 3ms/step - accuracy: 0.9120 - loss: 0.2994 - val_accuracy: 0.9109 - val_loss: 0.3060
Epoch 7/100
1750/1750  4s 2ms/step - accuracy: 0.9147 - loss: 0.2884 - val_accuracy: 0.9074 - val_loss: 0.3233
Epoch 8/100
1750/1750  3s 2ms/step - accuracy: 0.9195 - loss: 0.2772 - val_accuracy: 0.9176 - val_loss: 0.2864
Epoch 9/100
1750/1750  3s 2ms/step - accuracy: 0.9199 - loss: 0.2670 - val_accuracy: 0.9230 - val_loss: 0.2740
Epoch 10/100
1750/1750  4s 3ms/step - accuracy: 0.9239 - loss: 0.2550 - val_accuracy: 0.9257 - val_loss: 0.2627
Epoch 11/100
1750/1750  4s 2ms/step - accuracy: 0.9252 - loss: 0.2477 - val_accuracy: 0.9237 - val_loss: 0.2596
Epoch 12/100
1750/1750  3s 2ms/step - accuracy: 0.9281 - loss: 0.2439 - val_accuracy: 0.9304 - val_loss: 0.2455
Epoch 13/100
1750/1750  4s 2ms/step - accuracy: 0.9280 - loss: 0.2379 - val_accuracy: 0.9224 - val_loss: 0.2721
Epoch 14/100
1750/1750  5s 2ms/step - accuracy: 0.9296 - loss: 0.2420 - val_accuracy: 0.9241 - val_loss: 0.2594
Epoch 15/100
1750/1750  5s 2ms/step - accuracy: 0.9300 - loss: 0.2330 - val_accuracy: 0.9209 - val_loss: 0.2623
Epoch 16/100
1750/1750  6s 2ms/step - accuracy: 0.9308 - loss: 0.2288 - val_accuracy: 0.9270 - val_loss: 0.2498
Epoch 17/100
1750/1750  4s 2ms/step - accuracy: 0.9336 - loss: 0.2185 - val_accuracy: 0.9207 - val_loss: 0.2798
219/219  0s 1ms/step - accuracy: 0.9319 - loss: 0.2338
Test loss=0.24552959203720093 Test accuracy = 0.9307143092155457
Time taken to train the model is 73.31176495552063 seconds

Loss vs epochs for sigmoid [16, 32]



Accuracy vs epochs for sigmoid [16, 32]



In [49]: `## Model 6 --> Adam, lr =0.001, CCE loss, hidden_neurons=[16,32,64]`

```
hidden_neurons = [16,32,64]
activation_function = 'sigmoid'

model, history, duration = train_model(
    activation_func=activation_function,
    hidden_neurons=hidden_neurons,
    optimizer=Adam,
    loss_func=CategoricalCrossentropy,
    epochs=100,
    dropout_rate=None,
    learning_rate=0.001
)

test_loss, test_acc = model.evaluate(X_test,y_test)

print(f"Test loss={test_loss} Test accuracy = {test_acc}")
```

```

print(f"Time taken to train the model is {duration} seconds")

result_df.loc[len(result_df.index)] = [
    len(hidden_neurons),
    activation_function,
    str(hidden_neurons),
    'Adam',
    duration,
    test_loss,
    test_acc
]

plot_metrics(history, activation_function, hidden_neurons)

```

Model: "sequential_21"

Layer (type)	Output Shape	Param #
flatten_21 (Flatten)	(None, 1024)	0
dense_70 (Dense)	(None, 16)	16,400
dense_71 (Dense)	(None, 32)	544
dense_72 (Dense)	(None, 64)	2,112
dense_73 (Dense)	(None, 10)	650

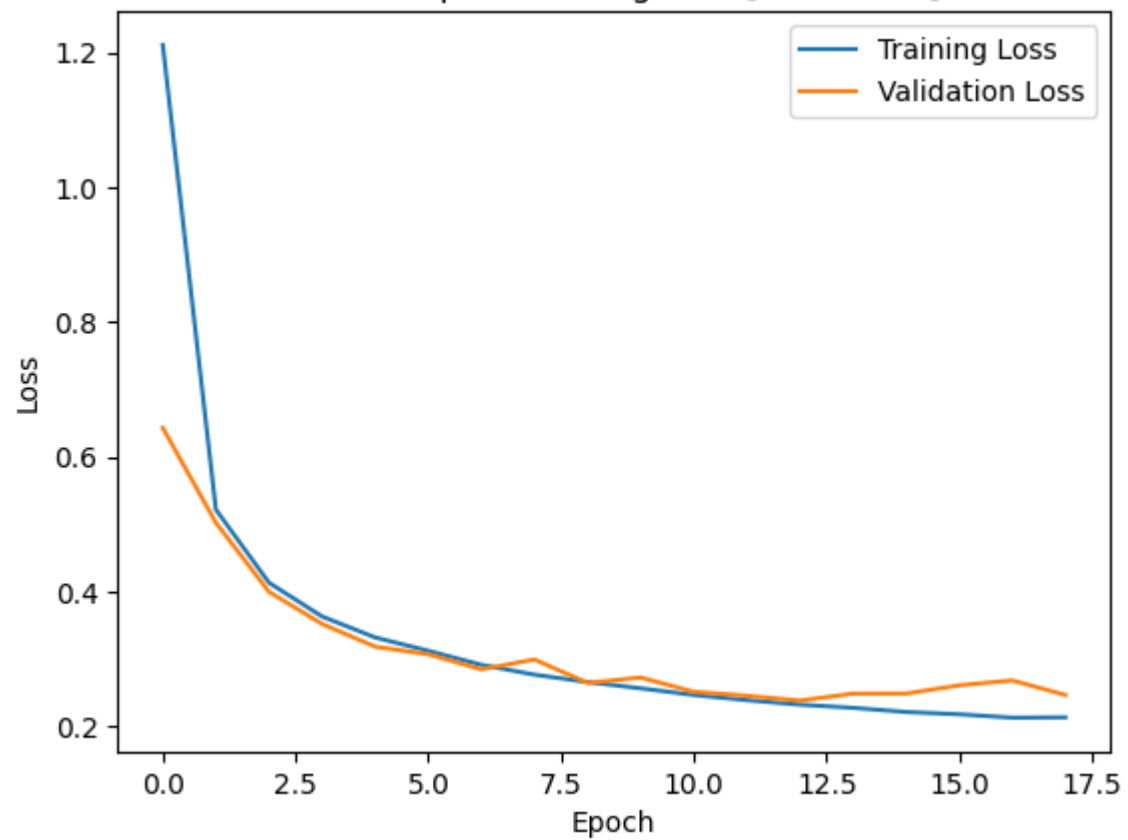
Total params: 19,706 (76.98 KB)

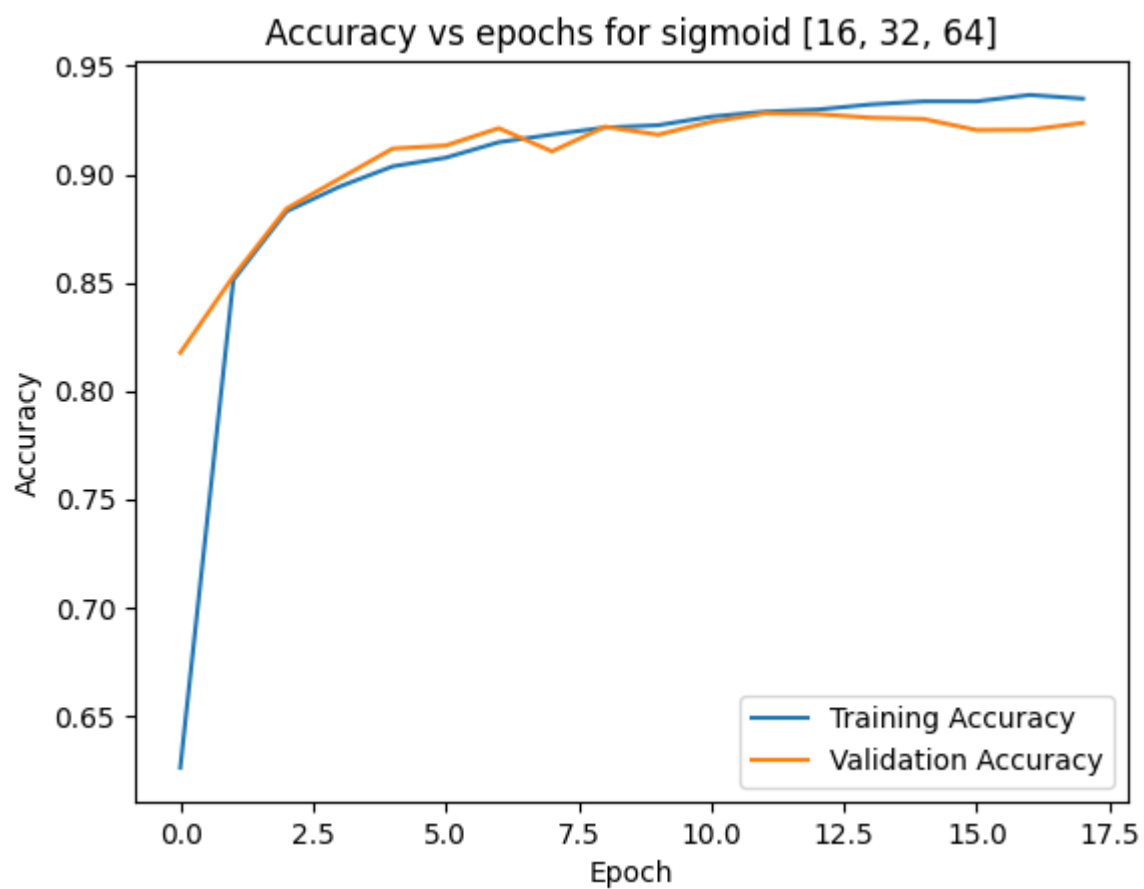
Trainable params: 19,706 (76.98 KB)

Non-trainable params: 0 (0.00 B)

Epoch 1/100
1750/1750 ————— 7s 3ms/step - accuracy: 0.4173 - loss: 1.7356 - val_accuracy: 0.8176 - val_loss: 0.6427
Epoch 2/100
1750/1750 ————— 4s 2ms/step - accuracy: 0.8375 - loss: 0.5732 - val_accuracy: 0.8529 - val_loss: 0.5020
Epoch 3/100
1750/1750 ————— 6s 2ms/step - accuracy: 0.8826 - loss: 0.4201 - val_accuracy: 0.8841 - val_loss: 0.3995
Epoch 4/100
1750/1750 ————— 4s 3ms/step - accuracy: 0.8932 - loss: 0.3659 - val_accuracy: 0.8980 - val_loss: 0.3518
Epoch 5/100
1750/1750 ————— 4s 2ms/step - accuracy: 0.9025 - loss: 0.3347 - val_accuracy: 0.9119 - val_loss: 0.3181
Epoch 6/100
1750/1750 ————— 3s 2ms/step - accuracy: 0.9079 - loss: 0.3120 - val_accuracy: 0.9133 - val_loss: 0.3068
Epoch 7/100
1750/1750 ————— 6s 3ms/step - accuracy: 0.9131 - loss: 0.2963 - val_accuracy: 0.9213 - val_loss: 0.2845
Epoch 8/100
1750/1750 ————— 4s 2ms/step - accuracy: 0.9166 - loss: 0.2805 - val_accuracy: 0.9106 - val_loss: 0.2991
Epoch 9/100
1750/1750 ————— 6s 2ms/step - accuracy: 0.9242 - loss: 0.2604 - val_accuracy: 0.9220 - val_loss: 0.2643
Epoch 10/100
1750/1750 ————— 5s 2ms/step - accuracy: 0.9233 - loss: 0.2518 - val_accuracy: 0.9183 - val_loss: 0.2723
Epoch 11/100
1750/1750 ————— 4s 2ms/step - accuracy: 0.9293 - loss: 0.2408 - val_accuracy: 0.9243 - val_loss: 0.2512
Epoch 12/100
1750/1750 ————— 6s 3ms/step - accuracy: 0.9294 - loss: 0.2372 - val_accuracy: 0.9281 - val_loss: 0.2452
Epoch 13/100
1750/1750 ————— 4s 2ms/step - accuracy: 0.9304 - loss: 0.2347 - val_accuracy: 0.9279 - val_loss: 0.2381
Epoch 14/100
1750/1750 ————— 4s 2ms/step - accuracy: 0.9350 - loss: 0.2225 - val_accuracy: 0.9261 - val_loss: 0.2485
Epoch 15/100
1750/1750 ————— 6s 3ms/step - accuracy: 0.9333 - loss: 0.2207 - val_accuracy: 0.9256 - val_loss: 0.2486
Epoch 16/100
1750/1750 ————— 4s 2ms/step - accuracy: 0.9348 - loss: 0.2179 - val_accuracy: 0.9204 - val_loss: 0.2607
Epoch 17/100
1750/1750 ————— 4s 2ms/step - accuracy: 0.9366 - loss: 0.2104 - val_accuracy: 0.9206 - val_loss: 0.2680
Epoch 18/100
1750/1750 ————— 4s 2ms/step - accuracy: 0.9356 - loss: 0.2119 - val_accuracy: 0.9237 - val_loss: 0.2464
219/219 ————— 0s 2ms/step - accuracy: 0.9357 - loss: 0.2305
Test loss=0.24331416189670563 Test accuracy = 0.9314285516738892
Time taken to train the model is 84.13690066337585 seconds

Loss vs epochs for sigmoid [16, 32, 64]





In [50]: result_df

Out[50]:

	Hidden Layers	Activation Function	Hidden Neurons	Optimizer	Training Time(in seconds)	Test Loss	Test Accuracy
0	1	sigmoid	[16]	SGD	428.502323	0.085315	0.497714
1	2	sigmoid	[16, 32]	SGD	436.430634	0.089831	0.157286
2	3	sigmoid	[16, 32, 64]	SGD	428.479271	0.089967	0.119429
3	1	sigmoid	[16]	Adam	63.880293	0.282636	0.919714
4	2	sigmoid	[16, 32]	Adam	73.311765	0.245530	0.930714
5	3	sigmoid	[16, 32, 64]	Adam	84.136901	0.243314	0.931429

Question 05

- v. For the following few tasks, use Adam optimizer (learning rate = 0.001) and Categorical Cross Entropy Loss. Run the network by changing the activation function hyper-parameter:

Hidden Layers	Activation Function	Hidden Neurons
3	Sigmoid	[16, 32, 64]
3	Tanh	[16, 32, 64]
3	ReLU	[16, 32, 64]

```
In [31]: ## Model 1 --> Adam, lr=0.001, CCE loss, hidden_neurons=[16,32,64]. Sigmoid
```

```
result_df = pd.DataFrame(  
    columns=[  
        'Hidden Layers',  
        'Activation Function',  
        'Hidden Neurons',  
        'Training Time (in seconds)',  
        'Test Loss',  
        'Test Accuracy'  
    ]  
)  
  
hidden_neurons = [16,32,64]  
activation_function = 'sigmoid'  
  
model, history, duration = train_model(  
    activation_func=activation_function,  
    hidden_neurons=hidden_neurons,  
    optimizer=Adam,  
    loss_func=CategoricalCrossentropy,  
    epochs=100,  
    dropout_rate=None,  
    learning_rate=0.001  
)  
  
test_loss, test_acc = model.evaluate(X_test,y_test)
```

```

print(f"Test loss={test_loss} Test accuracy = {test_acc}")
print(f"Time taken to train the model is {duration} seconds")

result_df.loc[len(result_df.index)] = [
    len(hidden_neurons),
    activation_function,
    str(hidden_neurons),
    duration,
    test_loss,
    test_acc
]

plot_metrics(history, activation_function, hidden_neurons)

```

Model: "sequential_8"

Layer (type)	Output Shape	Param #
flatten_8 (Flatten)	(None, 1024)	0
dense_24 (Dense)	(None, 16)	16,400
dense_25 (Dense)	(None, 32)	544
dense_26 (Dense)	(None, 64)	2,112
dense_27 (Dense)	(None, 10)	650

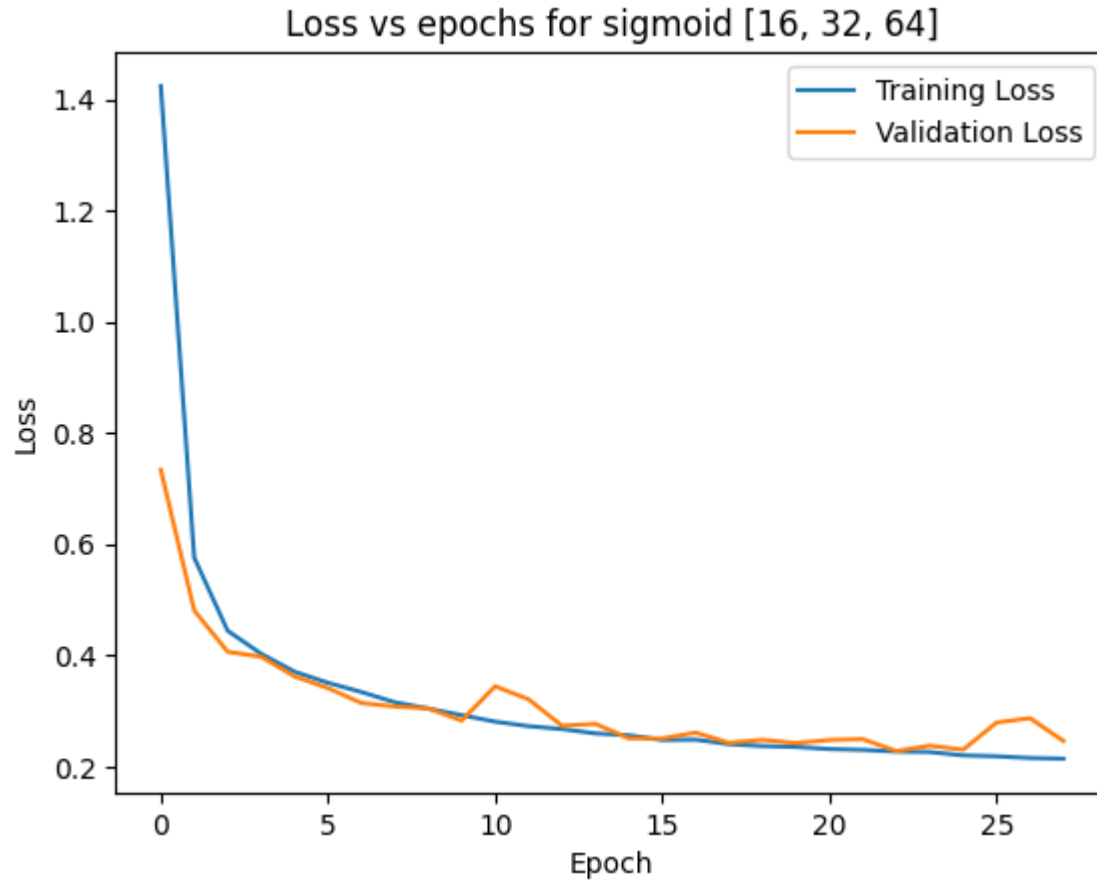
Total params: 19,706 (76.98 KB)

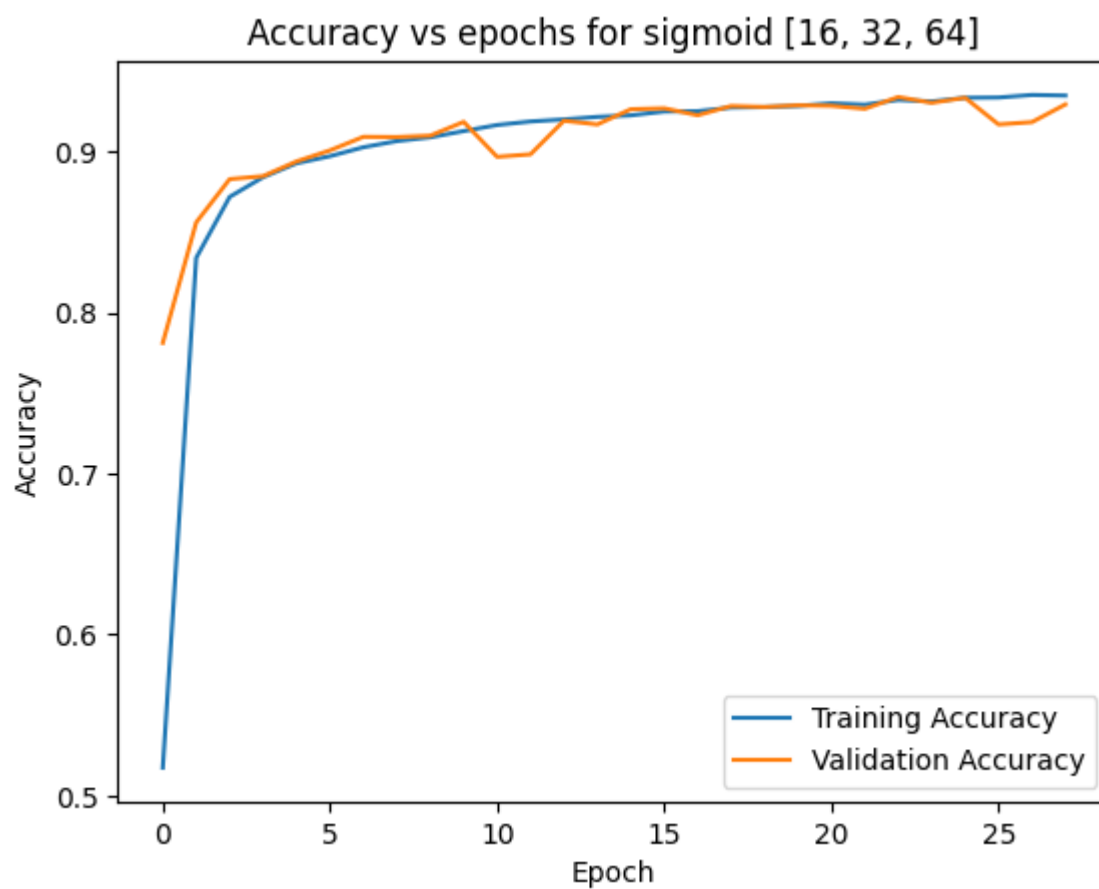
Trainable params: 19,706 (76.98 KB)

Non-trainable params: 0 (0.00 B)

Epoch 1/100	1750/1750	13s	6ms/step	- accuracy: 0.3228	- loss: 1.9027	- val_accuracy: 0.7813	- val_loss: 0.7335
Epoch 2/100	1750/1750	4s	2ms/step	- accuracy: 0.8098	- loss: 0.6485	- val_accuracy: 0.8561	- val_loss: 0.4806
Epoch 3/100	1750/1750	5s	3ms/step	- accuracy: 0.8665	- loss: 0.4655	- val_accuracy: 0.8830	- val_loss: 0.4062
Epoch 4/100	1750/1750	4s	2ms/step	- accuracy: 0.8819	- loss: 0.4082	- val_accuracy: 0.8849	- val_loss: 0.3973
Epoch 5/100	1750/1750	5s	2ms/step	- accuracy: 0.8927	- loss: 0.3756	- val_accuracy: 0.8940	- val_loss: 0.3619
Epoch 6/100	1750/1750	5s	3ms/step	- accuracy: 0.8940	- loss: 0.3618	- val_accuracy: 0.9009	- val_loss: 0.3409
Epoch 7/100	1750/1750	4s	2ms/step	- accuracy: 0.9006	- loss: 0.3429	- val_accuracy: 0.9093	- val_loss: 0.3141
Epoch 8/100	1750/1750	5s	2ms/step	- accuracy: 0.9025	- loss: 0.3270	- val_accuracy: 0.9091	- val_loss: 0.3078
Epoch 9/100	1750/1750	6s	3ms/step	- accuracy: 0.9090	- loss: 0.3059	- val_accuracy: 0.9100	- val_loss: 0.3043
Epoch 10/100	1750/1750	4s	2ms/step	- accuracy: 0.9127	- loss: 0.2916	- val_accuracy: 0.9186	- val_loss: 0.2824
Epoch 11/100	1750/1750	4s	2ms/step	- accuracy: 0.9154	- loss: 0.2855	- val_accuracy: 0.8969	- val_loss: 0.3442
Epoch 12/100	1750/1750	6s	2ms/step	- accuracy: 0.9202	- loss: 0.2690	- val_accuracy: 0.8984	- val_loss: 0.3208
Epoch 13/100	1750/1750	4s	2ms/step	- accuracy: 0.9209	- loss: 0.2629	- val_accuracy: 0.9193	- val_loss: 0.2733
Epoch 14/100	1750/1750	4s	2ms/step	- accuracy: 0.9212	- loss: 0.2627	- val_accuracy: 0.9170	- val_loss: 0.2766
Epoch 15/100	1750/1750	5s	3ms/step	- accuracy: 0.9225	- loss: 0.2560	- val_accuracy: 0.9264	- val_loss: 0.2505
Epoch 16/100	1750/1750	4s	2ms/step	- accuracy: 0.9250	- loss: 0.2474	- val_accuracy: 0.9270	- val_loss: 0.2503
Epoch 17/100	1750/1750	4s	2ms/step	- accuracy: 0.9263	- loss: 0.2429	- val_accuracy: 0.9227	- val_loss: 0.2609
Epoch 18/100	1750/1750	7s	3ms/step	- accuracy: 0.9271	- loss: 0.2426	- val_accuracy: 0.9284	- val_loss: 0.2425
Epoch 19/100	1750/1750	4s	2ms/step	- accuracy: 0.9285	- loss: 0.2337	- val_accuracy: 0.9279	- val_loss: 0.2479
Epoch 20/100	1750/1750	4s	2ms/step	- accuracy: 0.9293	- loss: 0.2318	- val_accuracy: 0.9289	- val_loss: 0.2420
Epoch 21/100	1750/1750	4s	2ms/step	- accuracy: 0.9301	- loss: 0.2302	- val_accuracy: 0.9286	- val_loss: 0.2476
Epoch 22/100	1750/1750	4s	2ms/step	- accuracy: 0.9297	- loss: 0.2279	- val_accuracy: 0.9267	- val_loss: 0.2491
Epoch 23/100	1750/1750	5s	2ms/step	- accuracy: 0.9335	- loss: 0.2233	- val_accuracy: 0.9340	- val_loss: 0.2277

Epoch 24/100
1750/1750 — 4s 2ms/step - accuracy: 0.9331 - loss: 0.2206 - val_accuracy: 0.9304 - val_loss: 0.2373
Epoch 25/100
1750/1750 — 5s 3ms/step - accuracy: 0.9328 - loss: 0.2211 - val_accuracy: 0.9336 - val_loss: 0.2301
Epoch 26/100
1750/1750 — 4s 2ms/step - accuracy: 0.9352 - loss: 0.2139 - val_accuracy: 0.9170 - val_loss: 0.2789
Epoch 27/100
1750/1750 — 5s 2ms/step - accuracy: 0.9349 - loss: 0.2177 - val_accuracy: 0.9184 - val_loss: 0.2869
Epoch 28/100
1750/1750 — 5s 2ms/step - accuracy: 0.9361 - loss: 0.2082 - val_accuracy: 0.9294 - val_loss: 0.2458
219/219 — 0s 1ms/step - accuracy: 0.9338 - loss: 0.2278
Test loss=0.22948826849460602 Test accuracy = 0.9328571557998657
Time taken to train the model is 134.11440420150757 seconds





In [32]: `## Model 2 --> Adam, lr =0.001, CCE loss, hidden_neurons=[16,32,64]. tanh`

```
hidden_neurons = [16,32,64]
activation_function = 'tanh'

model, history, duration = train_model(
    activation_func=activation_function,
    hidden_neurons=hidden_neurons,
    optimizer=Adam,
    loss_func=CategoricalCrossentropy,
    epochs=100,
    dropout_rate=None,
    learning_rate=0.001
)

test_loss, test_acc = model.evaluate(X_test,y_test)

print(f"Test loss={test_loss} Test accuracy = {test_acc}")
```

```

print(f"Time taken to train the model is {duration} seconds")

result_df.loc[len(result_df.index)] = [
    len(hidden_neurons),
    activation_function,
    str(hidden_neurons),
    duration,
    test_loss,
    test_acc
]

plot_metrics(history, activation_function, hidden_neurons)

```











Model: "sequential_9"

Layer (type)	Output Shape	Param #
flatten_9 (Flatten)	(None , 1024)	0
dense_28 (Dense)	(None , 16)	16,400
dense_29 (Dense)	(None , 32)	544
dense_30 (Dense)	(None , 64)	2,112
dense_31 (Dense)	(None , 10)	650

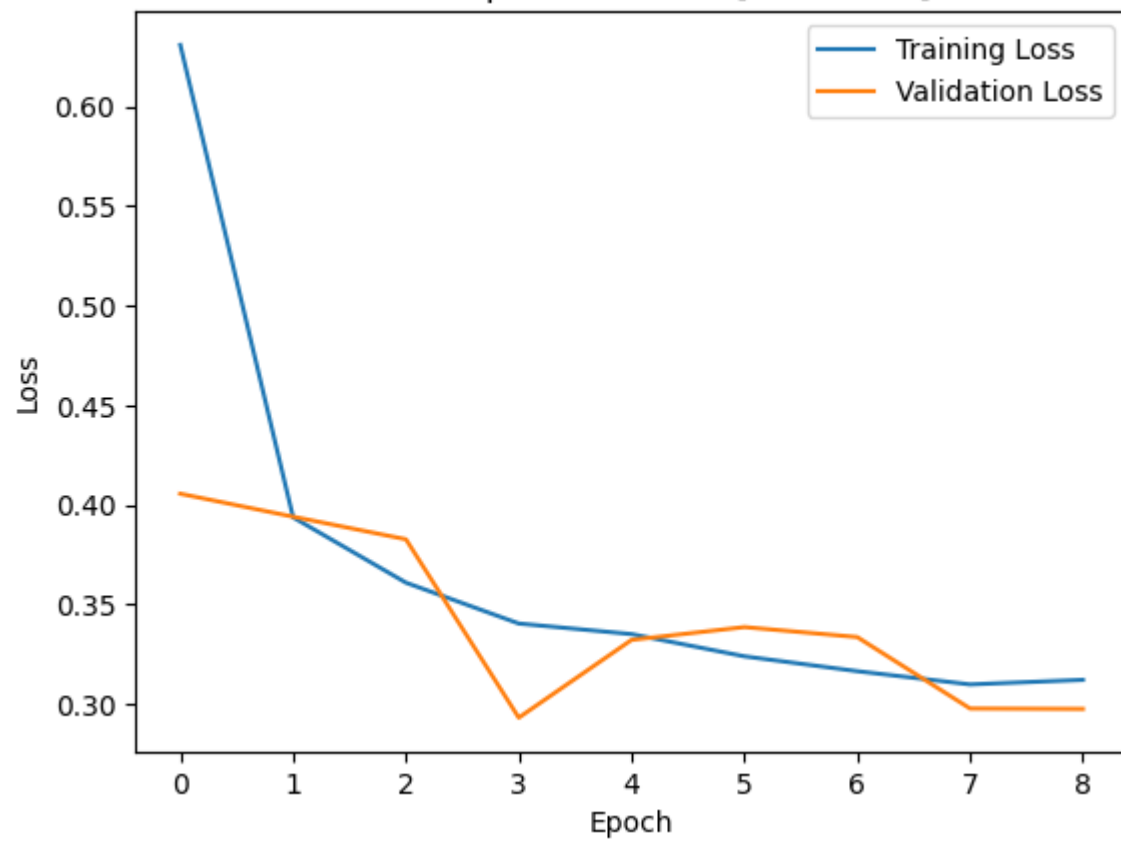
Total params: [19,706](#) (76.98 KB)

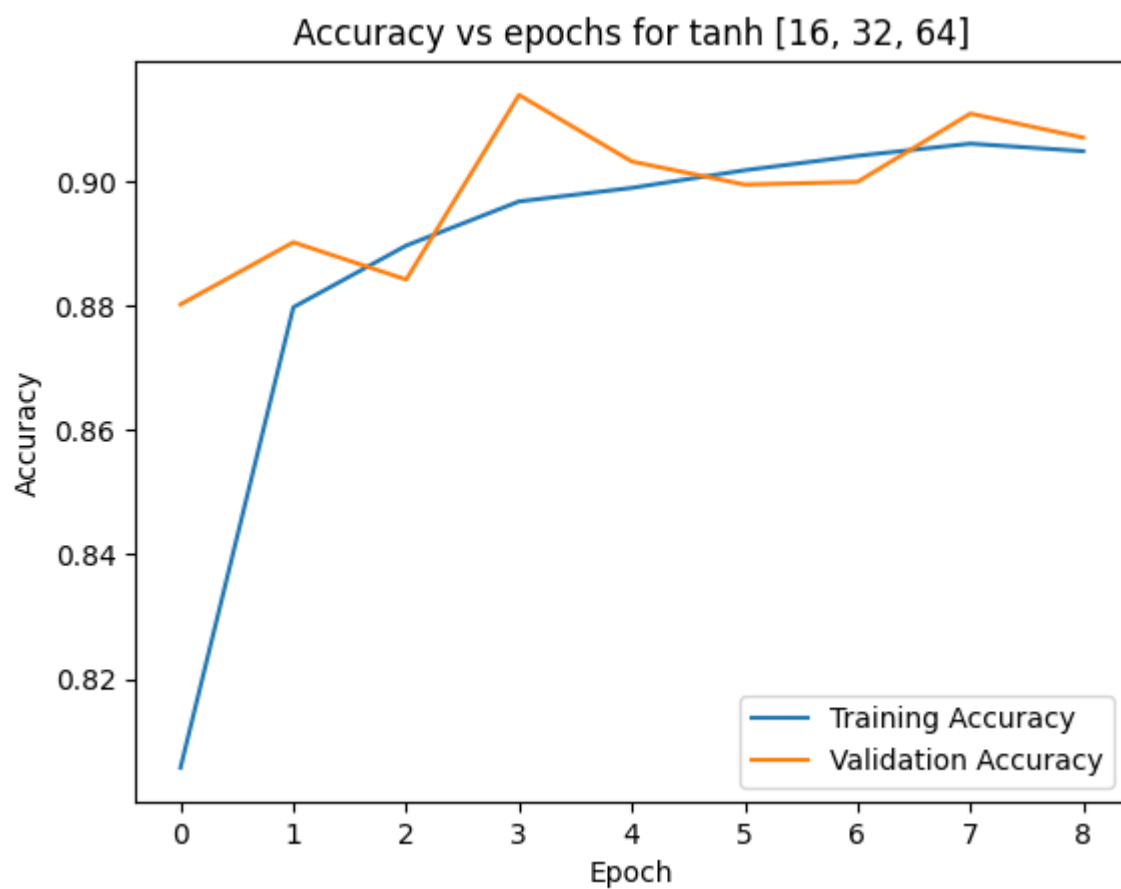
Trainable params: [19,706](#) (76.98 KB)

Non-trainable params: [0](#) (0.00 B)

Epoch 1/100
1750/1750  5s 2ms/step - accuracy: 0.6881 - loss: 0.9802 - val_accuracy: 0.8803 - val_loss: 0.4057
Epoch 2/100
1750/1750  6s 2ms/step - accuracy: 0.8780 - loss: 0.3985 - val_accuracy: 0.8903 - val_loss: 0.3941
Epoch 3/100
1750/1750  3s 2ms/step - accuracy: 0.8899 - loss: 0.3602 - val_accuracy: 0.8843 - val_loss: 0.3828
Epoch 4/100
1750/1750  6s 2ms/step - accuracy: 0.8964 - loss: 0.3406 - val_accuracy: 0.9140 - val_loss: 0.2933
Epoch 5/100
1750/1750  4s 2ms/step - accuracy: 0.8982 - loss: 0.3361 - val_accuracy: 0.9033 - val_loss: 0.3323
Epoch 6/100
1750/1750  4s 2ms/step - accuracy: 0.9023 - loss: 0.3262 - val_accuracy: 0.8996 - val_loss: 0.3387
Epoch 7/100
1750/1750  3s 2ms/step - accuracy: 0.9052 - loss: 0.3155 - val_accuracy: 0.9000 - val_loss: 0.3337
Epoch 8/100
1750/1750  5s 3ms/step - accuracy: 0.9037 - loss: 0.3185 - val_accuracy: 0.9110 - val_loss: 0.2978
Epoch 9/100
1750/1750  3s 2ms/step - accuracy: 0.9060 - loss: 0.3082 - val_accuracy: 0.9071 - val_loss: 0.2976
219/219  0s 1ms/step - accuracy: 0.9142 - loss: 0.2919
Test loss=0.3065372705459595 Test accuracy = 0.9105714559555054
Time taken to train the model is 41.307180643081665 seconds

Loss vs epochs for tanh [16, 32, 64]





In [33]: `## Model 3 --> Adam, lr =0.001, CCE loss, hidden_neurons=[16,32,64]. relu`

```
hidden_neurons = [16,32,64]
activation_function = 'relu'

model, history, duration = train_model(
    activation_func=activation_function,
    hidden_neurons=hidden_neurons,
    optimizer=Adam,
    loss_func=CategoricalCrossentropy,
    epochs=100,
    dropout_rate=None,
    learning_rate=0.001
)

test_loss, test_acc = model.evaluate(X_test,y_test)

print(f"Test loss={test_loss} Test accuracy = {test_acc}")
```

```

print(f"Time taken to train the model is {duration} seconds")

result_df.loc[len(result_df.index)] = [
    len(hidden_neurons),
    activation_function,
    str(hidden_neurons),
    duration,
    test_loss,
    test_acc
]

plot_metrics(history, activation_function, hidden_neurons)

```

Model: "sequential_10"

Layer (type)	Output Shape	Param #
flatten_10 (Flatten)	(None, 1024)	0
dense_32 (Dense)	(None, 16)	16,400
dense_33 (Dense)	(None, 32)	544
dense_34 (Dense)	(None, 64)	2,112
dense_35 (Dense)	(None, 10)	650

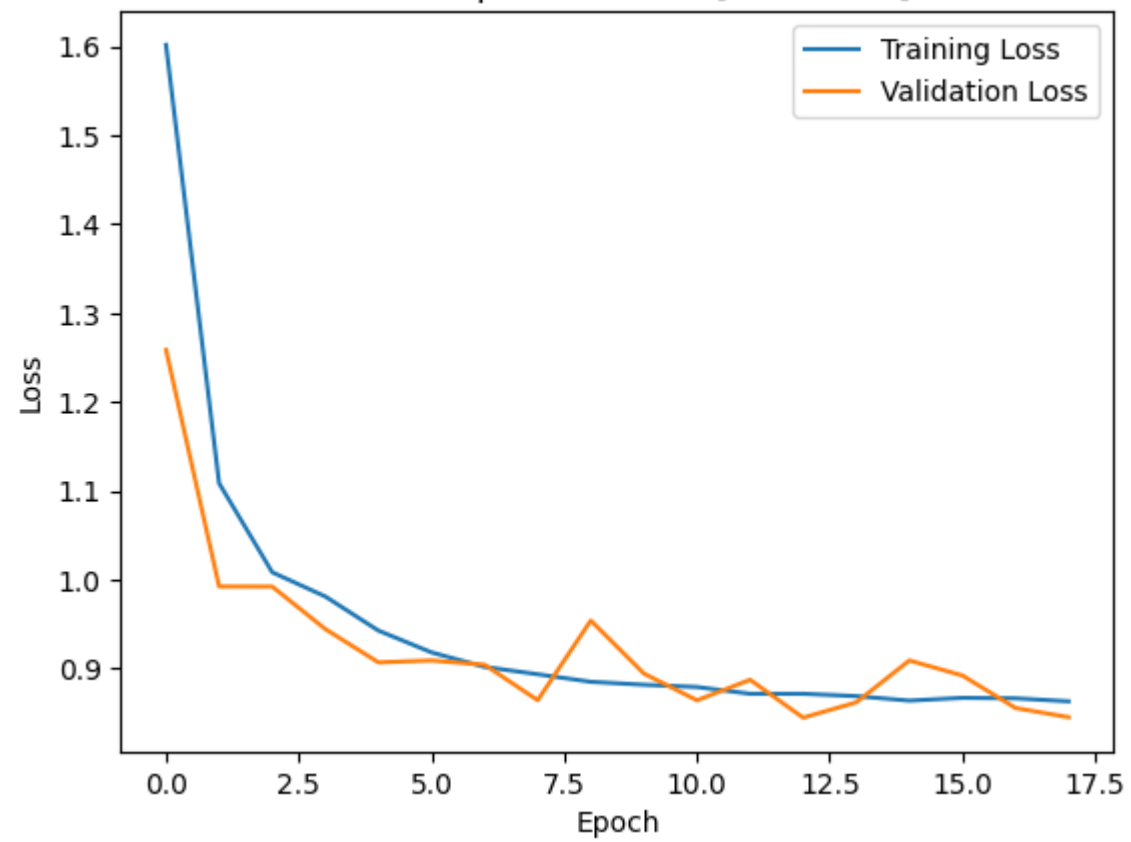
Total params: 19,706 (76.98 KB)

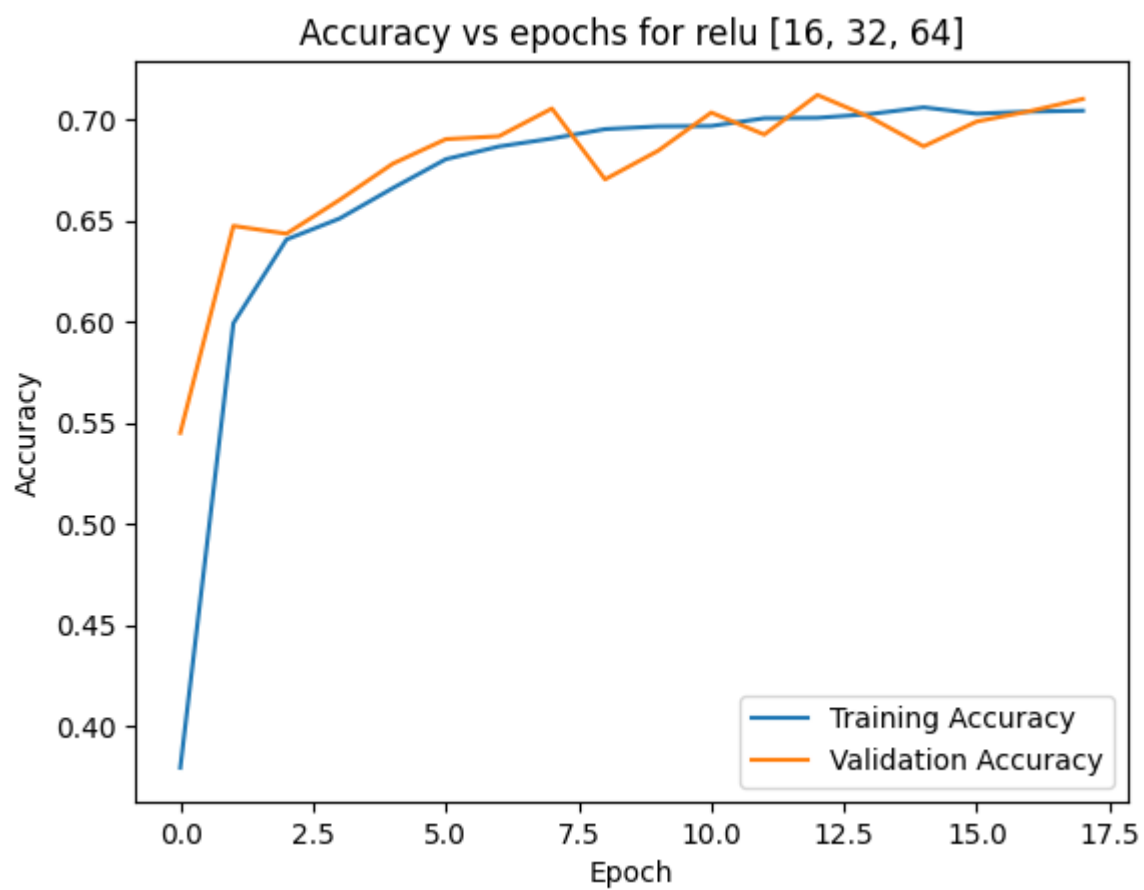
Trainable params: 19,706 (76.98 KB)

Non-trainable params: 0 (0.00 B)

Epoch 1/100
1750/1750 ————— 6s 3ms/step - accuracy: 0.2840 - loss: 1.8179 - val_accuracy: 0.5450 - val_loss: 1.2585
Epoch 2/100
1750/1750 ————— 4s 2ms/step - accuracy: 0.5777 - loss: 1.1678 - val_accuracy: 0.6473 - val_loss: 0.9924
Epoch 3/100
1750/1750 ————— 3s 2ms/step - accuracy: 0.6378 - loss: 1.0118 - val_accuracy: 0.6434 - val_loss: 0.9923
Epoch 4/100
1750/1750 ————— 6s 3ms/step - accuracy: 0.6501 - loss: 0.9843 - val_accuracy: 0.6601 - val_loss: 0.9445
Epoch 5/100
1750/1750 ————— 4s 2ms/step - accuracy: 0.6603 - loss: 0.9599 - val_accuracy: 0.6780 - val_loss: 0.9070
Epoch 6/100
1750/1750 ————— 3s 2ms/step - accuracy: 0.6788 - loss: 0.9196 - val_accuracy: 0.6901 - val_loss: 0.9092
Epoch 7/100
1750/1750 ————— 4s 2ms/step - accuracy: 0.6863 - loss: 0.9003 - val_accuracy: 0.6916 - val_loss: 0.9044
Epoch 8/100
1750/1750 ————— 5s 2ms/step - accuracy: 0.6861 - loss: 0.8993 - val_accuracy: 0.7053 - val_loss: 0.8645
Epoch 9/100
1750/1750 ————— 5s 2ms/step - accuracy: 0.6986 - loss: 0.8743 - val_accuracy: 0.6703 - val_loss: 0.9541
Epoch 10/100
1750/1750 ————— 5s 3ms/step - accuracy: 0.6951 - loss: 0.8849 - val_accuracy: 0.6844 - val_loss: 0.8945
Epoch 11/100
1750/1750 ————— 4s 2ms/step - accuracy: 0.6954 - loss: 0.8820 - val_accuracy: 0.7033 - val_loss: 0.8643
Epoch 12/100
1750/1750 ————— 5s 2ms/step - accuracy: 0.6994 - loss: 0.8737 - val_accuracy: 0.6924 - val_loss: 0.8875
Epoch 13/100
1750/1750 ————— 5s 3ms/step - accuracy: 0.7044 - loss: 0.8581 - val_accuracy: 0.7120 - val_loss: 0.8448
Epoch 14/100
1750/1750 ————— 4s 2ms/step - accuracy: 0.6999 - loss: 0.8727 - val_accuracy: 0.7009 - val_loss: 0.8620
Epoch 15/100
1750/1750 ————— 5s 2ms/step - accuracy: 0.7063 - loss: 0.8580 - val_accuracy: 0.6866 - val_loss: 0.9092
Epoch 16/100
1750/1750 ————— 5s 3ms/step - accuracy: 0.7027 - loss: 0.8695 - val_accuracy: 0.6989 - val_loss: 0.8923
Epoch 17/100
1750/1750 ————— 3s 2ms/step - accuracy: 0.7070 - loss: 0.8619 - val_accuracy: 0.7040 - val_loss: 0.8557
Epoch 18/100
1750/1750 ————— 3s 2ms/step - accuracy: 0.7056 - loss: 0.8605 - val_accuracy: 0.7100 - val_loss: 0.8454
219/219 ————— 0s 1ms/step - accuracy: 0.7070 - loss: 0.8464
Test loss=0.849205493927002 Test accuracy = 0.7095714211463928
Time taken to train the model is 81.0392382144928 seconds

Loss vs epochs for relu [16, 32, 64]





In [34]: result_df

Out[34]:

	Hidden Layers	Activation Function	Hidden Neurons	Training Time (in seconds)	Test Loss	Test Accuracy
0	3	sigmoid	[16, 32, 64]	134.114404	0.229488	0.932857
1	3	tanh	[16, 32, 64]	41.307181	0.306537	0.910571
2	3	relu	[16, 32, 64]	81.039238	0.849205	0.709571

From the above table, we can infer that the best activation funtion is `sigmoid`

Question 06

vi. Run the network by changing the number of Dropout hyper-parameters:

Hidden Layers	Activation Function	Hidden Neurons	Dropout
3	ReLU	[16, 32, 64]	0.9
3	ReLU	[16, 32, 64]	0.75
3	ReLU	[16, 32, 64]	0.5
3	ReLU	[16, 32, 64]	0.25
3	ReLU	[16, 32, 64]	0.10

```
In [35]: result_df_dropout = pd.DataFrame(  
    columns=[  
        'Hidden Layers',  
        'Activation Function',  
        'Hidden Neurons',  
        'Dropout Rate',  
        'Training Time (in seconds)',  
        'Test Loss',  
        'Test Accuracy'  
    ]  
)
```

```
In [36]: ## Model 1 --> Adam, lr =0.001, CCE loss, hidden_neurons=[16,32,64]. relu, dropout=0.9
```

```
hidden_neurons = [16,32,64]  
activation_function = 'relu'  
dropout_rate = 0.9  
  
model, history,duration= train_model(  
    activation_func=activation_function,  
    hidden_neurons=hidden_neurons,  
    optimizer=Adam,  
    loss_func=CategoricalCrossentropy,  
    epochs=100,  
    dropout_rate=dropout_rate,  
    learning_rate=0.001
```

```

)

test_loss, test_acc = model.evaluate(X_test,y_test)

print(f"Test loss={test_loss} Test accuracy = {test_acc}")
print(f"Time required to train the model is {duration} seconds")

result_df_dropout.loc[len(result_df_dropout.index)] = [
    len(hidden_neurons),
    activation_function,
    str(hidden_neurons),
    dropout_rate,
    duration,
    test_loss,
    test_acc
]

plot_metrics(history, activation_function, hidden_neurons)

```

Model: "sequential_11"

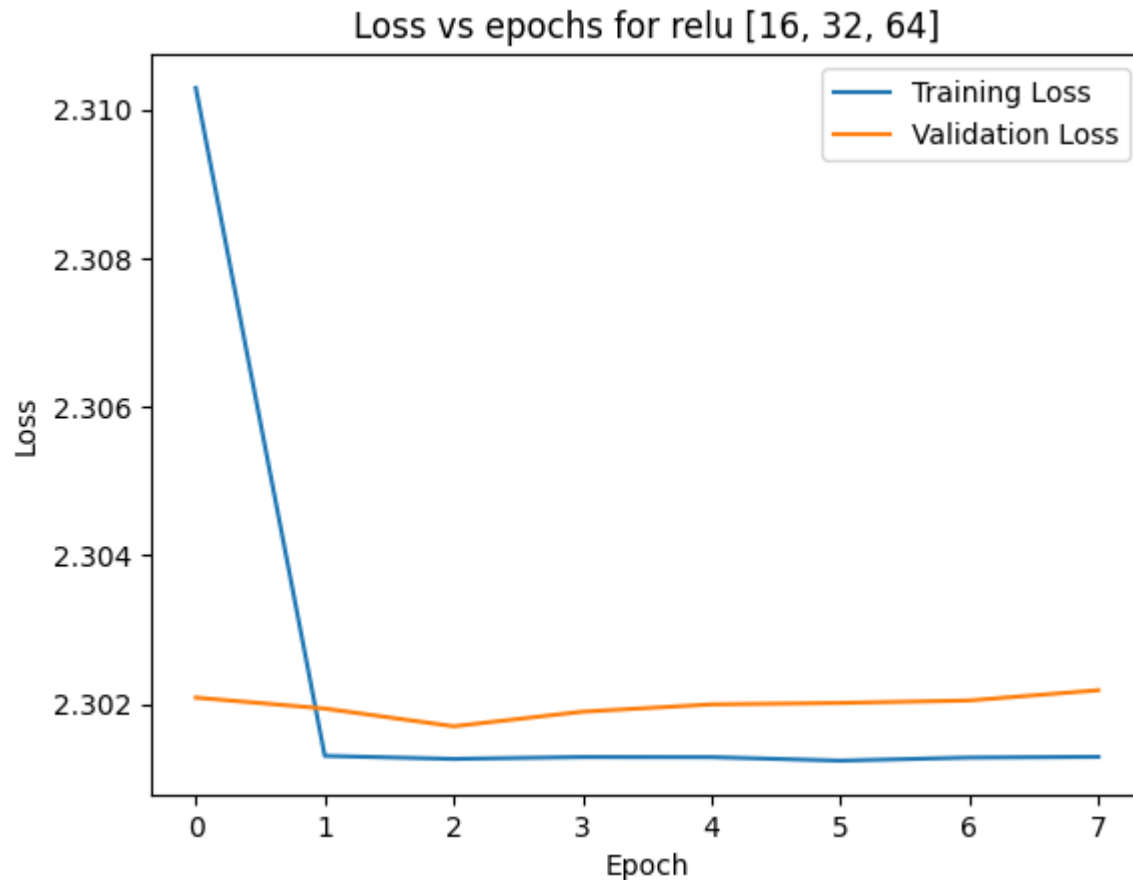
Layer (type)	Output Shape	Param #
flatten_11 (Flatten)	(None, 1024)	0
dense_36 (Dense)	(None, 16)	16,400
dropout (Dropout)	(None, 16)	0
dense_37 (Dense)	(None, 32)	544
dropout_1 (Dropout)	(None, 32)	0
dense_38 (Dense)	(None, 64)	2,112
dropout_2 (Dropout)	(None, 64)	0
dense_39 (Dense)	(None, 10)	650

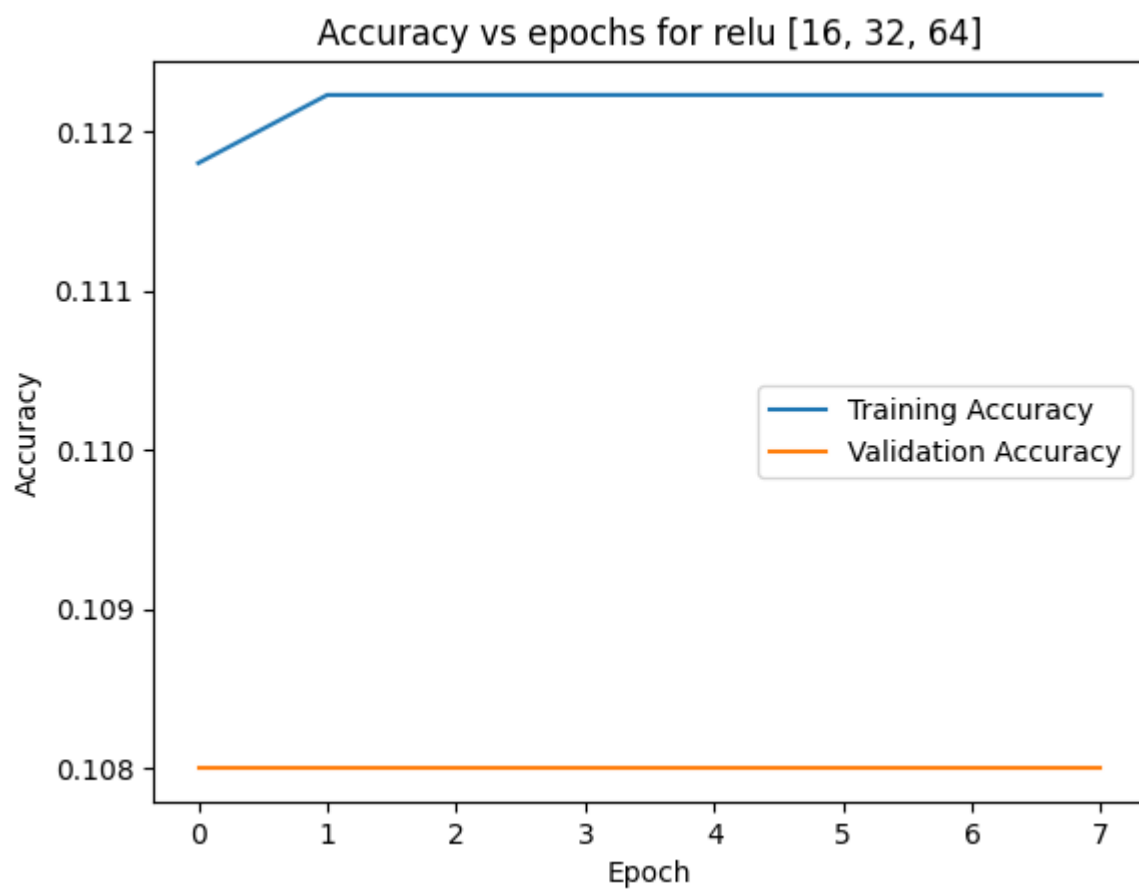
Total params: 19,706 (76.98 KB)

Trainable params: 19,706 (76.98 KB)

Non-trainable params: 0 (0.00 B)

Epoch 1/100
1750/1750 ————— 7s 3ms/step - accuracy: 0.1100 - loss: 2.3640 - val_accuracy: 0.1080 - val_loss: 2.3021
Epoch 2/100
1750/1750 ————— 4s 2ms/step - accuracy: 0.1111 - loss: 2.3014 - val_accuracy: 0.1080 - val_loss: 2.3019
Epoch 3/100
1750/1750 ————— 4s 2ms/step - accuracy: 0.1120 - loss: 2.3011 - val_accuracy: 0.1080 - val_loss: 2.3017
Epoch 4/100
1750/1750 ————— 5s 3ms/step - accuracy: 0.1106 - loss: 2.3016 - val_accuracy: 0.1080 - val_loss: 2.3019
Epoch 5/100
1750/1750 ————— 4s 2ms/step - accuracy: 0.1117 - loss: 2.3015 - val_accuracy: 0.1080 - val_loss: 2.3020
Epoch 6/100
1750/1750 ————— 5s 2ms/step - accuracy: 0.1130 - loss: 2.3014 - val_accuracy: 0.1080 - val_loss: 2.3020
Epoch 7/100
1750/1750 ————— 5s 3ms/step - accuracy: 0.1118 - loss: 2.3016 - val_accuracy: 0.1080 - val_loss: 2.3021
Epoch 8/100
1750/1750 ————— 4s 2ms/step - accuracy: 0.1116 - loss: 2.3015 - val_accuracy: 0.1080 - val_loss: 2.3022
219/219 ————— 0s 1ms/step - accuracy: 0.1219 - loss: 2.3007
Test loss=2.300971031188965 Test accuracy = 0.11942857503890991
Time required to train the model is 37.40566968917847 seconds





In [37]: `## Model 2 --> Adam, lr =0.001, CCE loss, hidden_neurons=[16,32,64]. relu, dropout=0.75`

```
hidden_neurons = [16,32,64]
activation_function = 'relu'
dropout_rate = 0.75

model, history,duration = train_model(
    activation_func=activation_function,
    hidden_neurons=hidden_neurons,
    optimizer=Adam,
    loss_func=CategoricalCrossentropy,
    epochs=100,
    dropout_rate=dropout_rate,
    learning_rate=0.001
)

test_loss, test_acc = model.evaluate(X_test,y_test)
```

```

print(f"Test loss={test_loss} Test accuracy = {test_acc}")
print(f"Time required to train the model is {duration} seconds")

result_df_dropout.loc[len(result_df_dropout.index)] = [
    len(hidden_neurons),
    activation_function,
    str(hidden_neurons),
    dropout_rate,
    duration,
    test_loss,
    test_acc
]

plot_metrics(history, activation_function, hidden_neurons)

```














Model: "sequential_12"

Layer (type)	Output Shape	Param #
flatten_12 (Flatten)	(None, 1024)	0
dense_40 (Dense)	(None, 16)	16,400
dropout_3 (Dropout)	(None, 16)	0
dense_41 (Dense)	(None, 32)	544
dropout_4 (Dropout)	(None, 32)	0
dense_42 (Dense)	(None, 64)	2,112
dropout_5 (Dropout)	(None, 64)	0
dense_43 (Dense)	(None, 10)	650

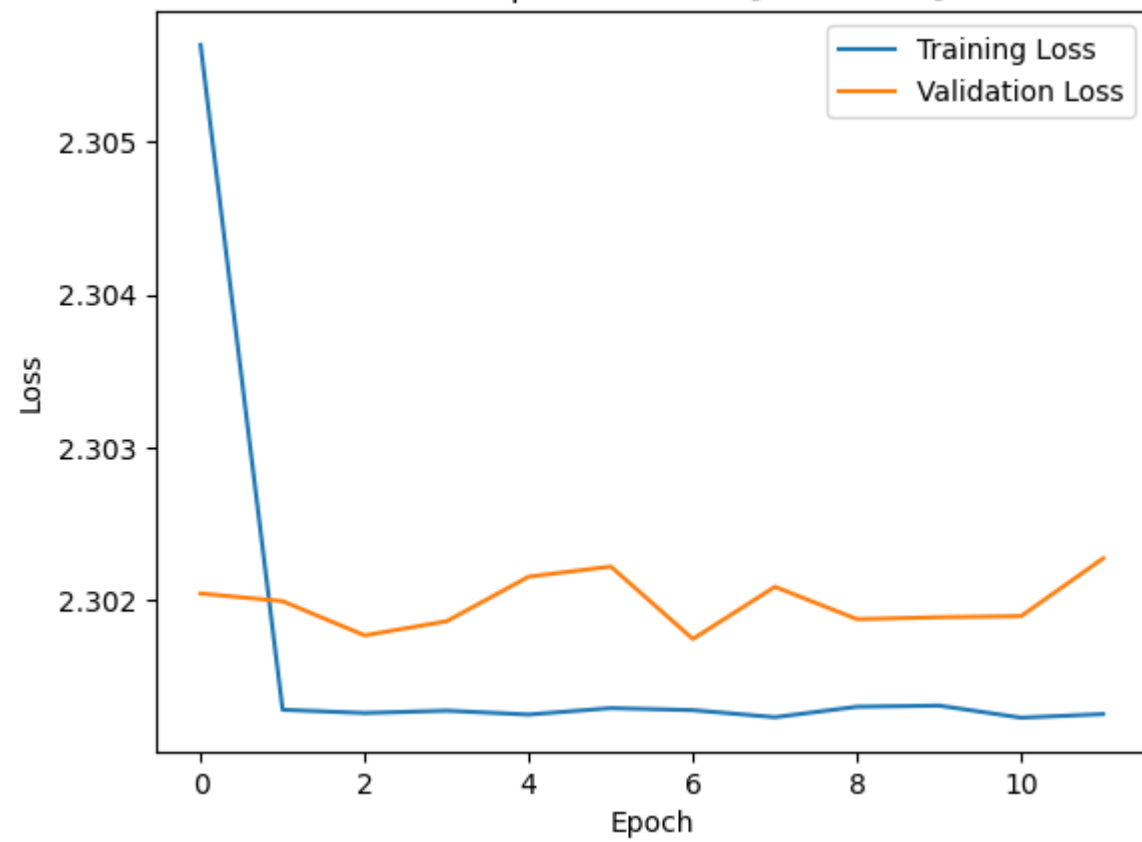
Total params: 19,706 (76.98 KB)

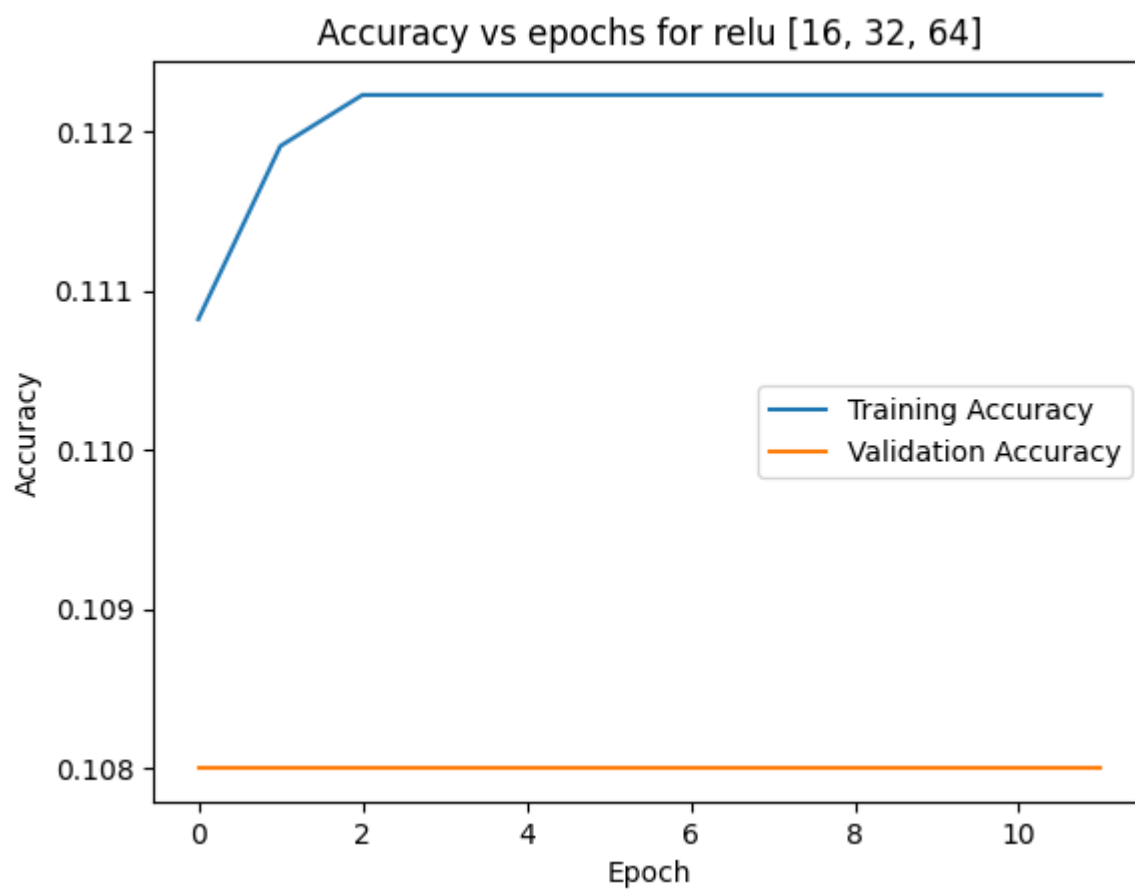
Trainable params: 19,706 (76.98 KB)

Non-trainable params: 0 (0.00 B)

Epoch 1/100
1750/1750  6s 3ms/step - accuracy: 0.1090 - loss: 2.3317 - val_accuracy: 0.1080 - val_loss: 2.3020
Epoch 2/100
1750/1750  5s 2ms/step - accuracy: 0.1107 - loss: 2.3014 - val_accuracy: 0.1080 - val_loss: 2.3020
Epoch 3/100
1750/1750  4s 2ms/step - accuracy: 0.1131 - loss: 2.3012 - val_accuracy: 0.1080 - val_loss: 2.3018
Epoch 4/100
1750/1750  4s 2ms/step - accuracy: 0.1128 - loss: 2.3011 - val_accuracy: 0.1080 - val_loss: 2.3019
Epoch 5/100
1750/1750  5s 2ms/step - accuracy: 0.1131 - loss: 2.3012 - val_accuracy: 0.1080 - val_loss: 2.3022
Epoch 6/100
1750/1750  4s 2ms/step - accuracy: 0.1109 - loss: 2.3014 - val_accuracy: 0.1080 - val_loss: 2.3022
Epoch 7/100
1750/1750  4s 3ms/step - accuracy: 0.1135 - loss: 2.3009 - val_accuracy: 0.1080 - val_loss: 2.3017
Epoch 8/100
1750/1750  5s 2ms/step - accuracy: 0.1131 - loss: 2.3012 - val_accuracy: 0.1080 - val_loss: 2.3021
Epoch 9/100
1750/1750  4s 2ms/step - accuracy: 0.1136 - loss: 2.3011 - val_accuracy: 0.1080 - val_loss: 2.3019
Epoch 10/100
1750/1750  6s 3ms/step - accuracy: 0.1130 - loss: 2.3011 - val_accuracy: 0.1080 - val_loss: 2.3019
Epoch 11/100
1750/1750  9s 2ms/step - accuracy: 0.1140 - loss: 2.3008 - val_accuracy: 0.1080 - val_loss: 2.3019
Epoch 12/100
1750/1750  5s 3ms/step - accuracy: 0.1122 - loss: 2.3015 - val_accuracy: 0.1080 - val_loss: 2.3023
219/219  0s 1ms/step - accuracy: 0.1219 - loss: 2.3006
Test loss=2.3008530139923096 Test accuracy = 0.11942857503890991
Time required to train the model is 61.32969856262207 seconds

Loss vs epochs for relu [16, 32, 64]





In [38]: `## Model 3 --> Adam, lr =0.001, CCE loss, hidden_neurons=[16,32,64]. relu, dropout=0.5`

```
hidden_neurons = [16,32,64]
activation_function = 'relu'
dropout_rate = 0.5

model, history, duration = train_model(
    activation_func=activation_function,
    hidden_neurons=hidden_neurons,
    optimizer=Adam,
    loss_func=CategoricalCrossentropy,
    epochs=100,
    dropout_rate=dropout_rate,
    learning_rate=0.001
)

test_loss, test_acc = model.evaluate(X_test,y_test)
```

```

print(f"Test loss={test_loss} Test accuracy = {test_acc}")
print(f"Time taken to train the model is {duration} seconds")

result_df_dropout.loc[len(result_df_dropout.index)] = [
    len(hidden_neurons),
    activation_function,
    str(hidden_neurons),
    dropout_rate,
    duration,
    test_loss,
    test_acc
]

plot_metrics(history, activation_function, hidden_neurons)

```












Model: "sequential_13"

Layer (type)	Output Shape	Param #
flatten_13 (Flatten)	(None, 1024)	0
dense_44 (Dense)	(None, 16)	16,400
dropout_6 (Dropout)	(None, 16)	0
dense_45 (Dense)	(None, 32)	544
dropout_7 (Dropout)	(None, 32)	0
dense_46 (Dense)	(None, 64)	2,112
dropout_8 (Dropout)	(None, 64)	0
dense_47 (Dense)	(None, 10)	650

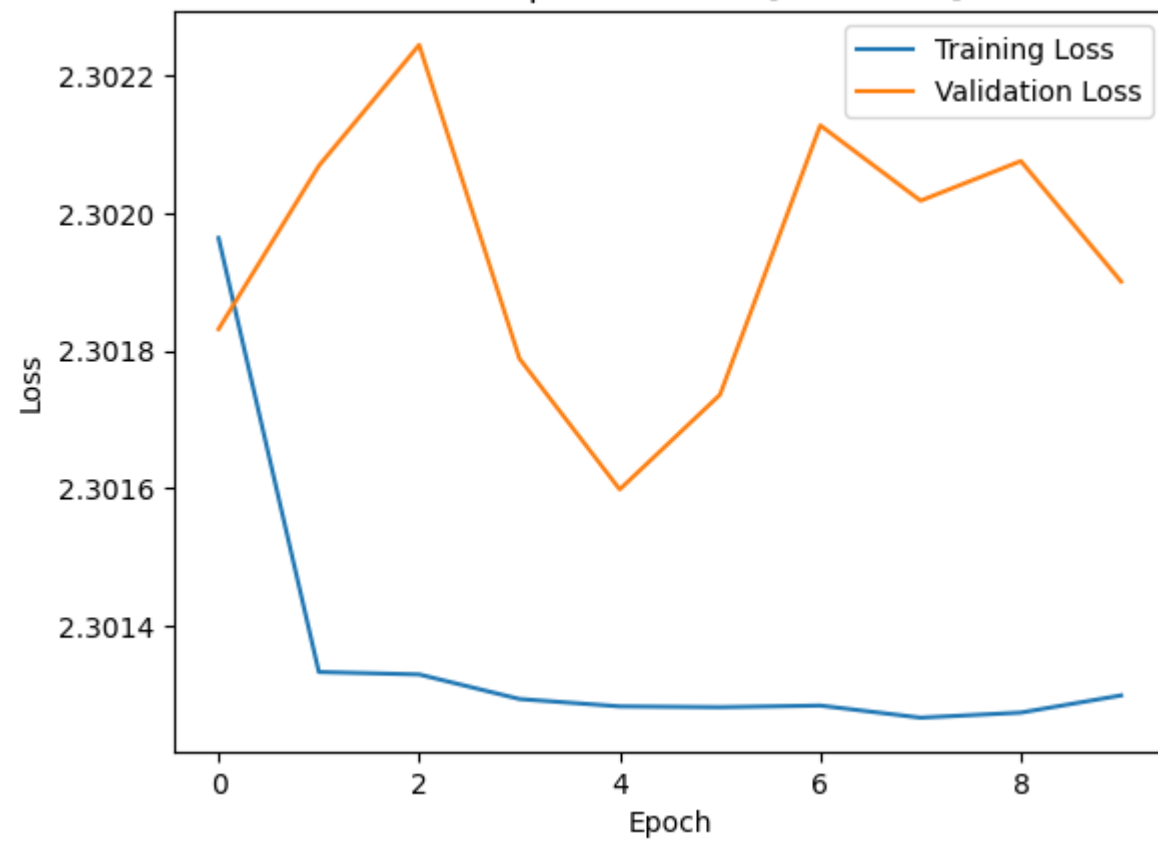
Total params: 19,706 (76.98 KB)

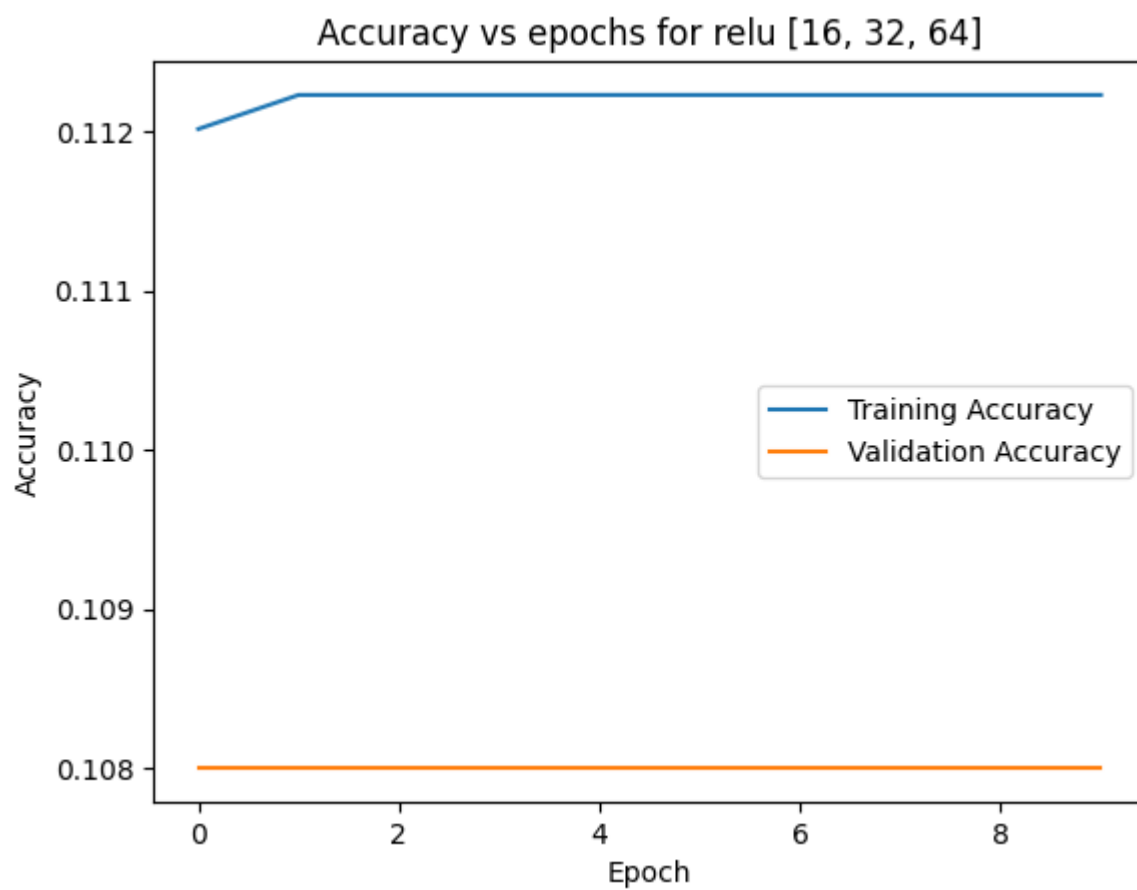
Trainable params: 19,706 (76.98 KB)

Non-trainable params: 0 (0.00 B)

Epoch 1/100
1750/1750  5s 2ms/step - accuracy: 0.1118 - loss: 2.3042 - val_accuracy: 0.1080 - val_loss: 2.3018
Epoch 2/100
1750/1750  6s 3ms/step - accuracy: 0.1134 - loss: 2.3013 - val_accuracy: 0.1080 - val_loss: 2.3021
Epoch 3/100
1750/1750  4s 2ms/step - accuracy: 0.1101 - loss: 2.3016 - val_accuracy: 0.1080 - val_loss: 2.3022
Epoch 4/100
1750/1750  5s 2ms/step - accuracy: 0.1120 - loss: 2.3012 - val_accuracy: 0.1080 - val_loss: 2.3018
Epoch 5/100
1750/1750  6s 3ms/step - accuracy: 0.1124 - loss: 2.3010 - val_accuracy: 0.1080 - val_loss: 2.3016
Epoch 6/100
1750/1750  4s 2ms/step - accuracy: 0.1155 - loss: 2.3010 - val_accuracy: 0.1080 - val_loss: 2.3017
Epoch 7/100
1750/1750  5s 2ms/step - accuracy: 0.1139 - loss: 2.3011 - val_accuracy: 0.1080 - val_loss: 2.3021
Epoch 8/100
1750/1750  5s 2ms/step - accuracy: 0.1125 - loss: 2.3012 - val_accuracy: 0.1080 - val_loss: 2.3020
Epoch 9/100
1750/1750  4s 2ms/step - accuracy: 0.1094 - loss: 2.3016 - val_accuracy: 0.1080 - val_loss: 2.3021
Epoch 10/100
1750/1750  6s 3ms/step - accuracy: 0.1111 - loss: 2.3014 - val_accuracy: 0.1080 - val_loss: 2.3019
219/219  0s 1ms/step - accuracy: 0.1219 - loss: 2.3008
Test loss=2.3010246753692627 Test accuracy = 0.11942857503890991
Time taken to train the model is 51.37219834327698 seconds

Loss vs epochs for relu [16, 32, 64]





In [39]: `## Model 4 --> Adam, lr =0.001, CCE loss, hidden_neurons=[16,32,64]. relu, dropout=0.25`

```
hidden_neurons = [16,32,64]
activation_function = 'relu'
dropout_rate = 0.25

model, history, duration = train_model(
    activation_func=activation_function,
    hidden_neurons=hidden_neurons,
    optimizer=Adam,
    loss_func=CategoricalCrossentropy,
    epochs=100,
    dropout_rate=dropout_rate,
    learning_rate=0.001
)

test_loss, test_acc = model.evaluate(X_test,y_test)
```

```

print(f"Test loss={test_loss} Test accuracy = {test_acc}")
print(f"Time taken to train the model is {duration} seconds")

result_df_dropout.loc[len(result_df_dropout.index)] = [
    len(hidden_neurons),
    activation_function,
    str(hidden_neurons),
    dropout_rate,
    duration,
    test_loss,
    test_acc
]

plot_metrics(history, activation_function, hidden_neurons)

```














Model: "sequential_14"

Layer (type)	Output Shape	Param #
flatten_14 (Flatten)	(None, 1024)	0
dense_48 (Dense)	(None, 16)	16,400
dropout_9 (Dropout)	(None, 16)	0
dense_49 (Dense)	(None, 32)	544
dropout_10 (Dropout)	(None, 32)	0
dense_50 (Dense)	(None, 64)	2,112
dropout_11 (Dropout)	(None, 64)	0
dense_51 (Dense)	(None, 10)	650

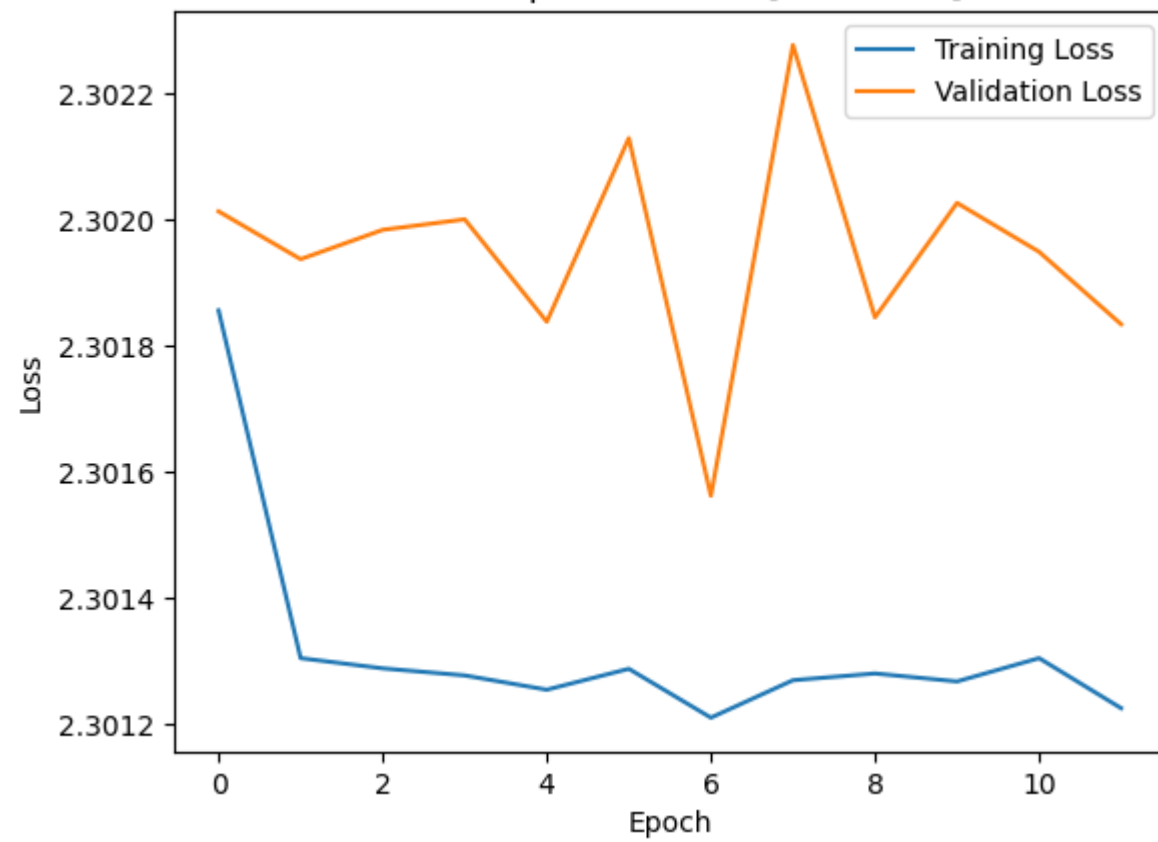
Total params: 19,706 (76.98 KB)

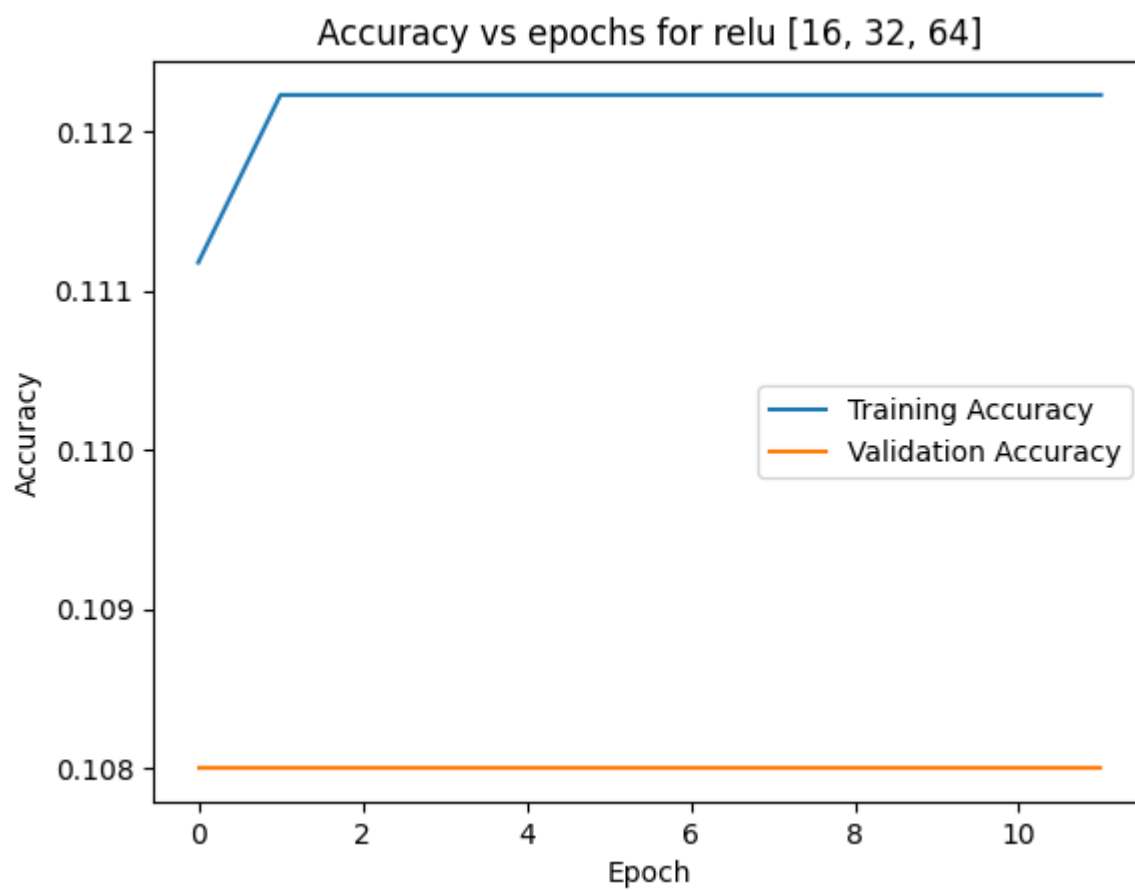
Trainable params: 19,706 (76.98 KB)

Non-trainable params: 0 (0.00 B)

Epoch 1/100
1750/1750  6s 2ms/step - accuracy: 0.1099 - loss: 2.3045 - val_accuracy: 0.1080 - val_loss: 2.3020
Epoch 2/100
1750/1750  6s 3ms/step - accuracy: 0.1137 - loss: 2.3010 - val_accuracy: 0.1080 - val_loss: 2.3019
Epoch 3/100
1750/1750  4s 2ms/step - accuracy: 0.1126 - loss: 2.3008 - val_accuracy: 0.1080 - val_loss: 2.3020
Epoch 4/100
1750/1750  4s 2ms/step - accuracy: 0.1137 - loss: 2.3011 - val_accuracy: 0.1080 - val_loss: 2.3020
Epoch 5/100
1750/1750  6s 3ms/step - accuracy: 0.1100 - loss: 2.3015 - val_accuracy: 0.1080 - val_loss: 2.3018
Epoch 6/100
1750/1750  4s 2ms/step - accuracy: 0.1107 - loss: 2.3015 - val_accuracy: 0.1080 - val_loss: 2.3021
Epoch 7/100
1750/1750  5s 2ms/step - accuracy: 0.1142 - loss: 2.3010 - val_accuracy: 0.1080 - val_loss: 2.3016
Epoch 8/100
1750/1750  6s 3ms/step - accuracy: 0.1111 - loss: 2.3014 - val_accuracy: 0.1080 - val_loss: 2.3023
Epoch 9/100
1750/1750  4s 2ms/step - accuracy: 0.1136 - loss: 2.3014 - val_accuracy: 0.1080 - val_loss: 2.3018
Epoch 10/100
1750/1750  6s 2ms/step - accuracy: 0.1127 - loss: 2.3013 - val_accuracy: 0.1080 - val_loss: 2.3020
Epoch 11/100
1750/1750  5s 3ms/step - accuracy: 0.1120 - loss: 2.3012 - val_accuracy: 0.1080 - val_loss: 2.3019
Epoch 12/100
1750/1750  4s 2ms/step - accuracy: 0.1126 - loss: 2.3013 - val_accuracy: 0.1080 - val_loss: 2.3018
219/219  0s 1ms/step - accuracy: 0.1219 - loss: 2.3009
Test loss=2.3010637760162354 Test accuracy = 0.11942857503890991
Time taken to train the model is 59.987740993499756 seconds

Loss vs epochs for relu [16, 32, 64]





In [40]: `## Model 5 --> Adam, lr =0.001, CCE loss, hidden_neurons=[16,32,64]. relu, dropout=0.1`

```
hidden_neurons = [16,32,64]
activation_function = 'relu'
dropout_rate = 0.1

model, history, duration = train_model(
    activation_func=activation_function,
    hidden_neurons=hidden_neurons,
    optimizer=Adam,
    loss_func=CategoricalCrossentropy,
    epochs=100,
    dropout_rate=dropout_rate,
    learning_rate=0.001
)

test_loss, test_acc = model.evaluate(X_test,y_test)
```

```

print(f"Test loss={test_loss} Test accuracy = {test_acc}")
print(f"Time taken to train the model is {duration} seconds")

result_df_dropout.loc[len(result_df_dropout.index)] = [
    len(hidden_neurons),
    activation_function,
    str(hidden_neurons),
    dropout_rate,
    duration,
    test_loss,
    test_acc
]

plot_metrics(history, activation_function, hidden_neurons)

```

Model: "sequential_15"

Layer (type)	Output Shape	Param #
flatten_15 (Flatten)	(None, 1024)	0
dense_52 (Dense)	(None, 16)	16,400
dropout_12 (Dropout)	(None, 16)	0
dense_53 (Dense)	(None, 32)	544
dropout_13 (Dropout)	(None, 32)	0
dense_54 (Dense)	(None, 64)	2,112
dropout_14 (Dropout)	(None, 64)	0
dense_55 (Dense)	(None, 10)	650

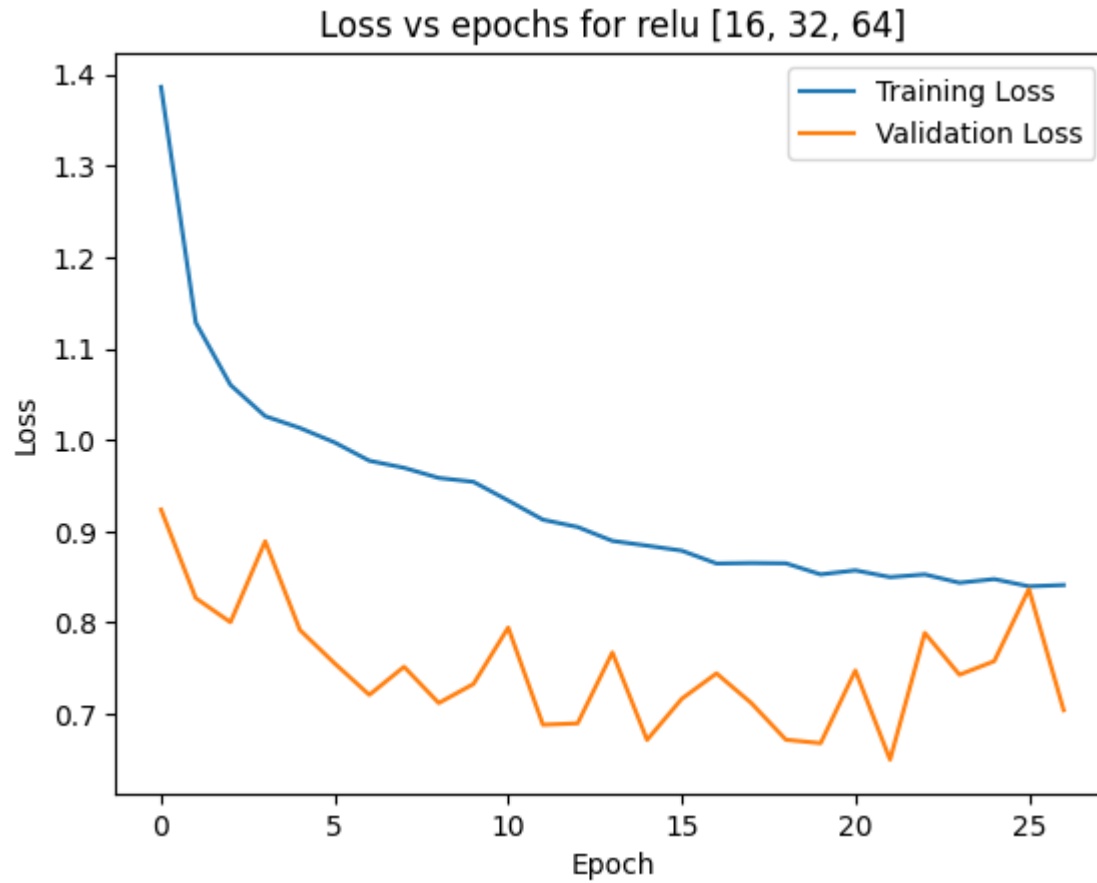
Total params: 19,706 (76.98 KB)

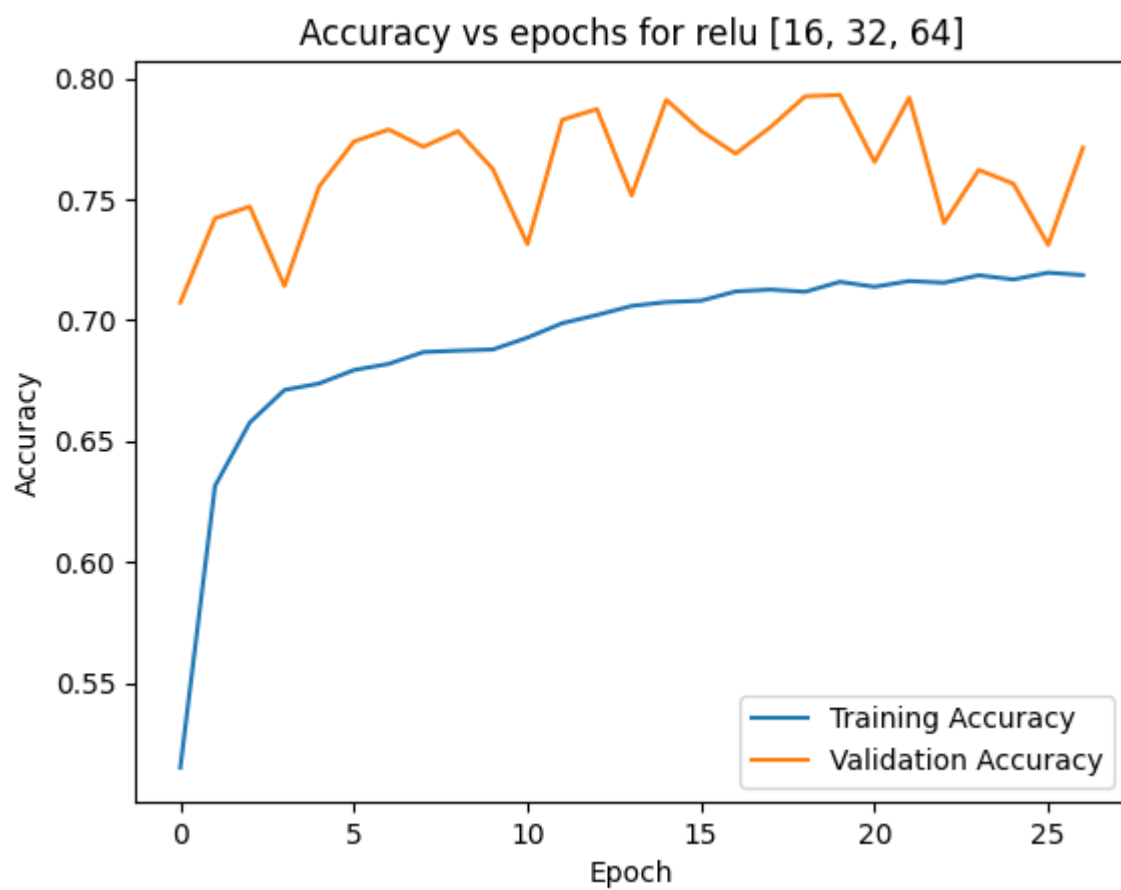
Trainable params: 19,706 (76.98 KB)

Non-trainable params: 0 (0.00 B)

Epoch 1/100	1750/1750	7s	3ms/step	- accuracy: 0.4060	- loss: 1.6379	- val_accuracy: 0.7073	- val_loss: 0.9236
Epoch 2/100	1750/1750	9s	2ms/step	- accuracy: 0.6211	- loss: 1.1505	- val_accuracy: 0.7421	- val_loss: 0.8267
Epoch 3/100	1750/1750	5s	3ms/step	- accuracy: 0.6544	- loss: 1.0685	- val_accuracy: 0.7470	- val_loss: 0.8004
Epoch 4/100	1750/1750	4s	2ms/step	- accuracy: 0.6727	- loss: 1.0284	- val_accuracy: 0.7141	- val_loss: 0.8889
Epoch 5/100	1750/1750	5s	2ms/step	- accuracy: 0.6723	- loss: 1.0184	- val_accuracy: 0.7553	- val_loss: 0.7918
Epoch 6/100	1750/1750	6s	3ms/step	- accuracy: 0.6791	- loss: 0.9962	- val_accuracy: 0.7739	- val_loss: 0.7553
Epoch 7/100	1750/1750	4s	2ms/step	- accuracy: 0.6836	- loss: 0.9714	- val_accuracy: 0.7789	- val_loss: 0.7208
Epoch 8/100	1750/1750	4s	2ms/step	- accuracy: 0.6905	- loss: 0.9607	- val_accuracy: 0.7717	- val_loss: 0.7515
Epoch 9/100	1750/1750	6s	2ms/step	- accuracy: 0.6872	- loss: 0.9594	- val_accuracy: 0.7781	- val_loss: 0.7119
Epoch 10/100	1750/1750	4s	2ms/step	- accuracy: 0.6858	- loss: 0.9545	- val_accuracy: 0.7626	- val_loss: 0.7327
Epoch 11/100	1750/1750	7s	3ms/step	- accuracy: 0.6922	- loss: 0.9382	- val_accuracy: 0.7314	- val_loss: 0.7945
Epoch 12/100	1750/1750	4s	2ms/step	- accuracy: 0.6985	- loss: 0.9171	- val_accuracy: 0.7829	- val_loss: 0.6881
Epoch 13/100	1750/1750	5s	2ms/step	- accuracy: 0.7034	- loss: 0.8968	- val_accuracy: 0.7873	- val_loss: 0.6896
Epoch 14/100	1750/1750	6s	3ms/step	- accuracy: 0.7059	- loss: 0.8926	- val_accuracy: 0.7516	- val_loss: 0.7672
Epoch 15/100	1750/1750	4s	2ms/step	- accuracy: 0.7079	- loss: 0.8829	- val_accuracy: 0.7911	- val_loss: 0.6715
Epoch 16/100	1750/1750	5s	2ms/step	- accuracy: 0.7068	- loss: 0.8835	- val_accuracy: 0.7783	- val_loss: 0.7165
Epoch 17/100	1750/1750	5s	2ms/step	- accuracy: 0.7160	- loss: 0.8552	- val_accuracy: 0.7689	- val_loss: 0.7443
Epoch 18/100	1750/1750	5s	2ms/step	- accuracy: 0.7110	- loss: 0.8725	- val_accuracy: 0.7799	- val_loss: 0.7120
Epoch 19/100	1750/1750	5s	3ms/step	- accuracy: 0.7116	- loss: 0.8656	- val_accuracy: 0.7926	- val_loss: 0.6717
Epoch 20/100	1750/1750	4s	2ms/step	- accuracy: 0.7148	- loss: 0.8525	- val_accuracy: 0.7931	- val_loss: 0.6679
Epoch 21/100	1750/1750	5s	2ms/step	- accuracy: 0.7135	- loss: 0.8574	- val_accuracy: 0.7654	- val_loss: 0.7475
Epoch 22/100	1750/1750	6s	3ms/step	- accuracy: 0.7156	- loss: 0.8492	- val_accuracy: 0.7920	- val_loss: 0.6498
Epoch 23/100	1750/1750	4s	2ms/step	- accuracy: 0.7127	- loss: 0.8636	- val_accuracy: 0.7401	- val_loss: 0.7886

Epoch 24/100 **1750/1750** 4s 2ms/step - accuracy: 0.7181 - loss: 0.8411 - val_accuracy: 0.7621 - val_loss: 0.7428
Epoch 25/100 **1750/1750** 6s 3ms/step - accuracy: 0.7129 - loss: 0.8549 - val_accuracy: 0.7564 - val_loss: 0.7577
Epoch 26/100 **1750/1750** 4s 2ms/step - accuracy: 0.7206 - loss: 0.8388 - val_accuracy: 0.7311 - val_loss: 0.8366
Epoch 27/100 **1750/1750** 6s 3ms/step - accuracy: 0.7158 - loss: 0.8482 - val_accuracy: 0.7713 - val_loss: 0.7041
219/219 0s 2ms/step - accuracy: 0.7998 - loss: 0.6454
Test loss=0.6585350632667542 Test accuracy = 0.8001428842544556
Time taken to train the model is 138.4717516899109 seconds





In [41]: result_df_dropout

Out[41]:	Hidden Layers	Activation Function	Hidden Neurons	Dropout Rate	Training Time (in seconds)	Test Loss	Test Accuracy
0	3	relu	[16, 32, 64]	0.90	37.405670	2.300971	0.119429
1	3	relu	[16, 32, 64]	0.75	61.329699	2.300853	0.119429
2	3	relu	[16, 32, 64]	0.50	51.372198	2.301025	0.119429
3	3	relu	[16, 32, 64]	0.25	59.987741	2.301064	0.119429
4	3	relu	[16, 32, 64]	0.10	138.471752	0.658535	0.800143

The best dropout rate is 0.1

Question 07

- vii. Plot the graph for loss vs epoch and accuracy (train and validation accuracy) vs epoch for all the above cases. Point out the logic in the report.

Logic

1. As we can see from the above reports for less number of hidden layers, the model performs better since the training dataset is of small size and the number of trainable parameters for more neurons in each layer exceeds the expectation of the dataset. So less neurons and less layers is the way to go for small datasets like MNIST.
2. Sigmoid performs better than tanh and relu, although if dataset size is increased, we will more likely see that relu performs better, owing to its ability to prevent over-fitting.
3. Adam performs better optimization(smooth decrease in loss) as compared to SGD.
4. CategoricalCrossentropy loss is appropriate as compared to MSE since this is a classification task not regression.
5. Low dropout rate is favourable for small datasets.
6. For tanh and relu, early stopping occurs leading to fast convergence due to small dataset but lower accuracy than sigmoid due to early stopping criterion

Question 08

- viii. With the best set hyper-parameters from above run vary the Adam Optimizer learning rate {0.01, 0.001, 0.005, 0.0001, 0.0005}. Print the time to achieve the best validation accuracy (as reported before from all runs) from all these five run.

```
In [53]: best_dropout_rate = 0.1
best_activation_function = 'sigmoid'
hidden_neurons = [16,32,64]

learning_rates = [0.01,0.001,0.005,0.0001,0.0005]

result_df_lr = pd.DataFrame(
    columns=[
        'Hidden Layers',
```

```

        'Hidden Neurons',
        'Learning Rate',
        'Dropout Rate',
        'Activation Function',
        'Training Time (in seconds)',
        'Test Loss',
        'Test Accuracy'
    ]
)

for lr in learning_rates:
    model, history, duration = train_model(
        hidden_neurons=hidden_neurons,
        learning_rate=lr,
        dropout_rate=best_dropout_rate,
        activation_func=best_activation_function,
        loss_func=CategoricalCrossentropy,
        optimizer=Adam,
        epochs=100
    )

    test_loss, test_acc = model.evaluate(X_test,y_test)

    print(f"Learning rate = {lr}")
    print(f"Test loss={test_loss} Test accuracy = {test_acc}")
    print(f"Time taken to train the model is {duration} seconds")

    result_df_lr.loc[len(result_df_lr.index)] = [
        len(hidden_neurons),
        str(hidden_neurons),
        lr,
        best_dropout_rate,
        best_activation_function,
        duration,
        test_loss,
        test_acc
    ]

    plot_metrics(history, activation_function, hidden_neurons)

```

Model: "sequential_22"

Layer (type)	Output Shape	Param #
flatten_22 (Flatten)	(None, 1024)	0
dense_74 (Dense)	(None, 16)	16,400
dropout_15 (Dropout)	(None, 16)	0
dense_75 (Dense)	(None, 32)	544
dropout_16 (Dropout)	(None, 32)	0
dense_76 (Dense)	(None, 64)	2,112
dropout_17 (Dropout)	(None, 64)	0
dense_77 (Dense)	(None, 10)	650

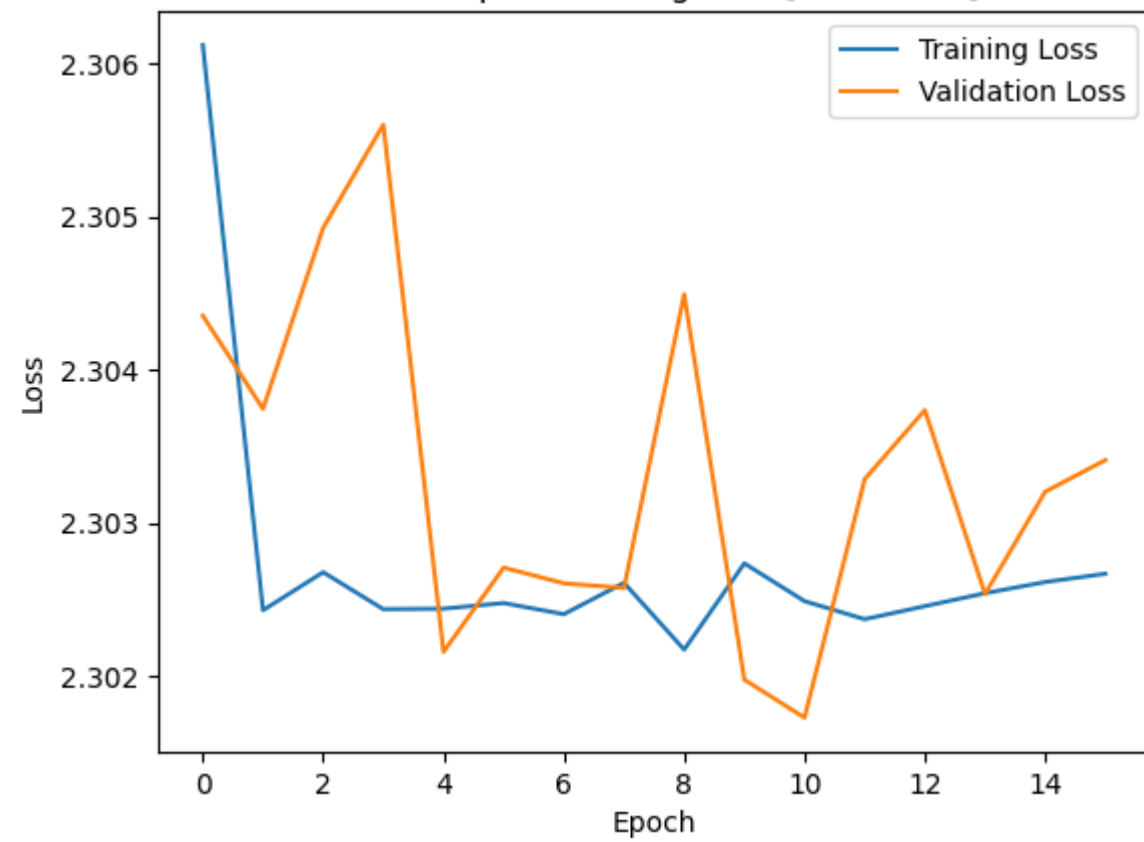
Total params: 19,706 (76.98 KB)

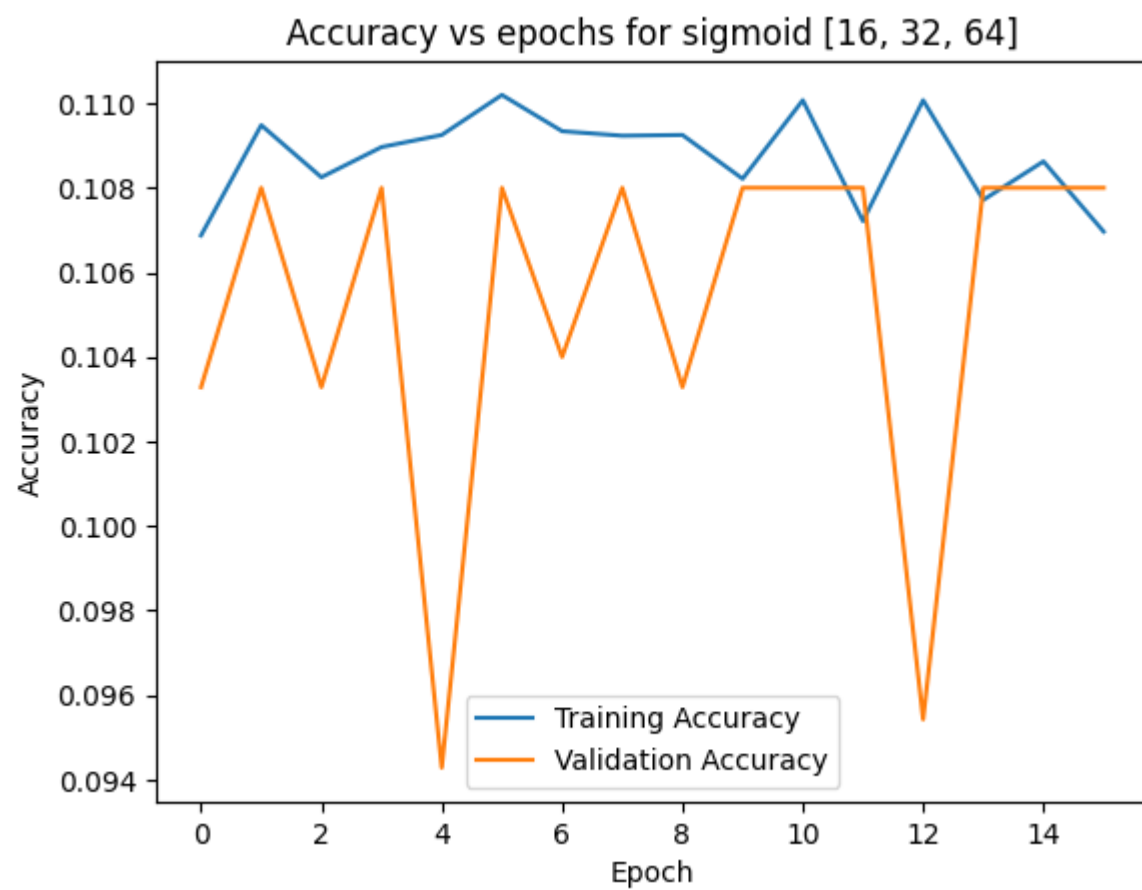
Trainable params: 19,706 (76.98 KB)

Non-trainable params: 0 (0.00 B)

Epoch 1/100
1750/1750 ————— 5s 2ms/step - accuracy: 0.1033 - loss: 2.3183 - val_accuracy: 0.1033 - val_loss: 2.3044
Epoch 2/100
1750/1750 ————— 5s 3ms/step - accuracy: 0.1074 - loss: 2.3025 - val_accuracy: 0.1080 - val_loss: 2.3037
Epoch 3/100
1750/1750 ————— 4s 2ms/step - accuracy: 0.1110 - loss: 2.3022 - val_accuracy: 0.1033 - val_loss: 2.3049
Epoch 4/100
1750/1750 ————— 5s 2ms/step - accuracy: 0.1100 - loss: 2.3025 - val_accuracy: 0.1080 - val_loss: 2.3056
Epoch 5/100
1750/1750 ————— 6s 3ms/step - accuracy: 0.1096 - loss: 2.3026 - val_accuracy: 0.0943 - val_loss: 2.3022
Epoch 6/100
1750/1750 ————— 4s 2ms/step - accuracy: 0.1111 - loss: 2.3021 - val_accuracy: 0.1080 - val_loss: 2.3027
Epoch 7/100
1750/1750 ————— 6s 2ms/step - accuracy: 0.1097 - loss: 2.3024 - val_accuracy: 0.1040 - val_loss: 2.3026
Epoch 8/100
1750/1750 ————— 5s 2ms/step - accuracy: 0.1071 - loss: 2.3027 - val_accuracy: 0.1080 - val_loss: 2.3026
Epoch 9/100
1750/1750 ————— 4s 2ms/step - accuracy: 0.1062 - loss: 2.3025 - val_accuracy: 0.1033 - val_loss: 2.3045
Epoch 10/100
1750/1750 ————— 4s 2ms/step - accuracy: 0.1064 - loss: 2.3029 - val_accuracy: 0.1080 - val_loss: 2.3020
Epoch 11/100
1750/1750 ————— 5s 3ms/step - accuracy: 0.1110 - loss: 2.3023 - val_accuracy: 0.1080 - val_loss: 2.3017
Epoch 12/100
1750/1750 ————— 4s 2ms/step - accuracy: 0.1103 - loss: 2.3021 - val_accuracy: 0.1080 - val_loss: 2.3033
Epoch 13/100
1750/1750 ————— 4s 2ms/step - accuracy: 0.1116 - loss: 2.3024 - val_accuracy: 0.0954 - val_loss: 2.3037
Epoch 14/100
1750/1750 ————— 5s 3ms/step - accuracy: 0.1070 - loss: 2.3026 - val_accuracy: 0.1080 - val_loss: 2.3025
Epoch 15/100
1750/1750 ————— 4s 2ms/step - accuracy: 0.1080 - loss: 2.3027 - val_accuracy: 0.1080 - val_loss: 2.3032
Epoch 16/100
1750/1750 ————— 4s 2ms/step - accuracy: 0.1093 - loss: 2.3024 - val_accuracy: 0.1080 - val_loss: 2.3034
219/219 ————— 0s 1ms/step - accuracy: 0.1219 - loss: 2.3007
Learning rate = 0.01
Test loss=2.301191568374634 Test accuracy = 0.11942857503890991
Time taken to train the model is 73.58484530448914 seconds

Loss vs epochs for sigmoid [16, 32, 64]





Model: "sequential_23"















Layer (type)	Output Shape	Param #
flatten_23 (Flatten)	(None, 1024)	0
dense_78 (Dense)	(None, 16)	16,400
dropout_18 (Dropout)	(None, 16)	0
dense_79 (Dense)	(None, 32)	544
dropout_19 (Dropout)	(None, 32)	0
dense_80 (Dense)	(None, 64)	2,112
dropout_20 (Dropout)	(None, 64)	0
dense_81 (Dense)	(None, 10)	650

Total params: 19,706 (76.98 KB)

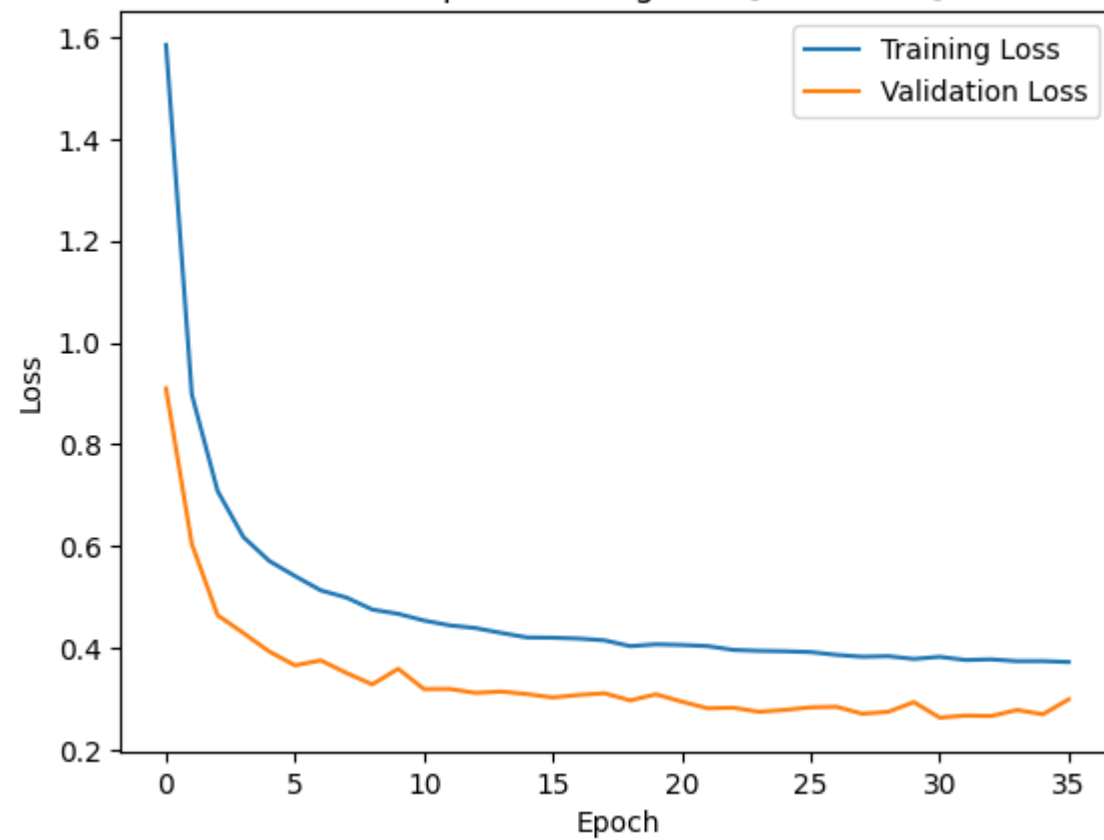
Trainable params: 19,706 (76.98 KB)

Non-trainable params: 0 (0.00 B)

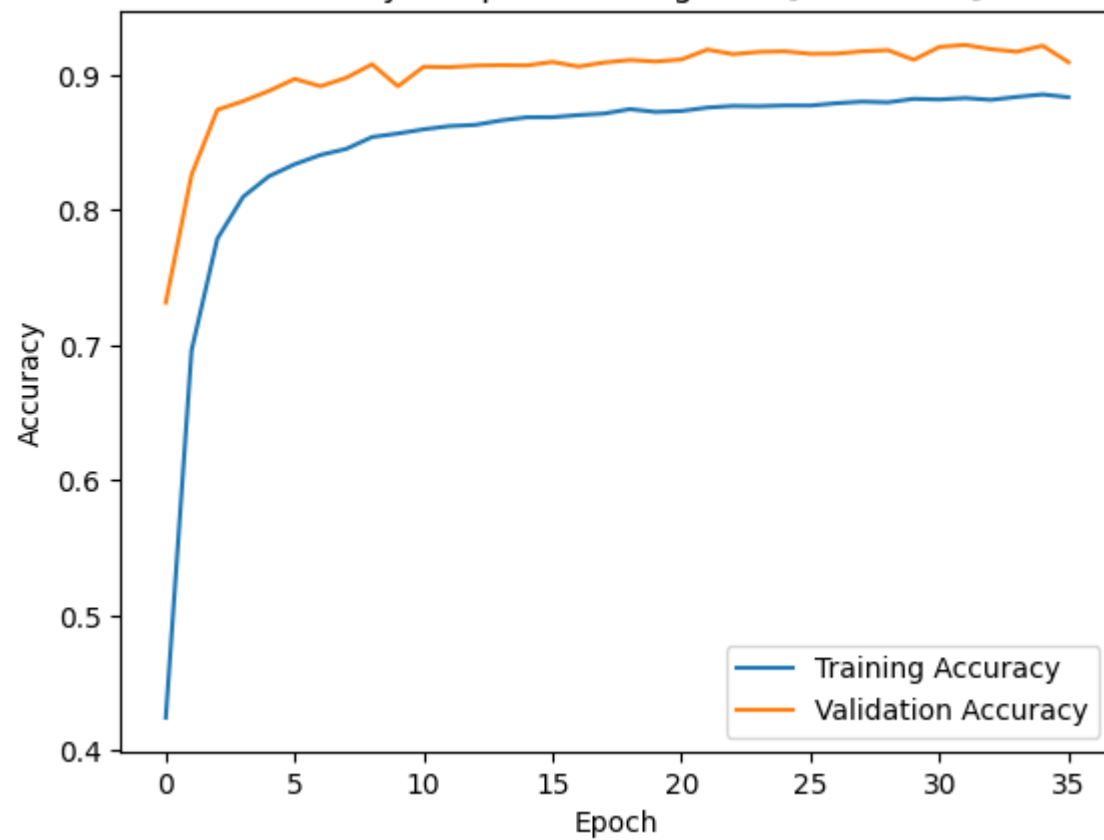
Epoch 1/100	1750/1750	6s	3ms/step	- accuracy: 0.2745	- loss: 1.9652	- val_accuracy: 0.7316	- val_loss: 0.9098
Epoch 2/100	1750/1750	4s	2ms/step	- accuracy: 0.6651	- loss: 0.9659	- val_accuracy: 0.8261	- val_loss: 0.6046
Epoch 3/100	1750/1750	5s	3ms/step	- accuracy: 0.7695	- loss: 0.7357	- val_accuracy: 0.8741	- val_loss: 0.4648
Epoch 4/100	1750/1750	5s	3ms/step	- accuracy: 0.8049	- loss: 0.6330	- val_accuracy: 0.8806	- val_loss: 0.4299
Epoch 5/100	1750/1750	4s	2ms/step	- accuracy: 0.8249	- loss: 0.5765	- val_accuracy: 0.8881	- val_loss: 0.3937
Epoch 6/100	1750/1750	4s	2ms/step	- accuracy: 0.8315	- loss: 0.5479	- val_accuracy: 0.8970	- val_loss: 0.3664
Epoch 7/100	1750/1750	5s	3ms/step	- accuracy: 0.8389	- loss: 0.5166	- val_accuracy: 0.8916	- val_loss: 0.3762
Epoch 8/100	1750/1750	4s	2ms/step	- accuracy: 0.8428	- loss: 0.5054	- val_accuracy: 0.8979	- val_loss: 0.3514
Epoch 9/100	1750/1750	4s	2ms/step	- accuracy: 0.8500	- loss: 0.4850	- val_accuracy: 0.9079	- val_loss: 0.3290
Epoch 10/100	1750/1750	5s	3ms/step	- accuracy: 0.8574	- loss: 0.4618	- val_accuracy: 0.8916	- val_loss: 0.3597
Epoch 11/100	1750/1750	4s	2ms/step	- accuracy: 0.8597	- loss: 0.4482	- val_accuracy: 0.9060	- val_loss: 0.3199
Epoch 12/100	1750/1750	4s	2ms/step	- accuracy: 0.8637	- loss: 0.4390	- val_accuracy: 0.9057	- val_loss: 0.3201
Epoch 13/100	1750/1750	7s	3ms/step	- accuracy: 0.8621	- loss: 0.4413	- val_accuracy: 0.9069	- val_loss: 0.3123
Epoch 14/100	1750/1750	4s	2ms/step	- accuracy: 0.8655	- loss: 0.4300	- val_accuracy: 0.9071	- val_loss: 0.3150
Epoch 15/100	1750/1750	5s	2ms/step	- accuracy: 0.8708	- loss: 0.4170	- val_accuracy: 0.9070	- val_loss: 0.3100
Epoch 16/100	1750/1750	5s	3ms/step	- accuracy: 0.8706	- loss: 0.4131	- val_accuracy: 0.9096	- val_loss: 0.3034
Epoch 17/100	1750/1750	4s	2ms/step	- accuracy: 0.8711	- loss: 0.4192	- val_accuracy: 0.9061	- val_loss: 0.3083
Epoch 18/100	1750/1750	4s	2ms/step	- accuracy: 0.8699	- loss: 0.4205	- val_accuracy: 0.9091	- val_loss: 0.3117
Epoch 19/100	1750/1750	7s	3ms/step	- accuracy: 0.8740	- loss: 0.4013	- val_accuracy: 0.9110	- val_loss: 0.2975
Epoch 20/100	1750/1750	9s	2ms/step	- accuracy: 0.8728	- loss: 0.4069	- val_accuracy: 0.9099	- val_loss: 0.3096
Epoch 21/100	1750/1750	5s	3ms/step	- accuracy: 0.8740	- loss: 0.4061	- val_accuracy: 0.9114	- val_loss: 0.2951
Epoch 22/100	1750/1750	9s	2ms/step	- accuracy: 0.8774	- loss: 0.3996	- val_accuracy: 0.9187	- val_loss: 0.2821
Epoch 23/100	1750/1750	7s	3ms/step	- accuracy: 0.8766	- loss: 0.4017	- val_accuracy: 0.9153	- val_loss: 0.2832

Epoch 24/100
1750/1750  9s 2ms/step - accuracy: 0.8753 - loss: 0.3954 - val_accuracy: 0.9170 - val_loss: 0.2752
Epoch 25/100
1750/1750  6s 3ms/step - accuracy: 0.8738 - loss: 0.4073 - val_accuracy: 0.9174 - val_loss: 0.2789
Epoch 26/100
1750/1750  4s 2ms/step - accuracy: 0.8786 - loss: 0.3915 - val_accuracy: 0.9156 - val_loss: 0.2840
Epoch 27/100
1750/1750  6s 3ms/step - accuracy: 0.8803 - loss: 0.3831 - val_accuracy: 0.9157 - val_loss: 0.2850
Epoch 28/100
1750/1750  5s 3ms/step - accuracy: 0.8795 - loss: 0.3879 - val_accuracy: 0.9174 - val_loss: 0.2714
Epoch 29/100
1750/1750  4s 2ms/step - accuracy: 0.8771 - loss: 0.3892 - val_accuracy: 0.9183 - val_loss: 0.2752
Epoch 30/100
1750/1750  6s 3ms/step - accuracy: 0.8819 - loss: 0.3805 - val_accuracy: 0.9111 - val_loss: 0.2942
Epoch 31/100
1750/1750  4s 2ms/step - accuracy: 0.8829 - loss: 0.3784 - val_accuracy: 0.9207 - val_loss: 0.2637
Epoch 32/100
1750/1750  4s 2ms/step - accuracy: 0.8824 - loss: 0.3797 - val_accuracy: 0.9221 - val_loss: 0.2678
Epoch 33/100
1750/1750  5s 3ms/step - accuracy: 0.8801 - loss: 0.3805 - val_accuracy: 0.9190 - val_loss: 0.2670
Epoch 34/100
1750/1750  5s 3ms/step - accuracy: 0.8848 - loss: 0.3703 - val_accuracy: 0.9171 - val_loss: 0.2788
Epoch 35/100
1750/1750  4s 2ms/step - accuracy: 0.8864 - loss: 0.3721 - val_accuracy: 0.9216 - val_loss: 0.2703
Epoch 36/100
1750/1750  6s 3ms/step - accuracy: 0.8823 - loss: 0.3753 - val_accuracy: 0.9094 - val_loss: 0.2998
219/219  0s 2ms/step - accuracy: 0.9215 - loss: 0.2569
Learning rate = 0.001
Test loss=0.2631760239601135 Test accuracy = 0.920285701751709
Time taken to train the model is 186.7718403339386 seconds

Loss vs epochs for sigmoid [16, 32, 64]



Accuracy vs epochs for sigmoid [16, 32, 64]



Model: "sequential_24"

Layer (type)	Output Shape	Param #
flatten_24 (Flatten)	(None, 1024)	0
dense_82 (Dense)	(None, 16)	16,400
dropout_21 (Dropout)	(None, 16)	0
dense_83 (Dense)	(None, 32)	544
dropout_22 (Dropout)	(None, 32)	0
dense_84 (Dense)	(None, 64)	2,112
dropout_23 (Dropout)	(None, 64)	0
dense_85 (Dense)	(None, 10)	650

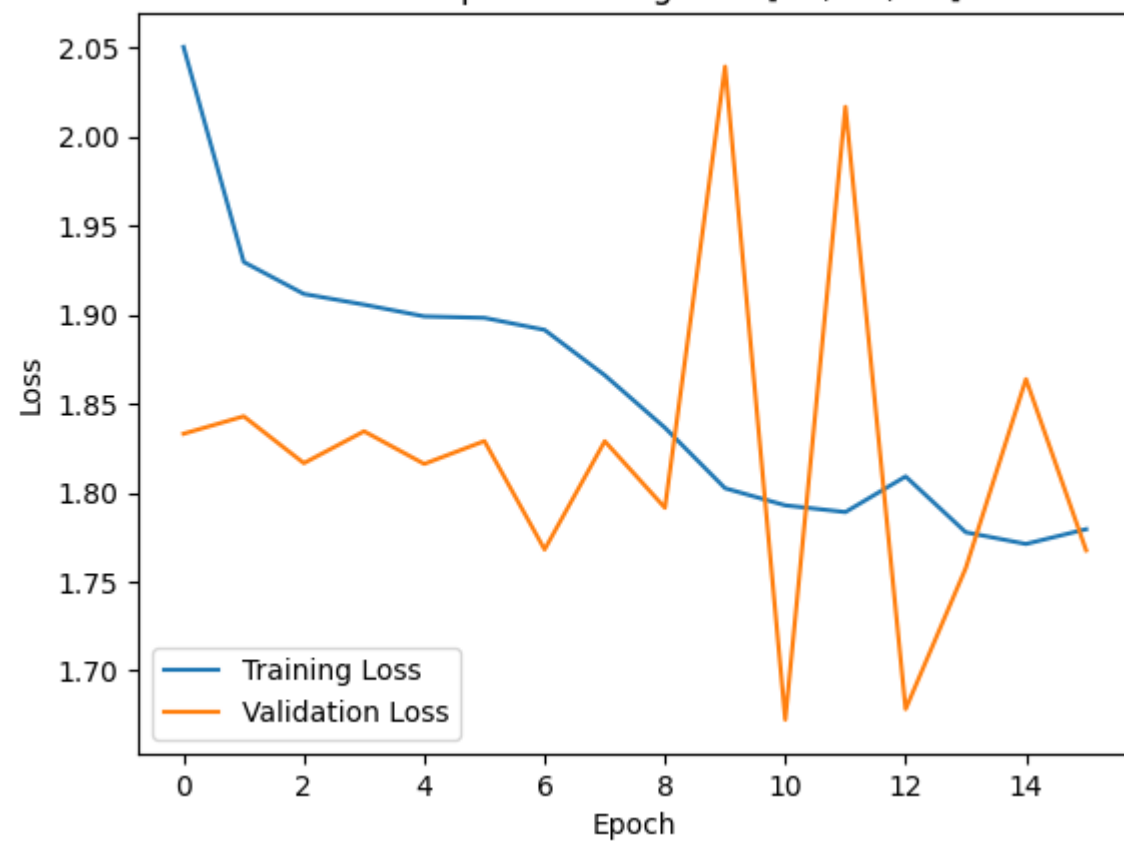
Total params: 19,706 (76.98 KB)

Trainable params: 19,706 (76.98 KB)

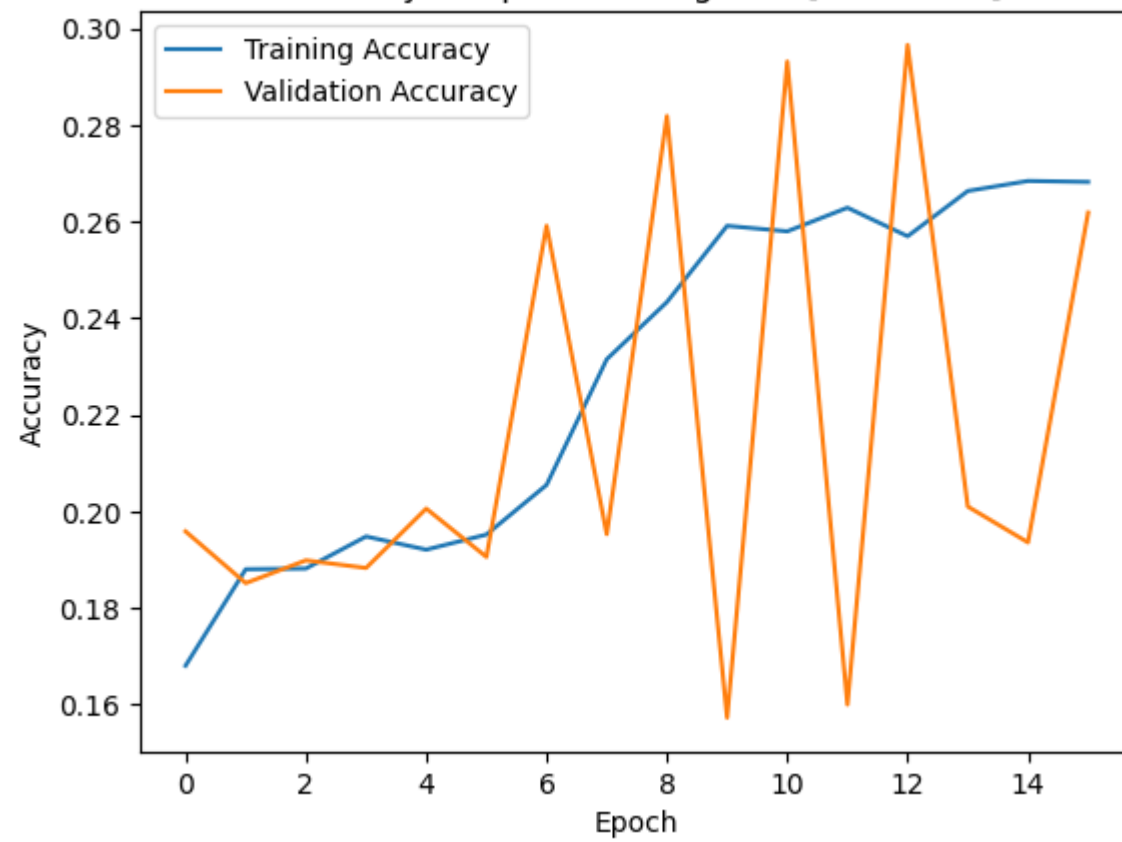
Non-trainable params: 0 (0.00 B)

Epoch 1/100
1750/1750 ————— 6s 2ms/step - accuracy: 0.1384 - loss: 2.1863 - val_accuracy: 0.1959 - val_loss: 1.8331
Epoch 2/100
1750/1750 ————— 6s 3ms/step - accuracy: 0.1881 - loss: 1.9369 - val_accuracy: 0.1851 - val_loss: 1.8429
Epoch 3/100
1750/1750 ————— 5s 2ms/step - accuracy: 0.1908 - loss: 1.9149 - val_accuracy: 0.1899 - val_loss: 1.8165
Epoch 4/100
1750/1750 ————— 4s 2ms/step - accuracy: 0.1948 - loss: 1.9022 - val_accuracy: 0.1883 - val_loss: 1.8345
Epoch 5/100
1750/1750 ————— 5s 3ms/step - accuracy: 0.1919 - loss: 1.8962 - val_accuracy: 0.2006 - val_loss: 1.8161
Epoch 6/100
1750/1750 ————— 5s 3ms/step - accuracy: 0.1952 - loss: 1.8931 - val_accuracy: 0.1906 - val_loss: 1.8290
Epoch 7/100
1750/1750 ————— 4s 2ms/step - accuracy: 0.2011 - loss: 1.8921 - val_accuracy: 0.2591 - val_loss: 1.7680
Epoch 8/100
1750/1750 ————— 4s 3ms/step - accuracy: 0.2250 - loss: 1.8747 - val_accuracy: 0.1953 - val_loss: 1.8289
Epoch 9/100
1750/1750 ————— 5s 3ms/step - accuracy: 0.2396 - loss: 1.8514 - val_accuracy: 0.2819 - val_loss: 1.7915
Epoch 10/100
1750/1750 ————— 4s 2ms/step - accuracy: 0.2550 - loss: 1.8020 - val_accuracy: 0.1573 - val_loss: 2.0394
Epoch 11/100
1750/1750 ————— 4s 2ms/step - accuracy: 0.2573 - loss: 1.7974 - val_accuracy: 0.2931 - val_loss: 1.6724
Epoch 12/100
1750/1750 ————— 5s 3ms/step - accuracy: 0.2638 - loss: 1.7868 - val_accuracy: 0.1600 - val_loss: 2.0168
Epoch 13/100
1750/1750 ————— 4s 2ms/step - accuracy: 0.2562 - loss: 1.8041 - val_accuracy: 0.2966 - val_loss: 1.6783
Epoch 14/100
1750/1750 ————— 6s 3ms/step - accuracy: 0.2663 - loss: 1.7814 - val_accuracy: 0.2010 - val_loss: 1.7576
Epoch 15/100
1750/1750 ————— 5s 3ms/step - accuracy: 0.2718 - loss: 1.7680 - val_accuracy: 0.1936 - val_loss: 1.8638
Epoch 16/100
1750/1750 ————— 4s 2ms/step - accuracy: 0.2632 - loss: 1.7932 - val_accuracy: 0.2619 - val_loss: 1.7675
219/219 ————— 0s 1ms/step - accuracy: 0.3042 - loss: 1.6644
Learning rate = 0.005
Test loss=1.6627888679504395 Test accuracy = 0.29671427607536316
Time taken to train the model is 75.20755457878113 seconds

Loss vs epochs for sigmoid [16, 32, 64]



Accuracy vs epochs for sigmoid [16, 32, 64]



Model: "sequential_25"

Layer (type)	Output Shape	Param #
flatten_25 (Flatten)	(None, 1024)	0
dense_86 (Dense)	(None, 16)	16,400
dropout_24 (Dropout)	(None, 16)	0
dense_87 (Dense)	(None, 32)	544
dropout_25 (Dropout)	(None, 32)	0
dense_88 (Dense)	(None, 64)	2,112
dropout_26 (Dropout)	(None, 64)	0
dense_89 (Dense)	(None, 10)	650

Total params: 19,706 (76.98 KB)

Trainable params: 19,706 (76.98 KB)










Non-trainable params: 0 (0.00 B)

Epoch 1/100	1750/1750	7s	3ms/step	- accuracy: 0.1175	- loss: 2.3349	- val_accuracy: 0.2574	- val_loss: 2.1762
Epoch 2/100	1750/1750	8s	2ms/step	- accuracy: 0.2592	- loss: 2.1093	- val_accuracy: 0.4377	- val_loss: 1.8098
Epoch 3/100	1750/1750	5s	3ms/step	- accuracy: 0.3984	- loss: 1.7545	- val_accuracy: 0.5507	- val_loss: 1.4426
Epoch 4/100	1750/1750	9s	2ms/step	- accuracy: 0.5074	- loss: 1.4313	- val_accuracy: 0.6306	- val_loss: 1.1801
Epoch 5/100	1750/1750	7s	3ms/step	- accuracy: 0.5710	- loss: 1.2213	- val_accuracy: 0.7029	- val_loss: 1.0208
Epoch 6/100	1750/1750	9s	2ms/step	- accuracy: 0.6147	- loss: 1.0800	- val_accuracy: 0.7191	- val_loss: 0.9271
Epoch 7/100	1750/1750	5s	3ms/step	- accuracy: 0.6410	- loss: 1.0077	- val_accuracy: 0.7576	- val_loss: 0.8654
Epoch 8/100	1750/1750	9s	2ms/step	- accuracy: 0.6644	- loss: 0.9574	- val_accuracy: 0.7600	- val_loss: 0.8153
Epoch 9/100	1750/1750	5s	3ms/step	- accuracy: 0.6880	- loss: 0.9034	- val_accuracy: 0.7760	- val_loss: 0.7686
Epoch 10/100	1750/1750	4s	2ms/step	- accuracy: 0.7108	- loss: 0.8617	- val_accuracy: 0.7976	- val_loss: 0.7145
Epoch 11/100	1750/1750	5s	2ms/step	- accuracy: 0.7307	- loss: 0.8184	- val_accuracy: 0.8049	- val_loss: 0.6726
Epoch 12/100	1750/1750	6s	3ms/step	- accuracy: 0.7474	- loss: 0.7769	- val_accuracy: 0.8257	- val_loss: 0.6288
Epoch 13/100	1750/1750	4s	2ms/step	- accuracy: 0.7546	- loss: 0.7600	- val_accuracy: 0.8310	- val_loss: 0.6005
Epoch 14/100	1750/1750	5s	3ms/step	- accuracy: 0.7724	- loss: 0.7123	- val_accuracy: 0.8420	- val_loss: 0.5682
Epoch 15/100	1750/1750	5s	3ms/step	- accuracy: 0.7837	- loss: 0.6930	- val_accuracy: 0.8503	- val_loss: 0.5440
Epoch 16/100	1750/1750	4s	2ms/step	- accuracy: 0.7941	- loss: 0.6685	- val_accuracy: 0.8567	- val_loss: 0.5216
Epoch 17/100	1750/1750	6s	3ms/step	- accuracy: 0.7984	- loss: 0.6453	- val_accuracy: 0.8611	- val_loss: 0.5028
Epoch 18/100	1750/1750	4s	2ms/step	- accuracy: 0.8069	- loss: 0.6324	- val_accuracy: 0.8667	- val_loss: 0.4815
Epoch 19/100	1750/1750	5s	2ms/step	- accuracy: 0.8155	- loss: 0.6002	- val_accuracy: 0.8720	- val_loss: 0.4660
Epoch 20/100	1750/1750	6s	3ms/step	- accuracy: 0.8203	- loss: 0.5920	- val_accuracy: 0.8754	- val_loss: 0.4526
Epoch 21/100	1750/1750	9s	2ms/step	- accuracy: 0.8237	- loss: 0.5782	- val_accuracy: 0.8770	- val_loss: 0.4447
Epoch 22/100	1750/1750	7s	3ms/step	- accuracy: 0.8304	- loss: 0.5657	- val_accuracy: 0.8814	- val_loss: 0.4356
Epoch 23/100	1750/1750	9s	2ms/step	- accuracy: 0.8315	- loss: 0.5588	- val_accuracy: 0.8853	- val_loss: 0.4225

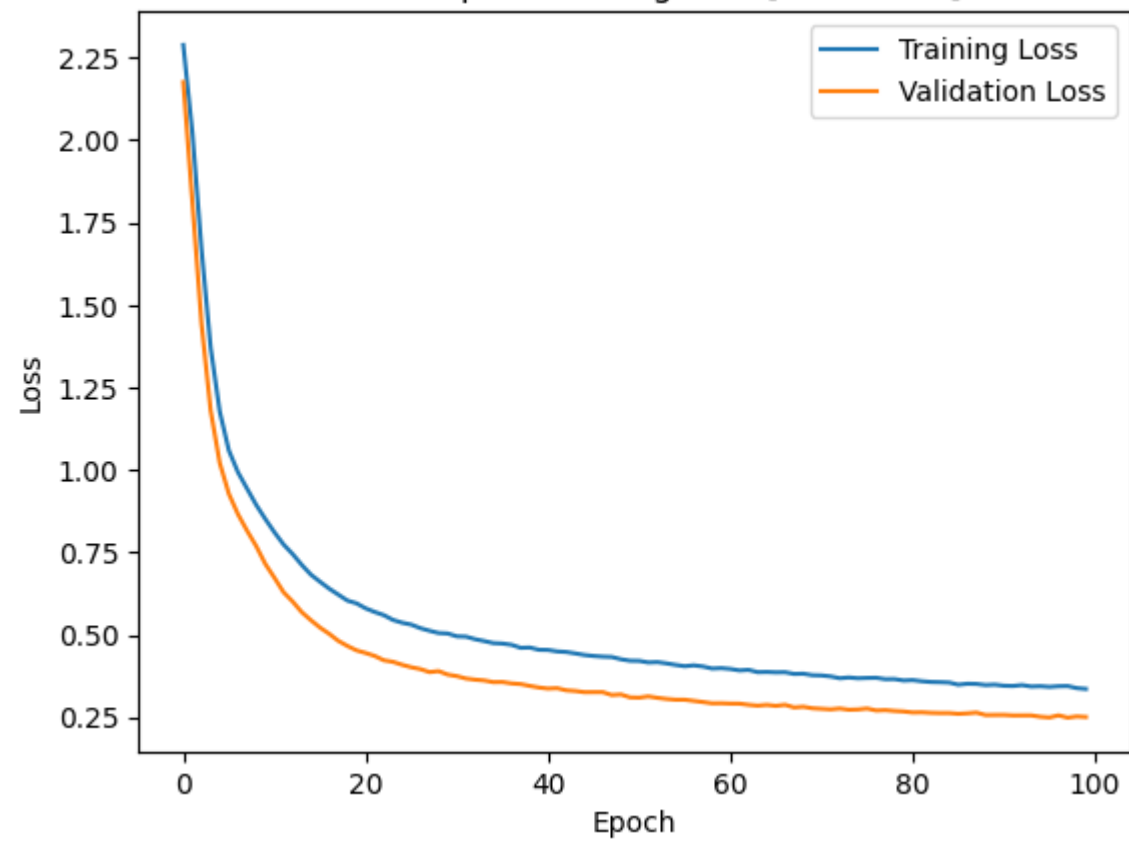
Epoch 24/100	1750/1750	5s	3ms/step	- accuracy: 0.8370	- loss: 0.5414	- val_accuracy: 0.8856	- val_loss: 0.4178
Epoch 25/100	1750/1750	4s	2ms/step	- accuracy: 0.8388	- loss: 0.5352	- val_accuracy: 0.8876	- val_loss: 0.4089
Epoch 26/100	1750/1750	5s	2ms/step	- accuracy: 0.8412	- loss: 0.5313	- val_accuracy: 0.8900	- val_loss: 0.4015
Epoch 27/100	1750/1750	5s	3ms/step	- accuracy: 0.8435	- loss: 0.5237	- val_accuracy: 0.8914	- val_loss: 0.3959
Epoch 28/100	1750/1750	4s	2ms/step	- accuracy: 0.8506	- loss: 0.5052	- val_accuracy: 0.8930	- val_loss: 0.3867
Epoch 29/100	1750/1750	4s	2ms/step	- accuracy: 0.8528	- loss: 0.4992	- val_accuracy: 0.8917	- val_loss: 0.3896
Epoch 30/100	1750/1750	6s	3ms/step	- accuracy: 0.8482	- loss: 0.5073	- val_accuracy: 0.8941	- val_loss: 0.3792
Epoch 31/100	1750/1750	4s	2ms/step	- accuracy: 0.8550	- loss: 0.4946	- val_accuracy: 0.8951	- val_loss: 0.3745
Epoch 32/100	1750/1750	6s	3ms/step	- accuracy: 0.8488	- loss: 0.5020	- val_accuracy: 0.8991	- val_loss: 0.3673
Epoch 33/100	1750/1750	5s	3ms/step	- accuracy: 0.8534	- loss: 0.4922	- val_accuracy: 0.8990	- val_loss: 0.3638
Epoch 34/100	1750/1750	5s	2ms/step	- accuracy: 0.8579	- loss: 0.4831	- val_accuracy: 0.8993	- val_loss: 0.3614
Epoch 35/100	1750/1750	5s	3ms/step	- accuracy: 0.8589	- loss: 0.4773	- val_accuracy: 0.9007	- val_loss: 0.3563
Epoch 36/100	1750/1750	5s	3ms/step	- accuracy: 0.8591	- loss: 0.4743	- val_accuracy: 0.9019	- val_loss: 0.3567
Epoch 37/100	1750/1750	5s	2ms/step	- accuracy: 0.8605	- loss: 0.4721	- val_accuracy: 0.9033	- val_loss: 0.3524
Epoch 38/100	1750/1750	5s	3ms/step	- accuracy: 0.8644	- loss: 0.4612	- val_accuracy: 0.9029	- val_loss: 0.3504
Epoch 39/100	1750/1750	5s	3ms/step	- accuracy: 0.8641	- loss: 0.4638	- val_accuracy: 0.9029	- val_loss: 0.3449
Epoch 40/100	1750/1750	4s	2ms/step	- accuracy: 0.8674	- loss: 0.4502	- val_accuracy: 0.9070	- val_loss: 0.3400
Epoch 41/100	1750/1750	6s	3ms/step	- accuracy: 0.8636	- loss: 0.4621	- val_accuracy: 0.9066	- val_loss: 0.3368
Epoch 42/100	1750/1750	5s	3ms/step	- accuracy: 0.8645	- loss: 0.4489	- val_accuracy: 0.9061	- val_loss: 0.3380
Epoch 43/100	1750/1750	4s	2ms/step	- accuracy: 0.8680	- loss: 0.4469	- val_accuracy: 0.9086	- val_loss: 0.3315
Epoch 44/100	1750/1750	5s	3ms/step	- accuracy: 0.8687	- loss: 0.4385	- val_accuracy: 0.9071	- val_loss: 0.3294
Epoch 45/100	1750/1750	5s	3ms/step	- accuracy: 0.8710	- loss: 0.4347	- val_accuracy: 0.9076	- val_loss: 0.3258
Epoch 46/100	1750/1750	4s	2ms/step	- accuracy: 0.8711	- loss: 0.4353	- val_accuracy: 0.9129	- val_loss: 0.3260

Epoch 47/100	1750/1750	5s	3ms/step	- accuracy: 0.8699	- loss: 0.4375	- val_accuracy: 0.9063	- val_loss: 0.3260
Epoch 48/100	1750/1750	5s	3ms/step	- accuracy: 0.8709	- loss: 0.4336	- val_accuracy: 0.9124	- val_loss: 0.3171
Epoch 49/100	1750/1750	5s	2ms/step	- accuracy: 0.8760	- loss: 0.4211	- val_accuracy: 0.9131	- val_loss: 0.3189
Epoch 50/100	1750/1750	6s	3ms/step	- accuracy: 0.8771	- loss: 0.4150	- val_accuracy: 0.9134	- val_loss: 0.3100
Epoch 51/100	1750/1750	4s	2ms/step	- accuracy: 0.8735	- loss: 0.4255	- val_accuracy: 0.9129	- val_loss: 0.3092
Epoch 52/100	1750/1750	5s	2ms/step	- accuracy: 0.8764	- loss: 0.4174	- val_accuracy: 0.9131	- val_loss: 0.3131
Epoch 53/100	1750/1750	6s	3ms/step	- accuracy: 0.8734	- loss: 0.4239	- val_accuracy: 0.9133	- val_loss: 0.3087
Epoch 54/100	1750/1750	4s	2ms/step	- accuracy: 0.8809	- loss: 0.4133	- val_accuracy: 0.9151	- val_loss: 0.3053
Epoch 55/100	1750/1750	4s	2ms/step	- accuracy: 0.8793	- loss: 0.4091	- val_accuracy: 0.9166	- val_loss: 0.3030
Epoch 56/100	1750/1750	7s	3ms/step	- accuracy: 0.8808	- loss: 0.4080	- val_accuracy: 0.9171	- val_loss: 0.3030
Epoch 57/100	1750/1750	9s	2ms/step	- accuracy: 0.8771	- loss: 0.4116	- val_accuracy: 0.9177	- val_loss: 0.2994
Epoch 58/100	1750/1750	6s	3ms/step	- accuracy: 0.8812	- loss: 0.4054	- val_accuracy: 0.9191	- val_loss: 0.2959
Epoch 59/100	1750/1750	4s	2ms/step	- accuracy: 0.8810	- loss: 0.4009	- val_accuracy: 0.9193	- val_loss: 0.2920
Epoch 60/100	1750/1750	4s	2ms/step	- accuracy: 0.8817	- loss: 0.3977	- val_accuracy: 0.9194	- val_loss: 0.2922
Epoch 61/100	1750/1750	7s	3ms/step	- accuracy: 0.8847	- loss: 0.3874	- val_accuracy: 0.9199	- val_loss: 0.2913
Epoch 62/100	1750/1750	4s	2ms/step	- accuracy: 0.8840	- loss: 0.3886	- val_accuracy: 0.9184	- val_loss: 0.2911
Epoch 63/100	1750/1750	4s	2ms/step	- accuracy: 0.8839	- loss: 0.3867	- val_accuracy: 0.9213	- val_loss: 0.2875
Epoch 64/100	1750/1750	7s	3ms/step	- accuracy: 0.8843	- loss: 0.3846	- val_accuracy: 0.9211	- val_loss: 0.2850
Epoch 65/100	1750/1750	4s	2ms/step	- accuracy: 0.8864	- loss: 0.3846	- val_accuracy: 0.9180	- val_loss: 0.2872
Epoch 66/100	1750/1750	4s	2ms/step	- accuracy: 0.8843	- loss: 0.3860	- val_accuracy: 0.9217	- val_loss: 0.2844
Epoch 67/100	1750/1750	6s	3ms/step	- accuracy: 0.8836	- loss: 0.3932	- val_accuracy: 0.9183	- val_loss: 0.2875
Epoch 68/100	1750/1750	9s	2ms/step	- accuracy: 0.8862	- loss: 0.3799	- val_accuracy: 0.9223	- val_loss: 0.2795
Epoch 69/100	1750/1750	6s	3ms/step	- accuracy: 0.8858	- loss: 0.3847	- val_accuracy: 0.9213	- val_loss: 0.2816

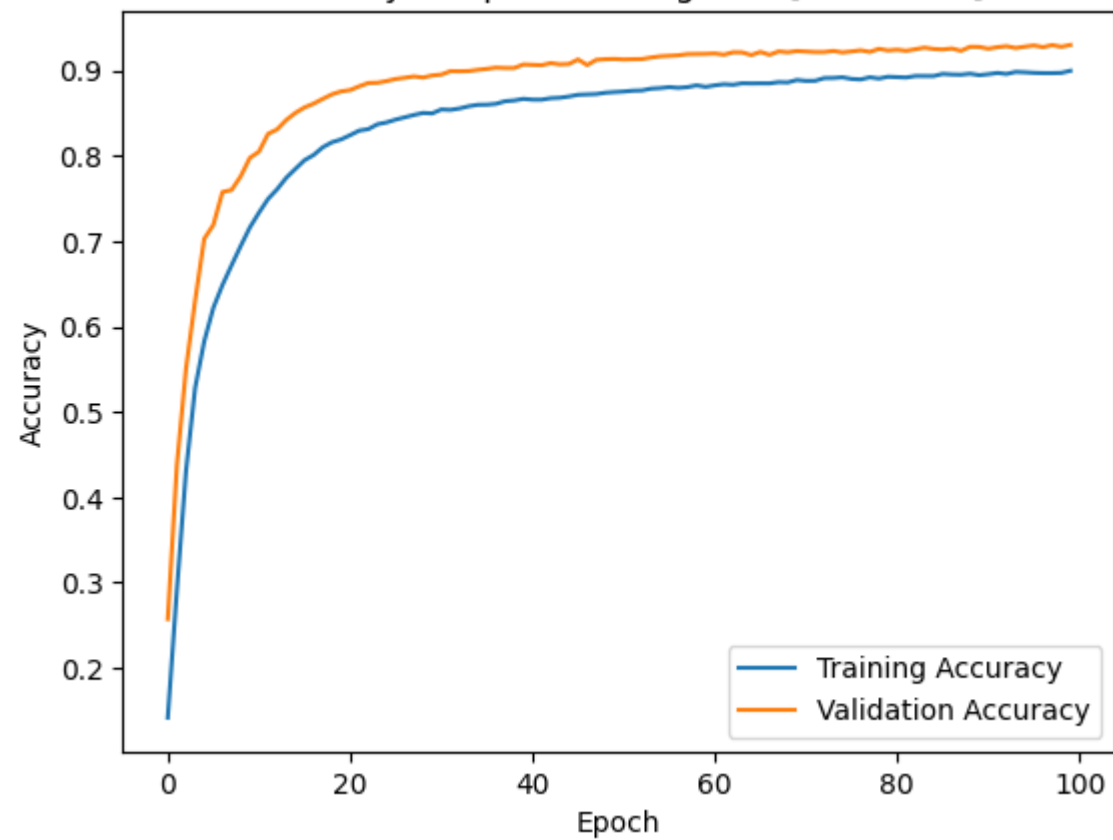
Epoch 70/100	1750/1750	4s	2ms/step	- accuracy: 0.8888	- loss: 0.3739	- val_accuracy: 0.9227	- val_loss: 0.2770
Epoch 71/100	1750/1750	4s	2ms/step	- accuracy: 0.8877	- loss: 0.3788	- val_accuracy: 0.9221	- val_loss: 0.2754
Epoch 72/100	1750/1750	5s	3ms/step	- accuracy: 0.8874	- loss: 0.3751	- val_accuracy: 0.9216	- val_loss: 0.2737
Epoch 73/100	1750/1750	4s	2ms/step	- accuracy: 0.8924	- loss: 0.3658	- val_accuracy: 0.9216	- val_loss: 0.2763
Epoch 74/100	1750/1750	5s	2ms/step	- accuracy: 0.8909	- loss: 0.3745	- val_accuracy: 0.9229	- val_loss: 0.2726
Epoch 75/100	1750/1750	6s	3ms/step	- accuracy: 0.8923	- loss: 0.3673	- val_accuracy: 0.9211	- val_loss: 0.2734
Epoch 76/100	1750/1750	4s	2ms/step	- accuracy: 0.8904	- loss: 0.3679	- val_accuracy: 0.9223	- val_loss: 0.2763
Epoch 77/100	1750/1750	5s	2ms/step	- accuracy: 0.8893	- loss: 0.3699	- val_accuracy: 0.9236	- val_loss: 0.2703
Epoch 78/100	1750/1750	6s	3ms/step	- accuracy: 0.8914	- loss: 0.3639	- val_accuracy: 0.9217	- val_loss: 0.2714
Epoch 79/100	1750/1750	8s	2ms/step	- accuracy: 0.8906	- loss: 0.3622	- val_accuracy: 0.9253	- val_loss: 0.2689
Epoch 80/100	1750/1750	6s	3ms/step	- accuracy: 0.8917	- loss: 0.3632	- val_accuracy: 0.9234	- val_loss: 0.2674
Epoch 81/100	1750/1750	4s	2ms/step	- accuracy: 0.8888	- loss: 0.3740	- val_accuracy: 0.9243	- val_loss: 0.2643
Epoch 82/100	1750/1750	5s	2ms/step	- accuracy: 0.8930	- loss: 0.3551	- val_accuracy: 0.9230	- val_loss: 0.2648
Epoch 83/100	1750/1750	6s	3ms/step	- accuracy: 0.8950	- loss: 0.3568	- val_accuracy: 0.9249	- val_loss: 0.2631
Epoch 84/100	1750/1750	4s	2ms/step	- accuracy: 0.8937	- loss: 0.3539	- val_accuracy: 0.9267	- val_loss: 0.2624
Epoch 85/100	1750/1750	5s	2ms/step	- accuracy: 0.8942	- loss: 0.3498	- val_accuracy: 0.9251	- val_loss: 0.2625
Epoch 86/100	1750/1750	5s	3ms/step	- accuracy: 0.8949	- loss: 0.3487	- val_accuracy: 0.9246	- val_loss: 0.2603
Epoch 87/100	1750/1750	9s	2ms/step	- accuracy: 0.8970	- loss: 0.3461	- val_accuracy: 0.9257	- val_loss: 0.2616
Epoch 88/100	1750/1750	5s	3ms/step	- accuracy: 0.8946	- loss: 0.3510	- val_accuracy: 0.9230	- val_loss: 0.2639
Epoch 89/100	1750/1750	9s	2ms/step	- accuracy: 0.8967	- loss: 0.3489	- val_accuracy: 0.9277	- val_loss: 0.2564
Epoch 90/100	1750/1750	6s	3ms/step	- accuracy: 0.8938	- loss: 0.3551	- val_accuracy: 0.9276	- val_loss: 0.2567
Epoch 91/100	1750/1750	4s	2ms/step	- accuracy: 0.8954	- loss: 0.3448	- val_accuracy: 0.9254	- val_loss: 0.2568
Epoch 92/100	1750/1750	4s	2ms/step	- accuracy: 0.8967	- loss: 0.3483	- val_accuracy: 0.9273	- val_loss: 0.2553

Epoch 93/100
1750/1750  6s 3ms/step - accuracy: 0.8946 - loss: 0.3482 - val_accuracy: 0.9284 - val_loss: 0.2554
Epoch 94/100
1750/1750  4s 2ms/step - accuracy: 0.8988 - loss: 0.3428 - val_accuracy: 0.9263 - val_loss: 0.2554
Epoch 95/100
1750/1750  6s 3ms/step - accuracy: 0.8988 - loss: 0.3381 - val_accuracy: 0.9279 - val_loss: 0.2512
Epoch 96/100
1750/1750  9s 2ms/step - accuracy: 0.8955 - loss: 0.3461 - val_accuracy: 0.9293 - val_loss: 0.2491
Epoch 97/100
1750/1750  5s 3ms/step - accuracy: 0.8988 - loss: 0.3398 - val_accuracy: 0.9273 - val_loss: 0.2556
Epoch 98/100
1750/1750  4s 2ms/step - accuracy: 0.8973 - loss: 0.3409 - val_accuracy: 0.9297 - val_loss: 0.2485
Epoch 99/100
1750/1750  5s 2ms/step - accuracy: 0.8979 - loss: 0.3399 - val_accuracy: 0.9277 - val_loss: 0.2522
Epoch 100/100
1750/1750  6s 3ms/step - accuracy: 0.8989 - loss: 0.3356 - val_accuracy: 0.9299 - val_loss: 0.2505
219/219  0s 1ms/step - accuracy: 0.9338 - loss: 0.2457
Learning rate = 0.0001
Test loss=0.25098052620887756 Test accuracy = 0.9292857050895691
Time taken to train the model is 554.5450575351715 seconds

Loss vs epochs for sigmoid [16, 32, 64]



Accuracy vs epochs for sigmoid [16, 32, 64]



Model: "sequential_26"























Layer (type)	Output Shape	Param #
flatten_26 (Flatten)	(None, 1024)	0
dense_90 (Dense)	(None, 16)	16,400
dropout_27 (Dropout)	(None, 16)	0
dense_91 (Dense)	(None, 32)	544
dropout_28 (Dropout)	(None, 32)	0
dense_92 (Dense)	(None, 64)	2,112
dropout_29 (Dropout)	(None, 64)	0
dense_93 (Dense)	(None, 10)	650

Total params: 19,706 (76.98 KB)

Trainable params: 19,706 (76.98 KB)

Non-trainable params: 0 (0.00 B)

Epoch 1/100	1750/1750	7s	3ms/step	- accuracy: 0.2199	- loss: 2.1256	- val_accuracy: 0.6759	- val_loss: 1.1909
Epoch 2/100	1750/1750	4s	2ms/step	- accuracy: 0.6131	- loss: 1.1406	- val_accuracy: 0.8139	- val_loss: 0.7003
Epoch 3/100	1750/1750	4s	2ms/step	- accuracy: 0.7564	- loss: 0.7731	- val_accuracy: 0.8556	- val_loss: 0.5296
Epoch 4/100	1750/1750	5s	3ms/step	- accuracy: 0.8025	- loss: 0.6327	- val_accuracy: 0.8709	- val_loss: 0.4540
Epoch 5/100	1750/1750	5s	3ms/step	- accuracy: 0.8264	- loss: 0.5732	- val_accuracy: 0.8879	- val_loss: 0.4100
Epoch 6/100	1750/1750	5s	2ms/step	- accuracy: 0.8400	- loss: 0.5258	- val_accuracy: 0.8980	- val_loss: 0.3741
Epoch 7/100	1750/1750	5s	3ms/step	- accuracy: 0.8512	- loss: 0.4889	- val_accuracy: 0.9034	- val_loss: 0.3468
Epoch 8/100	1750/1750	5s	3ms/step	- accuracy: 0.8541	- loss: 0.4719	- val_accuracy: 0.9037	- val_loss: 0.3426
Epoch 9/100	1750/1750	5s	2ms/step	- accuracy: 0.8629	- loss: 0.4522	- val_accuracy: 0.9076	- val_loss: 0.3327
Epoch 10/100	1750/1750	7s	3ms/step	- accuracy: 0.8645	- loss: 0.4338	- val_accuracy: 0.9121	- val_loss: 0.3021
Epoch 11/100	1750/1750	9s	2ms/step	- accuracy: 0.8678	- loss: 0.4283	- val_accuracy: 0.9190	- val_loss: 0.2859
Epoch 12/100	1750/1750	6s	3ms/step	- accuracy: 0.8733	- loss: 0.4042	- val_accuracy: 0.9187	- val_loss: 0.2782
Epoch 13/100	1750/1750	4s	2ms/step	- accuracy: 0.8734	- loss: 0.4071	- val_accuracy: 0.9217	- val_loss: 0.2698
Epoch 14/100	1750/1750	5s	2ms/step	- accuracy: 0.8752	- loss: 0.3954	- val_accuracy: 0.9234	- val_loss: 0.2668
Epoch 15/100	1750/1750	6s	3ms/step	- accuracy: 0.8791	- loss: 0.3873	- val_accuracy: 0.9226	- val_loss: 0.2658
Epoch 16/100	1750/1750	9s	2ms/step	- accuracy: 0.8840	- loss: 0.3724	- val_accuracy: 0.9257	- val_loss: 0.2598
Epoch 17/100	1750/1750	7s	3ms/step	- accuracy: 0.8837	- loss: 0.3638	- val_accuracy: 0.9199	- val_loss: 0.2696
Epoch 18/100	1750/1750	9s	2ms/step	- accuracy: 0.8854	- loss: 0.3687	- val_accuracy: 0.9197	- val_loss: 0.2665
Epoch 19/100	1750/1750	7s	3ms/step	- accuracy: 0.8872	- loss: 0.3584	- val_accuracy: 0.9267	- val_loss: 0.2503
Epoch 20/100	1750/1750	4s	2ms/step	- accuracy: 0.8886	- loss: 0.3580	- val_accuracy: 0.9261	- val_loss: 0.2517
Epoch 21/100	1750/1750	5s	2ms/step	- accuracy: 0.8905	- loss: 0.3442	- val_accuracy: 0.9276	- val_loss: 0.2502
Epoch 22/100	1750/1750	6s	3ms/step	- accuracy: 0.8941	- loss: 0.3399	- val_accuracy: 0.9256	- val_loss: 0.2446
Epoch 23/100	1750/1750	4s	2ms/step	- accuracy: 0.8914	- loss: 0.3409	- val_accuracy: 0.9201	- val_loss: 0.2700

Epoch 24/100		4s	2ms/step	- accuracy: 0.8940	- loss: 0.3410	- val_accuracy: 0.9266	- val_loss: 0.2502
Epoch 25/100		6s	3ms/step	- accuracy: 0.8988	- loss: 0.3321	- val_accuracy: 0.9243	- val_loss: 0.2446
Epoch 26/100		9s	2ms/step	- accuracy: 0.8968	- loss: 0.3344	- val_accuracy: 0.9340	- val_loss: 0.2266
Epoch 27/100		5s	3ms/step	- accuracy: 0.8988	- loss: 0.3231	- val_accuracy: 0.9329	- val_loss: 0.2313
Epoch 28/100		9s	2ms/step	- accuracy: 0.8989	- loss: 0.3236	- val_accuracy: 0.9290	- val_loss: 0.2315
Epoch 29/100		5s	3ms/step	- accuracy: 0.9008	- loss: 0.3189	- val_accuracy: 0.9320	- val_loss: 0.2342
Epoch 30/100		4s	2ms/step	- accuracy: 0.8971	- loss: 0.3266	- val_accuracy: 0.9309	- val_loss: 0.2300
Epoch 31/100		5s	2ms/step	- accuracy: 0.9009	- loss: 0.3171	- val_accuracy: 0.9333	- val_loss: 0.2240
Epoch 32/100		7s	3ms/step	- accuracy: 0.9024	- loss: 0.3088	- val_accuracy: 0.9270	- val_loss: 0.2352
Epoch 33/100		4s	2ms/step	- accuracy: 0.9023	- loss: 0.3124	- val_accuracy: 0.9353	- val_loss: 0.2250
Epoch 34/100		4s	2ms/step	- accuracy: 0.9034	- loss: 0.3058	- val_accuracy: 0.9337	- val_loss: 0.2198
Epoch 35/100		5s	3ms/step	- accuracy: 0.9056	- loss: 0.3053	- val_accuracy: 0.9320	- val_loss: 0.2239
Epoch 36/100		9s	2ms/step	- accuracy: 0.9068	- loss: 0.3001	- val_accuracy: 0.9317	- val_loss: 0.2270
Epoch 37/100		5s	3ms/step	- accuracy: 0.9029	- loss: 0.3093	- val_accuracy: 0.9336	- val_loss: 0.2115
Epoch 38/100		4s	2ms/step	- accuracy: 0.9069	- loss: 0.3009	- val_accuracy: 0.9314	- val_loss: 0.2251
Epoch 39/100		4s	2ms/step	- accuracy: 0.9065	- loss: 0.2977	- val_accuracy: 0.9336	- val_loss: 0.2229
Epoch 40/100		7s	3ms/step	- accuracy: 0.9086	- loss: 0.2919	- val_accuracy: 0.9380	- val_loss: 0.2087
Epoch 41/100		4s	2ms/step	- accuracy: 0.9083	- loss: 0.2994	- val_accuracy: 0.9361	- val_loss: 0.2163
Epoch 42/100		4s	2ms/step	- accuracy: 0.9071	- loss: 0.2999	- val_accuracy: 0.9377	- val_loss: 0.2070
Epoch 43/100		6s	3ms/step	- accuracy: 0.9091	- loss: 0.2935	- val_accuracy: 0.9376	- val_loss: 0.2165
Epoch 44/100		4s	2ms/step	- accuracy: 0.9100	- loss: 0.2945	- val_accuracy: 0.9330	- val_loss: 0.2227
Epoch 45/100		5s	2ms/step	- accuracy: 0.9101	- loss: 0.2899	- val_accuracy: 0.9337	- val_loss: 0.2153
Epoch 46/100		7s	3ms/step	- accuracy: 0.9108	- loss: 0.2875	- val_accuracy: 0.9363	- val_loss: 0.2097

Epoch 47/100

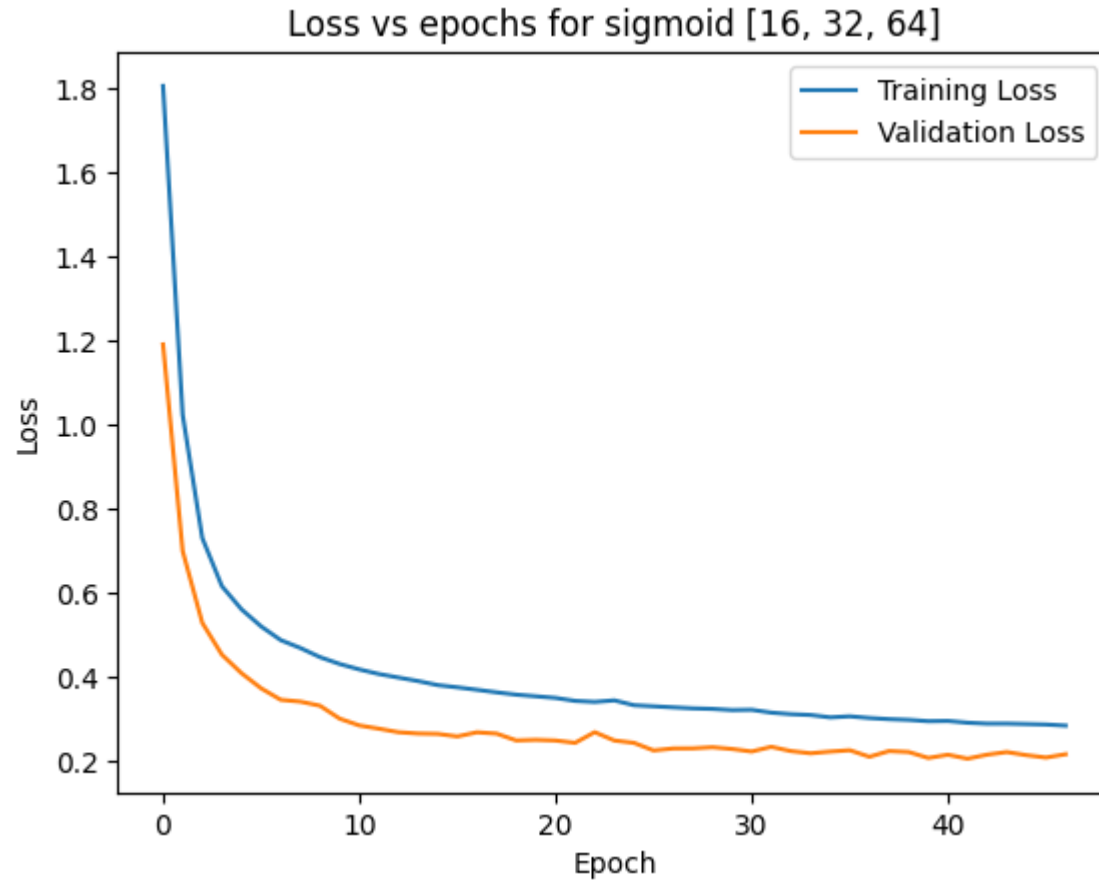
1750/1750 ————— **4s** 2ms/step - accuracy: 0.9142 - loss: 0.2818 - val_accuracy: 0.9346 - val_loss: 0.2176

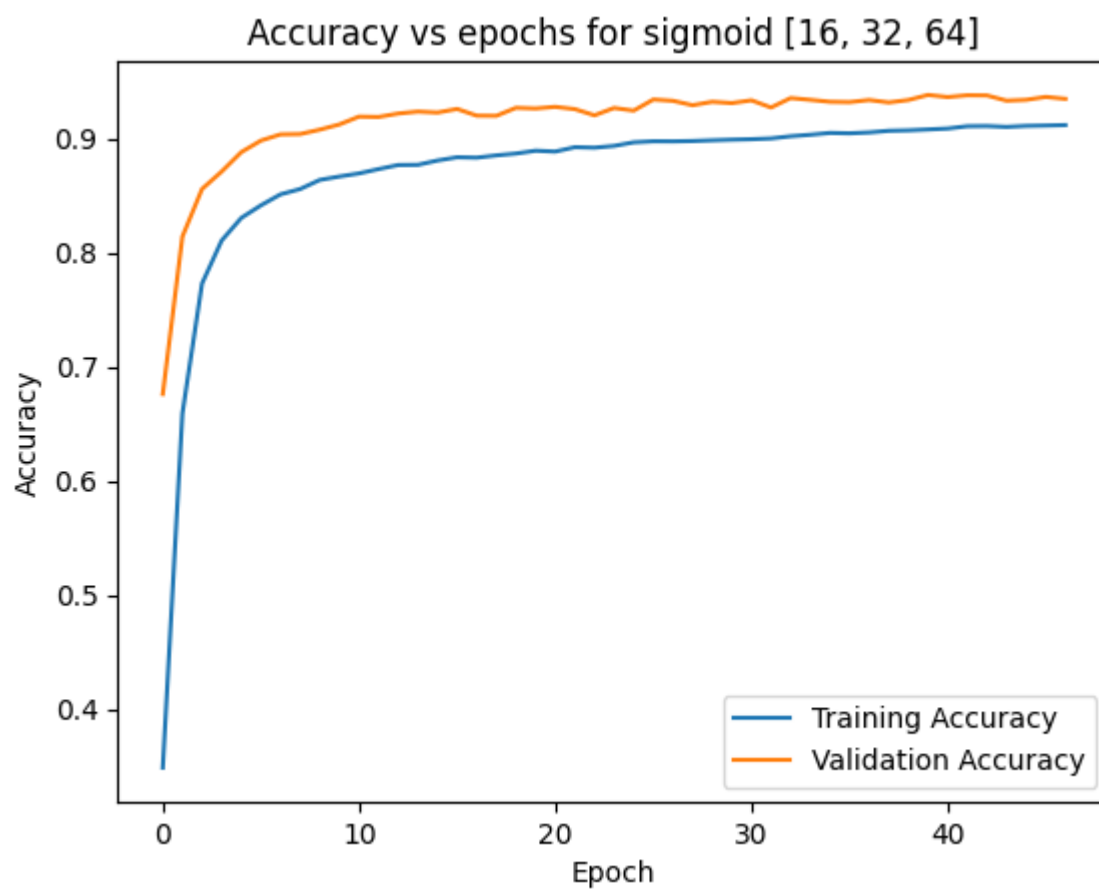
219/219 ————— **0s** 1ms/step - accuracy: 0.9422 - loss: 0.2004

Learning rate = 0.0005

Test loss=0.20578503608703613 Test accuracy = 0.9392856955528259

Time taken to train the model is 262.3124701976776 seconds





In [54]: `result_df_lr`

Out[54]:

	Hidden Layers	Hidden Neurons	Learning Rate	Dropout Rate	Activation Function	Training Time (in seconds)	Test Loss	Test Accuracy
0	3	[16, 32, 64]	0.0100	0.1	sigmoid	73.584845	2.301192	0.119429
1	3	[16, 32, 64]	0.0010	0.1	sigmoid	186.771840	0.263176	0.920286
2	3	[16, 32, 64]	0.0050	0.1	sigmoid	75.207555	1.662789	0.296714
3	3	[16, 32, 64]	0.0001	0.1	sigmoid	554.545058	0.250981	0.929286
4	3	[16, 32, 64]	0.0005	0.1	sigmoid	262.312470	0.205785	0.939286

As we can see from above the best learning rate is `0.0005`

Question 09

- ix. Create five images (of size 28×28) containing a digit of your own handwriting and test whether your trained classifier can predict it or not.

```
In [55]: # Install necessary libraries
```

```
!pip install opencv-python
!pip install Pillow
```

Requirement already satisfied: opencv-python in /usr/local/lib/python3.10/dist-packages (4.10.0.84)

Requirement already satisfied: numpy>=1.21.2 in /usr/local/lib/python3.10/dist-packages (from opencv-python) (1.26.4)

Requirement already satisfied: Pillow in /usr/local/lib/python3.10/dist-packages (10.4.0)

```
In [56]: # Import the libraries
```

```
import os
import numpy as np
from PIL import Image
import cv2
```

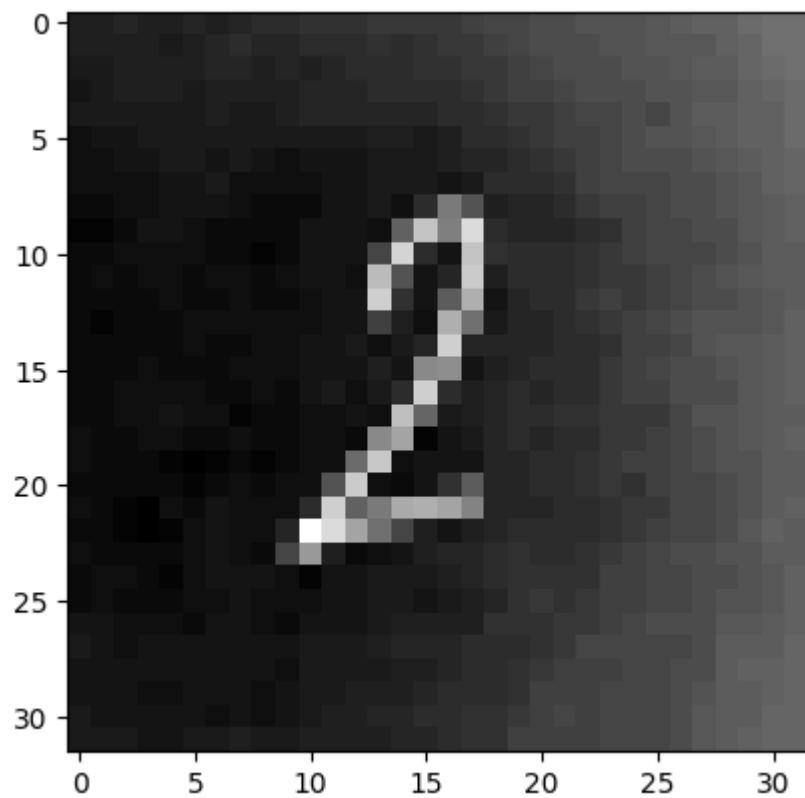
```
In [112]: def process_image(image_path):
            img = Image.open(image_path).convert('L') # Convert to gray-scale
            img = img.resize((32,32))
            img_array = np.array(img)

            return img_array

            # Testing it out
            img_array = process_image('/content/sample_images/two.png')
            print(img_array)
```

```
[[164 164 163 ... 151 150 150]
 [164 164 164 ... 152 150 150]
 [165 165 164 ... 152 151 150]
 ...
 [166 166 166 ... 154 153 152]
 [165 166 166 ... 154 152 152]
 [165 165 165 ... 153 152 152]]
```

```
In [113]: plt.imshow(img_array, cmap='Greys')
plt.show()
```



The resolution of the image is not upto the mark

```
In [114... IMG_DIR = '/content/sample_images'

images = []

image_files = [f for f in os.listdir(IMG_DIR) if os.path.isfile(os.path.join(IMG_DIR,f))]

images = []

for image_file in image_files:
    image_path = os.path.join(IMG_DIR,image_file)
    image_array = process_image(image_path)
    images.append(image_array)

images = np.array(images)
images.shape
```

```
Out[114... (5, 32, 32)
```

In [115... # Train the model (lr=0.0005), Adam optimizer, sigmoid function, dropout rate=0.1, CCE

```
model, history, duration = train_model(
    hidden_neurons=hidden_neurons,
    learning_rate=0.0005,
    dropout_rate=best_dropout_rate,
    activation_func=best_activation_function,
    loss_func=CategoricalCrossentropy,
    optimizer=Adam,
    epochs=100
)

test_loss, test_acc = model.evaluate(X_test,y_test)

print(f"Learning rate = {0.0005}")
print(f"Test loss={test_loss} Test accuracy = {test_acc}")
print(f"Time taken to train the model is {duration} seconds")

plot_metrics(history,best_activation_function,hidden_neurons,dropout_rate=best_dropout_rate)
```

Model: "sequential_27"

Layer (type)	Output Shape	Param #
flatten_27 (Flatten)	(None, 1024)	0
dense_94 (Dense)	(None, 16)	16,400
dropout_30 (Dropout)	(None, 16)	0
dense_95 (Dense)	(None, 32)	544
dropout_31 (Dropout)	(None, 32)	0
dense_96 (Dense)	(None, 64)	2,112
dropout_32 (Dropout)	(None, 64)	0
dense_97 (Dense)	(None, 10)	650











Total params: 19,706 (76.98 KB)

Trainable params: 19,706 (76.98 KB)

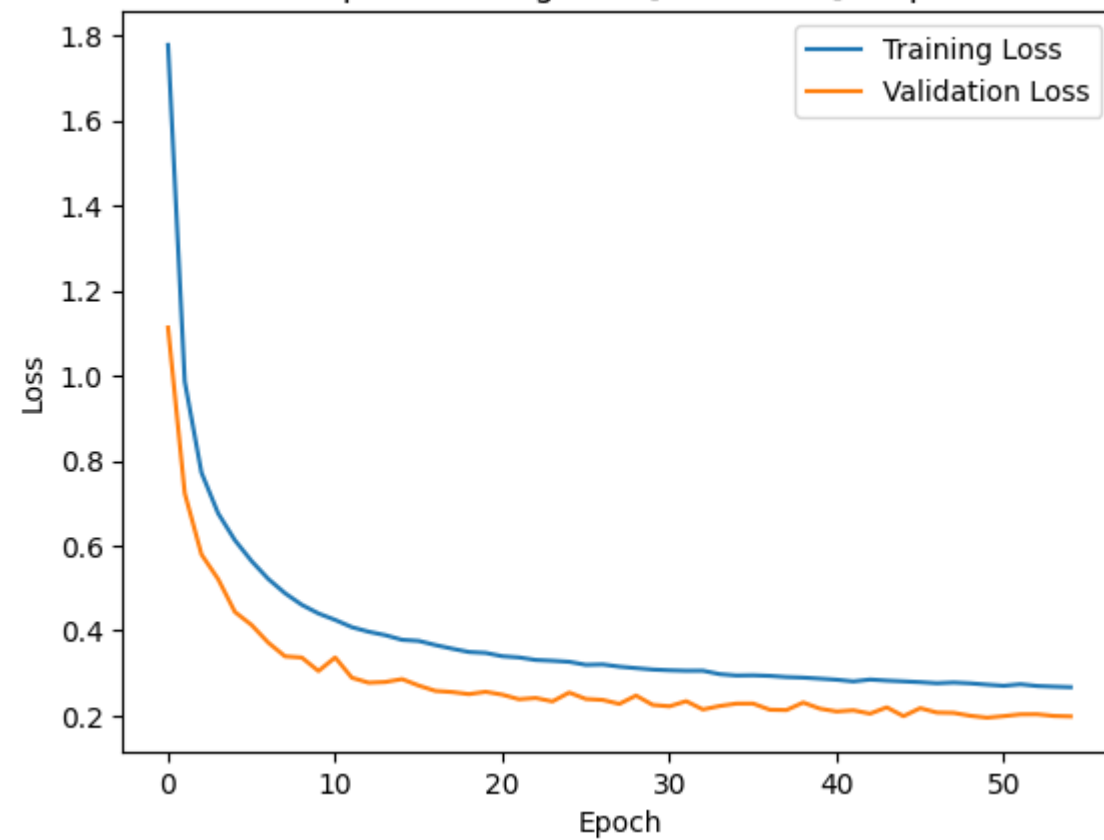
Non-trainable params: 0 (0.00 B)

Epoch 1/100	1750/1750	7s	3ms/step	- accuracy: 0.2258	- loss: 2.1062	- val_accuracy: 0.6690	- val_loss: 1.1130
Epoch 2/100	1750/1750	9s	2ms/step	- accuracy: 0.6298	- loss: 1.0791	- val_accuracy: 0.7670	- val_loss: 0.7234
Epoch 3/100	1750/1750	7s	3ms/step	- accuracy: 0.7355	- loss: 0.8068	- val_accuracy: 0.8324	- val_loss: 0.5794
Epoch 4/100	1750/1750	9s	2ms/step	- accuracy: 0.7855	- loss: 0.6915	- val_accuracy: 0.8546	- val_loss: 0.5212
Epoch 5/100	1750/1750	5s	3ms/step	- accuracy: 0.8123	- loss: 0.6292	- val_accuracy: 0.8790	- val_loss: 0.4441
Epoch 6/100	1750/1750	9s	2ms/step	- accuracy: 0.8284	- loss: 0.5773	- val_accuracy: 0.8866	- val_loss: 0.4133
Epoch 7/100	1750/1750	5s	3ms/step	- accuracy: 0.8434	- loss: 0.5277	- val_accuracy: 0.8971	- val_loss: 0.3718
Epoch 8/100	1750/1750	4s	2ms/step	- accuracy: 0.8563	- loss: 0.4882	- val_accuracy: 0.9051	- val_loss: 0.3396
Epoch 9/100	1750/1750	4s	2ms/step	- accuracy: 0.8594	- loss: 0.4711	- val_accuracy: 0.9033	- val_loss: 0.3369
Epoch 10/100	1750/1750	6s	3ms/step	- accuracy: 0.8665	- loss: 0.4468	- val_accuracy: 0.9131	- val_loss: 0.3052
Epoch 11/100	1750/1750	11s	4ms/step	- accuracy: 0.8770	- loss: 0.4149	- val_accuracy: 0.9010	- val_loss: 0.3370
Epoch 12/100	1750/1750	8s	2ms/step	- accuracy: 0.8773	- loss: 0.4079	- val_accuracy: 0.9170	- val_loss: 0.2894
Epoch 13/100	1750/1750	4s	2ms/step	- accuracy: 0.8788	- loss: 0.4024	- val_accuracy: 0.9199	- val_loss: 0.2777
Epoch 14/100	1750/1750	5s	2ms/step	- accuracy: 0.8842	- loss: 0.3862	- val_accuracy: 0.9201	- val_loss: 0.2793
Epoch 15/100	1750/1750	5s	2ms/step	- accuracy: 0.8887	- loss: 0.3777	- val_accuracy: 0.9170	- val_loss: 0.2861
Epoch 16/100	1750/1750	7s	3ms/step	- accuracy: 0.8909	- loss: 0.3669	- val_accuracy: 0.9233	- val_loss: 0.2709
Epoch 17/100	1750/1750	4s	2ms/step	- accuracy: 0.8931	- loss: 0.3629	- val_accuracy: 0.9231	- val_loss: 0.2583
Epoch 18/100	1750/1750	5s	2ms/step	- accuracy: 0.8911	- loss: 0.3584	- val_accuracy: 0.9271	- val_loss: 0.2557
Epoch 19/100	1750/1750	6s	3ms/step	- accuracy: 0.8937	- loss: 0.3470	- val_accuracy: 0.9263	- val_loss: 0.2508
Epoch 20/100	1750/1750	4s	2ms/step	- accuracy: 0.8928	- loss: 0.3534	- val_accuracy: 0.9237	- val_loss: 0.2563
Epoch 21/100	1750/1750	4s	2ms/step	- accuracy: 0.8999	- loss: 0.3419	- val_accuracy: 0.9249	- val_loss: 0.2495
Epoch 22/100	1750/1750	6s	3ms/step	- accuracy: 0.8977	- loss: 0.3417	- val_accuracy: 0.9290	- val_loss: 0.2387
Epoch 23/100	1750/1750	4s	2ms/step	- accuracy: 0.9043	- loss: 0.3229	- val_accuracy: 0.9289	- val_loss: 0.2417

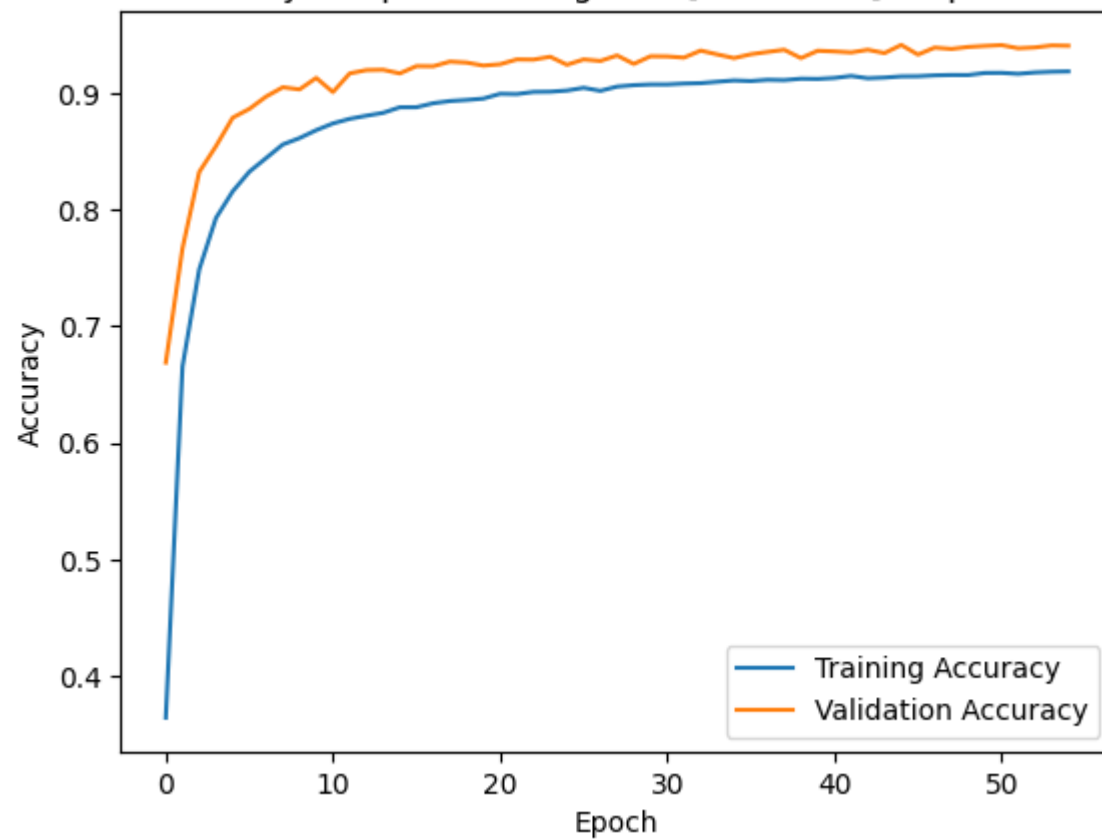
Epoch 24/100	1750/1750	4s	2ms/step	- accuracy: 0.9017	- loss: 0.3301	- val_accuracy: 0.9313	- val_loss: 0.2336
Epoch 25/100	1750/1750	6s	3ms/step	- accuracy: 0.9027	- loss: 0.3240	- val_accuracy: 0.9244	- val_loss: 0.2544
Epoch 26/100	1750/1750	4s	2ms/step	- accuracy: 0.9043	- loss: 0.3228	- val_accuracy: 0.9290	- val_loss: 0.2393
Epoch 27/100	1750/1750	6s	3ms/step	- accuracy: 0.9031	- loss: 0.3186	- val_accuracy: 0.9277	- val_loss: 0.2372
Epoch 28/100	1750/1750	9s	2ms/step	- accuracy: 0.9065	- loss: 0.3186	- val_accuracy: 0.9324	- val_loss: 0.2278
Epoch 29/100	1750/1750	5s	3ms/step	- accuracy: 0.9053	- loss: 0.3203	- val_accuracy: 0.9251	- val_loss: 0.2474
Epoch 30/100	1750/1750	4s	2ms/step	- accuracy: 0.9085	- loss: 0.3021	- val_accuracy: 0.9317	- val_loss: 0.2253
Epoch 31/100	1750/1750	4s	2ms/step	- accuracy: 0.9073	- loss: 0.3101	- val_accuracy: 0.9316	- val_loss: 0.2222
Epoch 32/100	1750/1750	5s	3ms/step	- accuracy: 0.9076	- loss: 0.3085	- val_accuracy: 0.9306	- val_loss: 0.2342
Epoch 33/100	1750/1750	4s	2ms/step	- accuracy: 0.9112	- loss: 0.2981	- val_accuracy: 0.9366	- val_loss: 0.2147
Epoch 34/100	1750/1750	5s	2ms/step	- accuracy: 0.9112	- loss: 0.2945	- val_accuracy: 0.9333	- val_loss: 0.2228
Epoch 35/100	1750/1750	7s	3ms/step	- accuracy: 0.9114	- loss: 0.2936	- val_accuracy: 0.9301	- val_loss: 0.2285
Epoch 36/100	1750/1750	9s	2ms/step	- accuracy: 0.9083	- loss: 0.2961	- val_accuracy: 0.9334	- val_loss: 0.2286
Epoch 37/100	1750/1750	5s	3ms/step	- accuracy: 0.9087	- loss: 0.3005	- val_accuracy: 0.9354	- val_loss: 0.2141
Epoch 38/100	1750/1750	4s	2ms/step	- accuracy: 0.9117	- loss: 0.2900	- val_accuracy: 0.9373	- val_loss: 0.2133
Epoch 39/100	1750/1750	5s	2ms/step	- accuracy: 0.9154	- loss: 0.2814	- val_accuracy: 0.9301	- val_loss: 0.2309
Epoch 40/100	1750/1750	6s	3ms/step	- accuracy: 0.9121	- loss: 0.2878	- val_accuracy: 0.9364	- val_loss: 0.2167
Epoch 41/100	1750/1750	4s	2ms/step	- accuracy: 0.9138	- loss: 0.2833	- val_accuracy: 0.9359	- val_loss: 0.2100
Epoch 42/100	1750/1750	6s	2ms/step	- accuracy: 0.9145	- loss: 0.2809	- val_accuracy: 0.9350	- val_loss: 0.2127
Epoch 43/100	1750/1750	5s	3ms/step	- accuracy: 0.9152	- loss: 0.2802	- val_accuracy: 0.9373	- val_loss: 0.2048
Epoch 44/100	1750/1750	4s	2ms/step	- accuracy: 0.9129	- loss: 0.2869	- val_accuracy: 0.9344	- val_loss: 0.2197
Epoch 45/100	1750/1750	4s	2ms/step	- accuracy: 0.9153	- loss: 0.2797	- val_accuracy: 0.9414	- val_loss: 0.1990
Epoch 46/100	1750/1750	5s	3ms/step	- accuracy: 0.9153	- loss: 0.2799	- val_accuracy: 0.9331	- val_loss: 0.2178

Epoch 47/100
1750/1750  4s 2ms/step - accuracy: 0.9178 - loss: 0.2710 - val_accuracy: 0.9391 - val_loss: 0.2075
Epoch 48/100
1750/1750  7s 3ms/step - accuracy: 0.9169 - loss: 0.2739 - val_accuracy: 0.9380 - val_loss: 0.2067
Epoch 49/100
1750/1750  9s 2ms/step - accuracy: 0.9155 - loss: 0.2751 - val_accuracy: 0.9397 - val_loss: 0.1998
Epoch 50/100
1750/1750  6s 3ms/step - accuracy: 0.9156 - loss: 0.2794 - val_accuracy: 0.9406 - val_loss: 0.1954
Epoch 51/100
1750/1750  4s 2ms/step - accuracy: 0.9171 - loss: 0.2717 - val_accuracy: 0.9413 - val_loss: 0.1991
Epoch 52/100
1750/1750  8s 4ms/step - accuracy: 0.9161 - loss: 0.2743 - val_accuracy: 0.9386 - val_loss: 0.2032
Epoch 53/100
1750/1750  6s 2ms/step - accuracy: 0.9206 - loss: 0.2597 - val_accuracy: 0.9393 - val_loss: 0.2034
Epoch 54/100
1750/1750  6s 2ms/step - accuracy: 0.9171 - loss: 0.2730 - val_accuracy: 0.9410 - val_loss: 0.1994
Epoch 55/100
1750/1750  5s 3ms/step - accuracy: 0.9183 - loss: 0.2669 - val_accuracy: 0.9407 - val_loss: 0.1984
219/219  0s 1ms/step - accuracy: 0.9454 - loss: 0.1963
Learning rate = 0.0005
Test loss=0.20049400627613068 Test accuracy = 0.9425714015960693
Time taken to train the model is 313.5573968887329 seconds

Loss vs epochs for sigmoid [16, 32, 64] dropout 0.1

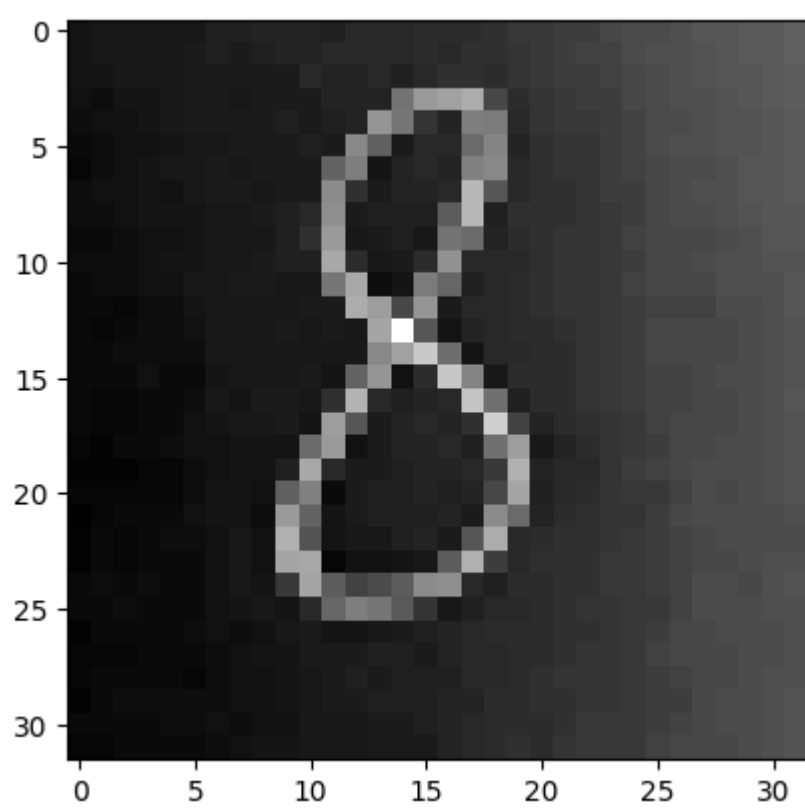


Accuracy vs epochs for sigmoid [16, 32, 64] dropout 0.1



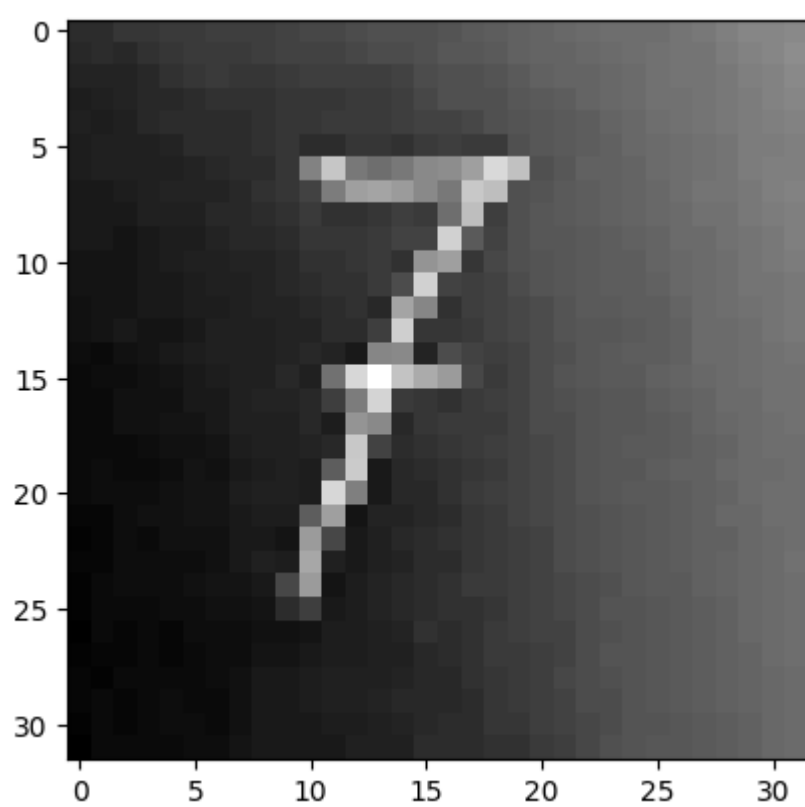
In [117... *# Prediction*

```
for image in images:
    plt.imshow(image, cmap='Greys')
    plt.show()
    prediction_array = model.predict(image.reshape(1,32,32))
    print(f"Prediction: {np.argmax(prediction_array)}")
```



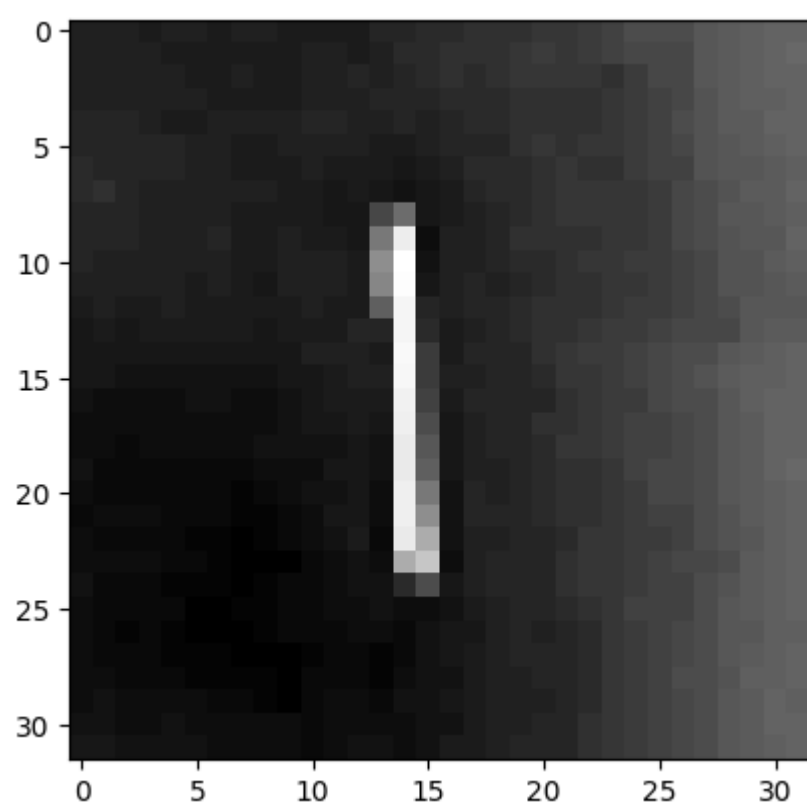
1/1 ————— 0s 246ms/step

Prediction: 0



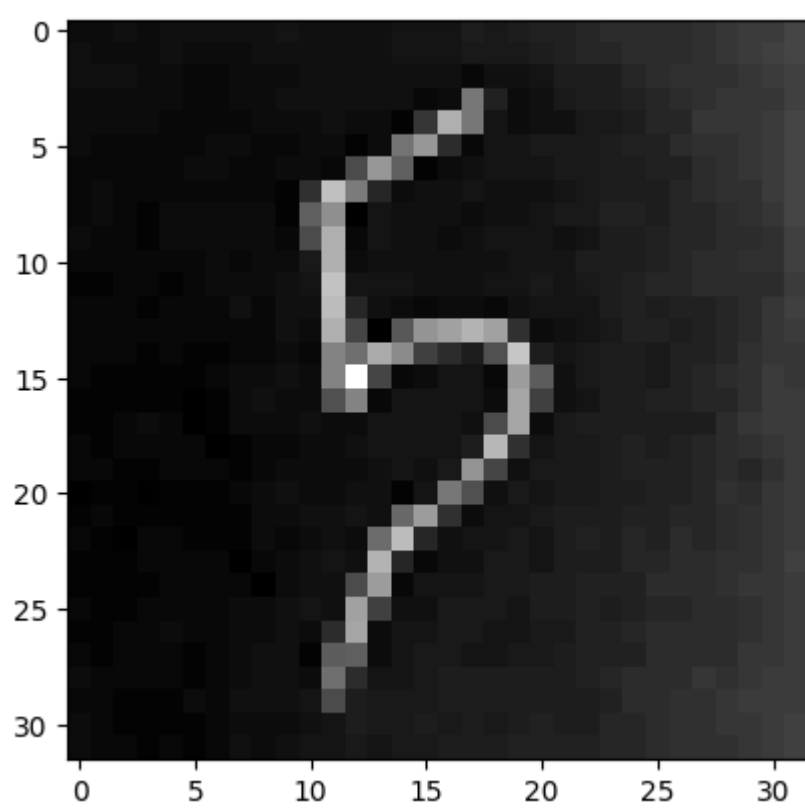
1/1 — 0s 34ms/step

Prediction: 0



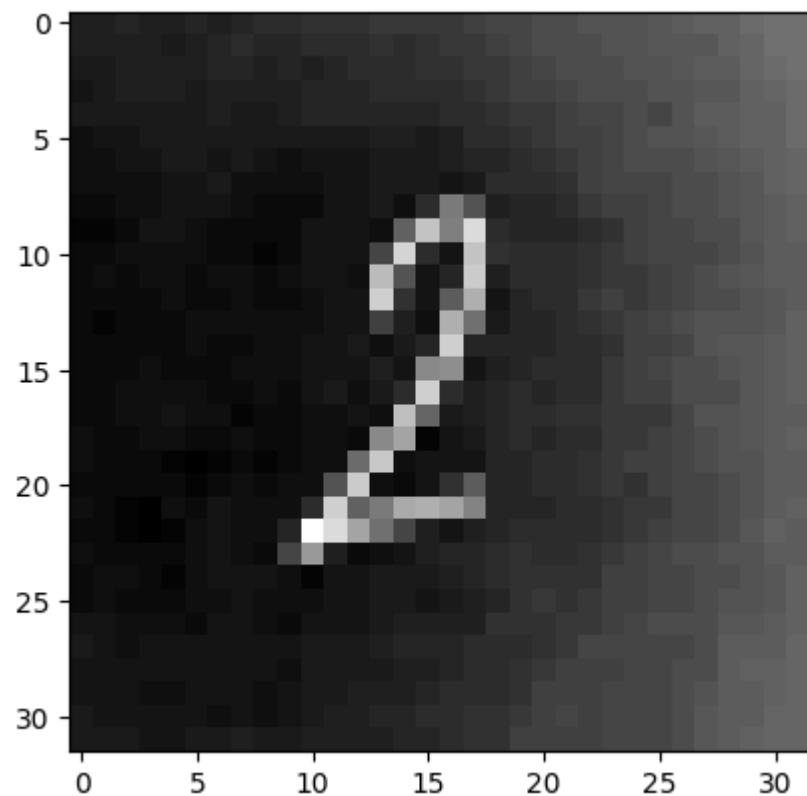
1/1 — 0s 31ms/step

Prediction: 5



1/1 — 0s 21ms/step

Prediction: 5



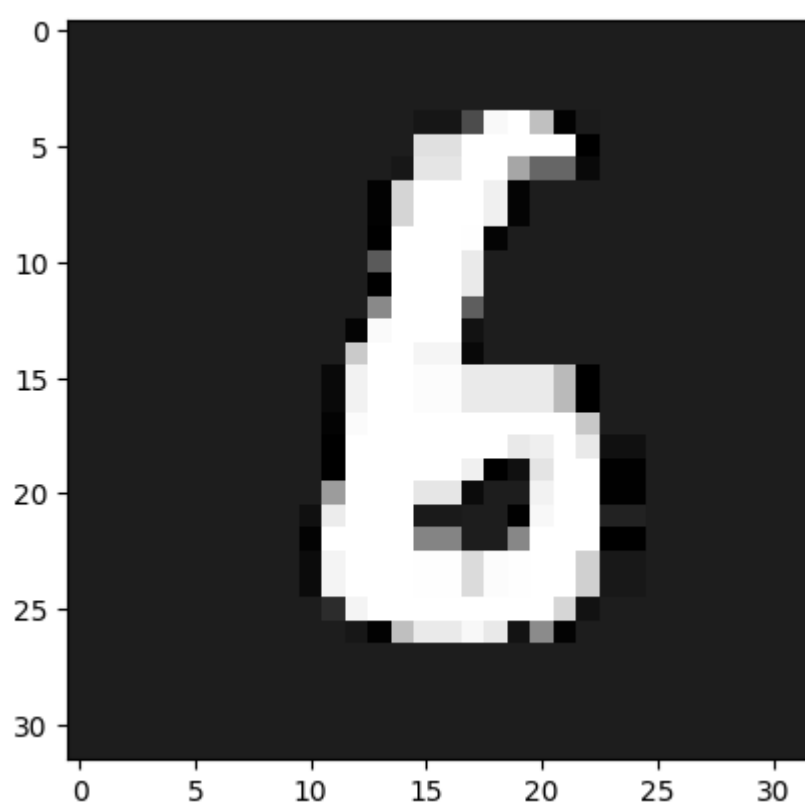
1/1 ————— 0s 20ms/step

Prediction: 5

```
In [118... # Testing on taking 10 random samples from MNIST dataset(testing portion)

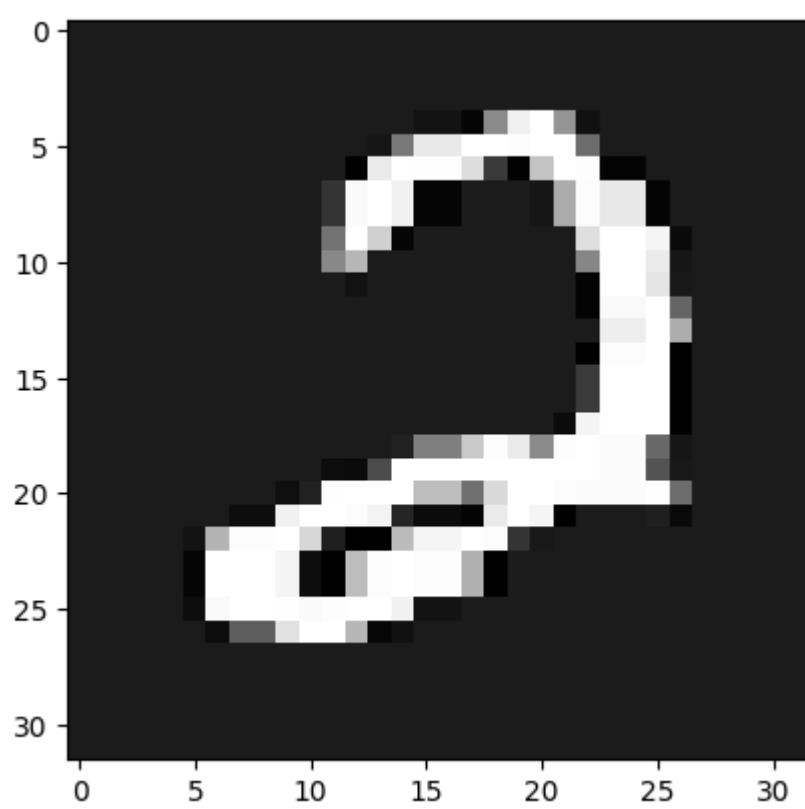
random_images = X_test[np.random.choice(X_test.shape[0], size=10, replace=False)]

for image in random_images:
    plt.imshow(image, cmap='Greys')
    plt.show()
    prediction_array = model.predict(image.reshape(1,32,32))
    print(f"Prediction: {np.argmax(prediction_array)}")
```



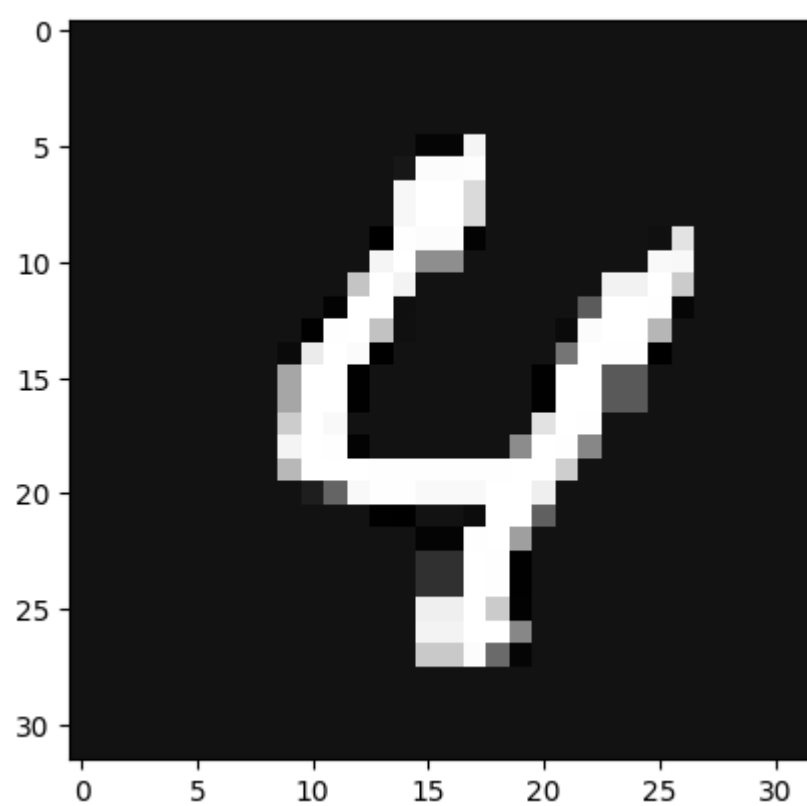
1/1 — 0s 116ms/step

Prediction: 6



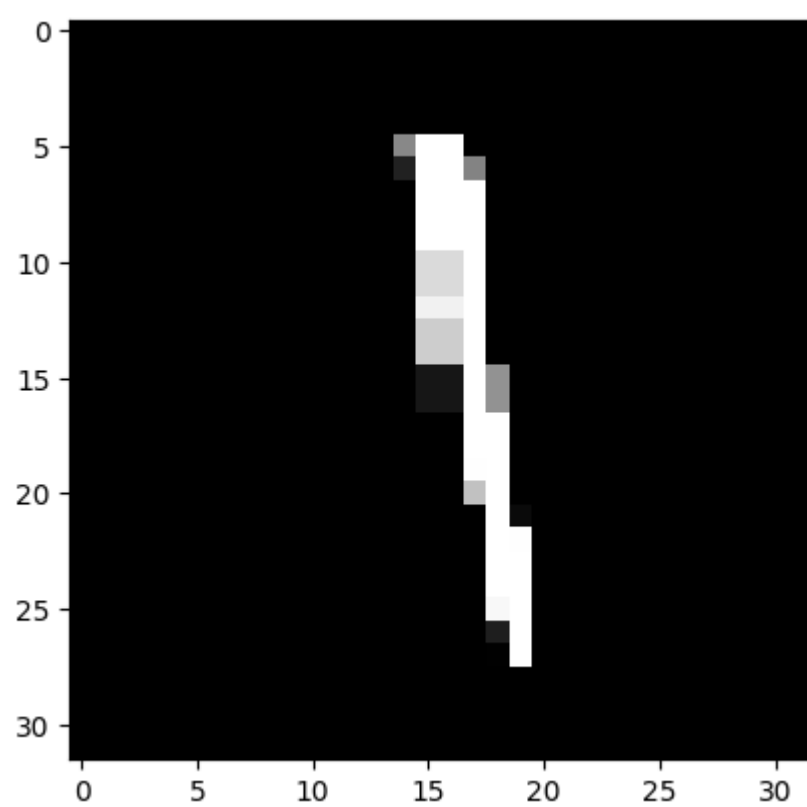
1/1 — 0s 31ms/step

Prediction: 2



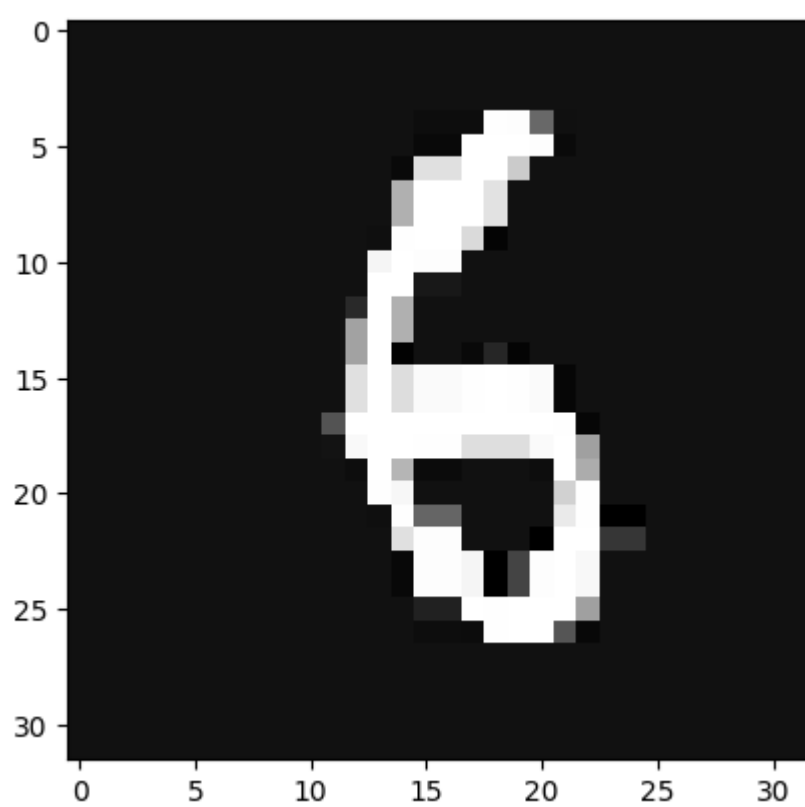
1/1 — 0s 35ms/step

Prediction: 4



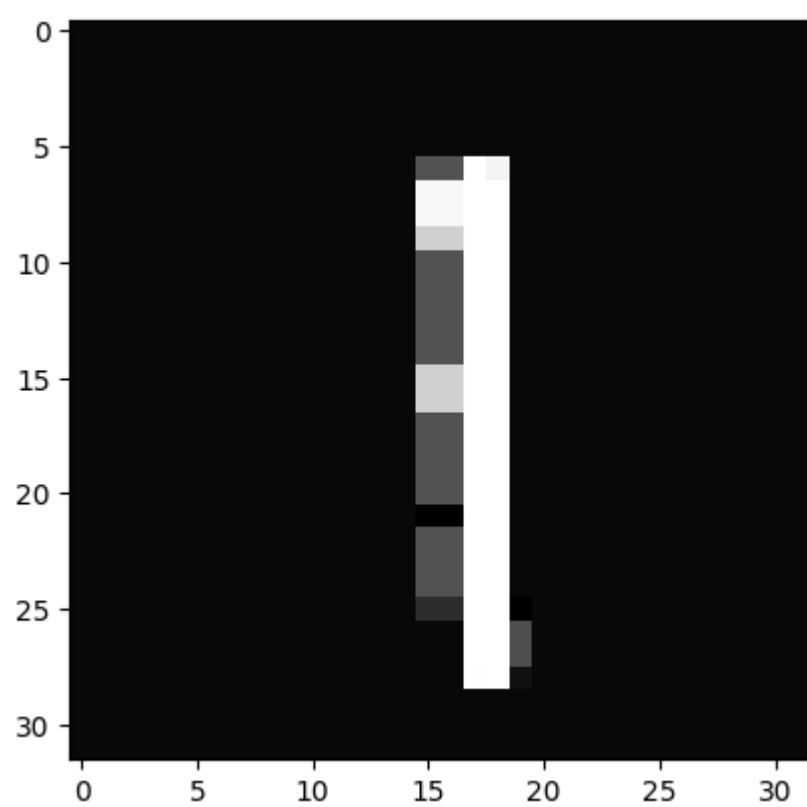
1/1 — 0s 162ms/step

Prediction: 1



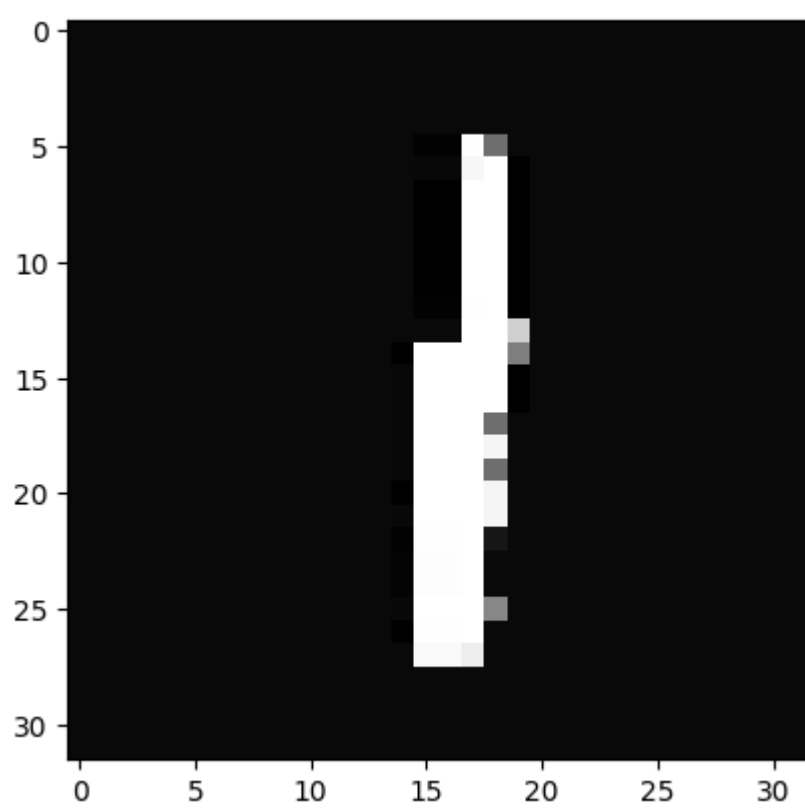
1/1 ————— 0s 23ms/step

Prediction: 6



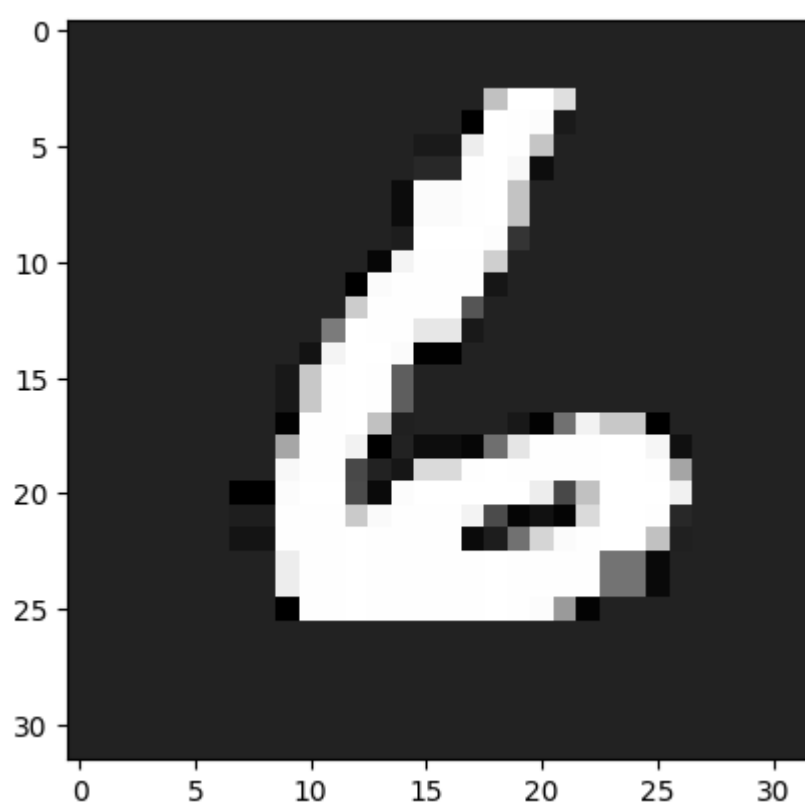
1/1 — 0s 22ms/step

Prediction: 1



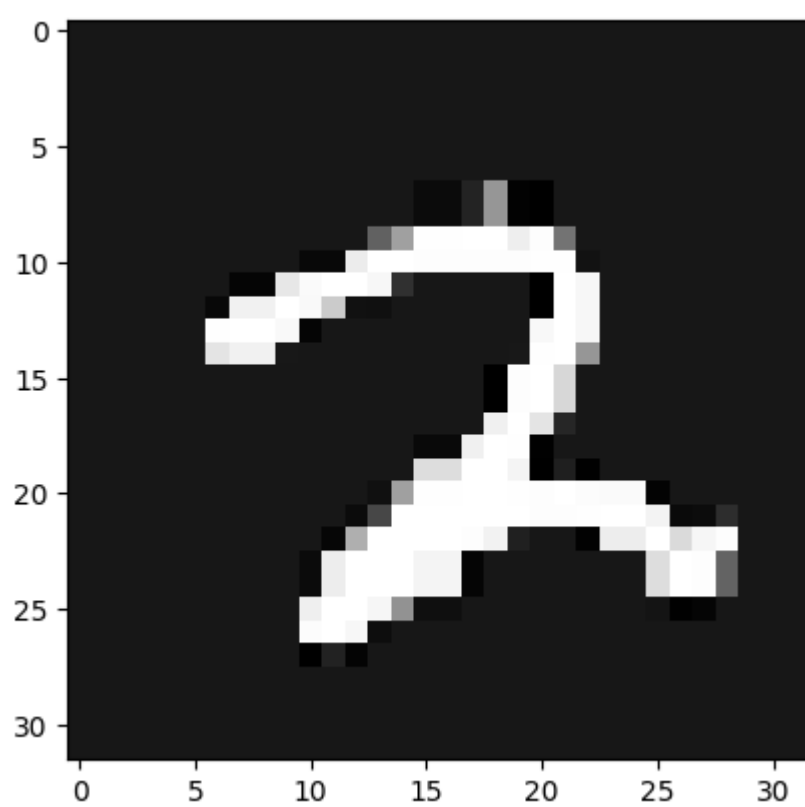
1/1 — 0s 20ms/step

Prediction: 1



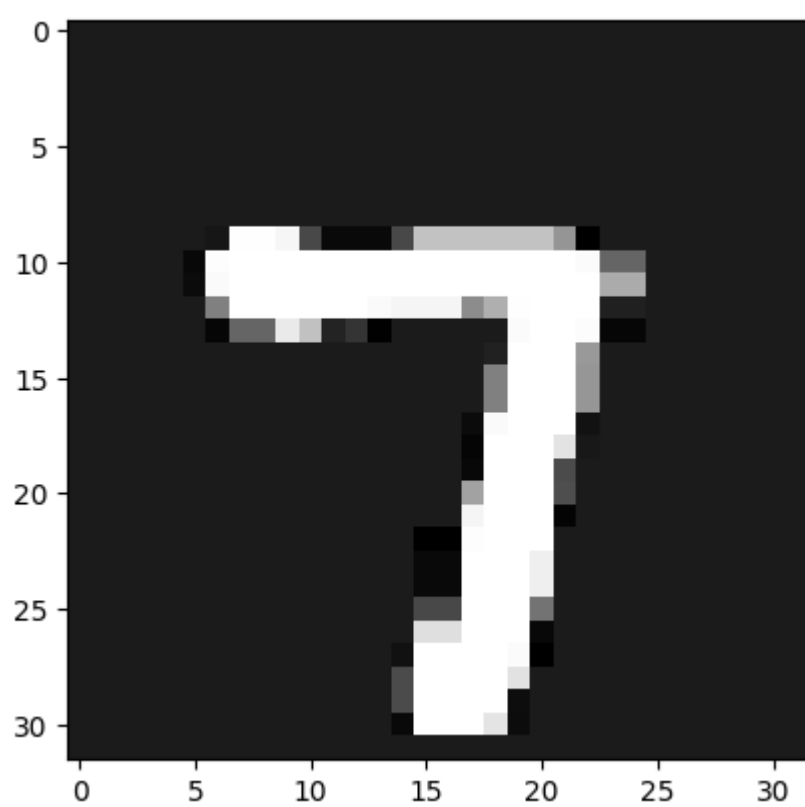
1/1 — 0s 21ms/step

Prediction: 6



1/1 — 0s 24ms/step

Prediction: 2



1/1 ————— 0s 20ms/step

Prediction: 7

As we can see that taking random samples from a standardised MNIST dataset, there is 100% accurate prediction but for non-standard images the accuracy is 20%