

Assignment 05

Name:- Sk Fardeen Hossain

Roll No. :- 2021CSB023

G-Suite Id:- 2021csb023.sk@students.iiests.ac.in

Department:- Computer Science and Technology

Question 01

Download and extract the flower image dataset from <https://www.kaggle.com/alxmamaev/flowers-recognition>.

```
In [ ]: # Download the dataset
import kagglehub

# Download latest version
path = kagglehub.dataset_download("alxmamaev/flowers-recognition")

path=path+'/flowers'

print("Path to dataset files:", path)
```

Downloading from https://www.kaggle.com/api/v1/datasets/download/alxmamaev/flowers-recognition?dataset_version_number=2...

100%|██████████| 225M/225M [00:12<00:00, 19.3MB/s]

Extracting files...

Path to dataset files: /root/.cache/kagglehub/datasets/alxmamaev/flowers-recognition/versions/2/flowers

```
In [ ]: import os

os.listdir(path)
```

```
Out[ ]: ['sunflower', 'daisy', 'rose', 'tulip', 'dandelion']
```

Question 02

The dataset contains five classes of flower images of variable size namely chamomile, tulip, rose, sunflower, and dandelion. Resize all images to 80×80

pixels and convert all colour images to grey images.

```
In [ ]: # Image Parameters
N_CLASSES = 5
IMG_SIZE = 80
DIR=path
FLOWER_DAISY_DIR=path+'/daisy'
FLOWER_SUNFLOWER_DIR=path+'/sunflower'
FLOWER_TULIP_DIR=path+'/tulip'
FLOWER_DANDI_DIR=path+'/dandelion'
FLOWER_ROSE_DIR=path+'/rose'
```

```
In [ ]: X=[] # Contains the image
Y=[] # Contains the labels
```

```
In [ ]: !pip install tqdm
```

Requirement already satisfied: tqdm in /usr/local/lib/python3.10/dist-packages (4.66.6)

```
In [ ]: from tqdm import tqdm
from PIL import Image
import numpy as np

def train_data(flower_type,path_dir):
    for img in tqdm(os.listdir(path_dir)):
        label=flower_type
        path = os.path.join(path_dir,img)
        img_array = Image.open(path).convert('L')
        img_array = img_array.resize((IMG_SIZE,IMG_SIZE))
        img_array = np.array(img_array)
        X.append(np.array(img_array))
        Y.append(str(label))
```

```
In [ ]: train_data('Daisy',FLOWER_DAISY_DIR)
train_data('Sunflower',FLOWER_SUNFLOWER_DIR)
train_data('Tulip',FLOWER_TULIP_DIR)
train_data('Dandelion',FLOWER_DANDI_DIR)
train_data('Rose',FLOWER_ROSE_DIR)
```

```
100%|██████████| 764/764 [00:01<00:00, 454.01it/s]
100%|██████████| 733/733 [00:01<00:00, 463.99it/s]
100%|██████████| 984/984 [00:02<00:00, 406.15it/s]
100%|██████████| 1052/1052 [00:02<00:00, 395.62it/s]
100%|██████████| 784/784 [00:01<00:00, 562.45it/s]
```


```
In [ ]: print(f"The number of samples in the dataset is {len(X)}")
```

The number of samples in the dataset is 4317

```
In [ ]: import matplotlib.pyplot as plt
```

```
plt.imshow(X[4235], cmap='gray')  
plt.show()
```

```
print(Y[4235])
```



Rose

Question 03

Randomly shuffle all images to create training, test set with ratio 90: 10, respectively. (Reduce the training size by 1/5, if computation resources are limited.)

```
In [ ]: from sklearn.model_selection import train_test_split
```

```
X_train, X_rem, y_train, y_rem = train_test_split(X, Y, test_size=0.2, random_state=7)
```

```
X_test,X_val,y_test,y_val=train_test_split(X_rem,y_rem,test_size=0.5,random_state=7)

print(len(X_train))
print(len(y_train))
print(len(X_val))
print(len(y_val))
print(len(X_test))
print(len(y_test))
```

3453

3453

432

432

432

432

In []: *# Convert to np array for tf processing*

```
X_train=np.array(X_train)
X_test=np.array(X_test)
X_val=np.array(X_val)
y_train=np.array(y_train)
y_val=np.array(y_val)
y_test=np.array(y_test)

# Reshape
X_train = X_train.reshape(-1,IMG_SIZE,IMG_SIZE,1)
X_test = X_test.reshape(-1,IMG_SIZE,IMG_SIZE,1)
X_val = X_val.reshape(-1,IMG_SIZE,IMG_SIZE,1)
```

In []: *# One hot encode the labels*

```
from sklearn.preprocessing import LabelEncoder

label_encoder = LabelEncoder()

label_encoder.fit(['Daisy','Sunflower','Tulip','Dandelion','Rose'])

y_train = label_encoder.transform(y_train)
y_test = label_encoder.transform(y_test)
y_val = label_encoder.transform(y_val)


import pandas as pd

y_train = pd.get_dummies(y_train,dtype='int').to_numpy()
y_test = pd.get_dummies(y_test,dtype='int').to_numpy()
y_val = pd.get_dummies(y_val,dtype='int').to_numpy()
```

```
In [ ]: import matplotlib.pyplot as plt

plt.imshow(X_train[135], cmap='gray')
plt.show()

print(y_train[135])
```



```
[0 1 0 0 0]
```

Question 04

- iv. Train a Convolutional Neural Network with max pooling and a fully connected layer at top, to classify the flower images. Now run the network by changing the following hyper-parameters:
- a. Analyze the performance of convolution window kernel size

Convolution Layers	Convolution Kernel Size	Convolution Filters Size	Pooling Layers	Activation	Fully Connected Layer (After Flatten)	Regularization
3	$(3 \times 3, 3 \times 3, 3 \times 3)$	[16,32,64]	Max Pooling	ReLU	1	Dropout of 0.1 after each layer
3	$(3 \times 3, 3 \times 3, 5 \times 5)$	[16,32,64]	Max Pooling	ReLU	1	Dropout of 0.1 after each layer
3	$(3 \times 3, 5 \times 5, 5 \times 5)$	[16,32,64]	Max Pooling	ReLU	1	Dropout of 0.1 after each layer
3	$(5 \times 5, 5 \times 5, 5 \times 5)$	[16,32,64]	Max Pooling	ReLU	1	Dropout of 0.1 after each layer

- b. For the best set of parameters obtained above, use two and three fully connected layers (After Flatten).
- c. For the best set of parameters obtained above, use average pooling instead of Max pooling.

```
In [ ]: # Import necessary libraries
import tensorflow as tf
```

```
In [ ]: from tensorflow.keras import Sequential
from tensorflow.keras.layers import Input, Conv2D, MaxPool2D, AveragePooling2D, Flatten, Dense, \
Dropout, BatchNormalization, LeakyReLU

from tensorflow.keras.metrics import F1Score
from tensorflow.math import confusion_matrix
from tensorflow.keras.losses import CategoricalCrossentropy
from tensorflow.keras.callbacks import EarlyStopping
from tensorflow.keras.optimizers import Adam
```

```
In [ ]: # Plot the accuracy and loss metrics of the model
```

```

def plot_model_metrics(
    history: 'tf.keras.callbacks.history',
    kernels: 'list(tuple(int,int))',
    filters: 'list(int)',
    activation_func: 'str',
    pool: 'str',
    num_dense_layers: 'int'
):
    plt.plot(history.history['loss'],label='Training loss')
    plt.plot(history.history['val_loss'],label='Validation loss')

    plt.ylabel('Loss')
    plt.xlabel('Epochs')
    plt.legend()
    plt.title(f"Filters: {filters}, Kernels: {kernels}, {pool} pool, {activation_func} activation function, No. of dense layers a

    plt.show()

    plt.plot(history.history['accuracy'],label='Training accuracy')
    plt.plot(history.history['val_accuracy'],label='Validation accuracy')
    plt.ylabel('Accuracy')
    plt.xlabel('Epochs')
    plt.legend()
    plt.title(f"Filters: {filters}, Kernels: {kernels}, {pool} pool, {activation_func} activation function, No. of dense layers a

    plt.show()

```

```

In [ ]: # Plot the confusion matrix of the model
from sklearn.metrics import ConfusionMatrixDisplay

def plot_confusion_matrix(
    y_test: 'list(int)',
    y_pred: 'list(int)'
):
    matrix = confusion_matrix(y_test,y_pred)
    # matrix = matrix.numpy()
    disp = ConfusionMatrixDisplay(confusion_matrix=matrix,display_labels=label_encoder.classes_)
    disp.plot(cmap=plt.cm.Blues)
    plt.title("Confusion Matrix for the above model")
    plt.show()

```

```

In [ ]: from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
from sklearn.preprocessing import LabelEncoder

def plot_confusion_matrix_mnist(y_test, y_pred):
    """Plots the confusion matrix for the given true and predicted labels, normalised"""

```

```

label_encoder = LabelEncoder()
label_encoder.fit(np.unique(np.concatenate((y_test, y_pred))))
matrix = confusion_matrix(y_test, y_pred, labels=label_encoder.classes_)
matrix = matrix.astype('float') / matrix.sum(axis=1)[:, np.newaxis] # Normalize the confusion matrix
disp = ConfusionMatrixDisplay(confusion_matrix=matrix,
                              display_labels=label_encoder.classes_)

disp.plot(cmap=plt.cm.Blues,
          xticks_rotation='vertical',
          include_values=True,
          values_format=".2f"
          )

plt.title("Confusion Matrix for the above model")
plt.tight_layout()
plt.show()

```

In []: # Train the model

```

import time

def train_model(
    kernels: 'list(tuple(int,int))',
    filters: 'list(int)',
    activation_func: 'str',
    pool: 'str',
    dropout_rate: 'float',
    num_dense_layers: 'int',
    X_train: 'numpy.array',
    y_train: 'numpy.array',
    X_test: 'numpy.array',
    y_test: 'numpy.array',
    add_batch_normalization=False,
    flatten_layer_size=64,
    num_epochs = 100,
    extra_conv_layers=0,
    is_rgb=False,
    is_mnist=False
):

    model = Sequential()
    input_shape=(80,80,1)
    num_classes=5
    val_data=(X_val,y_val)
    if is_rgb:
        input_shape=(80,80,3)
        val_data=(X_val_rgb,y_val_rgb)
    if is_mnist:

```



```

num_classes=10
for filter, kernel in zip(filters, kernels):
    if activation_func=='leaky_relu':
        model.add(Conv2D(
            filters=filter,
            kernel_size=kernel,
            activation=LeakyReLU(alpha=0.01)
        ))
    else:
        model.add(Conv2D(
            filters=filter,
            kernel_size=kernel,
            activation=activation_func
        ))

    if pool=='max':
        model.add(MaxPool2D())
    else:
        model.add(AveragePooling2D(pool_size=(2,2)))

    if add_batch_normalization:
        model.add(BatchNormalization())

    if dropout_rate>0:
        model.add(Dropout(rate=dropout_rate))

# Extra Conv Layers
for i in range(0, extra_conv_layers):
    filters.append(filters[-1]*2)
    kernels.append(kernels[-1])
    if activation_func=='leaky_relu':
        model.add(Conv2D(
            filters=filters[-1],
            kernel_size=kernels[-1],
            activation=LeakyReLU(alpha=0.01),
            padding='valid'
        ))
    else:
        model.add(Conv2D(
            filters=filters[-1],
            kernel_size=kernels[-1],
            activation=activation_func,
            padding='valid'
        ))

model.add(Flatten())

```

```

for i in range(num_dense_layers):
    if activation_func=='leaky_relu':
        model.add(Dense(
            units=flatten_layer_size,
            activation=LeakyReLU(alpha=0.01)
        ))
    else:
        model.add(Dense(
            units=flatten_layer_size,
            activation=activation_func
        ))

model.add(Dense(units=num_classes,activation='softmax'))

model.summary()

model.compile(
    optimizer=Adam(learning_rate=0.002),
    loss=CategoricalCrossentropy,
    metrics=['accuracy',F1Score(average='weighted')]
)

callback = [
    EarlyStopping(
        monitor = 'val_loss',
        patience = 10,
        restore_best_weights=True
    )
]

start_time = time.time()

history = model.fit(
    x=X_train,
    y=y_train,
    epochs=num_epochs,
    validation_data=val_data,
    callbacks=callback
)

end_time = time.time()

train_time=end_time-start_time

test_loss, test_accuracy ,test_f1 = model.evaluate(X_test,y_test)

```

```

print(f"Test Loss: {test_loss}, Test Accuracy: {test_accuracy}, Test F1 Score: {test_f1}")

print(f"Time required to train the model is {train_time} seconds")

plot_model_metrics(
    history=history,
    kernels=kernels,
    filters=filters,
    activation_func=activation_func,
    pool=pool,
    num_dense_layers=num_dense_layers
)

y_pred = model.predict(X_test)
y_pred_classes = np.argmax(y_pred,axis=1)
y_true_classes = np.argmax(y_test,axis=1)

if is_mnist:
    plot_confusion_matrix_mnist(
        y_test=y_true_classes,
        y_pred=y_pred_classes
    )
else:
    plot_confusion_matrix(
        y_test=y_true_classes,
        y_pred=y_pred_classes
    )

return test_loss,test_accuracy,test_f1,train_time,model

```

```

In [ ]: result_df_1 = pd.DataFrame(
    columns=[
        'Conv Kernel Size',
        'Conv Filter Size',
        'Pooling Layers',
        'Activation Function',
        'No. of Dense Layers after Flatten',
        'Dropout Rate',
        'Test Loss',
        'Test Accuracy',
        'Test F1 Score',
        'Training Time(in seconds)'
    ]
)

```

```

kernels = [
    [(3,3),(3,3),(3,3)],
    [(3,3),(3,3),(5,5)],
    [(3,3),(5,5),(5,5)],
    [(5,5),(5,5),(5,5)]
]

filters = [16,32,64]
activation='relu'
dropout_rate = 0.1
num_dense_layers = 0
pool='max'
epochs = 20

for kernel in kernels:
    test_loss,test_accuracy,test_f1,train_time,_ = train_model(
        kernels=kernel,
        filters=filters,
        activation_func=activation,
        pool=pool,
        dropout_rate=dropout_rate,
        num_dense_layers=num_dense_layers,
        X_train=X_train,
        y_train=y_train,
        X_test=X_test,
        y_test=y_test,
        num_epochs=epochs
    )

    result_df_1.loc[len(result_df_1.index)]=[
        kernel,
        filters,
        pool,
        activation,
        num_dense_layers,
        dropout_rate,
        test_loss,
        test_accuracy,
        test_f1,
        train_time
    ]

```















Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 78, 78, 16)	160
max_pooling2d (MaxPooling2D)	(None, 39, 39, 16)	0
dropout (Dropout)	(None, 39, 39, 16)	0
flatten (Flatten)	(None, 24336)	0
dense (Dense)	(None, 5)	121,685


Total params: 121,845 (475.96 KB)

Trainable params: 121,845 (475.96 KB)


Non-trainable params: 0 (0.00 B)

Epoch 1/20
108/108  **19s** 147ms/step - accuracy: 0.2377 - f1_score: 0.2347 - loss: 233.4431 - val_accuracy: 0.2477 - val_f1_score: 0.2424 - val_loss: 4.1256
Epoch 2/20
108/108  **14s** 93ms/step - accuracy: 0.4580 - f1_score: 0.4575 - loss: 1.9968 - val_accuracy: 0.2847 - val_f1_score: 0.2793 - val_loss: 3.5507
Epoch 3/20
108/108  **11s** 98ms/step - accuracy: 0.6470 - f1_score: 0.6464 - loss: 0.9763 - val_accuracy: 0.2870 - val_f1_score: 0.2894 - val_loss: 3.4204
Epoch 4/20
108/108  **10s** 90ms/step - accuracy: 0.7677 - f1_score: 0.7674 - loss: 0.6549 - val_accuracy: 0.3056 - val_f1_score: 0.2988 - val_loss: 3.6281
Epoch 5/20
108/108  **10s** 93ms/step - accuracy: 0.8737 - f1_score: 0.8738 - loss: 0.4159 - val_accuracy: 0.3148 - val_f1_score: 0.3075 - val_loss: 3.8635
Epoch 6/20
108/108  **11s** 100ms/step - accuracy: 0.8997 - f1_score: 0.8999 - loss: 0.3422 - val_accuracy: 0.3079 - val_f1_score: 0.2994 - val_loss: 4.2288
Epoch 7/20
108/108  **19s** 91ms/step - accuracy: 0.9322 - f1_score: 0.9324 - loss: 0.2716 - val_accuracy: 0.3194 - val_f1_score: 0.3144 - val_loss: 5.0555
Epoch 8/20
108/108  **10s** 88ms/step - accuracy: 0.9266 - f1_score: 0.9267 - loss: 0.3030 - val_accuracy: 0.3241 - val_f1_score: 0.3181 - val_loss: 4.9665
Epoch 9/20
108/108  **11s** 94ms/step - accuracy: 0.9281 - f1_score: 0.9282 - loss: 0.2895 - val_accuracy: 0.3449 - val_f1_score: 0.3406 - val_loss: 5.1423
Epoch 10/20
108/108  **11s** 102ms/step - accuracy: 0.9286 - f1_score: 0.9288 - loss: 0.2704 - val_accuracy: 0.3426 - val_f1_score: 0.3261 - val_loss: 5.6000
Epoch 11/20
108/108  **20s** 95ms/step - accuracy: 0.9363 - f1_score: 0.9363 - loss: 0.2558 - val_accuracy: 0.3333 - val_f1_score: 0.3300 - val_loss: 5.9610
Epoch 12/20
108/108  **21s** 99ms/step - accuracy: 0.9401 - f1_score: 0.9402 - loss: 0.2430 - val_accuracy: 0.3403 - val_f1_score: 0.3412 - val_loss: 5.9284
Epoch 13/20
108/108  **21s** 105ms/step - accuracy: 0.9434 - f1_score: 0.9434 - loss: 0.2411 - val_accuracy: 0.3356 - val_f1_score: 0.3386 - val_loss: 6.2893
14/14  **0s** 26ms/step - accuracy: 0.3013 - f1_score: 0.3062 - loss: 3.1925
Test Loss: 3.092764377593994, Test Accuracy: 0.32870370149612427, Test F1 Score: 0.3337514400482178
Time required to train the model is 197.24697875976562 seconds


Filters: [16, 32, 64], Kernels: [(3, 3), (3, 3), (3, 3)], max pool, relu activation function, No. of dense layers after flatten: 0



Filters: [16, 32, 64], Kernels: [(3, 3), (3, 3), (3, 3)], max pool, relu activation function, No. of dense layers after flatten: 0



14/14 0s 28ms/step

















Model: "sequential_1"

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 78, 78, 16)	160
max_pooling2d_1 (MaxPooling2D)	(None, 39, 39, 16)	0
dropout_1 (Dropout)	(None, 39, 39, 16)	0
flatten_1 (Flatten)	(None, 24336)	0
dense_1 (Dense)	(None, 5)	121,685


Total params: 121,845 (475.96 KB)

Trainable params: 121,845 (475.96 KB)


Non-trainable params: 0 (0.00 B)

Epoch 1/20
108/108  **21s** 100ms/step - accuracy: 0.2314 - f1_score: 0.2294 - loss: 237.6128 - val_accuracy: 0.2847 - val_f1_score: 0.2807 - val_loss: 3.0241
Epoch 2/20
108/108  **21s** 105ms/step - accuracy: 0.4779 - f1_score: 0.4776 - loss: 1.6270 - val_accuracy: 0.3056 - val_f1_score: 0.3026 - val_loss: 2.8588
Epoch 3/20
108/108  **19s** 88ms/step - accuracy: 0.6713 - f1_score: 0.6697 - loss: 0.8670 - val_accuracy: 0.3310 - val_f1_score: 0.3266 - val_loss: 2.7278
Epoch 4/20
108/108  **11s** 97ms/step - accuracy: 0.8005 - f1_score: 0.8005 - loss: 0.5713 - val_accuracy: 0.3241 - val_f1_score: 0.3199 - val_loss: 3.0708
Epoch 5/20
108/108  **11s** 101ms/step - accuracy: 0.8475 - f1_score: 0.8482 - loss: 0.4564 - val_accuracy: 0.3519 - val_f1_score: 0.3450 - val_loss: 3.3716
Epoch 6/20
108/108  **19s** 86ms/step - accuracy: 0.8912 - f1_score: 0.8917 - loss: 0.3845 - val_accuracy: 0.3495 - val_f1_score: 0.3378 - val_loss: 3.7420
Epoch 7/20
108/108  **12s** 100ms/step - accuracy: 0.9026 - f1_score: 0.9028 - loss: 0.3535 - val_accuracy: 0.3426 - val_f1_score: 0.3333 - val_loss: 3.9093
Epoch 8/20
108/108  **11s** 101ms/step - accuracy: 0.9235 - f1_score: 0.9240 - loss: 0.2933 - val_accuracy: 0.3264 - val_f1_score: 0.3184 - val_loss: 4.1339
Epoch 9/20
108/108  **20s** 94ms/step - accuracy: 0.9190 - f1_score: 0.9192 - loss: 0.2744 - val_accuracy: 0.3333 - val_f1_score: 0.3247 - val_loss: 4.4774
Epoch 10/20
108/108  **11s** 103ms/step - accuracy: 0.9277 - f1_score: 0.9282 - loss: 0.2795 - val_accuracy: 0.3380 - val_f1_score: 0.3361 - val_loss: 4.6995
Epoch 11/20
108/108  **11s** 98ms/step - accuracy: 0.9259 - f1_score: 0.9264 - loss: 0.3327 - val_accuracy: 0.3380 - val_f1_score: 0.3302 - val_loss: 4.6844
Epoch 12/20
108/108  **20s** 95ms/step - accuracy: 0.9408 - f1_score: 0.9415 - loss: 0.2403 - val_accuracy: 0.3519 - val_f1_score: 0.3288 - val_loss: 5.6084
Epoch 13/20
108/108  **21s** 99ms/step - accuracy: 0.9542 - f1_score: 0.9545 - loss: 0.1861 - val_accuracy: 0.3287 - val_f1_score: 0.3192 - val_loss: 5.4493
14/14  **0s** 24ms/step - accuracy: 0.2739 - f1_score: 0.2678 - loss: 3.1269
Test Loss: 2.8645572662353516, Test Accuracy: 0.31018519401550293, Test F1 Score: 0.30755409598350525
Time required to train the model is 198.43842554092407 seconds


Filters: [16, 32, 64], Kernels: [(3, 3), (3, 3), (5, 5)], max pool, relu activation function, No. of dense layers after flatten: 0



Filters: [16, 32, 64], Kernels: [(3, 3), (3, 3), (5, 5)], max pool, relu activation function, No. of dense layers after flatten: 0



14/14 0s 27ms/step
















Model: "sequential_2"

Layer (type)	Output Shape	Param #
conv2d_2 (Conv2D)	(None, 78, 78, 16)	160
max_pooling2d_2 (MaxPooling2D)	(None, 39, 39, 16)	0
dropout_2 (Dropout)	(None, 39, 39, 16)	0
flatten_2 (Flatten)	(None, 24336)	0
dense_2 (Dense)	(None, 5)	121,685


Total params: 121,845 (475.96 KB)

Trainable params: 121,845 (475.96 KB)


Non-trainable params: 0 (0.00 B)

Epoch 1/20
108/108  **13s** 107ms/step - accuracy: 0.2521 - f1_score: 0.2503 - loss: 492.9043 - val_accuracy: 0.3102 - val_f1_score: 0.2516 - val_loss: 8.3644
Epoch 2/20
108/108  **19s** 90ms/step - accuracy: 0.4321 - f1_score: 0.4323 - loss: 4.2141 - val_accuracy: 0.3380 - val_f1_score: 0.3099 - val_loss: 3.6129
Epoch 3/20
108/108  **11s** 104ms/step - accuracy: 0.6363 - f1_score: 0.6363 - loss: 1.1693 - val_accuracy: 0.3542 - val_f1_score: 0.3352 - val_loss: 3.7134
Epoch 4/20
108/108  **20s** 95ms/step - accuracy: 0.7514 - f1_score: 0.7510 - loss: 0.6698 - val_accuracy: 0.3495 - val_f1_score: 0.3391 - val_loss: 3.6550
Epoch 5/20
108/108  **10s** 94ms/step - accuracy: 0.8293 - f1_score: 0.8293 - loss: 0.4588 - val_accuracy: 0.3519 - val_f1_score: 0.3455 - val_loss: 3.8674
Epoch 6/20
108/108  **22s** 105ms/step - accuracy: 0.8783 - f1_score: 0.8784 - loss: 0.3393 - val_accuracy: 0.3519 - val_f1_score: 0.3518 - val_loss: 3.8306
Epoch 7/20
108/108  **19s** 90ms/step - accuracy: 0.9092 - f1_score: 0.9094 - loss: 0.2698 - val_accuracy: 0.3542 - val_f1_score: 0.3388 - val_loss: 4.2922
Epoch 8/20
108/108  **11s** 100ms/step - accuracy: 0.9270 - f1_score: 0.9271 - loss: 0.2313 - val_accuracy: 0.3796 - val_f1_score: 0.3774 - val_loss: 4.2985
Epoch 9/20
108/108  **21s** 109ms/step - accuracy: 0.9416 - f1_score: 0.9417 - loss: 0.1929 - val_accuracy: 0.3773 - val_f1_score: 0.3758 - val_loss: 4.6168
Epoch 10/20
108/108  **10s** 93ms/step - accuracy: 0.9514 - f1_score: 0.9515 - loss: 0.1693 - val_accuracy: 0.3866 - val_f1_score: 0.3813 - val_loss: 4.4255
Epoch 11/20
108/108  **10s** 94ms/step - accuracy: 0.9465 - f1_score: 0.9466 - loss: 0.1899 - val_accuracy: 0.3981 - val_f1_score: 0.3915 - val_loss: 4.8634
Epoch 12/20
108/108  **22s** 105ms/step - accuracy: 0.9627 - f1_score: 0.9627 - loss: 0.1619 - val_accuracy: 0.3773 - val_f1_score: 0.3745 - val_loss: 4.9914
14/14  **0s** 29ms/step - accuracy: 0.3556 - f1_score: 0.3423 - loss: 3.7078
Test Loss: 3.545902729034424, Test Accuracy: 0.3541666567325592, Test F1 Score: 0.33932924270629883
Time required to train the model is 187.31955242156982 seconds


Filters: [16, 32, 64], Kernels: [(3, 3), (5, 5), (5, 5)], max pool, relu activation function, No. of dense layers after flatten: 0



Filters: [16, 32, 64], Kernels: [(3, 3), (5, 5), (5, 5)], max pool, relu activation function, No. of dense layers after flatten: 0



14/14 0s 28ms/step















Model: "sequential_3"

Layer (type)	Output Shape	Param #
conv2d_3 (Conv2D)	(None, 76, 76, 16)	416
max_pooling2d_3 (MaxPooling2D)	(None, 38, 38, 16)	0
dropout_3 (Dropout)	(None, 38, 38, 16)	0
flatten_3 (Flatten)	(None, 23104)	0
dense_3 (Dense)	(None, 5)	115,525


Total params: 115,941 (452.89 KB)

Trainable params: 115,941 (452.89 KB)


Non-trainable params: 0 (0.00 B)

Epoch 1/20
108/108  **15s** 132ms/step - accuracy: 0.2181 - f1_score: 0.2116 - loss: 204.3684 - val_accuracy: 0.2407 - val_f1_score: 0.2167 - val_loss: 1.7584
Epoch 2/20
108/108  **14s** 126ms/step - accuracy: 0.3308 - f1_score: 0.3152 - loss: 1.4987 - val_accuracy: 0.2639 - val_f1_score: 0.2283 - val_loss: 1.8946
Epoch 3/20
108/108  **21s** 129ms/step - accuracy: 0.4827 - f1_score: 0.4761 - loss: 1.2595 - val_accuracy: 0.2731 - val_f1_score: 0.2455 - val_loss: 2.0592
Epoch 4/20
108/108  **14s** 130ms/step - accuracy: 0.5538 - f1_score: 0.5586 - loss: 1.0850 - val_accuracy: 0.2894 - val_f1_score: 0.2590 - val_loss: 2.2781
Epoch 5/20
108/108  **14s** 128ms/step - accuracy: 0.6392 - f1_score: 0.6459 - loss: 0.9273 - val_accuracy: 0.3056 - val_f1_score: 0.2865 - val_loss: 2.4814
Epoch 6/20
108/108  **21s** 132ms/step - accuracy: 0.7011 - f1_score: 0.7072 - loss: 0.8249 - val_accuracy: 0.2569 - val_f1_score: 0.2450 - val_loss: 3.3964
Epoch 7/20
108/108  **14s** 131ms/step - accuracy: 0.7224 - f1_score: 0.7279 - loss: 0.7498 - val_accuracy: 0.2963 - val_f1_score: 0.2802 - val_loss: 3.3382
Epoch 8/20
108/108  **15s** 135ms/step - accuracy: 0.7481 - f1_score: 0.7538 - loss: 0.7140 - val_accuracy: 0.2616 - val_f1_score: 0.2496 - val_loss: 3.7218
Epoch 9/20
108/108  **14s** 128ms/step - accuracy: 0.7450 - f1_score: 0.7521 - loss: 0.7188 - val_accuracy: 0.2824 - val_f1_score: 0.2427 - val_loss: 3.5226
Epoch 10/20
108/108  **21s** 135ms/step - accuracy: 0.6787 - f1_score: 0.6840 - loss: 0.8627 - val_accuracy: 0.2917 - val_f1_score: 0.2582 - val_loss: 4.4186
Epoch 11/20
108/108  **21s** 140ms/step - accuracy: 0.7338 - f1_score: 0.7402 - loss: 0.6956 - val_accuracy: 0.2731 - val_f1_score: 0.2418 - val_loss: 3.9892
14/14  **1s** 59ms/step - accuracy: 0.2363 - f1_score: 0.2113 - loss: 1.6868
Test Loss: 1.7128006219863892, Test Accuracy: 0.2569444477558136, Test F1 Score: 0.22423039376735687
Time required to train the model is 189.06018662452698 seconds


Filters: [16, 32, 64], Kernels: [(5, 5), (5, 5), (5, 5)], max pool, relu activation function, No. of dense layers after flatten: 0



Filters: [16, 32, 64], Kernels: [(5, 5), (5, 5), (5, 5)], max pool, relu activation function, No. of dense layers after flatten: 0



14/14 ————— 1s 45ms/step



In []: result_df_1

Out[]:	Conv Kernel Size	Conv Filter Size	Pooling Layers	Activation Function	No. of Dense Layers after Flatten	Dropout Rate	Test Loss	Test Accuracy	Test F1 Score	Training Time(in seconds)
0	[(3, 3), (3, 3), (3, 3)]	[16, 32, 64]	max	relu	0	0.1	3.092764	0.328704	0.333751	197.246979
1	[(3, 3), (3, 3), (5, 5)]	[16, 32, 64]	max	relu	0	0.1	2.864557	0.310185	0.307554	198.438426
2	[(3, 3), (5, 5), (5, 5)]	[16, 32, 64]	max	relu	0	0.1	3.545903	0.354167	0.339329	187.319552
3	[(5, 5), (5, 5), (5, 5)]	[16, 32, 64]	max	relu	0	0.1	1.712801	0.256944	0.224230	189.060187


In []: plt.figure(figsize=(20,5))
plt.bar(

```

[str(ker) for ker in result_df_1['Conv Kernel Size']],
result_df_1['Test Accuracy']
)

plt.ylabel('Test Accuracy')
plt.xlabel('Convolution Kernels')
plt.title('Validation Accuracy for various Convolution Kernels')
plt.show()

```



```

In [ ]: best_kernel = result_df_1.sort_values(
        by=['Test Accuracy', 'Test F1 Score'],
        ascending=[False, False]
        )['Conv Kernel Size'].iloc[0]

best_kernel

```

Out[]: [(3, 3), (5, 5), (5, 5)]

In []: ## Subtask 2

```

result_df_2 = pd.DataFrame(
    columns=[
        'Conv Kernel Size',
        'Conv Filter Size',
        'Pooling Layers',
        'Activation Function',
        'No. of Dense Layers after Flatten',
        'Dropout Rate',
    ]
)

```

```

        'Test Loss',
        'Test Accuracy',
        'Test F1 Score',
        'Training Time(in seconds)'
    ]
)

filters = [16,32,64]
activation='relu'
dropout_rate = 0.1
pool='max'
epochs = 20
num_dense_layers = [1,2,3]

for layer in num_dense_layers:
    test_loss,test_accuracy,test_f1,train_time,_ = train_model(
        kernels=best_kernel,
        filters=filters,
        activation_func=activation,
        pool=pool,
        dropout_rate=dropout_rate,
        num_dense_layers=layer,
        X_train=X_train,
        y_train=y_train,
        X_test=X_test,
        y_test=y_test,
        num_epochs=epochs
    )

    result_df_2.loc[len(result_df_2.index)]=[
        best_kernel,
        filters,
        pool,
        activation,
        layer,
        dropout_rate,
        test_loss,
        test_accuracy,
        test_f1,
        train_time
    ]

```














Model: "sequential_4"

Layer (type)	Output Shape	Param #
conv2d_4 (Conv2D)	(None, 78, 78, 16)	160
max_pooling2d_4 (MaxPooling2D)	(None, 39, 39, 16)	0
dropout_4 (Dropout)	(None, 39, 39, 16)	0
flatten_4 (Flatten)	(None, 24336)	0
dense_4 (Dense)	(None, 64)	1,557,568
dense_5 (Dense)	(None, 5)	325


Total params: 1,558,053 (5.94 MB)

Trainable params: 1,558,053 (5.94 MB)


Non-trainable params: 0 (0.00 B)

Epoch 1/20
108/108  **17s** 139ms/step - accuracy: 0.2455 - f1_score: 0.2424 - loss: 128.6618 - val_accuracy: 0.3148 - val_f1_score: 0.3127 - val_loss: 2.3806
Epoch 2/20
108/108  **18s** 119ms/step - accuracy: 0.5185 - f1_score: 0.5174 - loss: 1.2598 - val_accuracy: 0.3102 - val_f1_score: 0.3035 - val_loss: 2.3428
Epoch 3/20
108/108  **12s** 114ms/step - accuracy: 0.7168 - f1_score: 0.7161 - loss: 0.8011 - val_accuracy: 0.3449 - val_f1_score: 0.3353 - val_loss: 2.6991
Epoch 4/20
108/108  **12s** 115ms/step - accuracy: 0.8022 - f1_score: 0.8019 - loss: 0.6054 - val_accuracy: 0.3657 - val_f1_score: 0.3575 - val_loss: 3.0422
Epoch 5/20
108/108  **20s** 109ms/step - accuracy: 0.8647 - f1_score: 0.8642 - loss: 0.4511 - val_accuracy: 0.3657 - val_f1_score: 0.3642 - val_loss: 3.3770
Epoch 6/20
108/108  **22s** 123ms/step - accuracy: 0.8922 - f1_score: 0.8918 - loss: 0.3890 - val_accuracy: 0.3634 - val_f1_score: 0.3528 - val_loss: 3.7982
Epoch 7/20
108/108  **13s** 122ms/step - accuracy: 0.9092 - f1_score: 0.9089 - loss: 0.3430 - val_accuracy: 0.3611 - val_f1_score: 0.3421 - val_loss: 4.0634
Epoch 8/20
108/108  **20s** 115ms/step - accuracy: 0.9246 - f1_score: 0.9246 - loss: 0.2810 - val_accuracy: 0.3657 - val_f1_score: 0.3604 - val_loss: 4.4306
Epoch 9/20
108/108  **12s** 115ms/step - accuracy: 0.9293 - f1_score: 0.9290 - loss: 0.2849 - val_accuracy: 0.3796 - val_f1_score: 0.3662 - val_loss: 4.5676
Epoch 10/20
108/108  **12s** 113ms/step - accuracy: 0.9436 - f1_score: 0.9435 - loss: 0.2320 - val_accuracy: 0.3681 - val_f1_score: 0.3609 - val_loss: 4.9733
Epoch 11/20
108/108  **19s** 97ms/step - accuracy: 0.9429 - f1_score: 0.9426 - loss: 0.2837 - val_accuracy: 0.3796 - val_f1_score: 0.3762 - val_loss: 5.1438
Epoch 12/20
108/108  **22s** 114ms/step - accuracy: 0.9469 - f1_score: 0.9467 - loss: 0.2289 - val_accuracy: 0.3542 - val_f1_score: 0.3451 - val_loss: 5.9310
14/14  **0s** 27ms/step - accuracy: 0.2672 - f1_score: 0.2628 - loss: 2.1882
Test Loss: 2.231379985809326, Test Accuracy: 0.28703704476356506, Test F1 Score: 0.2833707928657532
Time required to train the model is 208.50289726257324 seconds


Filters: [16, 32, 64], Kernels: [(3, 3), (5, 5), (5, 5)], max pool, relu activation function, No. of dense layers after flatten: 1



Filters: [16, 32, 64], Kernels: [(3, 3), (5, 5), (5, 5)], max pool, relu activation function, No. of dense layers after flatten: 1



14/14 ————— 1s 33ms/step



Model: "sequential_5"


Layer (type)	Output Shape	Param #
conv2d_5 (Conv2D)	(None, 78, 78, 16)	160
max_pooling2d_5 (MaxPooling2D)	(None, 39, 39, 16)	0
dropout_5 (Dropout)	(None, 39, 39, 16)	0
flatten_5 (Flatten)	(None, 24336)	0
dense_6 (Dense)	(None, 64)	1,557,568
dense_7 (Dense)	(None, 64)	4,160
dense_8 (Dense)	(None, 5)	325

Total params: 1,562,213 (5.96 MB)


Trainable params: 1,562,213 (5.96 MB)

Non-trainable params: 0 (0.00 B)


Epoch 1/20

108/108  **14s** 117ms/step - accuracy: 0.2392 - f1_score: 0.2377 - loss: 159.1355 - val_accuracy: 0.3727 - val_f1_score: 0.3513 - val_loss: 1.9389


Epoch 2/20

108/108  **19s** 103ms/step - accuracy: 0.4622 - f1_score: 0.4577 - loss: 1.3670 - val_accuracy: 0.3866 - val_f1_score: 0.3790 - val_loss: 1.8688


Epoch 3/20

108/108  **22s** 117ms/step - accuracy: 0.6879 - f1_score: 0.6870 - loss: 0.8206 - val_accuracy: 0.3843 - val_f1_score: 0.3855 - val_loss: 2.0543


Epoch 4/20

108/108  **21s** 118ms/step - accuracy: 0.7927 - f1_score: 0.7936 - loss: 0.5902 - val_accuracy: 0.3611 - val_f1_score: 0.3582 - val_loss: 2.4087


Epoch 5/20

108/108  **13s** 123ms/step - accuracy: 0.8671 - f1_score: 0.8675 - loss: 0.4360 - val_accuracy: 0.3681 - val_f1_score: 0.3527 - val_loss: 2.7063


Epoch 6/20

108/108  **12s** 114ms/step - accuracy: 0.9006 - f1_score: 0.9011 - loss: 0.3367 - val_accuracy: 0.3843 - val_f1_score: 0.3803 - val_loss: 2.6917


Epoch 7/20

108/108  **21s** 116ms/step - accuracy: 0.9243 - f1_score: 0.9245 - loss: 0.2840 - val_accuracy: 0.3796 - val_f1_score: 0.3780 - val_loss: 3.2956


Epoch 8/20

108/108  **13s** 121ms/step - accuracy: 0.9312 - f1_score: 0.9311 - loss: 0.2383 - val_accuracy: 0.3773 - val_f1_score: 0.3839 - val_loss: 3.7464


Epoch 9/20

108/108  **20s** 114ms/step - accuracy: 0.9438 - f1_score: 0.9439 - loss: 0.2240 - val_accuracy: 0.3796 - val_f1_score: 0.3733 - val_loss: 3.6956


Epoch 10/20

108/108  **21s** 121ms/step - accuracy: 0.9513 - f1_score: 0.9514 - loss: 0.1860 - val_accuracy: 0.3981 - val_f1_score: 0.3870 - val_loss: 4.5019

Epoch 11/20

108/108  **13s** 121ms/step - accuracy: 0.9579 - f1_score: 0.9578 - loss: 0.1986 - val_accuracy: 0.3819 - val_f1_score: 0.3613 - val_loss: 4.2326

Epoch 12/20


108/108  **20s** 113ms/step - accuracy: 0.9680 - f1_score: 0.9681 - loss: 0.1501 - val_accuracy: 0.3704 - val_f1_score: 0.3652 - val_loss: 4.5444

14/14  **0s** 31ms/step - accuracy: 0.3901 - f1_score: 0.3797 - loss: 1.7396


Test Loss: 1.8127901554107666, Test Accuracy: 0.38425925374031067, Test F1 Score: 0.3741406202316284

Time required to train the model is 217.4804346561432 seconds

Filters: [16, 32, 64], Kernels: [(3, 3), (5, 5), (5, 5)], max pool, relu activation function, No. of dense layers after flatten: 2




Filters: [16, 32, 64], Kernels: [(3, 3), (5, 5), (5, 5)], max pool, relu activation function, No. of dense layers after flatten: 2



14/14 ————— 1s 36ms/step

Confusion Matrix for the above model
















Model: "sequential_6"

Layer (type)	Output Shape	Param #
conv2d_6 (Conv2D)	(None, 78, 78, 16)	160
max_pooling2d_6 (MaxPooling2D)	(None, 39, 39, 16)	0
dropout_6 (Dropout)	(None, 39, 39, 16)	0
flatten_6 (Flatten)	(None, 24336)	0
dense_9 (Dense)	(None, 64)	1,557,568
dense_10 (Dense)	(None, 64)	4,160
dense_11 (Dense)	(None, 64)	4,160
dense_12 (Dense)	(None, 5)	325


Total params: 1,566,373 (5.98 MB)

Trainable params: 1,566,373 (5.98 MB)


Non-trainable params: 0 (0.00 B)

Epoch 1/20
108/108  **16s** 124ms/step - accuracy: 0.2292 - f1_score: 0.2270 - loss: 110.7956 - val_accuracy: 0.2917 - val_f1_score: 0.2805 - val_loss: 1.8010
Epoch 2/20
108/108  **14s** 126ms/step - accuracy: 0.4593 - f1_score: 0.4495 - loss: 1.3546 - val_accuracy: 0.3125 - val_f1_score: 0.2987 - val_loss: 1.7974
Epoch 3/20
108/108  **13s** 124ms/step - accuracy: 0.6648 - f1_score: 0.6635 - loss: 0.9455 - val_accuracy: 0.3935 - val_f1_score: 0.3877 - val_loss: 1.8300
Epoch 4/20
108/108  **14s** 127ms/step - accuracy: 0.7984 - f1_score: 0.8000 - loss: 0.6543 - val_accuracy: 0.3750 - val_f1_score: 0.3600 - val_loss: 2.0824
Epoch 5/20
108/108  **19s** 114ms/step - accuracy: 0.8577 - f1_score: 0.8576 - loss: 0.4979 - val_accuracy: 0.3843 - val_f1_score: 0.3739 - val_loss: 2.2979
Epoch 6/20
108/108  **12s** 113ms/step - accuracy: 0.8861 - f1_score: 0.8868 - loss: 0.4069 - val_accuracy: 0.4051 - val_f1_score: 0.3954 - val_loss: 2.5747
Epoch 7/20
108/108  **20s** 114ms/step - accuracy: 0.9052 - f1_score: 0.9056 - loss: 0.3497 - val_accuracy: 0.3727 - val_f1_score: 0.3661 - val_loss: 3.0689
Epoch 8/20
108/108  **20s** 108ms/step - accuracy: 0.9304 - f1_score: 0.9306 - loss: 0.2558 - val_accuracy: 0.3819 - val_f1_score: 0.3723 - val_loss: 3.1783
Epoch 9/20
108/108  **21s** 117ms/step - accuracy: 0.9318 - f1_score: 0.9319 - loss: 0.2608 - val_accuracy: 0.3611 - val_f1_score: 0.3570 - val_loss: 3.6998
Epoch 10/20
108/108  **21s** 125ms/step - accuracy: 0.9557 - f1_score: 0.9558 - loss: 0.1828 - val_accuracy: 0.3704 - val_f1_score: 0.3655 - val_loss: 3.8902
Epoch 11/20
108/108  **20s** 118ms/step - accuracy: 0.9671 - f1_score: 0.9671 - loss: 0.1501 - val_accuracy: 0.3704 - val_f1_score: 0.3670 - val_loss: 4.4280
Epoch 12/20
108/108  **21s** 126ms/step - accuracy: 0.9656 - f1_score: 0.9656 - loss: 0.1274 - val_accuracy: 0.3449 - val_f1_score: 0.3397 - val_loss: 5.3462
14/14  **0s** 34ms/step - accuracy: 0.3728 - f1_score: 0.3530 - loss: 1.7272
Test Loss: 1.7116447687149048, Test Accuracy: 0.375, Test F1 Score: 0.3597155809402466
Time required to train the model is 219.30019235610962 seconds


Filters: [16, 32, 64], Kernels: [(3, 3), (5, 5), (5, 5)], max pool, relu activation function, No. of dense layers after flatten: 3



Filters: [16, 32, 64], Kernels: [(3, 3), (5, 5), (5, 5)], max pool, relu activation function, No. of dense layers after flatten: 3



14/14 ————— 1s 50ms/step



In []: result_df_2


Out[]:	Conv Kernel Size	Conv Filter Size	Pooling Layers	Activation Function	No. of Dense Layers after Flatten	Dropout Rate	Test Loss	Test Accuracy	Test F1 Score	Training Time(in seconds)
0	[(3, 3), (5, 5), (5, 5)]	[16, 32, 64]	max	relu	1	0.1	2.231380	0.287037	0.283371	208.502897
1	[(3, 3), (5, 5), (5, 5)]	[16, 32, 64]	max	relu	2	0.1	1.812790	0.384259	0.374141	217.480435
2	[(3, 3), (5, 5), (5, 5)]	[16, 32, 64]	max	relu	3	0.1	1.711645	0.375000	0.359716	219.300192

```

In [ ]: plt.figure(figsize=(20,5))
plt.bar(
    result_df_2['No. of Dense Layers after Flatten'],
    result_df_2['Test Accuracy']

```

```
)
plt.ylabel('Test Accuracy')
plt.xlabel('No. of Dense Layers after Flatten')
plt.title('Validation Accuracy for various number of dense layers after flatten')
plt.show()
```



```
In [ ]: best_num_dense = result_df_2.sort_values(
        by=['Test Accuracy', 'Test F1 Score'],
        ascending=[False, False]
    )['No. of Dense Layers after Flatten'].iloc[0]

best_num_dense
```

Out[]: 2

```
In [ ]: ## Subtask 3

result_df_3 = pd.DataFrame(
    columns=[
        'Conv Kernel Size',
        'Conv Filter Size',
        'Pooling Layers',
        'Activation Function',
        'No. of Dense Layers after Flatten',
        'Dropout Rate',
        'Test Loss',
        'Test Accuracy',
```

```

        'Test F1 Score',
        'Training Time(in seconds)'
    ]
)

filters = [16,32,64]
activation='relu'
dropout_rate = 0.1
pool=['max','avg']
epochs = 20

for p in pool:
    test_loss,test_accuracy,test_f1,train_time,_ = train_model(
        kernels=best_kernel,
        filters=filters,
        activation_func=activation,
        pool=p,
        dropout_rate=dropout_rate,
        num_dense_layers=best_num_dense,
        X_train=X_train,
        y_train=y_train,
        X_test=X_test,
        y_test=y_test,
        num_epochs=epochs
    )

    result_df_3.loc[len(result_df_3.index)]=[
        best_kernel,
        filters,
        p,
        activation,
        best_num_dense,
        dropout_rate,
        test_loss,
        test_accuracy,
        test_f1,
        train_time
    ]

```














Model: "sequential_7"

Layer (type)	Output Shape	Param #
conv2d_7 (Conv2D)	(None, 78, 78, 16)	160
max_pooling2d_7 (MaxPooling2D)	(None, 39, 39, 16)	0
dropout_7 (Dropout)	(None, 39, 39, 16)	0
flatten_7 (Flatten)	(None, 24336)	0
dense_13 (Dense)	(None, 64)	1,557,568
dense_14 (Dense)	(None, 64)	4,160
dense_15 (Dense)	(None, 5)	325


Total params: 1,562,213 (5.96 MB)

Trainable params: 1,562,213 (5.96 MB)


Non-trainable params: 0 (0.00 B)

Epoch 1/20
108/108  **16s** 121ms/step - accuracy: 0.2427 - f1_score: 0.2378 - loss: 160.6646 - val_accuracy: 0.3449 - val_f1_score: 0.3434 - val_loss: 1.8507
Epoch 2/20
108/108  **21s** 123ms/step - accuracy: 0.4452 - f1_score: 0.4431 - loss: 1.4091 - val_accuracy: 0.3472 - val_f1_score: 0.3418 - val_loss: 1.7605
Epoch 3/20
108/108  **17s** 153ms/step - accuracy: 0.6456 - f1_score: 0.6436 - loss: 1.0002 - val_accuracy: 0.3912 - val_f1_score: 0.3755 - val_loss: 1.9199
Epoch 4/20
108/108  **16s** 109ms/step - accuracy: 0.7856 - f1_score: 0.7845 - loss: 0.6875 - val_accuracy: 0.3750 - val_f1_score: 0.3702 - val_loss: 2.2030
Epoch 5/20
108/108  **21s** 118ms/step - accuracy: 0.8581 - f1_score: 0.8585 - loss: 0.4827 - val_accuracy: 0.3773 - val_f1_score: 0.3655 - val_loss: 2.6122
Epoch 6/20
108/108  **19s** 104ms/step - accuracy: 0.9084 - f1_score: 0.9088 - loss: 0.3535 - val_accuracy: 0.4028 - val_f1_score: 0.4009 - val_loss: 2.7527
Epoch 7/20
108/108  **22s** 121ms/step - accuracy: 0.9149 - f1_score: 0.9149 - loss: 0.3241 - val_accuracy: 0.3727 - val_f1_score: 0.3615 - val_loss: 3.3847
Epoch 8/20
108/108  **21s** 122ms/step - accuracy: 0.9254 - f1_score: 0.9254 - loss: 0.2878 - val_accuracy: 0.3796 - val_f1_score: 0.3774 - val_loss: 3.1862
Epoch 9/20
108/108  **20s** 121ms/step - accuracy: 0.9517 - f1_score: 0.9520 - loss: 0.1843 - val_accuracy: 0.3657 - val_f1_score: 0.3578 - val_loss: 3.6371
Epoch 10/20
108/108  **14s** 126ms/step - accuracy: 0.9589 - f1_score: 0.9590 - loss: 0.1612 - val_accuracy: 0.3565 - val_f1_score: 0.3495 - val_loss: 4.1129
Epoch 11/20
108/108  **14s** 125ms/step - accuracy: 0.9620 - f1_score: 0.9620 - loss: 0.1529 - val_accuracy: 0.3565 - val_f1_score: 0.3491 - val_loss: 3.8684
Epoch 12/20
108/108  **19s** 111ms/step - accuracy: 0.9588 - f1_score: 0.9589 - loss: 0.1641 - val_accuracy: 0.3472 - val_f1_score: 0.3284 - val_loss: 5.1755
14/14  **0s** 27ms/step - accuracy: 0.3671 - f1_score: 0.3718 - loss: 1.8074
Test Loss: 1.7586485147476196, Test Accuracy: 0.375, Test F1 Score: 0.36847788095474243
Time required to train the model is 227.50457501411438 seconds


Filters: [16, 32, 64], Kernels: [(3, 3), (5, 5), (5, 5)], max pool, relu activation function, No. of dense layers after flatten: 2



Filters: [16, 32, 64], Kernels: [(3, 3), (5, 5), (5, 5)], max pool, relu activation function, No. of dense layers after flatten: 2



14/14 ————— 1s 33ms/step



Model: "sequential_8"


Layer (type)	Output Shape	Param #
conv2d_8 (Conv2D)	(None, 78, 78, 16)	160
average_pooling2d (AveragePooling2D)	(None, 39, 39, 16)	0
dropout_8 (Dropout)	(None, 39, 39, 16)	0
flatten_8 (Flatten)	(None, 24336)	0
dense_16 (Dense)	(None, 64)	1,557,568
dense_17 (Dense)	(None, 64)	4,160
dense_18 (Dense)	(None, 5)	325

Total params: 1,562,213 (5.96 MB)


Trainable params: 1,562,213 (5.96 MB)

Non-trainable params: 0 (0.00 B)


Epoch 1/20

108/108  **15s** 116ms/step - accuracy: 0.2262 - f1_score: 0.2208 - loss: 61.0823 - val_accuracy: 0.2963 - val_f1_score: 0.2594 - val_loss: 1.5667


Epoch 2/20

108/108  **13s** 116ms/step - accuracy: 0.3516 - f1_score: 0.3497 - loss: 1.4920 - val_accuracy: 0.3588 - val_f1_score: 0.3571 - val_loss: 1.5527


Epoch 3/20

108/108  **21s** 119ms/step - accuracy: 0.4960 - f1_score: 0.4950 - loss: 1.2497 - val_accuracy: 0.3843 - val_f1_score: 0.3794 - val_loss: 1.6961


Epoch 4/20

108/108  **21s** 126ms/step - accuracy: 0.6142 - f1_score: 0.6156 - loss: 1.0053 - val_accuracy: 0.3681 - val_f1_score: 0.3668 - val_loss: 1.7546


Epoch 5/20

108/108  **13s** 120ms/step - accuracy: 0.7288 - f1_score: 0.7291 - loss: 0.7755 - val_accuracy: 0.3866 - val_f1_score: 0.3760 - val_loss: 2.1521


Epoch 6/20

108/108  **20s** 113ms/step - accuracy: 0.8187 - f1_score: 0.8197 - loss: 0.5467 - val_accuracy: 0.3727 - val_f1_score: 0.3683 - val_loss: 2.4708


Epoch 7/20

108/108  **12s** 111ms/step - accuracy: 0.8703 - f1_score: 0.8708 - loss: 0.4074 - val_accuracy: 0.3889 - val_f1_score: 0.3676 - val_loss: 2.8826


Epoch 8/20

108/108  **21s** 119ms/step - accuracy: 0.9069 - f1_score: 0.9071 - loss: 0.2996 - val_accuracy: 0.3426 - val_f1_score: 0.3429 - val_loss: 3.4772


Epoch 9/20

108/108  **20s** 118ms/step - accuracy: 0.9214 - f1_score: 0.9215 - loss: 0.2600 - val_accuracy: 0.3773 - val_f1_score: 0.3615 - val_loss: 3.9255


Epoch 10/20

108/108  **20s** 116ms/step - accuracy: 0.9221 - f1_score: 0.9222 - loss: 0.2445 - val_accuracy: 0.3681 - val_f1_score: 0.3384 - val_loss: 4.1631

Epoch 11/20

108/108  **13s** 119ms/step - accuracy: 0.9474 - f1_score: 0.9475 - loss: 0.1959 - val_accuracy: 0.3519 - val_f1_score: 0.3384 - val_loss: 4.7231

Epoch 12/20


108/108  **20s** 110ms/step - accuracy: 0.9438 - f1_score: 0.9438 - loss: 0.1878 - val_accuracy: 0.3426 - val_f1_score: 0.3288 - val_loss: 4.8987

14/14  **1s** 38ms/step - accuracy: 0.3504 - f1_score: 0.3421 - loss: 1.5816


Test Loss: 1.5979018211364746, Test Accuracy: 0.32175925374031067, Test F1 Score: 0.3118511438369751

Time required to train the model is 208.19212317466736 seconds


Filters: [16, 32, 64], Kernels: [(3, 3), (5, 5), (5, 5)], avg pool, relu activation function, No. of dense layers after flatten: 2



Filters: [16, 32, 64], Kernels: [(3, 3), (5, 5), (5, 5)], avg pool, relu activation function, No. of dense layers after flatten: 2



14/14 ————— 1s 61ms/step



In []: result_df_3


Out[]:	Conv Kernel Size	Conv Filter Size	Pooling Layers	Activation Function	No. of Dense Layers after Flatten	Dropout Rate	Test Loss	Test Accuracy	Test F1 Score	Training Time(in seconds)
0	[(3, 3), (5, 5), (5, 5)]	[16, 32, 64]	max	relu	2	0.1	1.758649	0.375000	0.368478	227.504575
1	[(3, 3), (5, 5), (5, 5)]	[16, 32, 64]	avg	relu	2	0.1	1.597902	0.321759	0.311851	208.192123

```
In [ ]: plt.figure(figsize=(20,5))
plt.bar(
    result_df_3['Pooling Layers'],
    result_df_3['Test Accuracy']
)

plt.ylabel('Test Accuracy')
```



```
plt.xlabel('Pooling Layers')
plt.title('Validation Accuracy for various pooling layers')
plt.show()
```



```
In [ ]: best_pool = result_df_3.sort_values(
        by=['Test Accuracy', 'Test F1 Score'],
        ascending=[False, False]
      )['Pooling Layers'].iloc[0]

best_pool
```

```
Out[ ]: 'max'
```

- d. For the best set of parameters obtained above, use the activation function: Sigmoid, ReLU, Leaky ReLU ($\alpha = 0.01$).
- e. For the best set of parameters from the above runs vary the regularization parameter:

Regularization
Dropout of 0.25 after each layer
Batch normalization after each layer (except the first)
Dropout of 0.1 after each layer along with Batch normalization after each layer (except the first)

```
In [ ]: # Subtask 4

result_df_4 = pd.DataFrame(
    columns=[
        'Conv Kernel Size',
        'Conv Filter Size',
        'Pooling Layers',
        'Activation Function',
        'No. of Dense Layers after Flatten',
        'Dropout Rate',
        'Test Loss',
        'Test Accuracy',
        'Test F1 Score',
        'Training Time(in seconds)'
    ]
)

filters = [16,32,64]
activation_list=['sigmoid','relu','leaky_relu']
dropout_rate = 0.1
epochs = 20

for activation in activation_list:
```

```

test_loss,test_accuracy,test_f1,train_time,_ = train_model(
    kernels=best_kernel,
    filters=filters,
    activation_func=activation,
    pool=best_pool,
    dropout_rate=dropout_rate,
    num_dense_layers=best_num_dense,
    X_train=X_train,
    y_train=y_train,
    X_test=X_test,
    y_test=y_test,
    num_epochs=epochs
)

result_df_4.loc[len(result_df_4.index)]=[
    best_kernel,
    filters,
    best_pool,
    activation,
    best_num_dense,
    dropout_rate,
    test_loss,
    test_accuracy,
    test_f1,
    train_time
]

```
















Model: "sequential_9"

Layer (type)	Output Shape	Param #
conv2d_9 (Conv2D)	(None, 78, 78, 16)	160
max_pooling2d_8 (MaxPooling2D)	(None, 39, 39, 16)	0
dropout_9 (Dropout)	(None, 39, 39, 16)	0
flatten_9 (Flatten)	(None, 24336)	0
dense_19 (Dense)	(None, 64)	1,557,568
dense_20 (Dense)	(None, 64)	4,160
dense_21 (Dense)	(None, 5)	325

Total params: 1,562,213 (5.96 MB)


Trainable params: 1,562,213 (5.96 MB)

Non-trainable params: 0 (0.00 B)


Epoch 1/20
108/108  **16s** 132ms/step - accuracy: 0.2421 - f1_score: 0.1551 - loss: 1.6267 - val_accuracy: 0.2083 - val_f1_score: 0.0718 - val_loss: 1.6189
Epoch 2/20
108/108  **14s** 127ms/step - accuracy: 0.2324 - f1_score: 0.1425 - loss: 1.6046 - val_accuracy: 0.2454 - val_f1_score: 0.0967 - val_loss: 1.6080
Epoch 3/20
108/108  **21s** 127ms/step - accuracy: 0.2201 - f1_score: 0.1440 - loss: 1.6107 - val_accuracy: 0.2454 - val_f1_score: 0.0967 - val_loss: 1.6078
Epoch 4/20
108/108  **21s** 128ms/step - accuracy: 0.2356 - f1_score: 0.1390 - loss: 1.6060 - val_accuracy: 0.2454 - val_f1_score: 0.0967 - val_loss: 1.6063
Epoch 5/20
108/108  **20s** 124ms/step - accuracy: 0.2345 - f1_score: 0.1568 - loss: 1.6080 - val_accuracy: 0.2083 - val_f1_score: 0.0718 - val_loss: 1.6118
Epoch 6/20
108/108  **21s** 131ms/step - accuracy: 0.2374 - f1_score: 0.1460 - loss: 1.6078 - val_accuracy: 0.2454 - val_f1_score: 0.0967 - val_loss: 1.6055
Epoch 7/20
108/108  **13s** 125ms/step - accuracy: 0.2348 - f1_score: 0.1685 - loss: 1.6052 - val_accuracy: 0.1667 - val_f1_score: 0.0476 - val_loss: 1.6297
Epoch 8/20
108/108  **21s** 132ms/step - accuracy: 0.2182 - f1_score: 0.1514 - loss: 1.6066 - val_accuracy: 0.2454 - val_f1_score: 0.0967 - val_loss: 1.6125
Epoch 9/20
108/108  **15s** 135ms/step - accuracy: 0.2389 - f1_score: 0.1484 - loss: 1.6066 - val_accuracy: 0.2454 - val_f1_score: 0.0967 - val_loss: 1.6064
Epoch 10/20
108/108  **19s** 125ms/step - accuracy: 0.2362 - f1_score: 0.1282 - loss: 1.6092 - val_accuracy: 0.2454 - val_f1_score: 0.0967 - val_loss: 1.6150
Epoch 11/20
108/108  **14s** 131ms/step - accuracy: 0.2188 - f1_score: 0.1356 - loss: 1.6072 - val_accuracy: 0.2454 - val_f1_score: 0.0967 - val_loss: 1.6031
Epoch 12/20
108/108  **14s** 127ms/step - accuracy: 0.2338 - f1_score: 0.1268 - loss: 1.6042 - val_accuracy: 0.2083 - val_f1_score: 0.0718 - val_loss: 1.6097
Epoch 13/20
108/108  **21s** 134ms/step - accuracy: 0.2351 - f1_score: 0.1485 - loss: 1.6040 - val_accuracy: 0.2083 - val_f1_score: 0.0718 - val_loss: 1.6052
Epoch 14/20
108/108  **20s** 130ms/step - accuracy: 0.2425 - f1_score: 0.1369 - loss: 1.5992 - val_accuracy: 0.2083 - val_f1_score: 0.0718 - val_loss: 1.6153
Epoch 15/20
108/108  **21s** 131ms/step - accuracy: 0.2336 - f1_score: 0.1402 - loss: 1.6065 - val_accuracy: 0.2454 - val_f1_score: 0.0967 - val_loss: 1.6096
Epoch 16/20

108/108 ████████████████████ **20s** 126ms/step - accuracy: 0.2414 - f1_score: 0.1285 - loss: 1.6009 - val_accuracy: 0.2083 - val_f1_score: 0.0718 - val_loss: 1.6127
Epoch 17/20
108/108 ████████████████████ **20s** 125ms/step - accuracy: 0.2380 - f1_score: 0.1398 - loss: 1.6027 - val_accuracy: 0.2454 - val_f1_score: 0.0967 - val_loss: 1.6114
Epoch 18/20
108/108 ████████████████████ **13s** 124ms/step - accuracy: 0.2363 - f1_score: 0.1471 - loss: 1.6021 - val_accuracy: 0.2454 - val_f1_score: 0.0967 - val_loss: 1.6056
Epoch 19/20
108/108 ████████████████████ **21s** 130ms/step - accuracy: 0.2358 - f1_score: 0.1442 - loss: 1.6027 - val_accuracy: 0.2083 - val_f1_score: 0.0718 - val_loss: 1.6070
Epoch 20/20
108/108 ████████████████████ **14s** 126ms/step - accuracy: 0.2248 - f1_score: 0.1388 - loss: 1.6016 - val_accuracy: 0.2454 - val_f1_score: 0.0967 - val_loss: 1.6102
14/14 ████████████████████ **0s** 27ms/step - accuracy: 0.1992 - f1_score: 0.0693 - loss: 1.6174
Test Loss: 1.6030747890472412, Test Accuracy: 0.24074074625968933, Test F1 Score: 0.09342177957296371
Time required to train the model is 359.34225511550903 seconds


Filters: [16, 32, 64], Kernels: [(3, 3), (5, 5), (5, 5)], max pool, sigmoid activation function, No. of dense layers after flatten: 2



Filters: [16, 32, 64], Kernels: [(3, 3), (5, 5), (5, 5)], max pool, sigmoid activation function, No. of dense layers after flatten: 2



14/14 ————— 1s 35ms/step



Model: "sequential_10"


Layer (type)	Output Shape	Param #
conv2d_10 (Conv2D)	(None, 78, 78, 16)	160
max_pooling2d_9 (MaxPooling2D)	(None, 39, 39, 16)	0
dropout_10 (Dropout)	(None, 39, 39, 16)	0
flatten_10 (Flatten)	(None, 24336)	0
dense_22 (Dense)	(None, 64)	1,557,568
dense_23 (Dense)	(None, 64)	4,160
dense_24 (Dense)	(None, 5)	325

Total params: 1,562,213 (5.96 MB)


Trainable params: 1,562,213 (5.96 MB)

Non-trainable params: 0 (0.00 B)


Epoch 1/20

108/108  **15s** 126ms/step - accuracy: 0.2408 - f1_score: 0.2380 - loss: 189.2480 - val_accuracy: 0.3681 - val_f1_score: 0.3521 - val_loss: 2.0062


Epoch 2/20

108/108  **13s** 121ms/step - accuracy: 0.4189 - f1_score: 0.4185 - loss: 1.7014 - val_accuracy: 0.3218 - val_f1_score: 0.2767 - val_loss: 1.9406


Epoch 3/20

108/108  **13s** 121ms/step - accuracy: 0.6098 - f1_score: 0.6088 - loss: 0.9774 - val_accuracy: 0.3819 - val_f1_score: 0.3697 - val_loss: 1.7696


Epoch 4/20

108/108  **21s** 125ms/step - accuracy: 0.7290 - f1_score: 0.7298 - loss: 0.7256 - val_accuracy: 0.3657 - val_f1_score: 0.3478 - val_loss: 2.3024


Epoch 5/20

108/108  **21s** 134ms/step - accuracy: 0.8288 - f1_score: 0.8298 - loss: 0.5196 - val_accuracy: 0.4074 - val_f1_score: 0.3933 - val_loss: 2.1891


Epoch 6/20

108/108  **20s** 128ms/step - accuracy: 0.8740 - f1_score: 0.8749 - loss: 0.4043 - val_accuracy: 0.4028 - val_f1_score: 0.3738 - val_loss: 2.5812


Epoch 7/20

108/108  **20s** 128ms/step - accuracy: 0.9124 - f1_score: 0.9129 - loss: 0.3020 - val_accuracy: 0.3750 - val_f1_score: 0.3651 - val_loss: 2.6319


Epoch 8/20

108/108  **14s** 127ms/step - accuracy: 0.9459 - f1_score: 0.9461 - loss: 0.2195 - val_accuracy: 0.4028 - val_f1_score: 0.3863 - val_loss: 2.9725


Epoch 9/20

108/108  **23s** 147ms/step - accuracy: 0.9454 - f1_score: 0.9455 - loss: 0.2061 - val_accuracy: 0.3981 - val_f1_score: 0.3871 - val_loss: 3.1585


Epoch 10/20

108/108  **17s** 118ms/step - accuracy: 0.9573 - f1_score: 0.9574 - loss: 0.2070 - val_accuracy: 0.3866 - val_f1_score: 0.3782 - val_loss: 3.3044


Epoch 11/20

108/108  **21s** 120ms/step - accuracy: 0.9590 - f1_score: 0.9589 - loss: 0.1825 - val_accuracy: 0.3796 - val_f1_score: 0.3555 - val_loss: 4.0280

Epoch 12/20

108/108  **14s** 128ms/step - accuracy: 0.9544 - f1_score: 0.9544 - loss: 0.1841 - val_accuracy: 0.3519 - val_f1_score: 0.3349 - val_loss: 3.9221

Epoch 13/20


108/108  **21s** 131ms/step - accuracy: 0.9625 - f1_score: 0.9625 - loss: 0.1585 - val_accuracy: 0.3750 - val_f1_score: 0.3705 - val_loss: 3.9933

14/14  **0s** 27ms/step - accuracy: 0.3617 - f1_score: 0.3417 - loss: 1.9335


Test Loss: 1.8960705995559692, Test Accuracy: 0.37731480598449707, Test F1 Score: 0.36335206031799316

Time required to train the model is 239.42326736450195 seconds


Filters: [16, 32, 64], Kernels: [(3, 3), (5, 5), (5, 5)], max pool, relu activation function, No. of dense layers after flatten: 2



Filters: [16, 32, 64], Kernels: [(3, 3), (5, 5), (5, 5)], max pool, relu activation function, No. of dense layers after flatten: 2



14/14 ————— 1s 31ms/step



Model: "sequential_11"


Layer (type)	Output Shape	Param #
conv2d_11 (Conv2D)	(None, 78, 78, 16)	160
max_pooling2d_10 (MaxPooling2D)	(None, 39, 39, 16)	0
dropout_11 (Dropout)	(None, 39, 39, 16)	0
flatten_11 (Flatten)	(None, 24336)	0
dense_25 (Dense)	(None, 64)	1,557,568
dense_26 (Dense)	(None, 64)	4,160
dense_27 (Dense)	(None, 5)	325

Total params: 1,562,213 (5.96 MB)


Trainable params: 1,562,213 (5.96 MB)

Non-trainable params: 0 (0.00 B)


Epoch 1/20

108/108  **17s** 136ms/step - accuracy: 0.2487 - f1_score: 0.2430 - loss: 226.8415 - val_accuracy: 0.3056 - val_f1_score: 0.3064 - val_loss: 2.1821


Epoch 2/20

108/108  **14s** 130ms/step - accuracy: 0.4977 - f1_score: 0.4959 - loss: 1.4525 - val_accuracy: 0.3194 - val_f1_score: 0.3240 - val_loss: 1.9707


Epoch 3/20

108/108  **21s** 131ms/step - accuracy: 0.6741 - f1_score: 0.6716 - loss: 0.8715 - val_accuracy: 0.3519 - val_f1_score: 0.3492 - val_loss: 2.0662


Epoch 4/20

108/108  **19s** 119ms/step - accuracy: 0.7846 - f1_score: 0.7818 - loss: 0.6272 - val_accuracy: 0.3634 - val_f1_score: 0.3487 - val_loss: 2.2012


Epoch 5/20

108/108  **21s** 128ms/step - accuracy: 0.8648 - f1_score: 0.8637 - loss: 0.4250 - val_accuracy: 0.3889 - val_f1_score: 0.3855 - val_loss: 2.3269


Epoch 6/20

108/108  **14s** 130ms/step - accuracy: 0.9251 - f1_score: 0.9250 - loss: 0.2935 - val_accuracy: 0.3773 - val_f1_score: 0.3679 - val_loss: 2.6274


Epoch 7/20

108/108  **14s** 131ms/step - accuracy: 0.9488 - f1_score: 0.9487 - loss: 0.2368 - val_accuracy: 0.3796 - val_f1_score: 0.3678 - val_loss: 2.7289


Epoch 8/20

108/108  **14s** 130ms/step - accuracy: 0.9492 - f1_score: 0.9491 - loss: 0.2063 - val_accuracy: 0.3472 - val_f1_score: 0.3328 - val_loss: 3.1098


Epoch 9/20

108/108  **14s** 131ms/step - accuracy: 0.9533 - f1_score: 0.9532 - loss: 0.2129 - val_accuracy: 0.3796 - val_f1_score: 0.3799 - val_loss: 3.3902


Epoch 10/20

108/108  **20s** 131ms/step - accuracy: 0.9430 - f1_score: 0.9430 - loss: 0.2216 - val_accuracy: 0.3495 - val_f1_score: 0.3459 - val_loss: 3.9098

Epoch 11/20

108/108  **21s** 133ms/step - accuracy: 0.9228 - f1_score: 0.9228 - loss: 0.3893 - val_accuracy: 0.3681 - val_f1_score: 0.3544 - val_loss: 3.7751

Epoch 12/20


108/108  **15s** 135ms/step - accuracy: 0.9641 - f1_score: 0.9640 - loss: 0.1480 - val_accuracy: 0.3773 - val_f1_score: 0.3675 - val_loss: 3.8326

14/14  **0s** 34ms/step - accuracy: 0.2961 - f1_score: 0.3012 - loss: 2.0987


Test Loss: 2.026577949523926, Test Accuracy: 0.31481480598449707, Test F1 Score: 0.32019156217575073

Time required to train the model is 210.24508500099182 seconds


Filters: [16, 32, 64], Kernels: [(3, 3), (5, 5), (5, 5)], max pool, leaky_relu activation function, No. of dense layers after flatten: 2



Filters: [16, 32, 64], Kernels: [(3, 3), (5, 5), (5, 5)], max pool, leaky_relu activation function, No. of dense layers after flatten: 2



14/14 ————— 1s 57ms/step




In []: result_df_4

Out[]:	Conv Kernel Size	Conv Filter Size	Pooling Layers	Activation Function	No. of Dense Layers after Flatten	Dropout Rate	Test Loss	Test Accuracy	Test F1 Score	Training Time(in seconds)
0	[(3, 3), (5, 5), (5, 5)]	[16, 32, 64]	max	sigmoid	2	0.1	1.603075	0.240741	0.093422	359.342255
1	[(3, 3), (5, 5), (5, 5)]	[16, 32, 64]	max	relu	2	0.1	1.896071	0.377315	0.363352	239.423267
2	[(3, 3), (5, 5), (5, 5)]	[16, 32, 64]	max	leaky_relu	2	0.1	2.026578	0.314815	0.320192	210.245085

```
In [ ]: plt.figure(figsize=(20,5))
plt.bar(
    result_df_4['Activation Function'],
    result_df_4['Test Accuracy']
```



```
)
plt.ylabel('Test Accuracy')
plt.xlabel('Activation Function')
plt.title('Validation Accuracy for various activation function')
plt.show()
```



```
In [ ]: best_activation_function = result_df_4.sort_values(
        by=['Test Accuracy', 'Test F1 Score'],
        ascending=[False, False]
    )['Activation Function'].iloc[0]

best_activation_function
```

Out[]: 'relu'

```
In [ ]: # Subtask 5

result_df_5 = pd.DataFrame(
    columns=[
        'Conv Kernel Size',
        'Conv Filter Size',
        'Pooling Layers',
        'Activation Function',
        'No. of Dense Layers after Flatten',
        'Dropout Rate',
        'Test Loss',
        'Test Accuracy',
```

```

        'Test F1 Score',
        'Batch Normalization Presence',
        'Training Time(in seconds)'
    ]
)

filters = [16,32,64]
dropout_rates = [0.1,0.25,0,0.1]
batch_norm=[False,False,True,True]
epochs = 20

for dropout_rate,add_batch in zip(dropout_rates,batch_norm):
    test_loss,test_accuracy,test_f1,train_time,_ = train_model(
        kernels=best_kernel,
        filters=filters,
        activation_func=best_activation_function,
        pool=best_pool,
        dropout_rate=dropout_rate,
        num_dense_layers=best_num_dense,
        X_train=X_train,
        y_train=y_train,
        X_test=X_test,
        y_test=y_test,
        num_epochs=epochs,
        add_batch_normalization=add_batch
    )

result_df_5.loc[len(result_df_5.index)]= [
    best_kernel,
    filters,
    best_pool,
    best_activation_function,
    best_num_dense,
    dropout_rate,
    test_loss,
    test_accuracy,
    test_f1,
    add_batch,
    train_time
]

```













Model: "sequential_14"

Layer (type)	Output Shape	Param #
conv2d_14 (Conv2D)	(None, 78, 78, 16)	160
max_pooling2d_13 (MaxPooling2D)	(None, 39, 39, 16)	0
dropout_14 (Dropout)	(None, 39, 39, 16)	0
flatten_14 (Flatten)	(None, 24336)	0
dense_34 (Dense)	(None, 64)	1,557,568
dense_35 (Dense)	(None, 64)	4,160
dense_36 (Dense)	(None, 5)	325


Total params: 1,562,213 (5.96 MB)

Trainable params: 1,562,213 (5.96 MB)


Non-trainable params: 0 (0.00 B)

Epoch 1/20
108/108  **15s** 113ms/step - accuracy: 0.2102 - f1_score: 0.2049 - loss: 9.4014 - val_accuracy: 0.2801 - val_f1_score: 0.2679 - val_loss: 1.6251
Epoch 2/20
108/108  **25s** 155ms/step - accuracy: 0.5163 - f1_score: 0.5125 - loss: 1.3066 - val_accuracy: 0.3194 - val_f1_score: 0.3024 - val_loss: 1.7811
Epoch 3/20
108/108  **20s** 180ms/step - accuracy: 0.6823 - f1_score: 0.6827 - loss: 0.9143 - val_accuracy: 0.3310 - val_f1_score: 0.3181 - val_loss: 2.1691
Epoch 4/20
108/108  **18s** 161ms/step - accuracy: 0.7920 - f1_score: 0.7919 - loss: 0.6391 - val_accuracy: 0.2801 - val_f1_score: 0.2662 - val_loss: 2.7814
Epoch 5/20
108/108  **18s** 167ms/step - accuracy: 0.8402 - f1_score: 0.8400 - loss: 0.5474 - val_accuracy: 0.3148 - val_f1_score: 0.2975 - val_loss: 3.6367
Epoch 6/20
108/108  **18s** 148ms/step - accuracy: 0.8896 - f1_score: 0.8899 - loss: 0.4420 - val_accuracy: 0.2593 - val_f1_score: 0.2531 - val_loss: 3.8784
Epoch 7/20
108/108  **19s** 136ms/step - accuracy: 0.9285 - f1_score: 0.9286 - loss: 0.2619 - val_accuracy: 0.2731 - val_f1_score: 0.2729 - val_loss: 4.4871
Epoch 8/20
108/108  **23s** 160ms/step - accuracy: 0.9189 - f1_score: 0.9192 - loss: 0.3187 - val_accuracy: 0.3009 - val_f1_score: 0.2960 - val_loss: 4.5305
Epoch 9/20
108/108  **18s** 138ms/step - accuracy: 0.9606 - f1_score: 0.9605 - loss: 0.1577 - val_accuracy: 0.2986 - val_f1_score: 0.2889 - val_loss: 6.0384
Epoch 10/20
108/108  **21s** 147ms/step - accuracy: 0.9742 - f1_score: 0.9742 - loss: 0.1252 - val_accuracy: 0.2870 - val_f1_score: 0.2835 - val_loss: 5.4415
Epoch 11/20
108/108  **18s** 121ms/step - accuracy: 0.9593 - f1_score: 0.9594 - loss: 0.1675 - val_accuracy: 0.3032 - val_f1_score: 0.2980 - val_loss: 5.8022
14/14  **0s** 27ms/step - accuracy: 0.3144 - f1_score: 0.3004 - loss: 1.5944
Test Loss: 1.597295880317688, Test Accuracy: 0.3263888955116272, Test F1 Score: 0.3075091540813446
Time required to train the model is 221.19082760810852 seconds


Filters: [16, 32, 64], Kernels: [(3, 3), (5, 5), (5, 5)], max pool, relu activation function, No. of dense layers after flatten: 2



Filters: [16, 32, 64], Kernels: [(3, 3), (5, 5), (5, 5)], max pool, relu activation function, No. of dense layers after flatten: 2



14/14 ————— 1s 32ms/step



Model: "sequential_15"


Layer (type)	Output Shape	Param #
conv2d_15 (Conv2D)	(None, 78, 78, 16)	160
max_pooling2d_14 (MaxPooling2D)	(None, 39, 39, 16)	0
dropout_15 (Dropout)	(None, 39, 39, 16)	0
flatten_15 (Flatten)	(None, 24336)	0
dense_37 (Dense)	(None, 64)	1,557,568
dense_38 (Dense)	(None, 64)	4,160
dense_39 (Dense)	(None, 5)	325

Total params: 1,562,213 (5.96 MB)


Trainable params: 1,562,213 (5.96 MB)

Non-trainable params: 0 (0.00 B)


Epoch 1/20

108/108  **15s** 115ms/step - accuracy: 0.2192 - f1_score: 0.2157 - loss: 89.8713 - val_accuracy: 0.2523 - val_f1_score: 0.2369 - val_loss: 1.6610


Epoch 2/20

108/108  **13s** 115ms/step - accuracy: 0.3869 - f1_score: 0.3713 - loss: 1.4517 - val_accuracy: 0.3356 - val_f1_score: 0.3129 - val_loss: 1.6257


Epoch 3/20

108/108  **21s** 125ms/step - accuracy: 0.5381 - f1_score: 0.5275 - loss: 1.1971 - val_accuracy: 0.3426 - val_f1_score: 0.3265 - val_loss: 1.7339


Epoch 4/20

108/108  **13s** 124ms/step - accuracy: 0.6532 - f1_score: 0.6511 - loss: 0.9490 - val_accuracy: 0.3634 - val_f1_score: 0.3326 - val_loss: 1.8115


Epoch 5/20

108/108  **20s** 121ms/step - accuracy: 0.7619 - f1_score: 0.7607 - loss: 0.7182 - val_accuracy: 0.3704 - val_f1_score: 0.3486 - val_loss: 1.8873


Epoch 6/20

108/108  **20s** 119ms/step - accuracy: 0.7859 - f1_score: 0.7842 - loss: 0.6197 - val_accuracy: 0.3565 - val_f1_score: 0.3513 - val_loss: 2.3846


Epoch 7/20

108/108  **20s** 111ms/step - accuracy: 0.8557 - f1_score: 0.8556 - loss: 0.4660 - val_accuracy: 0.3542 - val_f1_score: 0.3542 - val_loss: 2.8076


Epoch 8/20

108/108  **22s** 125ms/step - accuracy: 0.9103 - f1_score: 0.9102 - loss: 0.3105 - val_accuracy: 0.3681 - val_f1_score: 0.3700 - val_loss: 3.0242


Epoch 9/20

108/108  **20s** 117ms/step - accuracy: 0.9184 - f1_score: 0.9183 - loss: 0.3076 - val_accuracy: 0.3542 - val_f1_score: 0.3582 - val_loss: 3.1051


Epoch 10/20

108/108  **13s** 123ms/step - accuracy: 0.9397 - f1_score: 0.9396 - loss: 0.2515 - val_accuracy: 0.3657 - val_f1_score: 0.3426 - val_loss: 3.3166

Epoch 11/20

108/108  **13s** 124ms/step - accuracy: 0.9341 - f1_score: 0.9340 - loss: 0.2415 - val_accuracy: 0.3472 - val_f1_score: 0.3500 - val_loss: 3.7050

Epoch 12/20


108/108  **20s** 119ms/step - accuracy: 0.9575 - f1_score: 0.9574 - loss: 0.1611 - val_accuracy: 0.3681 - val_f1_score: 0.3677 - val_loss: 3.4058

14/14  **0s** 28ms/step - accuracy: 0.3418 - f1_score: 0.3253 - loss: 1.7100


Test Loss: 1.6537659168243408, Test Accuracy: 0.35648149251937866, Test F1 Score: 0.3351629376411438

Time required to train the model is 210.9004213809967 seconds

Filters: [16, 32, 64], Kernels: [(3, 3), (5, 5), (5, 5)], max pool, relu activation function, No. of dense layers after flatten: 2




Filters: [16, 32, 64], Kernels: [(3, 3), (5, 5), (5, 5)], max pool, relu activation function, No. of dense layers after flatten: 2



14/14 1s 35ms/step

Confusion Matrix for the above model



Model: "sequential_16"


Layer (type)	Output Shape	Param #
conv2d_16 (Conv2D)	(None, 78, 78, 16)	160
max_pooling2d_15 (MaxPooling2D)	(None, 39, 39, 16)	0
batch_normalization (BatchNormalization)	(None, 39, 39, 16)	64
flatten_16 (Flatten)	(None, 24336)	0
dense_40 (Dense)	(None, 64)	1,557,568
dense_41 (Dense)	(None, 64)	4,160
dense_42 (Dense)	(None, 5)	325

Total params: 1,562,277 (5.96 MB)


Trainable params: 1,562,245 (5.96 MB)

Non-trainable params: 32 (128.00 B)


Epoch 1/20

108/108  **16s** 130ms/step - accuracy: 0.2926 - f1_score: 0.2885 - loss: 3.2829 - val_accuracy: 0.3241 - val_f1_score: 0.2922 - val_loss: 2.3232


Epoch 2/20

108/108  **13s** 124ms/step - accuracy: 0.5564 - f1_score: 0.5523 - loss: 1.1791 - val_accuracy: 0.4213 - val_f1_score: 0.4066 - val_loss: 1.6195


Epoch 3/20

108/108  **20s** 119ms/step - accuracy: 0.7810 - f1_score: 0.7808 - loss: 0.6359 - val_accuracy: 0.4028 - val_f1_score: 0.3872 - val_loss: 2.0123


Epoch 4/20

108/108  **21s** 122ms/step - accuracy: 0.8960 - f1_score: 0.8963 - loss: 0.3274 - val_accuracy: 0.3935 - val_f1_score: 0.3668 - val_loss: 2.7691


Epoch 5/20

108/108  **20s** 116ms/step - accuracy: 0.9543 - f1_score: 0.9543 - loss: 0.1645 - val_accuracy: 0.4236 - val_f1_score: 0.4156 - val_loss: 2.3500


Epoch 6/20

108/108  **13s** 123ms/step - accuracy: 0.9813 - f1_score: 0.9813 - loss: 0.0838 - val_accuracy: 0.3889 - val_f1_score: 0.3752 - val_loss: 2.9027


Epoch 7/20

108/108  **13s** 124ms/step - accuracy: 0.9915 - f1_score: 0.9915 - loss: 0.0462 - val_accuracy: 0.4213 - val_f1_score: 0.4096 - val_loss: 2.7708


Epoch 8/20

108/108  **14s** 126ms/step - accuracy: 0.9845 - f1_score: 0.9846 - loss: 0.0615 - val_accuracy: 0.4028 - val_f1_score: 0.3928 - val_loss: 3.0520


Epoch 9/20

108/108  **21s** 127ms/step - accuracy: 0.9923 - f1_score: 0.9923 - loss: 0.0489 - val_accuracy: 0.3935 - val_f1_score: 0.3761 - val_loss: 3.5689


Epoch 10/20

108/108  **14s** 130ms/step - accuracy: 0.9885 - f1_score: 0.9885 - loss: 0.0520 - val_accuracy: 0.3843 - val_f1_score: 0.3690 - val_loss: 3.6576

Epoch 11/20

108/108  **20s** 123ms/step - accuracy: 0.9756 - f1_score: 0.9756 - loss: 0.0825 - val_accuracy: 0.3611 - val_f1_score: 0.3410 - val_loss: 4.5590

Epoch 12/20


108/108  **21s** 129ms/step - accuracy: 0.9242 - f1_score: 0.9240 - loss: 0.2676 - val_accuracy: 0.4005 - val_f1_score: 0.3929 - val_loss: 3.7008

14/14  **0s** 32ms/step - accuracy: 0.4347 - f1_score: 0.4237 - loss: 1.6590


Test Loss: 1.6189850568771362, Test Accuracy: 0.44907405972480774, Test F1 Score: 0.43202194571495056

Time required to train the model is 212.70346307754517 seconds

Filters: [16, 32, 64], Kernels: [(3, 3), (5, 5), (5, 5)], max pool, relu activation function, No. of dense layers after flatten: 2




Filters: [16, 32, 64], Kernels: [(3, 3), (5, 5), (5, 5)], max pool, relu activation function, No. of dense layers after flatten: 2



14/14 ————— 1s 39ms/step

Confusion Matrix for the above model
















Model: "sequential_17"

Layer (type)	Output Shape	Param #
conv2d_17 (Conv2D)	(None, 78, 78, 16)	160
max_pooling2d_16 (MaxPooling2D)	(None, 39, 39, 16)	0
batch_normalization_1 (BatchNormalization)	(None, 39, 39, 16)	64
dropout_16 (Dropout)	(None, 39, 39, 16)	0
flatten_17 (Flatten)	(None, 24336)	0
dense_43 (Dense)	(None, 64)	1,557,568
dense_44 (Dense)	(None, 64)	4,160
dense_45 (Dense)	(None, 5)	325


Total params: 1,562,277 (5.96 MB)

Trainable params: 1,562,245 (5.96 MB)


Non-trainable params: 32 (128.00 B)


Epoch 1/20
108/108  **19s** 147ms/step - accuracy: 0.3085 - f1_score: 0.3049 - loss: 2.8237 - val_accuracy: 0.3079 - val_f1_score: 0.2847 - val_loss: 1.8885
Epoch 2/20
108/108  **15s** 136ms/step - accuracy: 0.4960 - f1_score: 0.4902 - loss: 1.2529 - val_accuracy: 0.3449 - val_f1_score: 0.3296 - val_loss: 1.7837
Epoch 3/20
108/108  **15s** 140ms/step - accuracy: 0.6838 - f1_score: 0.6812 - loss: 0.8260 - val_accuracy: 0.2894 - val_f1_score: 0.2437 - val_loss: 2.1348
Epoch 4/20
108/108  **15s** 138ms/step - accuracy: 0.8160 - f1_score: 0.8152 - loss: 0.5066 - val_accuracy: 0.3542 - val_f1_score: 0.3437 - val_loss: 2.4368
Epoch 5/20
108/108  **20s** 135ms/step - accuracy: 0.9196 - f1_score: 0.9197 - loss: 0.2631 - val_accuracy: 0.3819 - val_f1_score: 0.3694 - val_loss: 2.2777
Epoch 6/20
108/108  **21s** 140ms/step - accuracy: 0.9489 - f1_score: 0.9488 - loss: 0.1810 - val_accuracy: 0.3773 - val_f1_score: 0.3709 - val_loss: 2.8069
Epoch 7/20
108/108  **15s** 137ms/step - accuracy: 0.9582 - f1_score: 0.9581 - loss: 0.1559 - val_accuracy: 0.3356 - val_f1_score: 0.3361 - val_loss: 3.4593
Epoch 8/20
108/108  **14s** 133ms/step - accuracy: 0.9447 - f1_score: 0.9447 - loss: 0.2027 - val_accuracy: 0.3750 - val_f1_score: 0.3591 - val_loss: 3.5186
Epoch 9/20
108/108  **21s** 140ms/step - accuracy: 0.9506 - f1_score: 0.9506 - loss: 0.1694 - val_accuracy: 0.3588 - val_f1_score: 0.3484 - val_loss: 3.4710
Epoch 10/20
108/108  **20s** 135ms/step - accuracy: 0.9510 - f1_score: 0.9509 - loss: 0.1960 - val_accuracy: 0.3310 - val_f1_score: 0.3114 - val_loss: 5.2249
Epoch 11/20
108/108  **20s** 135ms/step - accuracy: 0.9637 - f1_score: 0.9637 - loss: 0.1282 - val_accuracy: 0.3356 - val_f1_score: 0.3180 - val_loss: 4.2496
Epoch 12/20
108/108  **20s** 134ms/step - accuracy: 0.9637 - f1_score: 0.9637 - loss: 0.1365 - val_accuracy: 0.3866 - val_f1_score: 0.3737 - val_loss: 3.8974
14/14  **0s** 30ms/step - accuracy: 0.3305 - f1_score: 0.3218 - loss: 1.8348
Test Loss: 1.7746429443359375, Test Accuracy: 0.35185185074806213, Test F1 Score: 0.33424222469329834
Time required to train the model is 222.4370560646057 seconds

Filters: [16, 32, 64], Kernels: [(3, 3), (5, 5), (5, 5)], max pool, relu activation function, No. of dense layers after flatten: 2



Filters: [16, 32, 64], Kernels: [(3, 3), (5, 5), (5, 5)], max pool, relu activation function, No. of dense layers after flatten: 2





In []: result_df_5

Out[]:


	Conv Kernel Size	Conv Filter Size	Pooling Layers	Activation Function	No. of Dense Layers after Flatten	Dropout Rate	Test Loss	Test Accuracy	Test F1 Score	Batch Normalization Presence	Training Time(in seconds)
0	[(3, 3), (5, 5), (5, 5)]	[16, 32, 64]	max	relu	2	0.10	1.597296	0.326389	0.307509	False	221.190828
1	[(3, 3), (5, 5), (5, 5)]	[16, 32, 64]	max	relu	2	0.25	1.653766	0.356481	0.335163	False	210.900421
2	[(3, 3), (5, 5), (5, 5)]	[16, 32, 64]	max	relu	2	0.00	1.618985	0.449074	0.432022	True	212.703463
3	[(3, 3), (5, 5), (5, 5)]	[16, 32, 64]	max	relu	2	0.10	1.774643	0.351852	0.334242	True	222.437056

```

In [ ]: plt.figure(figsize=(20,5))
plt.bar(
    [f'Dropout rate: {drop} batch: {batch}' for drop,batch in zip(result_df_5['Dropout Rate'],result_df_5['Batch Normalization
    result_df_5['Test Accuracy']
    )

plt.ylabel('Test Accuracy')
plt.xlabel('Regularization')
plt.title('Validation Accuracy for various regularization')
plt.show()

```



```

In [ ]: best_dropout = result_df_5.sort_values(
    by=['Test Accuracy','Test F1 Score'],
    ascending=[False,False]
)['Dropout Rate'].iloc[0]

best_dropout

```

Out[]: 0.0

```

In [ ]: best_do_batch = result_df_5.sort_values(
    by=['Test Accuracy','Test F1 Score'],
    ascending=[False,False]
)['Batch Normalization Presence'].iloc[0]

best_do_batch

```

Out[]: True

- f. For the best set of parameters from the above runs, add [1,2,3] more convolution layers, and compare the size of trainable parameters and compare the time required to train each model for 10 epochs.
- g. For the best set of parameters obtained here repeat the experimentation for colour images and visualize the test result.

In []: # Subtask 6

```
result_df_6 = pd.DataFrame(
    columns=[
        'Conv Kernel Size',
        'Conv Filter Size',
        'Pooling Layers',
        'Activation Function',
        'No. of Dense Layers after Flatten',
        'Dropout Rate',
        'Test Loss',
        'Test Accuracy',
        'Test F1 Score',
        'Batch Normalization Presence',
        'No. of Extra Conv Layers',
        'Training Time(in seconds)'
    ]
)

filters = [16,32,64]
epochs = 10
extra_conv_layers = [0,1,2,3]

for c in extra_conv_layers:
    filter_copy=[]
    kernel_copy=[]

    for filt,ker in zip(filters,best_kernel):
        filter_copy.append(filt)
```

```

        kernel_copy.append(ker)

test_loss,test_accuracy,test_f1,train_time,_ = train_model(
    kernels=kernel_copy,
    filters=filter_copy,
    activation_func=best_activation_function,
    pool=best_pool,
    dropout_rate=best_dropout,
    num_dense_layers=best_num_dense,
    X_train=X_train,
    y_train=y_train,
    X_test=X_test,
    y_test=y_test,
    num_epochs=epochs,
    add_batch_normalization=best_do_batch,
    extra_conv_layers=c
)

result_df_6.loc[len(result_df_6.index)]=[
    kernel_copy,
    filter_copy,
    best_pool,
    best_activation_function,
    best_num_dense,
    best_dropout,
    test_loss,
    test_accuracy,
    test_f1,
    best_do_batch,
    c,
    train_time
]

```












Model: "sequential_6"

Layer (type)	Output Shape	Param #
conv2d_9 (Conv2D)	?	0 (unbuilt)
max_pooling2d_6 (MaxPooling2D)	?	0 (unbuilt)
batch_normalization_6 (BatchNormalization)	?	0 (unbuilt)
flatten_6 (Flatten)	?	0 (unbuilt)
dense_18 (Dense)	?	0 (unbuilt)
dense_19 (Dense)	?	0 (unbuilt)
dense_20 (Dense)	?	0 (unbuilt)


Total params: 0 (0.00 B)

Trainable params: 0 (0.00 B)


Non-trainable params: 0 (0.00 B)

Epoch 1/10
108/108  6s 28ms/step - accuracy: 0.2902 - f1_score: 0.2789 - loss: 2.7041 - val_accuracy: 0.3449 - val_f1_score: 0.2868 - val_loss: 1.6243
Epoch 2/10
108/108  2s 6ms/step - accuracy: 0.4680 - f1_score: 0.4586 - loss: 1.3390 - val_accuracy: 0.3403 - val_f1_score: 0.3248 - val_loss: 1.5946
Epoch 3/10
108/108  1s 4ms/step - accuracy: 0.6392 - f1_score: 0.6378 - loss: 0.9517 - val_accuracy: 0.2986 - val_f1_score: 0.2681 - val_loss: 2.0661
Epoch 4/10
108/108  1s 4ms/step - accuracy: 0.7875 - f1_score: 0.7874 - loss: 0.5885 - val_accuracy: 0.3148 - val_f1_score: 0.3014 - val_loss: 2.1523
Epoch 5/10
108/108  1s 4ms/step - accuracy: 0.8937 - f1_score: 0.8939 - loss: 0.3075 - val_accuracy: 0.3611 - val_f1_score: 0.3543 - val_loss: 2.3535
Epoch 6/10
108/108  1s 4ms/step - accuracy: 0.9372 - f1_score: 0.9372 - loss: 0.1979 - val_accuracy: 0.3565 - val_f1_score: 0.3462 - val_loss: 2.7796
Epoch 7/10
108/108  0s 3ms/step - accuracy: 0.9613 - f1_score: 0.9613 - loss: 0.1452 - val_accuracy: 0.3588 - val_f1_score: 0.3531 - val_loss: 2.9799
Epoch 8/10
108/108  0s 4ms/step - accuracy: 0.9716 - f1_score: 0.9716 - loss: 0.1070 - val_accuracy: 0.3912 - val_f1_score: 0.3877 - val_loss: 3.7011
Epoch 9/10
108/108  0s 3ms/step - accuracy: 0.9527 - f1_score: 0.9527 - loss: 0.1891 - val_accuracy: 0.3588 - val_f1_score: 0.3555 - val_loss: 3.4572
Epoch 10/10
108/108  0s 4ms/step - accuracy: 0.9521 - f1_score: 0.9521 - loss: 0.1813 - val_accuracy: 0.3634 - val_f1_score: 0.3502 - val_loss: 4.5450
14/14  0s 2ms/step - accuracy: 0.3482 - f1_score: 0.3309 - loss: 1.6105
Test Loss: 1.611054539680481, Test Accuracy: 0.33796295523643494, Test F1 Score: 0.3220447897911072
Time required to train the model is 12.167885541915894 seconds

Filters: [16, 32, 64], Kernels: [(3, 3), (5, 5), (5, 5)], max pool, relu activation function, No. of dense layers after flatten: 2




Filters: [16, 32, 64], Kernels: [(3, 3), (5, 5), (5, 5)], max pool, relu activation function, No. of dense layers after flatten: 2



14/14 0s 15ms/step

Confusion Matrix for the above model














Model: "sequential_7"

Layer (type)	Output Shape	Param #
conv2d_10 (Conv2D)	?	0 (unbuilt)
max_pooling2d_7 (MaxPooling2D)	?	0 (unbuilt)
batch_normalization_7 (BatchNormalization)	?	0 (unbuilt)
conv2d_11 (Conv2D)	?	0 (unbuilt)
flatten_7 (Flatten)	?	0 (unbuilt)
dense_21 (Dense)	?	0 (unbuilt)
dense_22 (Dense)	?	0 (unbuilt)
dense_23 (Dense)	?	0 (unbuilt)


Total params: 0 (0.00 B)

Trainable params: 0 (0.00 B)


Non-trainable params: 0 (0.00 B)

Epoch 1/10
108/108  **9s** 53ms/step - accuracy: 0.2472 - f1_score: 0.2428 - loss: 3.8436 - val_accuracy: 0.3264 - val_f1_score: 0.3001 - val_loss: 1.5401
Epoch 2/10
108/108  **4s** 10ms/step - accuracy: 0.4568 - f1_score: 0.4535 - loss: 1.3349 - val_accuracy: 0.3218 - val_f1_score: 0.2507 - val_loss: 2.2705
Epoch 3/10
108/108  **1s** 10ms/step - accuracy: 0.6738 - f1_score: 0.6730 - loss: 0.8941 - val_accuracy: 0.3171 - val_f1_score: 0.2621 - val_loss: 2.0527
Epoch 4/10
108/108  **1s** 9ms/step - accuracy: 0.8262 - f1_score: 0.8267 - loss: 0.5211 - val_accuracy: 0.3981 - val_f1_score: 0.3835 - val_loss: 2.3187
Epoch 5/10
108/108  **1s** 9ms/step - accuracy: 0.9072 - f1_score: 0.9072 - loss: 0.3232 - val_accuracy: 0.3542 - val_f1_score: 0.3458 - val_loss: 2.8635
Epoch 6/10
108/108  **1s** 10ms/step - accuracy: 0.9336 - f1_score: 0.9336 - loss: 0.2581 - val_accuracy: 0.3819 - val_f1_score: 0.3539 - val_loss: 3.9928
Epoch 7/10
108/108  **1s** 11ms/step - accuracy: 0.9620 - f1_score: 0.9620 - loss: 0.1399 - val_accuracy: 0.3634 - val_f1_score: 0.3541 - val_loss: 3.5446
Epoch 8/10
108/108  **1s** 11ms/step - accuracy: 0.9722 - f1_score: 0.9722 - loss: 0.1136 - val_accuracy: 0.3565 - val_f1_score: 0.3536 - val_loss: 4.2221
Epoch 9/10
108/108  **1s** 9ms/step - accuracy: 0.9624 - f1_score: 0.9624 - loss: 0.1448 - val_accuracy: 0.3773 - val_f1_score: 0.3833 - val_loss: 4.1511
Epoch 10/10
108/108  **1s** 9ms/step - accuracy: 0.9690 - f1_score: 0.9690 - loss: 0.1124 - val_accuracy: 0.3218 - val_f1_score: 0.2925 - val_loss: 6.7719
14/14  **0s** 3ms/step - accuracy: 0.2958 - f1_score: 0.2447 - loss: 1.5640
Test Loss: 1.5418537855148315, Test Accuracy: 0.3194444477558136, Test F1 Score: 0.2875923216342926
Time required to train the model is 22.891176462173462 seconds


Filters: [16, 32, 64, 128], Kernels: [(3, 3), (5, 5), (5, 5), (5, 5)], max pool, relu activation function, No. of dense layers after flatten: 2



Filters: [16, 32, 64, 128], Kernels: [(3, 3), (5, 5), (5, 5), (5, 5)], max pool, relu activation function, No. of dense layers after flatten: 2



Confusion Matrix for the above model












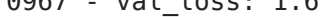

Model: "sequential_8"

Layer (type)	Output Shape	Param #
conv2d_12 (Conv2D)	?	0 (unbuilt)
max_pooling2d_8 (MaxPooling2D)	?	0 (unbuilt)
batch_normalization_8 (BatchNormalization)	?	0 (unbuilt)
conv2d_13 (Conv2D)	?	0 (unbuilt)
conv2d_14 (Conv2D)	?	0 (unbuilt)
flatten_8 (Flatten)	?	0 (unbuilt)
dense_24 (Dense)	?	0 (unbuilt)
dense_25 (Dense)	?	0 (unbuilt)
dense_26 (Dense)	?	0 (unbuilt)


Total params: 0 (0.00 B)

Trainable params: 0 (0.00 B)


Non-trainable params: 0 (0.00 B)

Epoch 1/10
108/108  **13s** 80ms/step - accuracy: 0.2309 - f1_score: 0.1277 - loss: 4.4110 - val_accuracy: 0.2454 - val_f1_score: 0.0967 - val_loss: 1.6031
Epoch 2/10
108/108  **3s** 32ms/step - accuracy: 0.2505 - f1_score: 0.1004 - loss: 1.5976 - val_accuracy: 0.2454 - val_f1_score: 0.0967 - val_loss: 1.6039
Epoch 3/10
108/108  **5s** 32ms/step - accuracy: 0.2345 - f1_score: 0.0892 - loss: 1.6025 - val_accuracy: 0.2454 - val_f1_score: 0.0967 - val_loss: 1.6050
Epoch 4/10
108/108  **3s** 31ms/step - accuracy: 0.2368 - f1_score: 0.0908 - loss: 1.6019 - val_accuracy: 0.2454 - val_f1_score: 0.0967 - val_loss: 1.6045
Epoch 5/10
108/108  **3s** 31ms/step - accuracy: 0.2544 - f1_score: 0.1033 - loss: 1.5904 - val_accuracy: 0.2454 - val_f1_score: 0.0967 - val_loss: 1.6043
Epoch 6/10
108/108  **5s** 33ms/step - accuracy: 0.2486 - f1_score: 0.0993 - loss: 1.5965 - val_accuracy: 0.2454 - val_f1_score: 0.0967 - val_loss: 1.6044
Epoch 7/10
108/108  **3s** 32ms/step - accuracy: 0.2507 - f1_score: 0.1006 - loss: 1.5965 - val_accuracy: 0.2454 - val_f1_score: 0.0967 - val_loss: 1.6042
Epoch 8/10
108/108  **5s** 32ms/step - accuracy: 0.2524 - f1_score: 0.1019 - loss: 1.5955 - val_accuracy: 0.2454 - val_f1_score: 0.0967 - val_loss: 1.6049
Epoch 9/10
108/108  **5s** 32ms/step - accuracy: 0.2477 - f1_score: 0.0985 - loss: 1.5987 - val_accuracy: 0.2454 - val_f1_score: 0.0967 - val_loss: 1.6052
Epoch 10/10
108/108  **5s** 31ms/step - accuracy: 0.2369 - f1_score: 0.0910 - loss: 1.6006 - val_accuracy: 0.2454 - val_f1_score: 0.0967 - val_loss: 1.6044
14/14  **0s** 10ms/step - accuracy: 0.1992 - f1_score: 0.0693 - loss: 1.6065
Test Loss: 1.598532795906067, Test Accuracy: 0.24074074625968933, Test F1 Score: 0.09342177957296371
Time required to train the model is 53.9805862903595 seconds


Filters: [16, 32, 64, 128, 256], Kernels: [(3, 3), (5, 5), (5, 5), (5, 5), (5, 5)], max pool, relu activation function, No. of dense layers after flatten: 2



Filters: [16, 32, 64, 128, 256], Kernels: [(3, 3), (5, 5), (5, 5), (5, 5), (5, 5)], max pool, relu activation function, No. of dense layers after flatten: 2



Confusion Matrix for the above model














Model: "sequential_9"

Layer (type)	Output Shape	Param #
conv2d_15 (Conv2D)	?	0 (unbuilt)
max_pooling2d_9 (MaxPooling2D)	?	0 (unbuilt)
batch_normalization_9 (BatchNormalization)	?	0 (unbuilt)
conv2d_16 (Conv2D)	?	0 (unbuilt)
conv2d_17 (Conv2D)	?	0 (unbuilt)
conv2d_18 (Conv2D)	?	0 (unbuilt)
flatten_9 (Flatten)	?	0 (unbuilt)
dense_27 (Dense)	?	0 (unbuilt)
dense_28 (Dense)	?	0 (unbuilt)
dense_29 (Dense)	?	0 (unbuilt)


Total params: 0 (0.00 B)

Trainable params: 0 (0.00 B)


Non-trainable params: 0 (0.00 B)


Epoch 1/10
108/108  **39s** 233ms/step - accuracy: 0.2140 - f1_score: 0.1530 - loss: 18.3655 - val_accuracy: 0.2083 - val_f1_score: 0.0718 - val_loss: 1.6065
Epoch 2/10
108/108  **16s** 82ms/step - accuracy: 0.2346 - f1_score: 0.1042 - loss: 1.6045 - val_accuracy: 0.2083 - val_f1_score: 0.0718 - val_loss: 1.6060
Epoch 3/10
108/108  **9s** 80ms/step - accuracy: 0.2551 - f1_score: 0.1448 - loss: 1.5932 - val_accuracy: 0.2454 - val_f1_score: 0.0967 - val_loss: 1.6032
Epoch 4/10
108/108  **9s** 79ms/step - accuracy: 0.2338 - f1_score: 0.0887 - loss: 1.6002 - val_accuracy: 0.2454 - val_f1_score: 0.0967 - val_loss: 1.6048
Epoch 5/10
108/108  **9s** 80ms/step - accuracy: 0.2420 - f1_score: 0.0944 - loss: 1.5970 - val_accuracy: 0.2454 - val_f1_score: 0.0967 - val_loss: 1.6047
Epoch 6/10
108/108  **10s** 78ms/step - accuracy: 0.2407 - f1_score: 0.0934 - loss: 1.6009 - val_accuracy: 0.2454 - val_f1_score: 0.0967 - val_loss: 1.6052
Epoch 7/10
108/108  **8s** 78ms/step - accuracy: 0.2403 - f1_score: 0.0932 - loss: 1.6002 - val_accuracy: 0.2454 - val_f1_score: 0.0967 - val_loss: 1.6053
Epoch 8/10
108/108  **9s** 79ms/step - accuracy: 0.2419 - f1_score: 0.0944 - loss: 1.5988 - val_accuracy: 0.2454 - val_f1_score: 0.0967 - val_loss: 1.6049
Epoch 9/10
108/108  **10s** 79ms/step - accuracy: 0.2359 - f1_score: 0.0901 - loss: 1.6003 - val_accuracy: 0.2454 - val_f1_score: 0.0967 - val_loss: 1.6052
Epoch 10/10
108/108  **10s** 77ms/step - accuracy: 0.2481 - f1_score: 0.0987 - loss: 1.5984 - val_accuracy: 0.2454 - val_f1_score: 0.0967 - val_loss: 1.6042
14/14  **0s** 12ms/step - accuracy: 0.1992 - f1_score: 0.0693 - loss: 1.6048
Test Loss: 1.5982946157455444, Test Accuracy: 0.24074074625968933, Test F1 Score: 0.09342177957296371
Time required to train the model is 128.19466972351074 seconds

Filters: [16, 32, 64, 128, 256, 512], Kernels: [(3, 3), (5, 5), (5, 5), (5, 5), (5, 5), (5, 5)], max pool, relu activation function, No. of dense layers after flatten: 2



Filters: [16, 32, 64, 128, 256, 512], Kernels: [(3, 3), (5, 5), (5, 5), (5, 5), (5, 5), (5, 5)], max pool, relu activation function, No. of dense layers after flatten: 2





In []: result_df_6


Out[]:

	Conv Kernel Size	Conv Filter Size	Pooling Layers	Activation Function	No. of Dense Layers after Flatten	Dropout Rate	Test Loss	Test Accuracy	Test F1 Score	Batch Normalization Presence	No. of Extra Conv Layers	Training Time(in seconds)
0	[(3, 3), (5, 5), (5, 5)]	[16, 32, 64]	max	relu	2	0.0	1.611055	0.337963	0.322045	True	0	12.167886
1	[(3, 3), (5, 5), (5, 5), (5, 5)]	[16, 32, 64, 128]	max	relu	2	0.0	1.541854	0.319444	0.287592	True	1	22.891176
2	[(3, 3), (5, 5), (5, 5), (5, 5), (5, 5)]	[16, 32, 64, 128, 256]	max	relu	2	0.0	1.598533	0.240741	0.093422	True	2	53.980586
3	[(3, 3), (5, 5), (5, 5), (5, 5), (5, 5), (5, 5)]	[16, 32, 64, 128, 256, 512]	max	relu	2	0.0	1.598295	0.240741	0.093422	True	3	128.194670

In []:

```
plt.figure(figsize=(20,5))
plt.bar(
    result_df_6['No. of Extra Conv Layers'],
    result_df_6['Test Accuracy']
)

plt.ylabel('Test Accuracy')
plt.xlabel('No. of Extra Conv Layers')
plt.title('Validation Accuracy for various extra conv layers')
plt.show()
```



```
In [ ]: best_conv = result_df_6.sort_values(  
        by=['Test Accuracy', 'Test F1 Score'],  
        ascending=[False, False]  
    )['No. of Extra Conv Layers'].iloc[0]  
  
best_conv
```

Out[]: 0

```
In [ ]: # Subtask 7  
  
# Prepare the RGB dataset  
  
X_rgb = [] # Contains the images  
Y_rgb = [] # Contains the labels
```

```
In [ ]: def train_data_rgb(flower_type, path_dir):  
    for img in tqdm(os.listdir(path_dir)):  
        label = flower_type  
        path = os.path.join(path_dir, img)  
        img_array = Image.open(path)  
        img_array = img_array.resize((IMG_SIZE, IMG_SIZE))  
        img_array = np.array(img_array)  
        X_rgb.append(np.array(img_array))  
        Y_rgb.append(str(label))
```

```
In [ ]: train_data_rgb('Daisy', FLOWER_DAISY_DIR)
```


```
train_data_rgb('Sunflower',FLOWER_SUNFLOWER_DIR)
train_data_rgb('Tulip',FLOWER_TULIP_DIR)
train_data_rgb('Dandelion',FLOWER_DANDI_DIR)
train_data_rgb('Rose',FLOWER_ROSE_DIR)
```

```
100%|██████████| 764/764 [00:03<00:00, 240.91it/s]
100%|██████████| 733/733 [00:03<00:00, 188.78it/s]
100%|██████████| 984/984 [00:04<00:00, 212.28it/s]
100%|██████████| 1052/1052 [00:04<00:00, 237.51it/s]
100%|██████████| 784/784 [00:03<00:00, 220.30it/s]
```

```
In [ ]: import matplotlib.image as mpimg
plt.imshow(X_rgb[223])

print(Y_rgb[223])
```

Daisy



```
In [ ]: from sklearn.model_selection import train_test_split

X_train_rgb,X_rem_rgb,y_train_rgb,y_rem_rgb = train_test_split(X_rgb,Y_rgb,test_size=0.2,random_state=5)
X_test_rgb,X_val_rgb,y_test_rgb,y_val_rgb = train_test_split(X_rem_rgb,y_rem_rgb,test_size=0.5,random_state=5)
```

```
print(len(X_train_rgb))
print(len(y_train_rgb))
print(len(X_val_rgb))
print(len(y_val_rgb))
print(len(X_test_rgb))
print(len(y_test_rgb))
```

3453

3453

432

432

432

432

In []: *# Convert to np array for tf processing*

```
X_train_rgb=np.array(X_train_rgb)
X_val_rgb=np.array(X_val_rgb)
X_test_rgb=np.array(X_test_rgb)
y_train_rgb=np.array(y_train_rgb)
y_val_rgb=np.array(y_val_rgb)
y_test_rgb=np.array(y_test_rgb)
```

Reshape

```
X_train_rgb = X_train_rgb.reshape(-1,IMG_SIZE,IMG_SIZE,3)
X_val_rgb = X_val_rgb.reshape(-1,IMG_SIZE,IMG_SIZE,3)
X_test_rgb = X_test_rgb.reshape(-1,IMG_SIZE,IMG_SIZE,3)
```

In []: *# One hot encode the labels*

```
label_encoder = LabelEncoder()
```


```
label_encoder.fit(['Daisy','Sunflower','Tulip','Dandelion','Rose'])
```

```
y_train_rgb = label_encoder.transform(y_train_rgb)
y_val_rgb = label_encoder.transform(y_val_rgb)
y_test_rgb = label_encoder.transform(y_test_rgb)
```

```
y_train_rgb = pd.get_dummies(y_train_rgb,dtype='int').to_numpy()
y_val_rgb = pd.get_dummies(y_val_rgb,dtype='int').to_numpy()
y_test_rgb = pd.get_dummies(y_test_rgb,dtype='int').to_numpy()
```

In []:

```
plt.imshow(X_train_rgb[24])
plt.show()
```



```
In [ ]: y_train_rgb[24]
```

```
Out[ ]: array([0, 0, 1, 0, 0])
```

```
In [ ]: result_df_7 = pd.DataFrame(  
    columns=[  
        'Conv Kernel Size',  
        'Conv Filter Size',  
        'Pooling Layers',  
        'Activation Function',  
        'No. of Dense Layers after Flatten',  
        'Dropout Rate',  
        'Test Loss',  
        'Test Accuracy',  
        'Test F1 Score',  
        'Batch Normalization Presence',  
        'No. of Extra Conv Layers',  
        'Color Mode',  
        'Training Time(in seconds)'  
    ]
```

```

)

filters = [16,32,64]
epochs = 20

test_loss,test_accuracy,test_f1,train_time,_ = train_model(
    kernels=best_kernel,
    filters=filters,
    activation_func=best_activation_function,
    pool=best_pool,
    dropout_rate=best_dropout,
    num_dense_layers=best_num_dense,
    X_train=X_train,
    y_train=y_train,
    X_test=X_test,
    y_test=y_test,
    num_epochs=epochs,
    add_batch_normalization=best_do_batch,
    extra_conv_layers=best_conv
)

result_df_7.loc[len(result_df_7.index)]=[
    best_kernel,
    filters,
    best_pool,
    best_activation_function,
    best_num_dense,
    best_dropout,
    test_loss,
    test_accuracy,
    test_f1,
    best_do_batch,
    best_conv,
    'grayscale',
    train_time
]

test_loss,test_accuracy,test_f1,train_time,_ = train_model(
    kernels=best_kernel,
    filters=filters,
    activation_func=best_activation_function,
    pool=best_pool,
    dropout_rate=best_dropout,
    num_dense_layers=best_num_dense,
    X_train=X_train_rgb,
    y_train=y_train_rgb,

```

```

X_test=X_test_rgb,
y_test=y_test_rgb,
num_epochs=epochs,
add_batch_normalization=best_do_batch,
extra_conv_layers=best_conv,
is_rgb=True
)

result_df_7.loc[len(result_df_7.index)]=[
    best_kernel,
    filters,
    best_pool,
    best_activation_function,
    best_num_dense,
    best_dropout,
    test_loss,
    test_accuracy,
    test_f1,
    best_do_batch,
    best_conv,
    'rgb',
    train_time
]

```

Model: "sequential_39"


Layer (type)	Output Shape	Param #
conv2d_58 (Conv2D)	?	0 (unbuilt)
max_pooling2d_38 (MaxPooling2D)	?	0 (unbuilt)
batch_normalization_23 (BatchNormalization)	?	0 (unbuilt)
flatten_39 (Flatten)	?	0 (unbuilt)
dense_109 (Dense)	?	0 (unbuilt)
dense_110 (Dense)	?	0 (unbuilt)
dense_111 (Dense)	?	0 (unbuilt)

Total params: 0 (0.00 B)


Trainable params: 0 (0.00 B)

Non-trainable params: 0 (0.00 B)


Epoch 1/20

108/108  **16s** 124ms/step - accuracy: 0.2898 - f1_score: 0.2828 - loss: 2.6201 - val_accuracy: 0.3750 - val_f1_score: 0.3660 - val_loss: 1.4675


Epoch 2/20

108/108  **20s** 121ms/step - accuracy: 0.5534 - f1_score: 0.5431 - loss: 1.2021 - val_accuracy: 0.3796 - val_f1_score: 0.3757 - val_loss: 1.5889


Epoch 3/20

108/108  **20s** 119ms/step - accuracy: 0.7582 - f1_score: 0.7554 - loss: 0.6858 - val_accuracy: 0.3264 - val_f1_score: 0.3189 - val_loss: 2.1540


Epoch 4/20

108/108  **21s** 125ms/step - accuracy: 0.8745 - f1_score: 0.8745 - loss: 0.4143 - val_accuracy: 0.3727 - val_f1_score: 0.3744 - val_loss: 2.4959


Epoch 5/20

108/108  **20s** 124ms/step - accuracy: 0.9413 - f1_score: 0.9412 - loss: 0.2176 - val_accuracy: 0.4259 - val_f1_score: 0.4236 - val_loss: 2.5334


Epoch 6/20

108/108  **20s** 115ms/step - accuracy: 0.9698 - f1_score: 0.9698 - loss: 0.1033 - val_accuracy: 0.3889 - val_f1_score: 0.3845 - val_loss: 2.9056


Epoch 7/20

108/108  **21s** 123ms/step - accuracy: 0.9838 - f1_score: 0.9838 - loss: 0.0656 - val_accuracy: 0.4213 - val_f1_score: 0.4097 - val_loss: 3.0003


Epoch 8/20

108/108  **13s** 119ms/step - accuracy: 0.9862 - f1_score: 0.9862 - loss: 0.0642 - val_accuracy: 0.4329 - val_f1_score: 0.4233 - val_loss: 3.0358


Epoch 9/20

108/108  **21s** 120ms/step - accuracy: 0.9961 - f1_score: 0.9961 - loss: 0.0331 - val_accuracy: 0.4005 - val_f1_score: 0.4002 - val_loss: 3.0924

Epoch 10/20

108/108  **14s** 125ms/step - accuracy: 0.9950 - f1_score: 0.9950 - loss: 0.0189 - val_accuracy: 0.3912 - val_f1_score: 0.3864 - val_loss: 3.0958

Epoch 11/20


108/108  **20s** 124ms/step - accuracy: 0.9962 - f1_score: 0.9962 - loss: 0.0175 - val_accuracy: 0.4005 - val_f1_score: 0.3901 - val_loss: 3.6262

14/14  **0s** 29ms/step - accuracy: 0.3190 - f1_score: 0.3046 - loss: 1.5945


Test Loss: 1.5915693044662476, Test Accuracy: 0.32870370149612427, Test F1 Score: 0.30711865425109863

Time required to train the model is 206.1395878791809 seconds


Filters: [16, 32, 64], Kernels: [(3, 3), (5, 5), (5, 5)], max pool, relu activation function, No. of dense layers after flatten: 2



Filters: [16, 32, 64], Kernels: [(3, 3), (5, 5), (5, 5)], max pool, relu activation function, No. of dense layers after flatten: 2



14/14 ————— 1s 36ms/step



Model: "sequential_40"


Layer (type)	Output Shape	Param #
conv2d_59 (Conv2D)	?	0 (unbuilt)
max_pooling2d_39 (MaxPooling2D)	?	0 (unbuilt)
batch_normalization_24 (BatchNormalization)	?	0 (unbuilt)
flatten_40 (Flatten)	?	0 (unbuilt)
dense_112 (Dense)	?	0 (unbuilt)
dense_113 (Dense)	?	0 (unbuilt)
dense_114 (Dense)	?	0 (unbuilt)

Total params: 0 (0.00 B)


Trainable params: 0 (0.00 B)

Non-trainable params: 0 (0.00 B)


Epoch 1/20

108/108  18s 143ms/step - accuracy: 0.3686 - f1_score: 0.3658 - loss: 3.1550 - val_accuracy: 0.4421 - val_f1_score: 0.3940 - val_loss: 2.2427


Epoch 2/20

108/108  15s 139ms/step - accuracy: 0.6656 - f1_score: 0.6655 - loss: 0.9375 - val_accuracy: 0.4444 - val_f1_score: 0.4186 - val_loss: 1.9477


Epoch 3/20

108/108  21s 142ms/step - accuracy: 0.8233 - f1_score: 0.8235 - loss: 0.5042 - val_accuracy: 0.5347 - val_f1_score: 0.5313 - val_loss: 1.7583


Epoch 4/20

108/108  15s 137ms/step - accuracy: 0.9275 - f1_score: 0.9275 - loss: 0.2410 - val_accuracy: 0.5301 - val_f1_score: 0.5176 - val_loss: 1.8613


Epoch 5/20

108/108  21s 138ms/step - accuracy: 0.9573 - f1_score: 0.9573 - loss: 0.1373 - val_accuracy: 0.5486 - val_f1_score: 0.5445 - val_loss: 2.2265


Epoch 6/20

108/108  22s 148ms/step - accuracy: 0.9773 - f1_score: 0.9773 - loss: 0.0895 - val_accuracy: 0.4884 - val_f1_score: 0.4884 - val_loss: 2.5013


Epoch 7/20

108/108  16s 152ms/step - accuracy: 0.9765 - f1_score: 0.9764 - loss: 0.0898 - val_accuracy: 0.5139 - val_f1_score: 0.4952 - val_loss: 2.6278


Epoch 8/20

108/108  15s 137ms/step - accuracy: 0.9753 - f1_score: 0.9753 - loss: 0.0752 - val_accuracy: 0.5417 - val_f1_score: 0.5298 - val_loss: 3.0524


Epoch 9/20

108/108  21s 139ms/step - accuracy: 0.9847 - f1_score: 0.9847 - loss: 0.0503 - val_accuracy: 0.5208 - val_f1_score: 0.5073 - val_loss: 3.1649


Epoch 10/20

108/108  20s 138ms/step - accuracy: 0.9891 - f1_score: 0.9890 - loss: 0.0498 - val_accuracy: 0.5278 - val_f1_score: 0.5217 - val_loss: 3.2112


Epoch 11/20


108/108  20s 136ms/step - accuracy: 0.9818 - f1_score: 0.9818 - loss: 0.0558 - val_accuracy: 0.5255 - val_f1_score: 0.5248 - val_loss: 3.4070

Epoch 12/20

108/108  21s 142ms/step - accuracy: 0.9493 - f1_score: 0.9493 - loss: 0.1980 - val_accuracy: 0.5486 - val_f1_score: 0.5482 - val_loss: 3.2304

Epoch 13/20


108/108  20s 135ms/step - accuracy: 0.9492 - f1_score: 0.9492 - loss: 0.1970 - val_accuracy: 0.5069 - val_f1_score: 0.4876 - val_loss: 4.1846

14/14  1s 64ms/step - accuracy: 0.5539 - f1_score: 0.5497 - loss: 1.8544


Test Loss: 1.6889742612838745, Test Accuracy: 0.5671296119689941, Test F1 Score: 0.5616947412490845

Time required to train the model is 250.11752653121948 seconds


Filters: [16, 32, 64], Kernels: [(3, 3), (5, 5), (5, 5)], max pool, relu activation function, No. of dense layers after flatten: 2



Filters: [16, 32, 64], Kernels: [(3, 3), (5, 5), (5, 5)], max pool, relu activation function, No. of dense layers after flatten: 2



14/14 1s 43ms/step



In []: result_df_7

Out[]:

	Conv Kernel Size	Conv Filter Size	Pooling Layers	Activation Function	No. of Dense Layers after Flatten	Dropout Rate	Test Loss	Test Accuracy	Test F1 Score	Batch Normalization Presence	No. of Extra Conv Layers	Color Mode	Training Time(in seconds)
0	[(3, 3), (5, 5), (5, 5)]	[16, 32, 64]	max	relu	2	0.0	1.591569	0.328704	0.307119	True	0	grayscale	206.139588
1	[(3, 3), (5, 5), (5, 5)]	[16, 32, 64]	max	relu	2	0.0	1.688974	0.567130	0.561695	True	0	rgb	250.117527


In []: plt.figure(figsize=(20,5))
plt.bar(


```

result_df_7['Color Mode'],
result_df_7['Test Accuracy']
)

plt.ylabel('Test Accuracy')
plt.xlabel('Color Mode')
plt.title('Validation Accuracy for various color modes')
plt.show()

```



Clearly the model performs better with rgb channel images

Question 05

- v. Plot the graph for the loss vs epoch and accuracy (train, test set) vs epoch for all the above cases. Also, plot the accuracy for all experimentation in a bar graph along with the confusion matrix and F1 Score.

Already done in continuation of question 4

Question 06

- vi. For the best model on the MNIST dataset in Assignment 4, train a model with MNIST data using the best set of parameters obtained in Question *iv*. Compare the test accuracy and the self-created images.

```
In [ ]: from keras.datasets import mnist
        (X_train, y_train), (X_test, y_test) = mnist.load_data()
```

```
In [ ]: X_train.shape, y_train.shape, X_test.shape, y_test.shape
```

```
Out[ ]: ((60000, 28, 28), (60000,)), (10000, 28, 28), (10000,))
```

```
In [ ]: X_train = np.concatenate((X_train, X_test),axis=0)
        y_train = np.concatenate((y_train, y_test),axis=0)
        X_train.shape
```

```
Out[ ]: (70000, 28, 28)
```

```
In [ ]: # Training Data -> 80% , Validation Data -> 10% , Testing Data -> 10%
```

```
from sklearn.model_selection import train_test_split
```

```
X_train, X_rem, y_train, y_rem = train_test_split(X_train, y_train, test_size=0.2, random_state=4)
X_val, X_test, y_val, y_test = train_test_split(X_rem, y_rem, test_size=0.5, random_state=4)
```

```
In [ ]: # Resizing to 80x80
```

```
from tensorflow.image import resize
```


```
X_train = np.reshape(X_train, (-1,28,28,1))
X_train=X_train/255
X_train = resize(X_train, (80,80))
```

```
X_val = np.reshape(X_val, (-1,28,28,1))
X_val=X_val/255
X_val = resize(X_val, (80,80))
```

```
X_test = np.reshape(X_test, (-1,28,28,1))
X_test=X_test/255
X_test = resize(X_test, (80,80))
```

```
In [ ]: plt.imshow(X_train[23], cmap='Greys')
```

```
Out[ ]: <matplotlib.image.AxesImage at 0x7a38cf234760>
```



```
In [ ]: y_train[23]
```

```
Out[ ]: 3
```

```
In [ ]: # Categorical encoding of y
y_train = pd.get_dummies(y_train, dtype='int').to_numpy()
y_val = pd.get_dummies(y_val, dtype='int').to_numpy()
y_test = pd.get_dummies(y_test, dtype='int').to_numpy()

y_train[23]
```

```
Out[ ]: array([0, 0, 0, 1, 0, 0, 0, 0, 0, 0])
```

```
In [ ]: result_df_8 = pd.DataFrame(
    columns=[
        'Conv Kernel Size',
        'Conv Filter Size',
```

```

        'Pooling Layers',
        'Activation Function',
        'No. of Dense Layers after Flatten',
        'Dropout Rate',
        'Test Loss',
        'Test Accuracy',
        'Test F1 Score',
        'Batch Normalization Presence',
        'No. of Extra Conv Layers',
        'Color Mode',
        'Training Time(in seconds)'
    ]
)

filters = [16,32,64]
epochs = 20

test_loss,test_accuracy,test_f1,train_time,model = train_model(
    kernels=best_kernel,
    filters=filters,
    activation_func=best_activation_function,
    pool=best_pool,
    dropout_rate=best_dropout,
    num_dense_layers=best_num_dense,
    X_train=X_train,
    y_train=y_train,
    X_test=X_test,
    y_test=y_test,
    num_epochs=epochs,
    add_batch_normalization=best_do_batch,
    extra_conv_layers=best_conv,
    is_mnist=True
)

result_df_8.loc[len(result_df_8.index)]= [
    best_kernel,
    filters,
    best_pool,
    best_activation_function,
    best_num_dense,
    best_dropout,
    test_loss,
    test_accuracy,
    test_f1,
    best_do_batch,
    best_conv,

```

```
'grayscale',  
train_time  
]
```

















Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	?	0 (unbuilt)
max_pooling2d (MaxPooling2D)	?	0 (unbuilt)
batch_normalization (BatchNormalization)	?	0 (unbuilt)
flatten (Flatten)	?	0 (unbuilt)
dense (Dense)	?	0 (unbuilt)
dense_1 (Dense)	?	0 (unbuilt)
dense_2 (Dense)	?	0 (unbuilt)

Total params: 0 (0.00 B)


Trainable params: 0 (0.00 B)

Non-trainable params: 0 (0.00 B)


Epoch 1/20
1750/1750  **14s** 4ms/step - accuracy: 0.9007 - f1_score: 0.9007 - loss: 0.3188 - val_accuracy: 0.9584 - val_f1_score: 0.9585 - val_loss: 0.1457
Epoch 2/20
1750/1750  **16s** 4ms/step - accuracy: 0.9731 - f1_score: 0.9731 - loss: 0.0883 - val_accuracy: 0.9664 - val_f1_score: 0.9664 - val_loss: 0.1161
Epoch 3/20
1750/1750  **6s** 3ms/step - accuracy: 0.9824 - f1_score: 0.9824 - loss: 0.0580 - val_accuracy: 0.9663 - val_f1_score: 0.9665 - val_loss: 0.1329
Epoch 4/20
1750/1750  **11s** 4ms/step - accuracy: 0.9855 - f1_score: 0.9855 - loss: 0.0465 - val_accuracy: 0.9684 - val_f1_score: 0.9685 - val_loss: 0.1333
Epoch 5/20
1750/1750  **6s** 4ms/step - accuracy: 0.9881 - f1_score: 0.9881 - loss: 0.0382 - val_accuracy: 0.9736 - val_f1_score: 0.9736 - val_loss: 0.1055
Epoch 6/20
1750/1750  **10s** 4ms/step - accuracy: 0.9909 - f1_score: 0.9909 - loss: 0.0275 - val_accuracy: 0.9657 - val_f1_score: 0.9657 - val_loss: 0.1633
Epoch 7/20
1750/1750  **10s** 4ms/step - accuracy: 0.9925 - f1_score: 0.9925 - loss: 0.0257 - val_accuracy: 0.9757 - val_f1_score: 0.9757 - val_loss: 0.1231
Epoch 8/20
1750/1750  **7s** 4ms/step - accuracy: 0.9927 - f1_score: 0.9927 - loss: 0.0239 - val_accuracy: 0.9697 - val_f1_score: 0.9696 - val_loss: 0.1456
Epoch 9/20
1750/1750  **7s** 4ms/step - accuracy: 0.9943 - f1_score: 0.9943 - loss: 0.0216 - val_accuracy: 0.9650 - val_f1_score: 0.9650 - val_loss: 0.2718
Epoch 10/20
1750/1750  **7s** 4ms/step - accuracy: 0.9934 - f1_score: 0.9934 - loss: 0.0234 - val_accuracy: 0.9723 - val_f1_score: 0.9723 - val_loss: 0.1539
Epoch 11/20
1750/1750  **10s** 4ms/step - accuracy: 0.9957 - f1_score: 0.9957 - loss: 0.0141 - val_accuracy: 0.9733 - val_f1_score: 0.9734 - val_loss: 0.1454
Epoch 12/20
1750/1750  **10s** 4ms/step - accuracy: 0.9956 - f1_score: 0.9956 - loss: 0.0159 - val_accuracy: 0.9743 - val_f1_score: 0.9743 - val_loss: 0.1380
Epoch 13/20
1750/1750  **10s** 3ms/step - accuracy: 0.9947 - f1_score: 0.9947 - loss: 0.0175 - val_accuracy: 0.9776 - val_f1_score: 0.9776 - val_loss: 0.1385
Epoch 14/20
1750/1750  **7s** 4ms/step - accuracy: 0.9962 - f1_score: 0.9962 - loss: 0.0151 - val_accuracy: 0.9664 - val_f1_score: 0.9664 - val_loss: 0.2354
Epoch 15/20
1750/1750  **6s** 3ms/step - accuracy: 0.9956 - f1_score: 0.9956 - loss: 0.0149 - val_accuracy: 0.9793 - val_f1_score: 0.9793 - val_loss: 0.1584
219/219  **0s** 2ms/step - accuracy: 0.9758 - f1_score: 0.9759 - loss: 0.0872

Test Loss: 0.09898536652326584, Test Accuracy: 0.9762856960296631, Test F1 Score: 0.9762869477272034
Time required to train the model is 142.12913298606873 seconds

Filters: [16, 32, 64], Kernels: [(3, 3), (5, 5), (5, 5)], max pool, relu activation function, No. of dense layers after flatten: 2




Filters: [16, 32, 64], Kernels: [(3, 3), (5, 5), (5, 5)], max pool, relu activation function, No. of dense layers after flatten: 2



219/219 ————— 1s 3ms/step

Confusion Matrix for the above model




```
In [ ]: import random

random_indices = random.sample(range(0,len(X_test)),10)

img_predict = np.array([X_test[ind] for ind in random_indices])
true_label = np.array([y_test[ind] for ind in random_indices])


cat_pred = model.predict(img_predict)

for img,cat_pred,true_val in zip(img_predict,cat_pred,true_label):
    plt.imshow(img,cmap="Greys")
    plt.show()
    prediction=np.argmax(cat_pred)
    print(f"Predicted value:- {prediction}")
    print(f"True value:- {np.argmax(true_val)}")
```




Predicted value:- 3

True value:- 3




Predicted value:- 9

True value:- 9




Predicted value:- 4

True value:- 4




Predicted value:- 6

True value:- 6




Predicted value:- 3

True value:- 3




Predicted value:- 7

True value:- 7




Predicted value:- 6

True value:- 6




Predicted value:- 3

True value:- 3



Predicted value:- 8

True value:- 8



Predicted value:- 8

True value:- 8

As we can see there is 100% accuracy achieved here, which was not the case for the earlier assignment's ANN model